



Flink Connector 分享

董亭亭 · 快手/高级开发工程师

Apache Flink Community China



Apache Flink

CONTENT

目录 >>

01 /

Flink Streaming Connectors

02 /

Flink Kafka Connector

03 /

Q&A

01

Flink Streaming Connectors

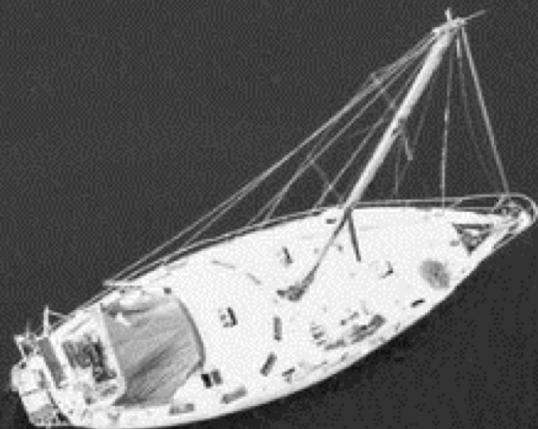
Streaming Connectors

预定义的Source和Sink

Bundled Connectors

Apache Bahir中的连接器

Async I/O





预定义的Source和Sink



基于文件的Source

`readTextFile(path)`

`readFile(fileInputFormat, path)`



基于Socket

`socketTextStream`



基于Collections、Iterators

`fromCollection`、`fromElements`

基于文件的Sink

`writeAsText`

`writeAsCsv`

基于Socket的Sink

`writeToSocket`

标准输出、标准错误

`Print`、`printToError`



Bundled Connectors

- [Apache Kafka](#) (source/sink)
 - [Apache Cassandra](#) (sink)
 - [Amazon Kinesis Streams](#) (source/sink)
 - [Elasticsearch](#) (sink)
 - [Hadoop FileSystem](#) (sink)
 - [RabbitMQ](#) (source/sink)
 - [Apache NiFi](#) (source/sink)
 - [Twitter Streaming API](#) (source)
- **注意：**以上流connector是Flink项目的一部分，但是不包括在二进制发布包中！



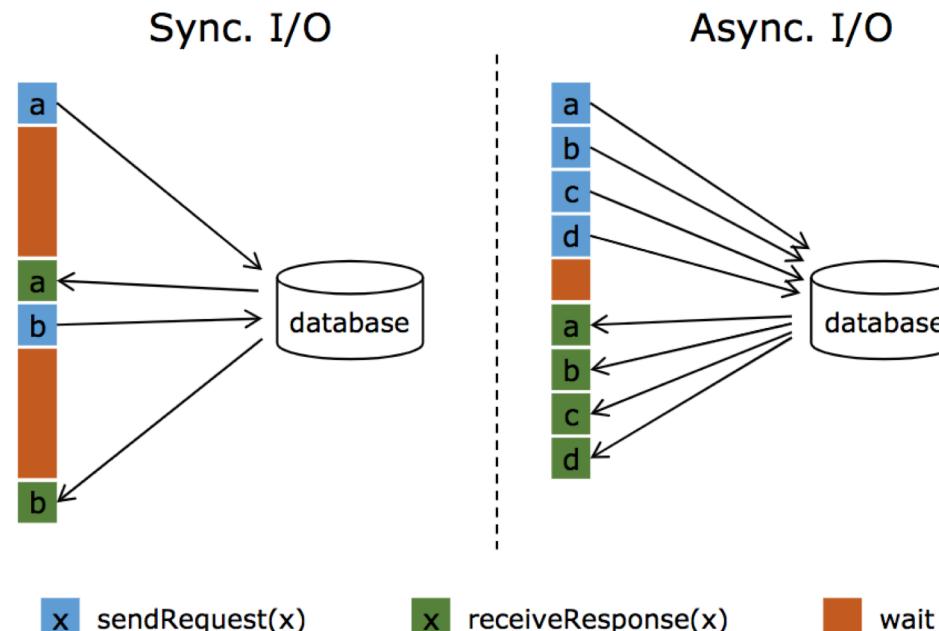
Apache Bahir中的连接器

- [Apache ActiveMQ](#)(source/sink)
- [Apache Flume](#)(sink)
- [Redis](#)(sink)
- [Akka](#)(sink)
- [Netty](#)(Source)



Async I/O

- 使用connector并不是数据输入输出Flink的唯一方式。
- 在Map、FlatMap中使用Async I/O方式读取外部数据库等。



<https://ci.apache.org/projects/flink/flink-docs-release-1.3/dev/stream/asyncio.html>

02

Flink Kafka Connector



Apache Flink

Flink Kafka Connector

Flink Kafka Consumer

- 反序列化数据
- 消费起始位置设置
- Topic和Partition动态发现
- Commit Offset方式
- Timestamp Extraction/Watermark 生成

Flink Kafka Producer

- Producer分区
- 容错



Flink Kafka Example

```
public class FlinkKafkaExample {
    public static void main(String[] args) throws Exception{
        final ParameterTool params = ParameterTool.fromArgs(args);

        // set up the execution environment
        final StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();
        env.getConfig().setGlobalJobParameters(params);
        env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);
        env.enableCheckpointing( interval: 60*1000, CheckpointingMode.EXACTLY_ONCE);

        String sourceTopic = "topic1";
        String bootStrapServers = "localhost:9092";source
        Properties properties = getConsumerProperties(bootStrapServers, groupId: "test-dtt-gi");
        FlinkKafkaConsumer010<String> consumer = new FlinkKafkaConsumer010<~>(sourceTopic, new SimpleStringSchema(), properties);
        //consumer.setStartFromEarliest();
        //consumer.setStartFromLatest();
        //consumer.setStartFromTimestamp(1561281792000L);
        //consumer.setStartFromGroupOffsets() //default

        String sinkTopic = "topic2";sink
        Properties prop = getProducerProperties(bootStrapServers);
        FlinkKafkaProducer010<String> producer = new FlinkKafkaProducer010<~>(sinkTopic, new SimpleStringSchema(), prop);

        env.addSource(consumer) DataStreamSource<String>
            .map(new MapFunction<String, Tuple2<Long, String>>() {
                @Override
                public Tuple2<Long, String> map(String s) throws Exception {...}
            }) SingleOutputStreamOperator<Tuple2<Long, String>>
            .filter(k -> k!=null) SingleOutputStreamOperator<Tuple2<Long, String>>
            .assignTimestampsAndWatermarks(new BoundedOutOfOrdernessTimestampExtractor<Tuple2<Long, String>>(Time.seconds(5)) {
                @Override
                public long extractTimestamp(Tuple2<Long, String> element) { return element.f0; }
            }) SingleOutputStreamOperator<Tuple2<Long, String>>
            .map(k -> k.toString()) SingleOutputStreamOperator<String>
            .addSink(producer);

        env.execute( jobName: "FlinkKafkaExample");
    }
}
```

Flink Kafka Consumer-反序列化数据

- 将kafka中二进制数据转化为具体的java、 scala 对象
- DeserializationSchema , T deserialize(byte[] message)
- KeyedDeserializationSchema , T deserialize(byte[] messageKey, byte[] message, String topic, int partition, long offset): 对于访问kafka key/value

常用

- SimpleStringSchema: 按字符串方式进行序列化、反序列化
- TypeInformationSerializationSchema: 基于Flink的TypeInformation来创建schema
- JsonDeserializationSchema: 使用jackson反序列化json格式消息，并返回
ObjectNode，可以使用.get(“property”)方法来访问字段。

Flink Kafka Consumer-消费起始位置

- **setStartFromGroupOffsets(默认)**
 - 从Kafka记录的group.id 的位置开始读取，如果没有根据auto.offset.reset设置
- **setStartFromEarliest**
 - 从kafka最早的位置读取
- **setStartFromLatest**
 - 从kafka最新数据开始读取
- **setStartFromTimestamp(long)**
 - 从时间戳大于或等于指定时间戳的位置开始读取
- **setStartFromSpecificOffsets**
 - 从指定的分区的offset位置开始读取，如指定的offsets中不存某个分区，该分区从group offset位置开始读取。

Flink Kafka Consumer-消费起始位置

```
public static void main(String[] args) throws Exception{
{
    final ParameterTool params = ParameterTool.fromArgs(args);

    // set up the execution environment
    final StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();
    env.getConfig().setGlobalJobParameters(params);
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);
    env.enableCheckpointing( interval: 60*1000, CheckpointingMode.EXACTLY_ONCE);

    String sourceTopic = "topic1";
    String bootStrapServers = "localhost:9092";
    Properties properties = getConsumerProperties(bootStrapServers, groupid: "test-gid");
    FlinkKafkaConsumer010<String> consumer = new FlinkKafkaConsumer010<~>(sourceTopic, new SimpleStringSchema(), properties);
    consumer.setStartFromEarliest();
    consumer.setStartFromLatest();
    consumer.setStartFromTimestamp(1561281792000L);
    consumer.setStartFromGroupOffsets();

    Map<KafkaTopicPartition, Long> specificStartOffsets = new HashMap<>();
    specificStartOffsets.put(new KafkaTopicPartition(sourceTopic, partition: 0), 0L);
    specificStartOffsets.put(new KafkaTopicPartition(sourceTopic, partition: 1), 1L);
    consumer.setStartFromSpecificOffsets(specificStartOffsets);
}
```

- 注意：作业故障从checkpoint自动恢复，以及手动做savepoint时，消费的位置从保存状态中恢复，与该配置无关！！！

Flink Kafka Consumer-topic partition自动发现

原理：内部单独的线程获取kafka meta信息进行更新

flink.partition-discovery.interval-millis:发现时间间隔。默认false，设置非负值开启。



分区发现

- 消费的Source kafka topic进行了partition 扩容
- 新发现的分区，从earliest位置开始读取



Topic发现

- 支持正则表达式描述topic名字

```
String sourceTopic = "topic1";
String bootStrapServers = "localhost:9092";
Properties properties = getConsumerProperties(bootStrapServers, groupId: "test-gid");
//FlinkKafkaConsumer010<String> consumer = new FlinkKafkaConsumer010<String>(sourceTopic, new SimpleStringSchema(), properties);
Pattern topicPattern = java.util.regex.Pattern.compile("topic[0-9]");
FlinkKafkaConsumer010<String> consumer = new FlinkKafkaConsumer010<~>(topicPattern, new SimpleStringSchema(), properties);
```



Flink Kafka Consumer-commit offset方式



Checkpoint关闭

- 依赖kafka客户端的auto commit 定期提交offset。
- 需设置enable.auto.commit, auto.commit.interval.ms 参数到consumer properties



Checkpoint开启

- Offset自己在checkpoint state中管理和容错。 提交kafka 仅作为外部监视消费进度。
- 通过setCommitOffsetsOnCheckpoints控制， Checkpoint成功之后， 是否提交 offset到kafka

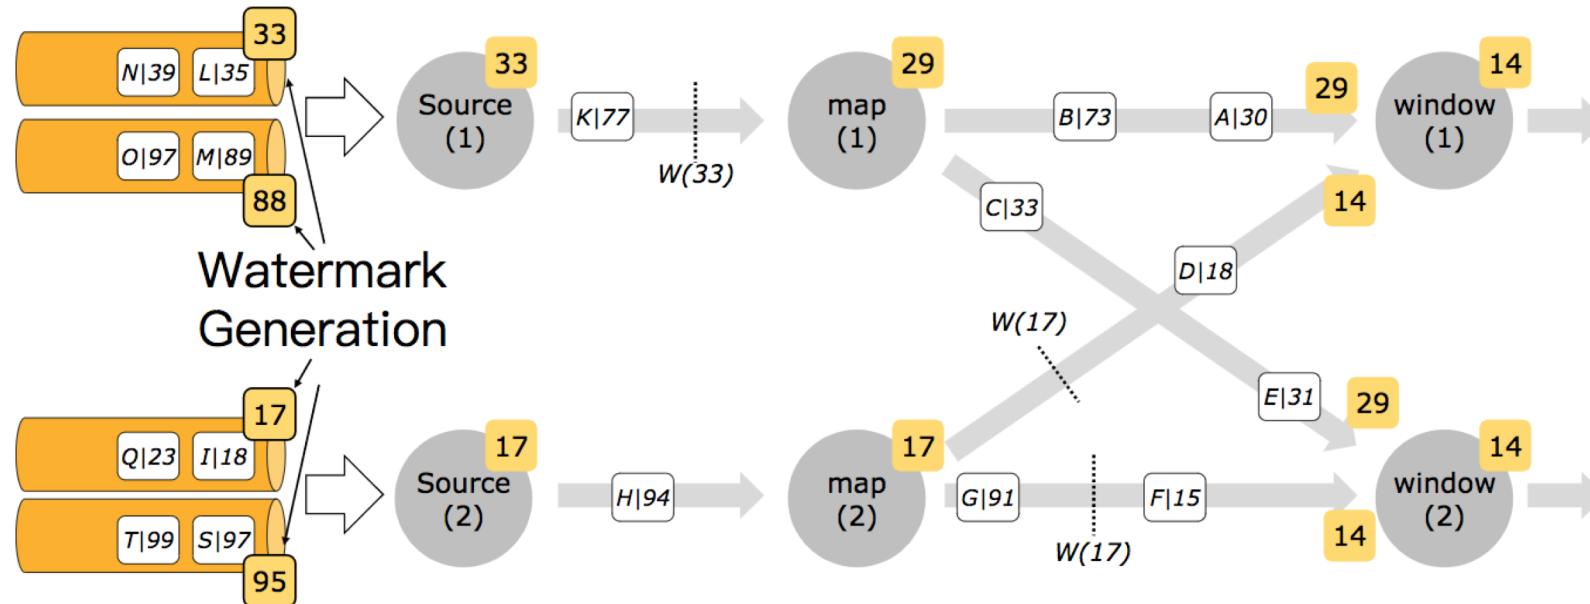


Flink Kafka Consumer-时戳提取/水位生成



per Kafka Partition watermark

- `assignTimestampsAndWatermarks`, 每个partition一个assigner, 水位为多个partition对齐后值
- 不在kafka source后生成watermark, 会出现扔掉部分数据情况。



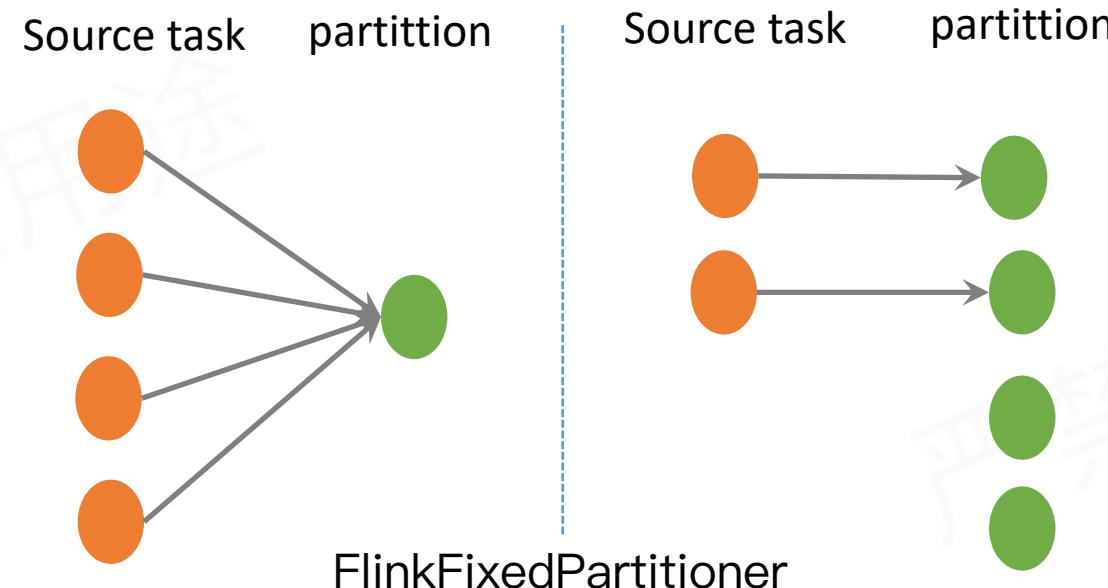


Flink Kafka Producer–Producer分区



Producer 分区

- FlinkFixedPartitioner (默认) : $\text{parallelInstanceID} \% \text{partitions.length}$
- Partitioner设置为null: round-robin kafka partitioner, 维持过多链接
- custom partitioner: 自定义分区



Flink Kafka Producer–Producer容错



Kafka 0.9 and 0.10

- setLogFailuresOnly: 默认false。写失败时，是否只打印失败log，不抛异常
- setFlushOnCheckpoint: 默认true。checkpoint时保证数据写到kafka。
- at-least-once语义: setLogFailuresOnly : flase + setFlushOnCheckpoint : true



Kafka 0.11

- FlinkKafkaProducer011，两阶段提交Sink结合kafka事务，可以保证端到端精准一次。

<https://www.ververica.com/blog/end-to-end-exactly-once-processing-apache-flink-apache-kafka>

03

Q&A



Apache Flink

THANKS

