

# Managing Flink on Kubernetes

Anand Swaminathan (@anand12100)  
Ketan Umare (@ketanumare)

April 2, 2019



# Agenda

1

## Kubernetes Primer

*Quick Introduction to concepts in Kubernetes*

2

## Background

*Summary of Lyft's legacy Flink Deployment*

3

## Solution

*Flink Kubernetes Operator*

4

## Demo

5

## Ecosystem

6

## Roadmap

# About us



# History

- **Google's** internal infrastructure is containerized and runs on **Borg/Omega**
- K8s was **open sourced** in **2014**, re-incarnation of the internal infrastructure
- Kubernetes **automates** - **deployment, scaling** and **management** of containerized apps.
- Containers are scheduled based on **CPU/GPU/Memory/Disk** etc

Production-Grade Container Scheduling and Management <https://kubernetes.io>

kubernetes

go

cncf

containers

 **76,018** commits


 **40** branches

 **489** releases

 **2,048** contributors

## Kubernetes Primer


# Pods



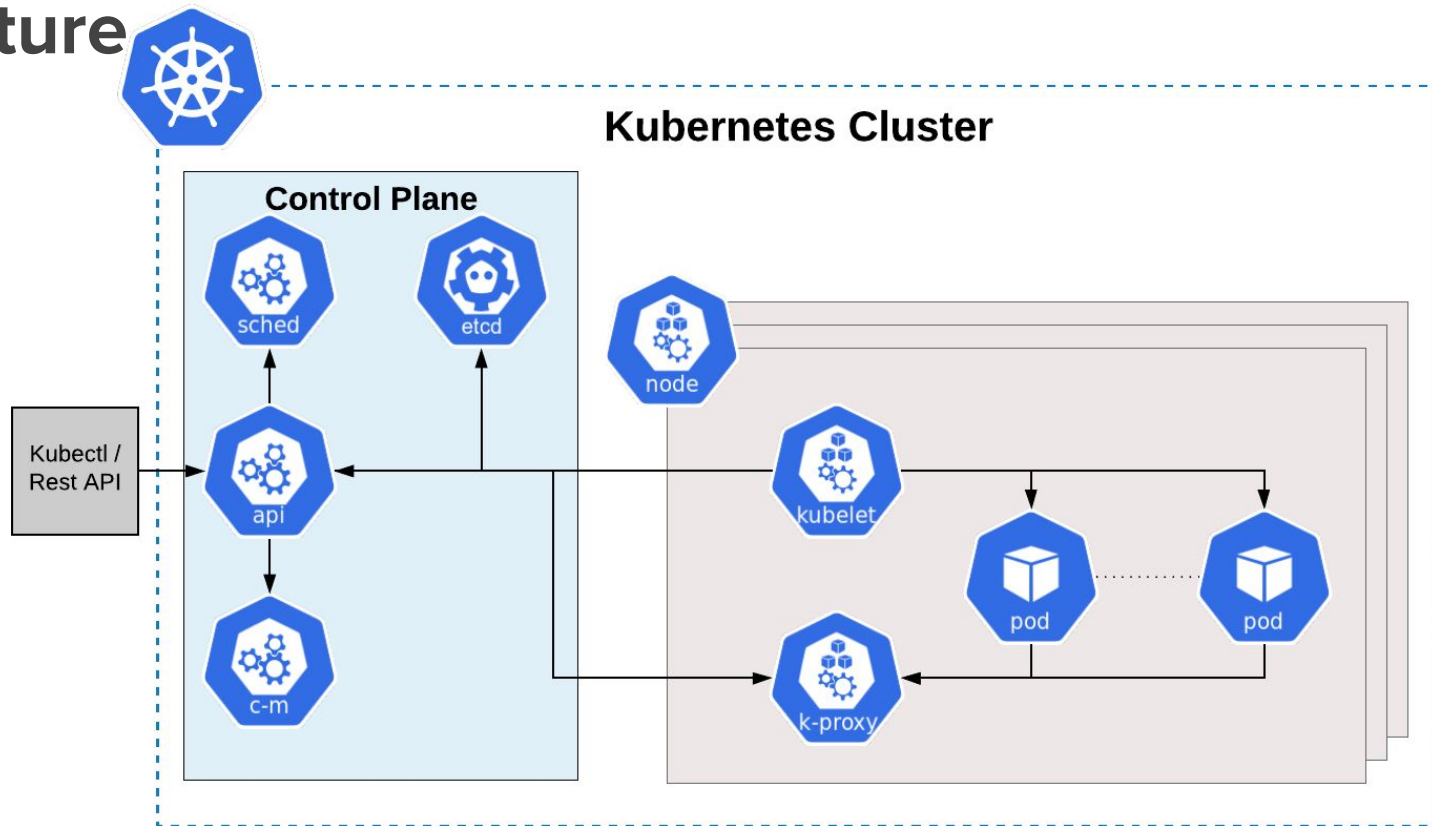
- A Pod is a group of **one or more Containers** as one unit
- Pods have **no durability** guarantees
- Each Pod has a **unique IP Address**
- **Containers** in a Pod can communicate using **localhost**
- **Multiple Pods** can be located on the **same node** - machine

# Other Concepts

- **Deployments** abstraction that enables rolling out changes to a set of pods
- **Service** abstraction to access sets of pods - like a load balancer within a k8s cluster
- **Ingress** abstraction to expose a service to the outside world (HTTP/HTTPS)
- **Controller** A reconciliation loop that drives current state towards desired state

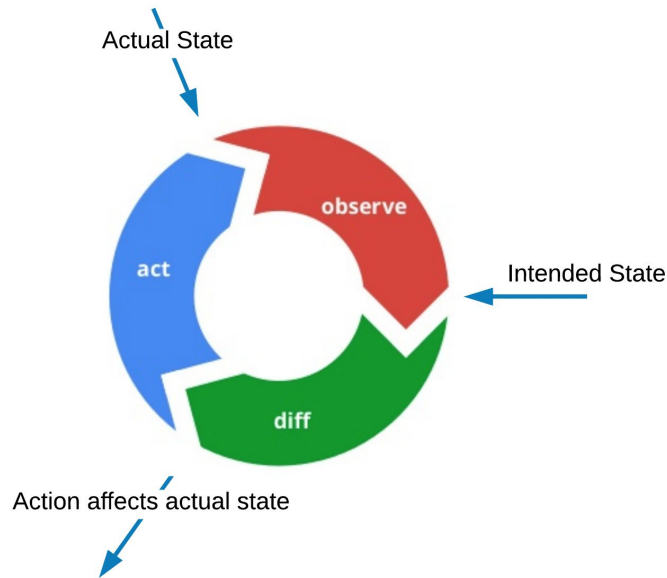


# Architecture



# Control Loops

- **Control loops** are fundamental building block industrial control systems
- **Desired State** refers to the intended state as requested
- **Current/Observed State** is the state of the system as **observed** by the controller
- Controller runs control loops
- Drive **Current State** -> **Desired State**
- This is the cornerstone of Kubernetes





# Custom Resources

- Custom Resource Definitions (**CRD**) allow **extending** Kubernetes API
- Custom resources are **optional** extensions
- Custom resources can be **added**/removed **dynamically**
- They can be manipulated using known tools - **kubectl** & kube clients
- State stored in **etcd**
- Custom control loops (**controllers**) are used to manage the state of the resource.
- CRD is essentially the **desired state**.

# Operators

- **Controller + CRD = Kubernetes Operator**
- Term coined by **CoreOS - 2017**
- **Manages** a **complex applications** lifecycle on Kubernetes.
- Core library to author operators @ [SIG/controller-runtime](https://github.com/kubernetes/community/tree/master/sig-controller-runtime)



## Background

# OK how does this relate

- @Lyft we started working on **Flyte** - a modern take at Pipelines/Workflows
- Orchestration is pervasive **throughout various sectors** of our Industry
  - Machine learning
  - Data engineering and processing
  - ETL
- Kubernetes has a solution to many of our problems
  - Deployment, Versioning, cluster management etc
- In parallel **Streaming Platform** started working on Flink for streaming applications

## Background

# Legacy deployment of Flink @Lyft

- Hosted on **AWS**
- Separate **AutoScalingGroups** for Task Managers and Job Managers
- Machines provisioned and bootstrapped by **SaltStack**
- Every deployment needs **provisioning of machines**
- Users started running multiple jobs in the same Flink Cluster
- **Multi-tenancy hell !**

# **Introducing Flink-k8s-operator**

# Goals

- Abstract out the complexity from application developers
  - **Hosting**
  - **Configuration**
  - **Management**
- Separate Flink cluster for each Flink application.
- **Deploy and rollback** support
- Support Flink application updates - **scaling**
- Simplified interface for instituting best practices
- **Scale to 100s** of flink applications

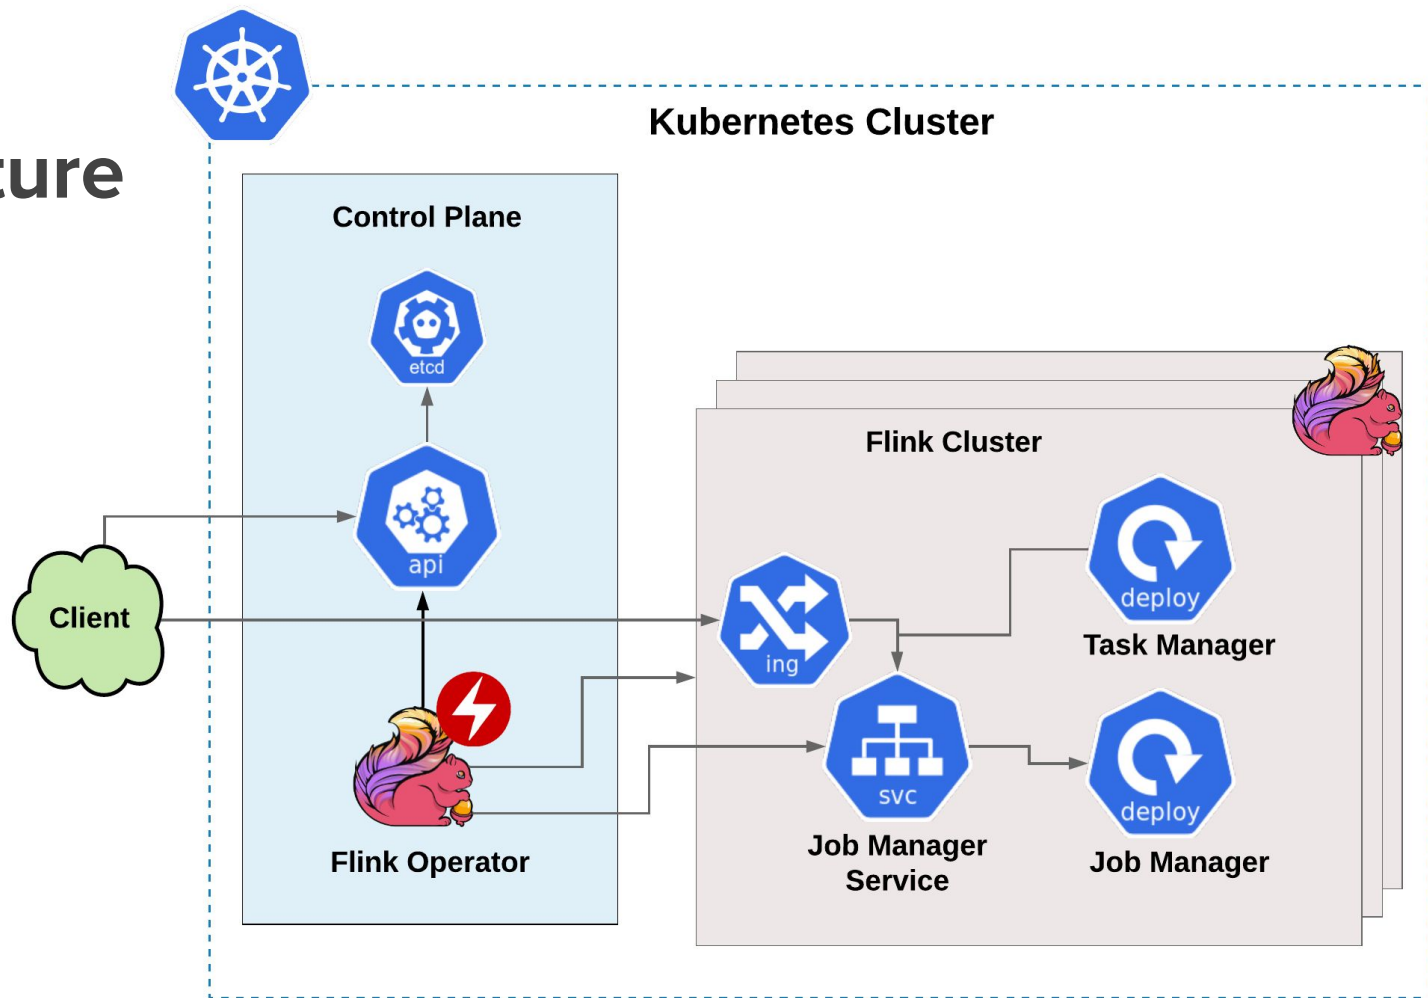
## Solution

# Flink Operator - CRD

- Each custom resource corresponds to a Flink application
- Each Flink application runs a single **Flink job**
- **Docker image** should be runnable

```
apiVersion: flink.k8s.io/v1alpha1
kind: FlinkApplication
metadata:
  name: flink-speeds-working-stats
  namespace: flink
  annotations:
    iam.amazonaws.com/role: 'arn:aws:iam::100:role/abc-iam'
  labels:
    app: app-name
    environment: staging
spec:
  image: '100.dkr.ecr.us-east-1.amazonaws.com/abc:xyz'
  flinkJob:
    jarName: name.jar
    parallelism: 10
  deploymentMode: Single
```

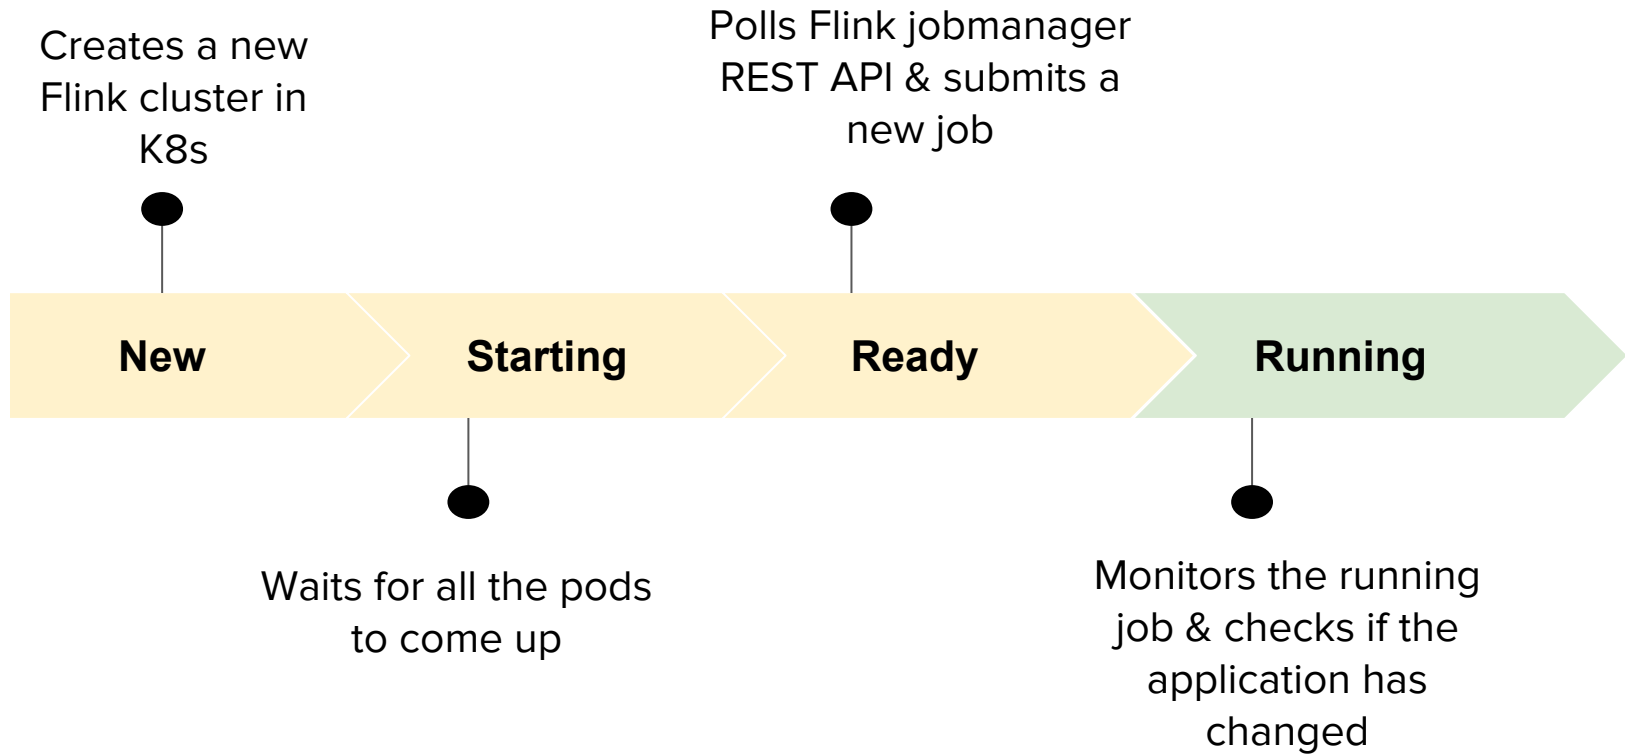
# Solution Architecture





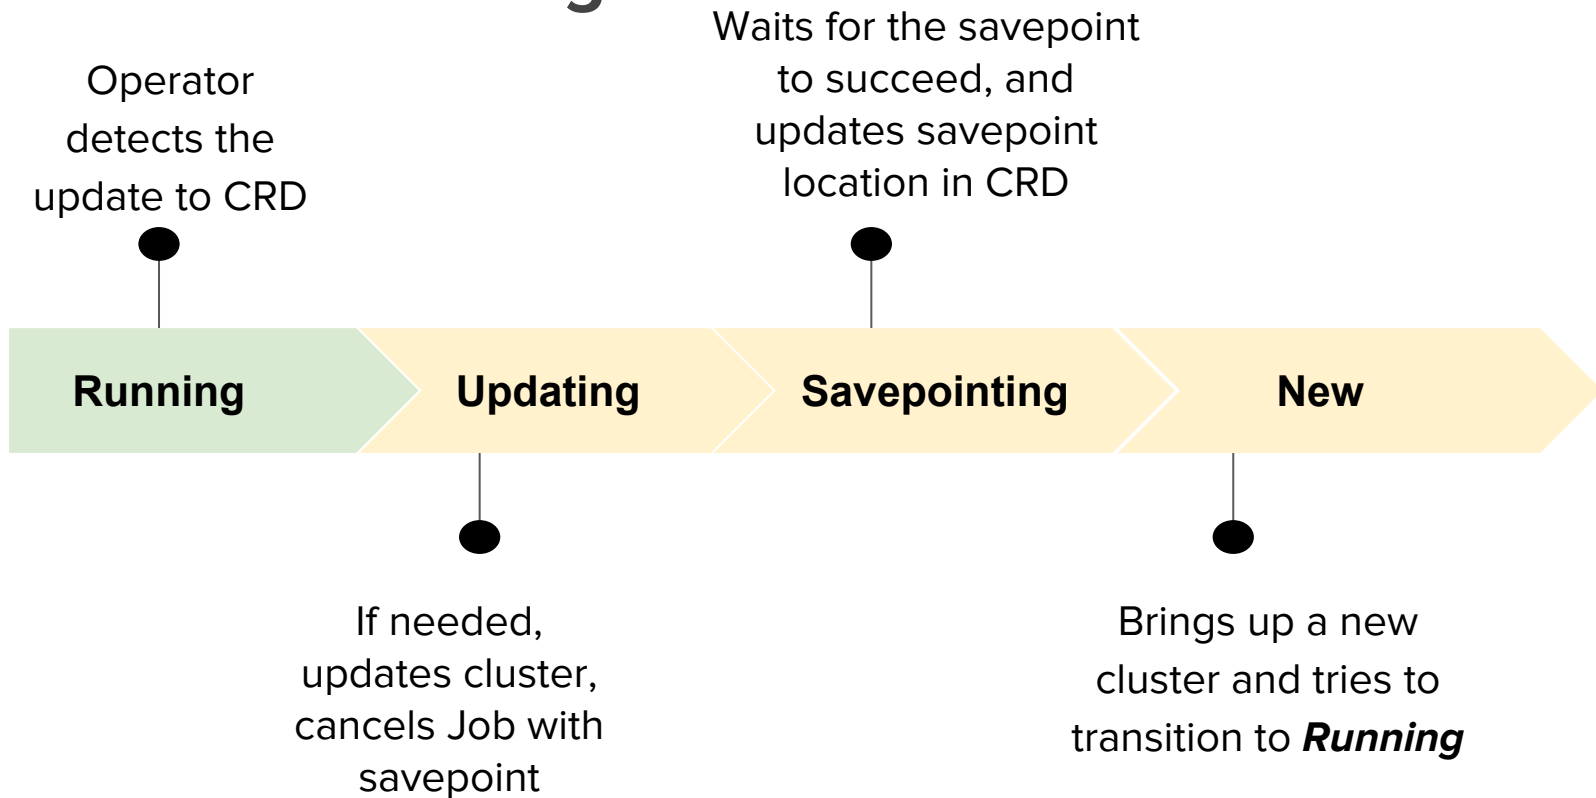
## Solution

# Operator Walkthrough



## Solution

# Operator Walkthrough

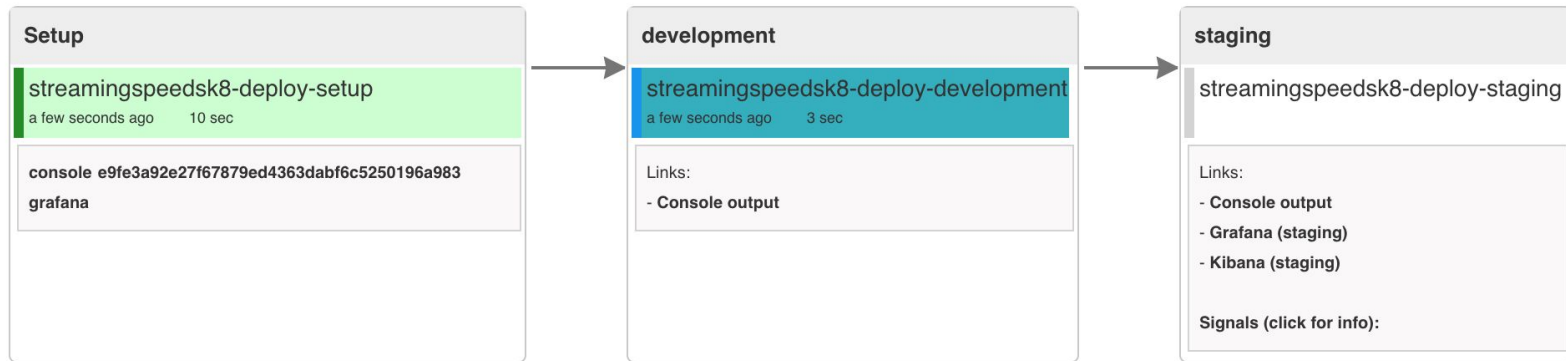


**Demo**

## Ecosystem

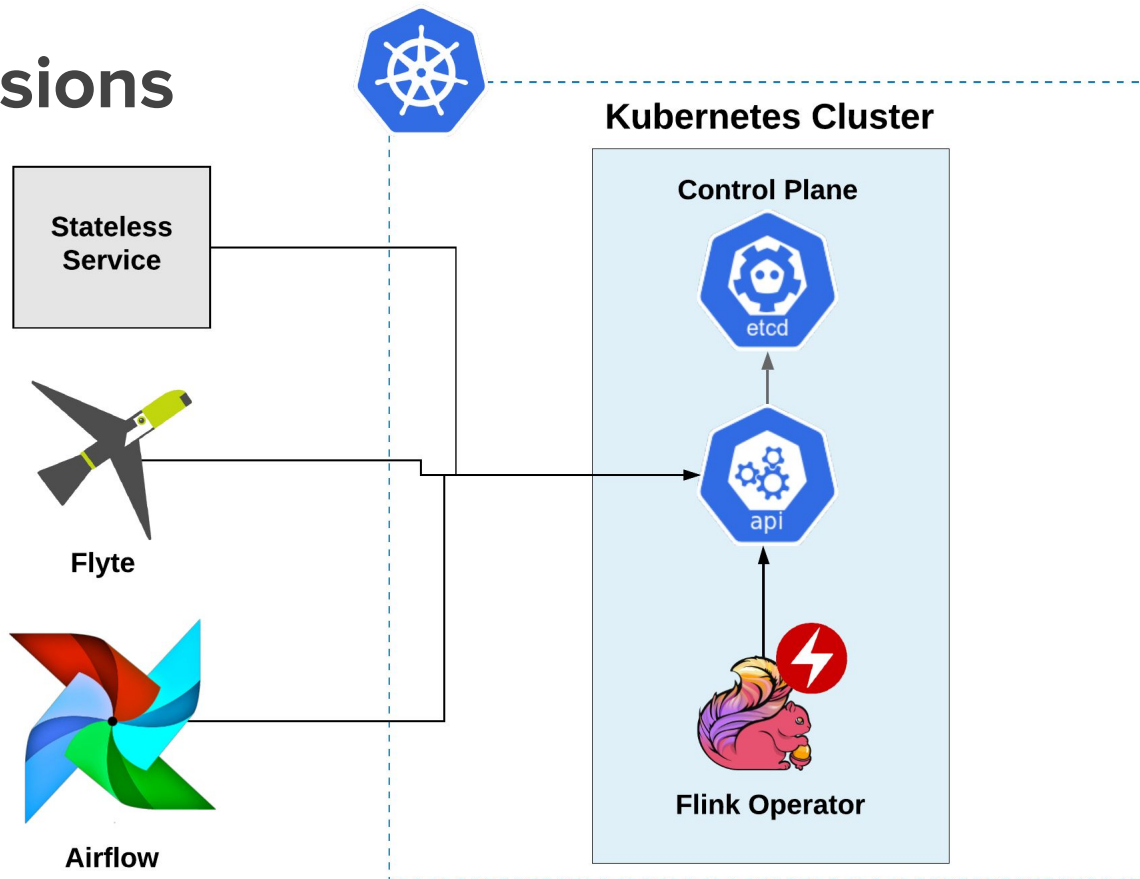
# Deployment @Lyft

- Jenkins based deployment
- Each stage creates or updates the resource in Kubernetes



Ecosystem

# Future Extensions



## Roadmap

# Open Source

- **Last week of April\***
- Project status: **Alpha**
- @Lyft:
  - Active development and testing in staging.
- Future
  - Flink Job failure handling
  - Tooling to manage CRD

Coming soon: <https://github.com/lyft/flinkk8soperator>

**We're Hiring!** Apply at [www.lyft.com/careers](https://www.lyft.com/careers)

**Data Engineering**

**Engineering Manager**  
San Francisco

**Software Engineer**  
San Francisco, Seattle, &  
New York City

**Data Infrastructure**

**Engineering Manager**  
San Francisco

**Software Engineer**  
San Francisco & Seattle

**Experimentation**

**Software Engineer**  
San Francisco

**Observability**

**Software Engineer**  
San Francisco

**Streaming**

**Software Engineer**  
San Francisco


**Thank you**  
**Questions please!**



## Background

# Example of Deployment

- User requests for a **Deployment** @ master
- Master accepts the request
- **Desired State:** 1 Pod running
- **Current State:** 0 Pods running



## Background

# Kubernetes 101

1. Master requests **Pod creation**
  - Current State: Deployment unhealthy
2. Master receives **pod created** event
  - Current State: Deployment healthy
3. Now if the pod crashes/dies etc
  - Current State: Deployment unhealthy
4. Goto 1

