

ADVENTURES IN SCALING FROM ZERO TO 5 BILLION DATA POINTS PER DAY

Dave Torok
Distinguished Architect
Comcast Corporation

2 April, 2019

Flink Forward – San Francisco 2019



COMCAST CUSTOMER RELATIONSHIPS

**30.3 MILLION OVERALL CUSTOMER
RELATIONSHIPS AT 2018 YEAR END**

**25.1 MILLION RESIDENTIAL HIGH-SPEED
INTERNET CUSTOMERS AT 2018 YEAR
END**

**1.2 MILLION RESIDENTIAL HIGH-SPEED
INTERNET CUSTOMER NET ADDITIONS IN
2018**



DELIVER THE ULTIMATE CUSTOMER EXPERIENCE

IS THE CUSTOMER HAVING A GOOD EXPERIENCE
FOR HIGH SPEED DATA (HSD) SERVICE?



IF THE CUSTOMER ENGAGES US DIGITALLY, CAN
WE OFFER A SELF-SERVICE SOLUTION?

IF THERE IS AN ISSUE CAN WE OFFER OUR AGENTS
AND TECHNICIANS A DIAGNOSIS TO HELP SOLVE
THE PROBLEM QUICKER?




REDUCE THE TIME TO RESOLVE ISSUES

REDUCE COST TO THE BUSINESS AND THE
CUSTOMER

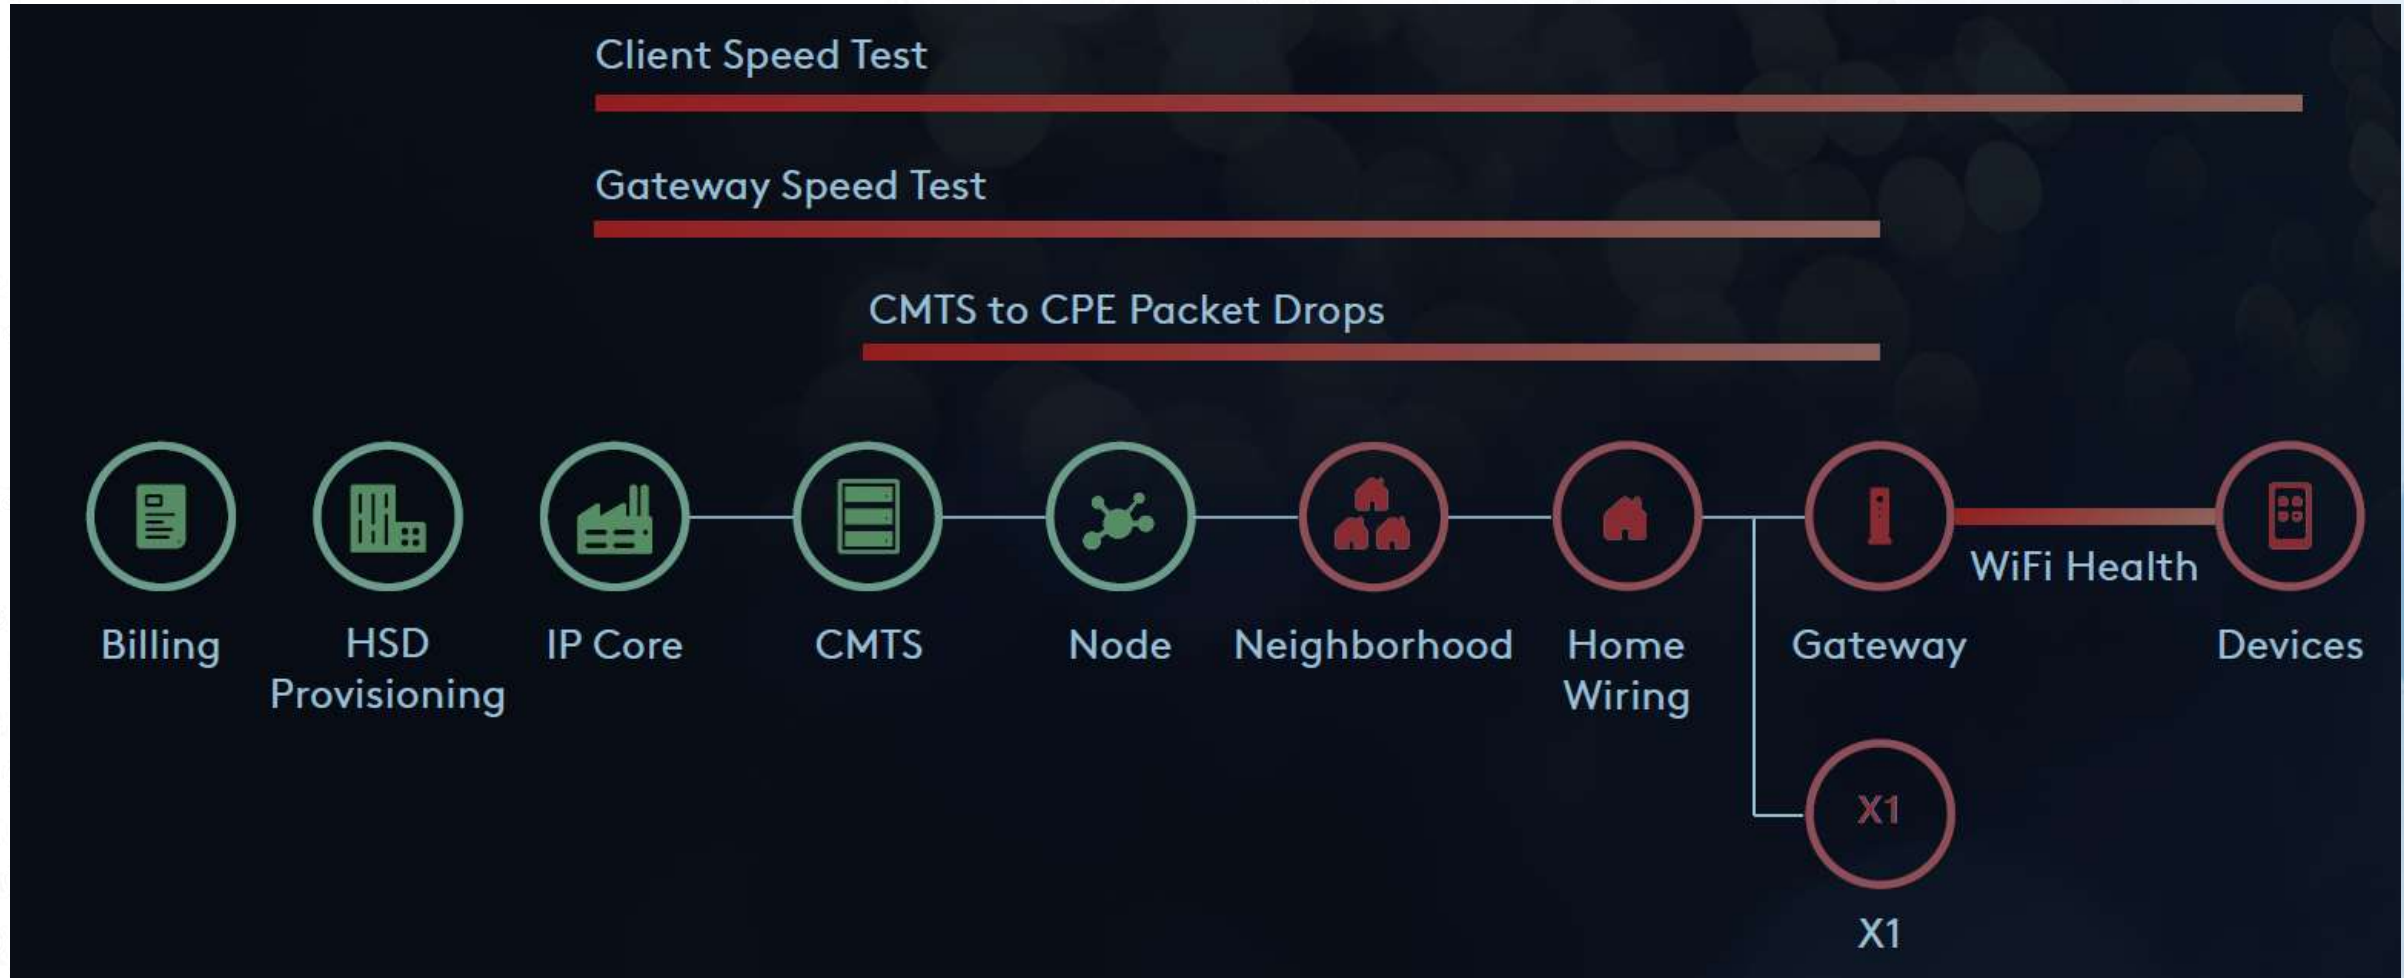


HIGH LEVEL CONCEPT



Comcast collects, stores, and uses all data in accordance with our privacy disclosures to users and applicable laws.

CUSTOMER EXPERIENCE INDICATORS



Comcast collects, stores, and uses all data in accordance with our privacy disclosures to users and applicable laws.

NINE ADVENTURES IN SCALING

THE
**TRIGGER AND
DIAGNOSIS**
PROBLEM

THE
REST
PROBLEM

THE
**INEFFICIENT
OBJECT HANDLING**
PROBLEM

THE
FEATURE STORE
PROBLEM

THE
VOLUME
PROBLEM

THE
CUSTOMER STATE
PROBLEM

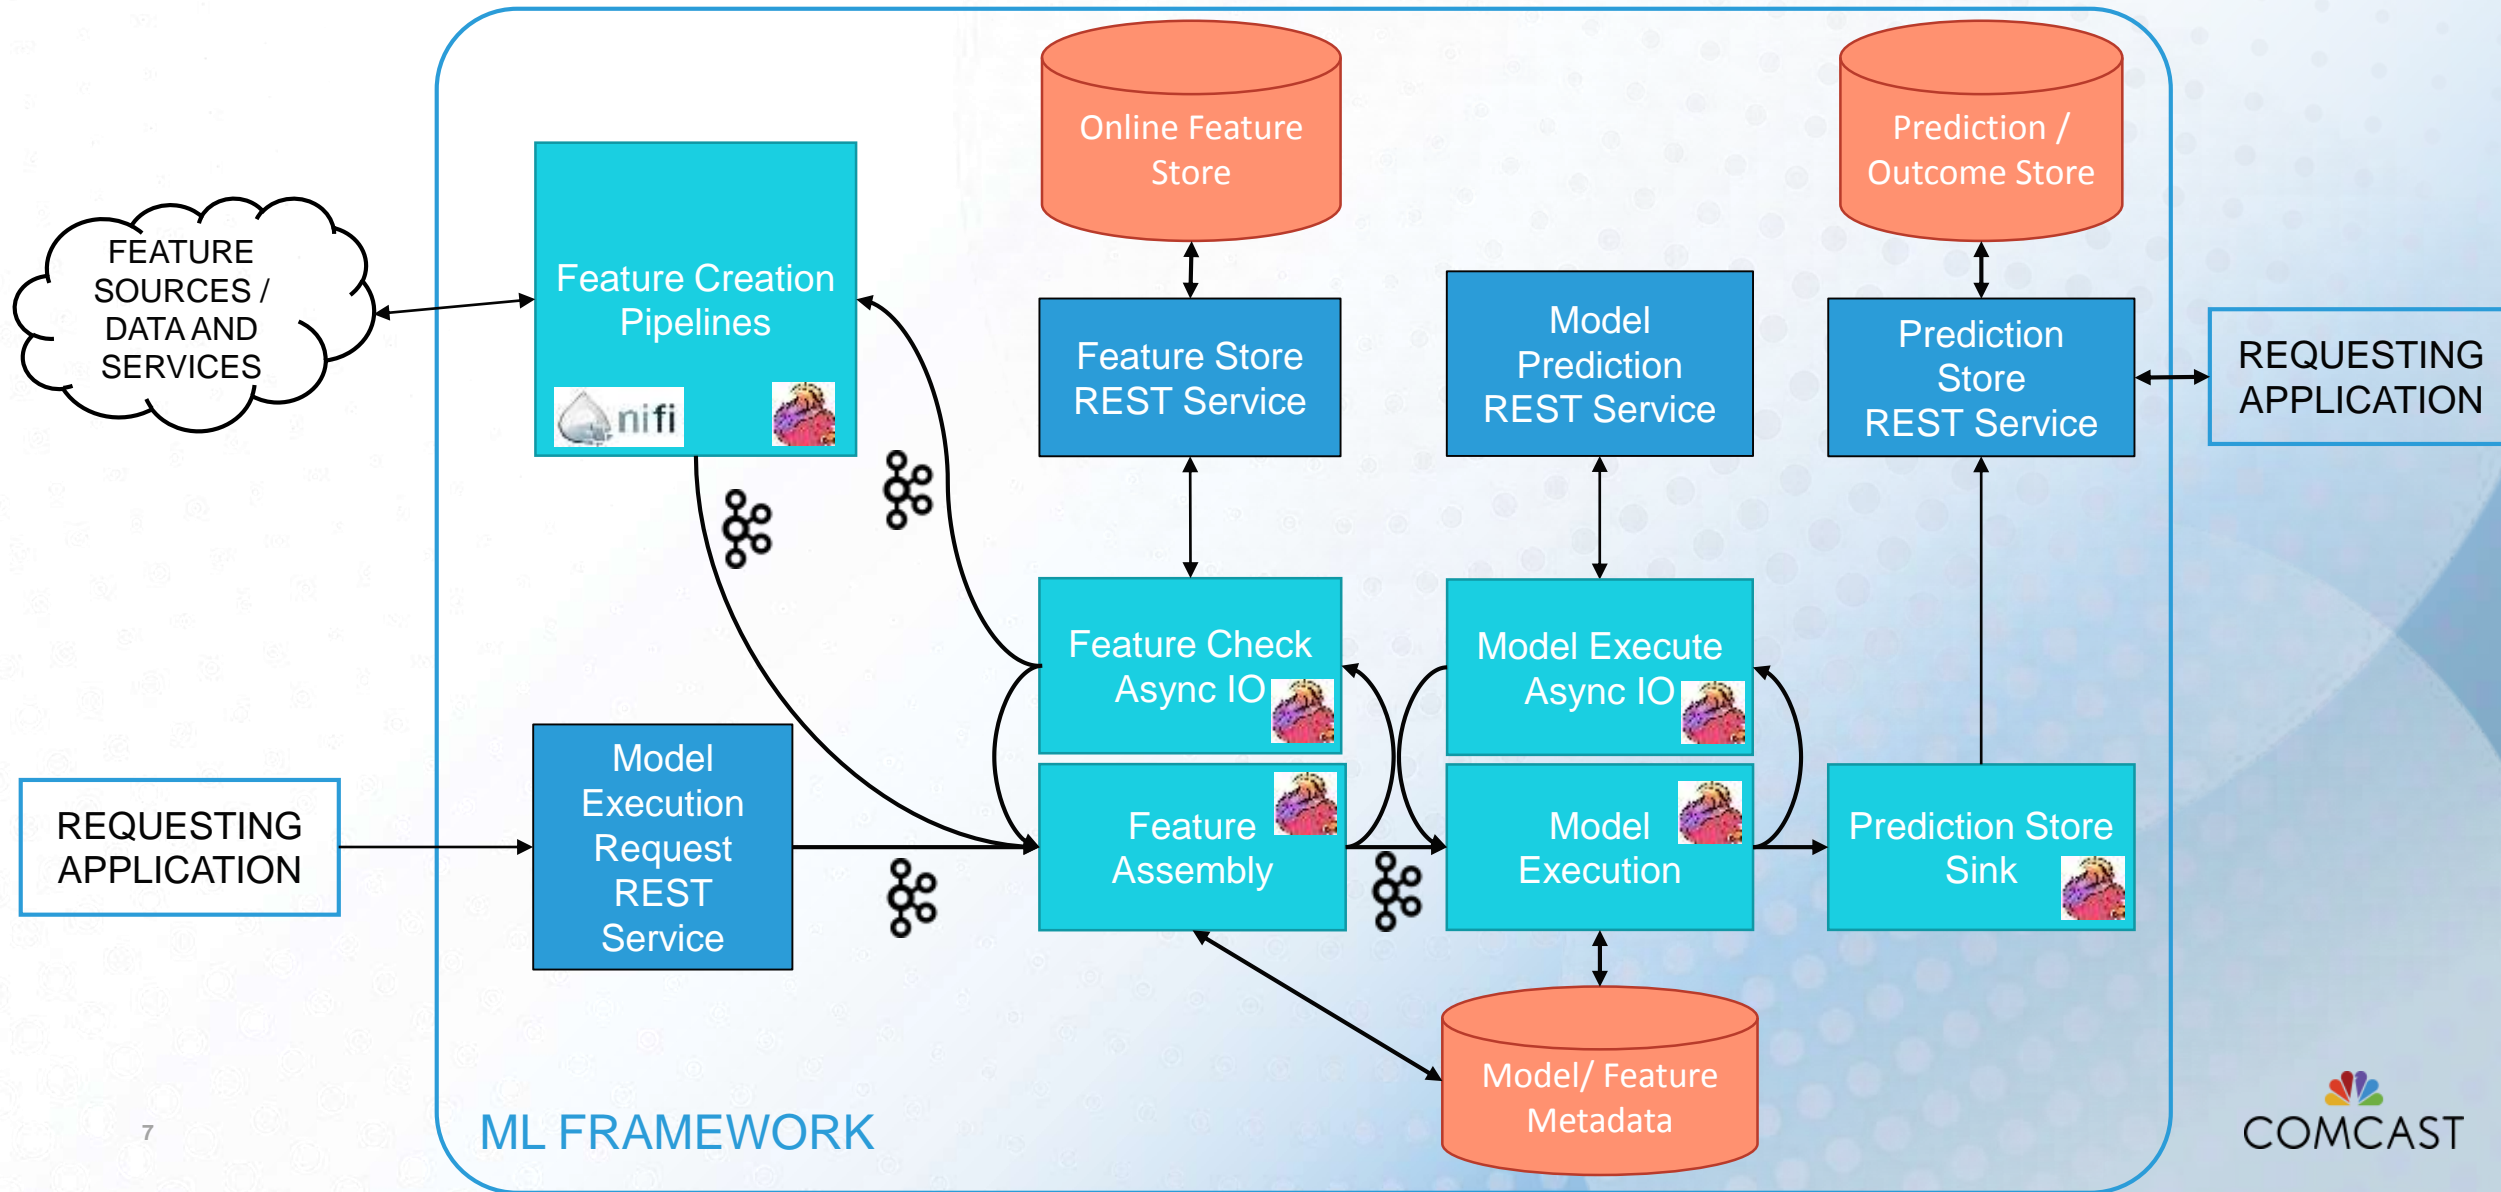
THE
CHECKPOINT
PROBLEM

THE
**TRIGGER AND
DIAGNOSIS**
PROBLEM

THE
**REALLY HIGH VOLUME
AND FEATURE STORE**
PROBLEM #2

ML FRAMEWORK ARCHITECTURE – 2018

"EMBEDDING FLINK THROUGHOUT AN OPERATIONALIZED STREAMING ML LIFECYCLE" – SF FLINK FORWARD 2018



THE TRIGGER AND DIAGNOSIS PROBLEM

MAY 2018 – INITIAL VOLUMES

INDICATOR #1

9 MILLION

INDICATOR #2

166 MILLION

INDICATOR #3

1.2 MILLION

INDICATOR #4

2500 (SMALL TRIAL)


TOTAL

175

MILLION / DAY

Comcast collects, stores, and uses all data in accordance with our privacy disclosures to users and applicable laws.


175 MILLION PREDICTIONS?




INTRODUCE A 'TRIGGER' CONDITION

	GREEN
	GREEN
	GREEN
	RED
	GREEN
	GREEN

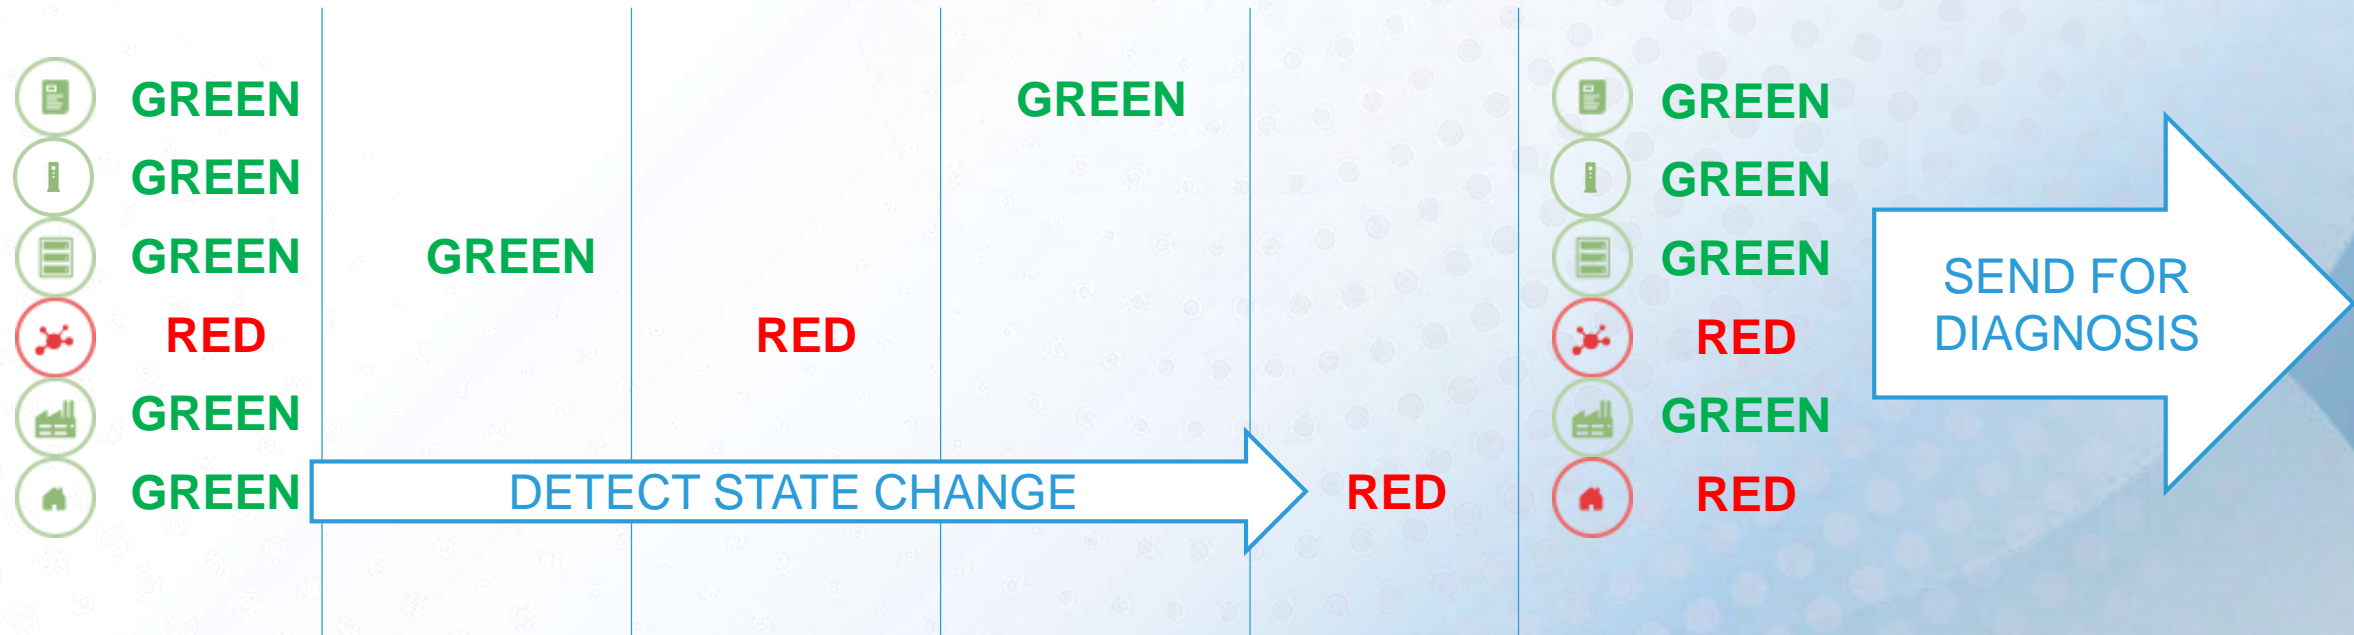
INTRODUCE A 'TRIGGER' CONDITION



INTRODUCE A 'TRIGGER' CONDITION

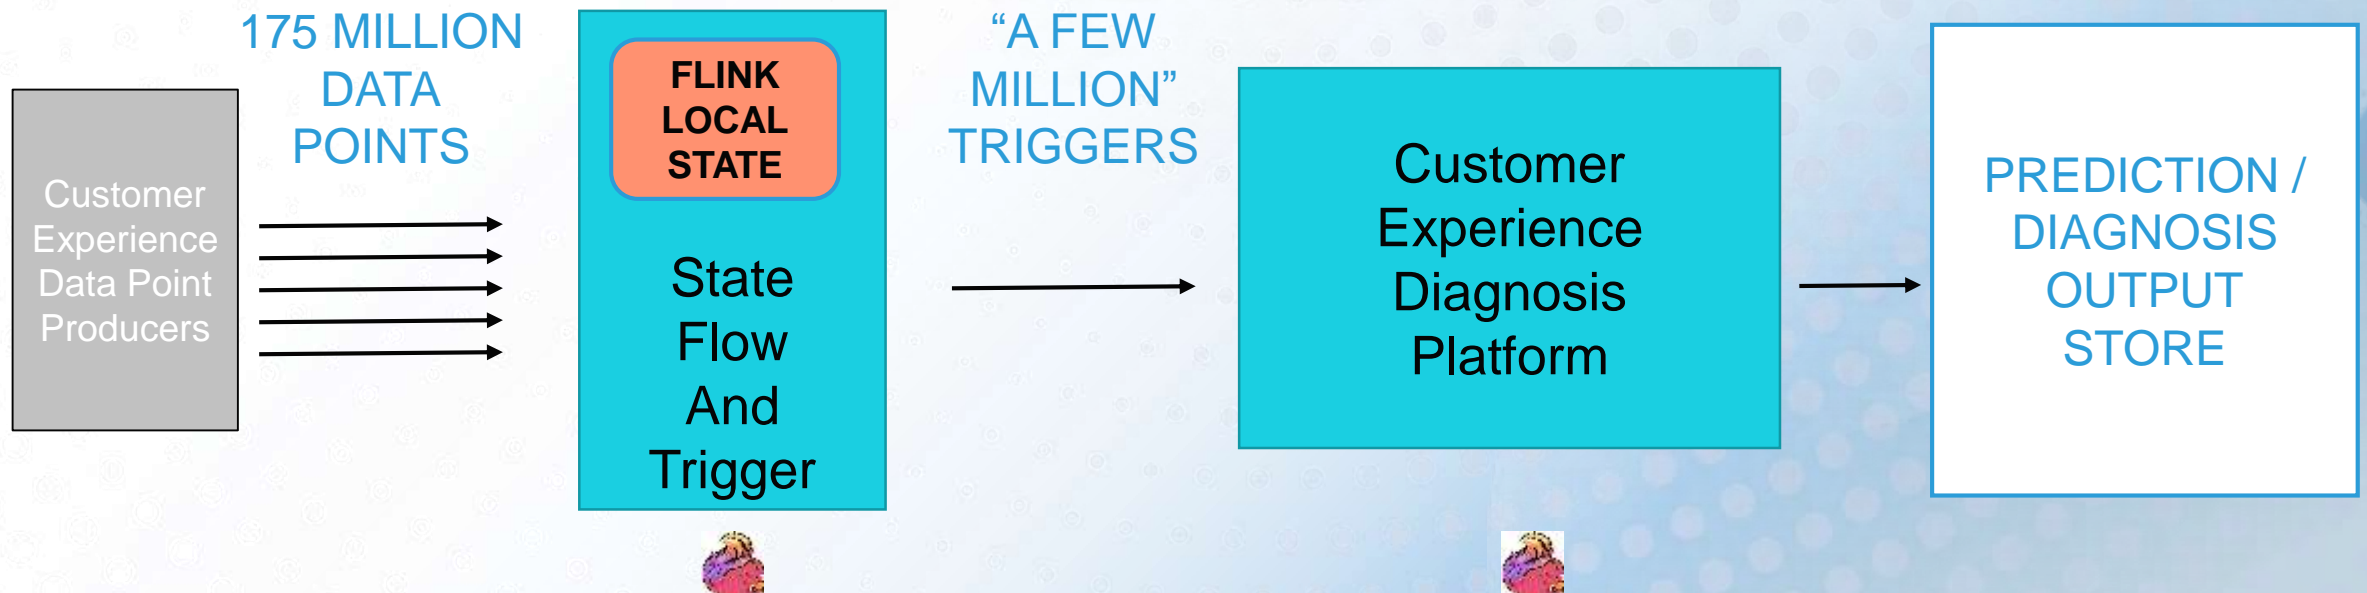


INTRODUCE A 'TRIGGER' CONDITION



STATE MANAGER AND TRIGGER

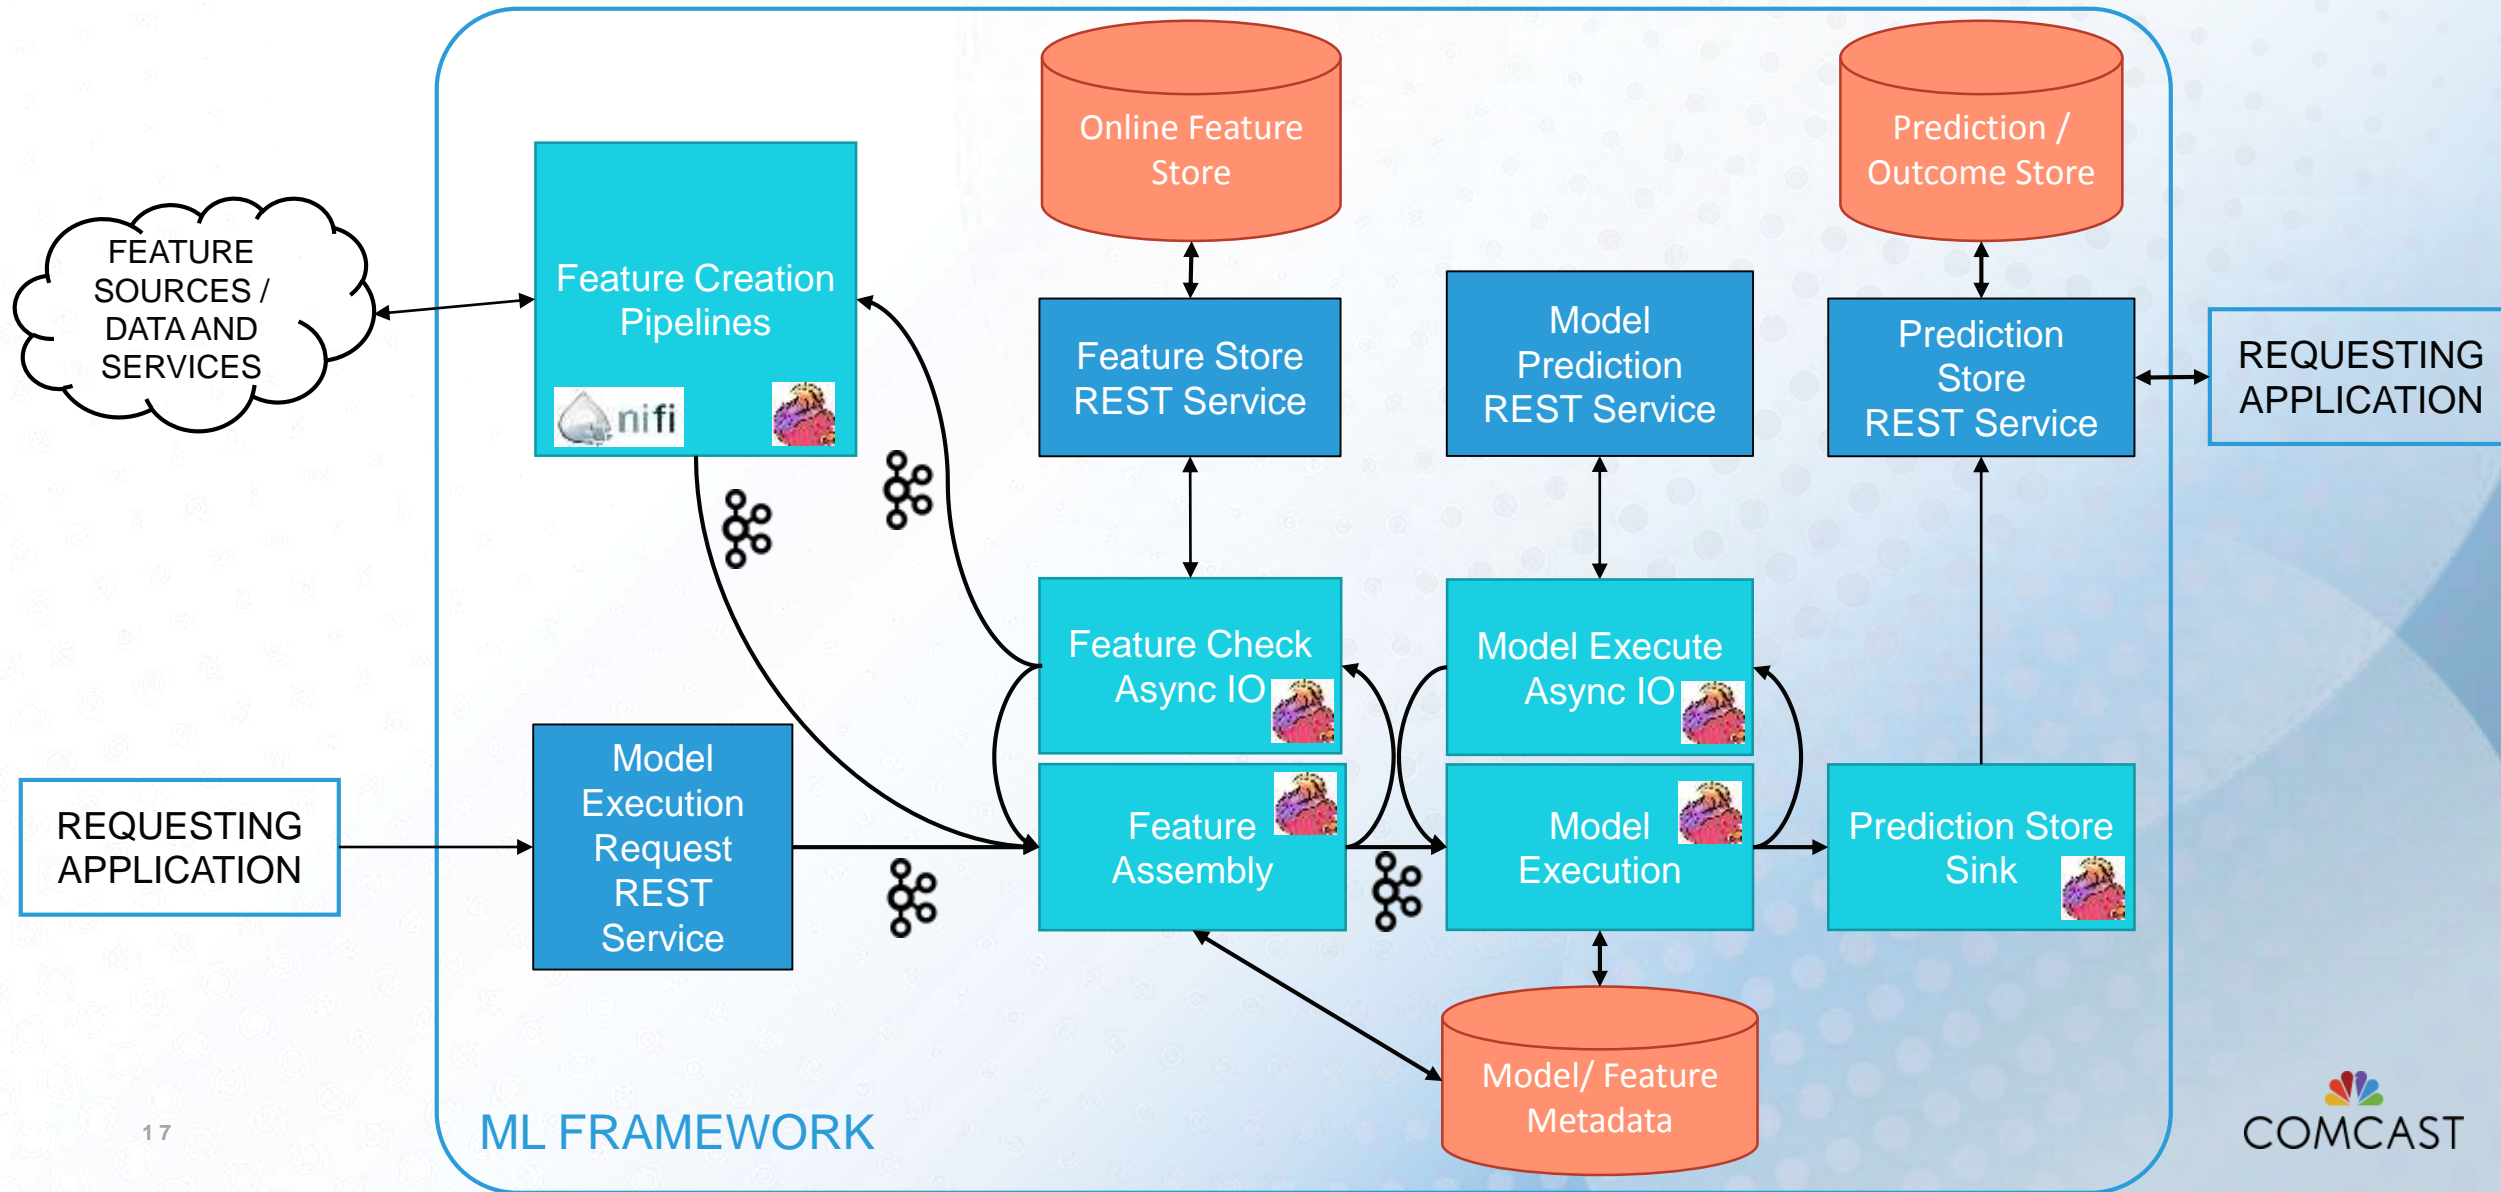
INTRODUCE FLINK LOCAL STATE



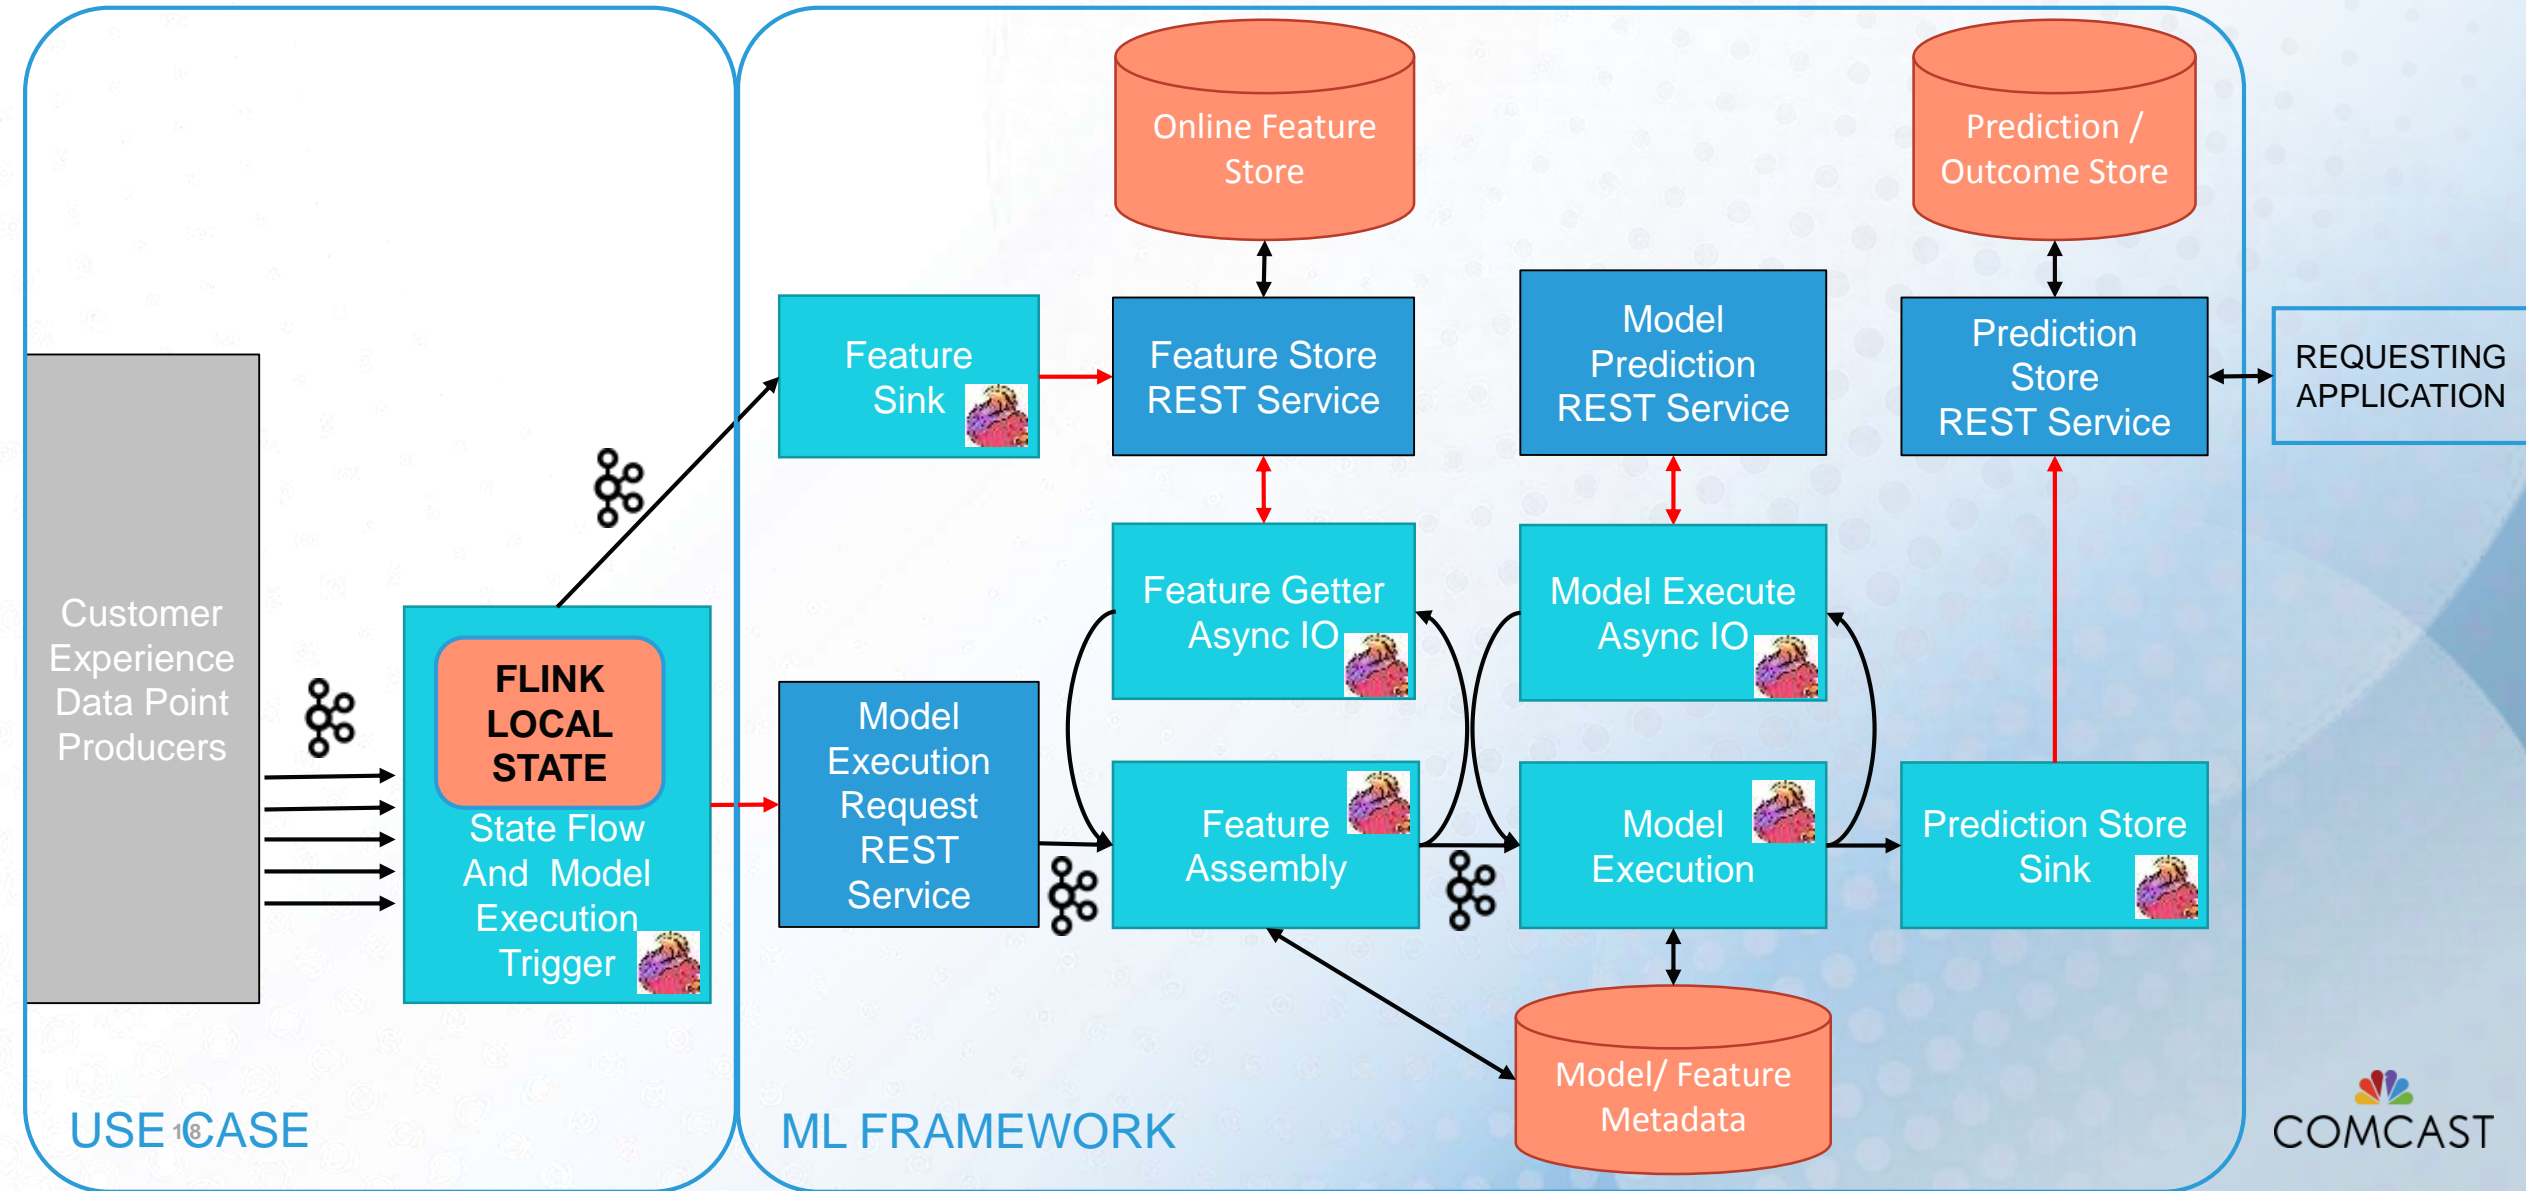
THE REST PROBLEM

ML FRAMEWORK ARCHITECTURE – 2018

"EMBEDDING FLINK THROUGHOUT AN OPERATIONALIZED STREAMING ML LIFECYCLE" – SF FLINK FORWARD 2018



INITIAL PLATFORM ARCHITECTURE



ORIGINAL ML FRAMEWORK ARCHITECTURE



**EVERY COMMUNICATION
WAS A REST SERVICE**



**ML FRAMEWORK
ISOLATED FROM DATA
POINT USE CASE
BY DESIGN**



**FLINK ASYNC + REST IS
DIFFICULT TO SCALE
ELASTICALLY**

TUNING ASYNC I/O – LATENCY AND VOLUME

MAX THROUGHPUT PER SECOND = (SLOTS) * (THREADS) * (1000 / TASK DURATION MSEC)
(12 SLOTS) * (20 THREADS) * (1000 / 300) = 800 / SECOND

APACHE HTTP CLIENT / EXECUTOR THREADS

`setMaxTotal`

`setDefaultMaxPerRoute (THREADS)`

FLINK ASYNC MAX OUTSTANDING

`RichAsyncFunction “Capacity”`

`AsyncDataStream.unorderedWait()` with max capacity parameter

RATE THROTTLING

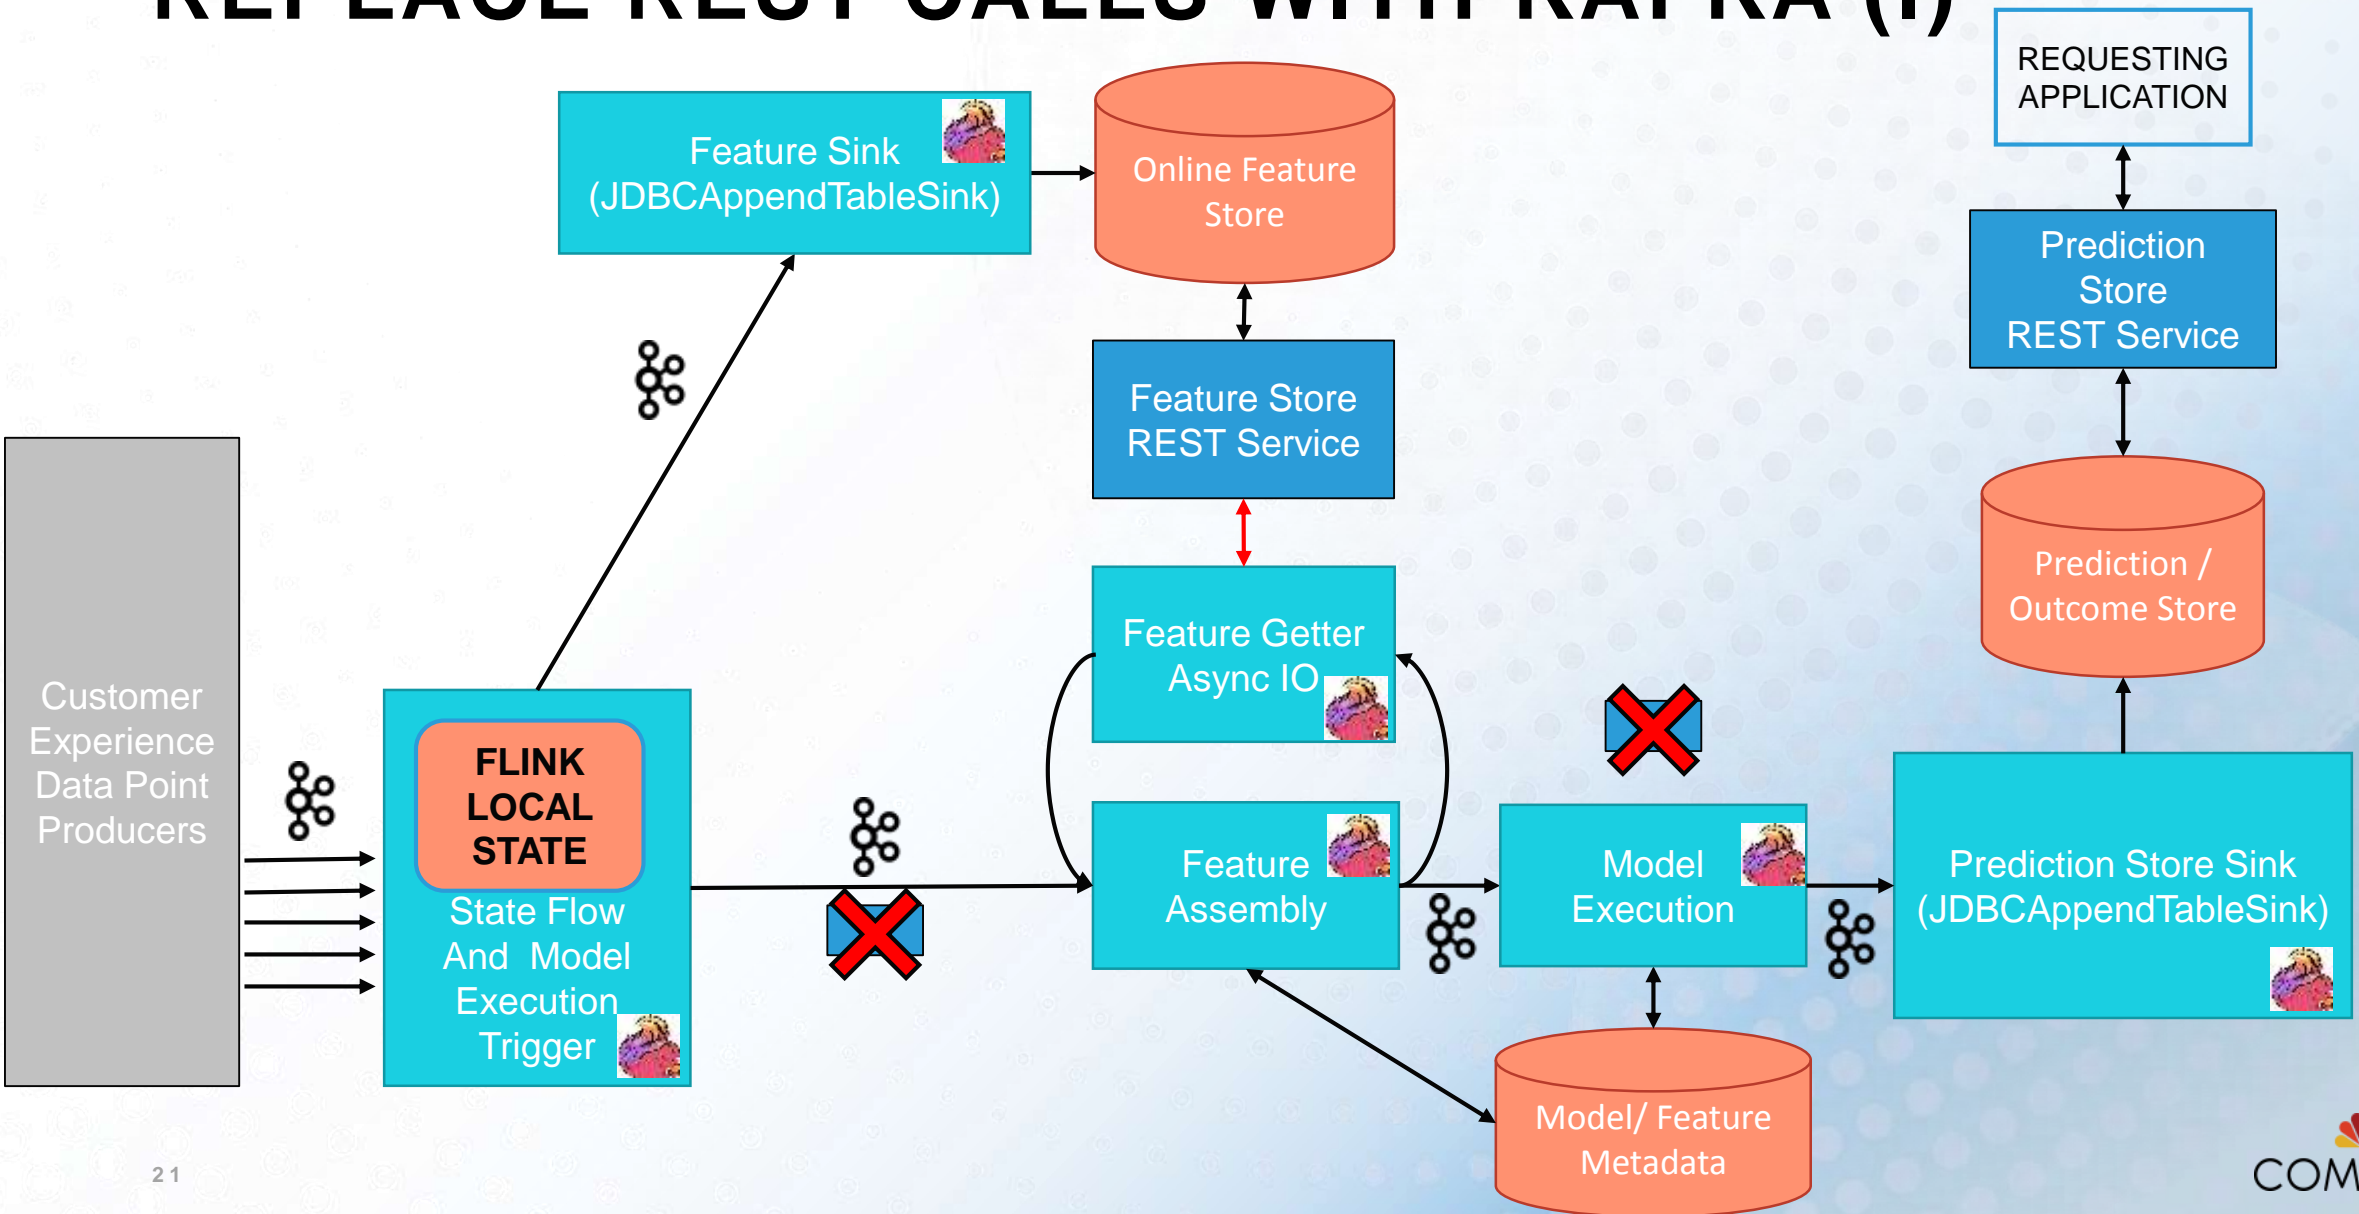
`if (numberOfRequestsPerPeriod > 0) {`

`int numberOfProcesses = getRuntimeContext()
.getNumberOfParallelSubtasks();`

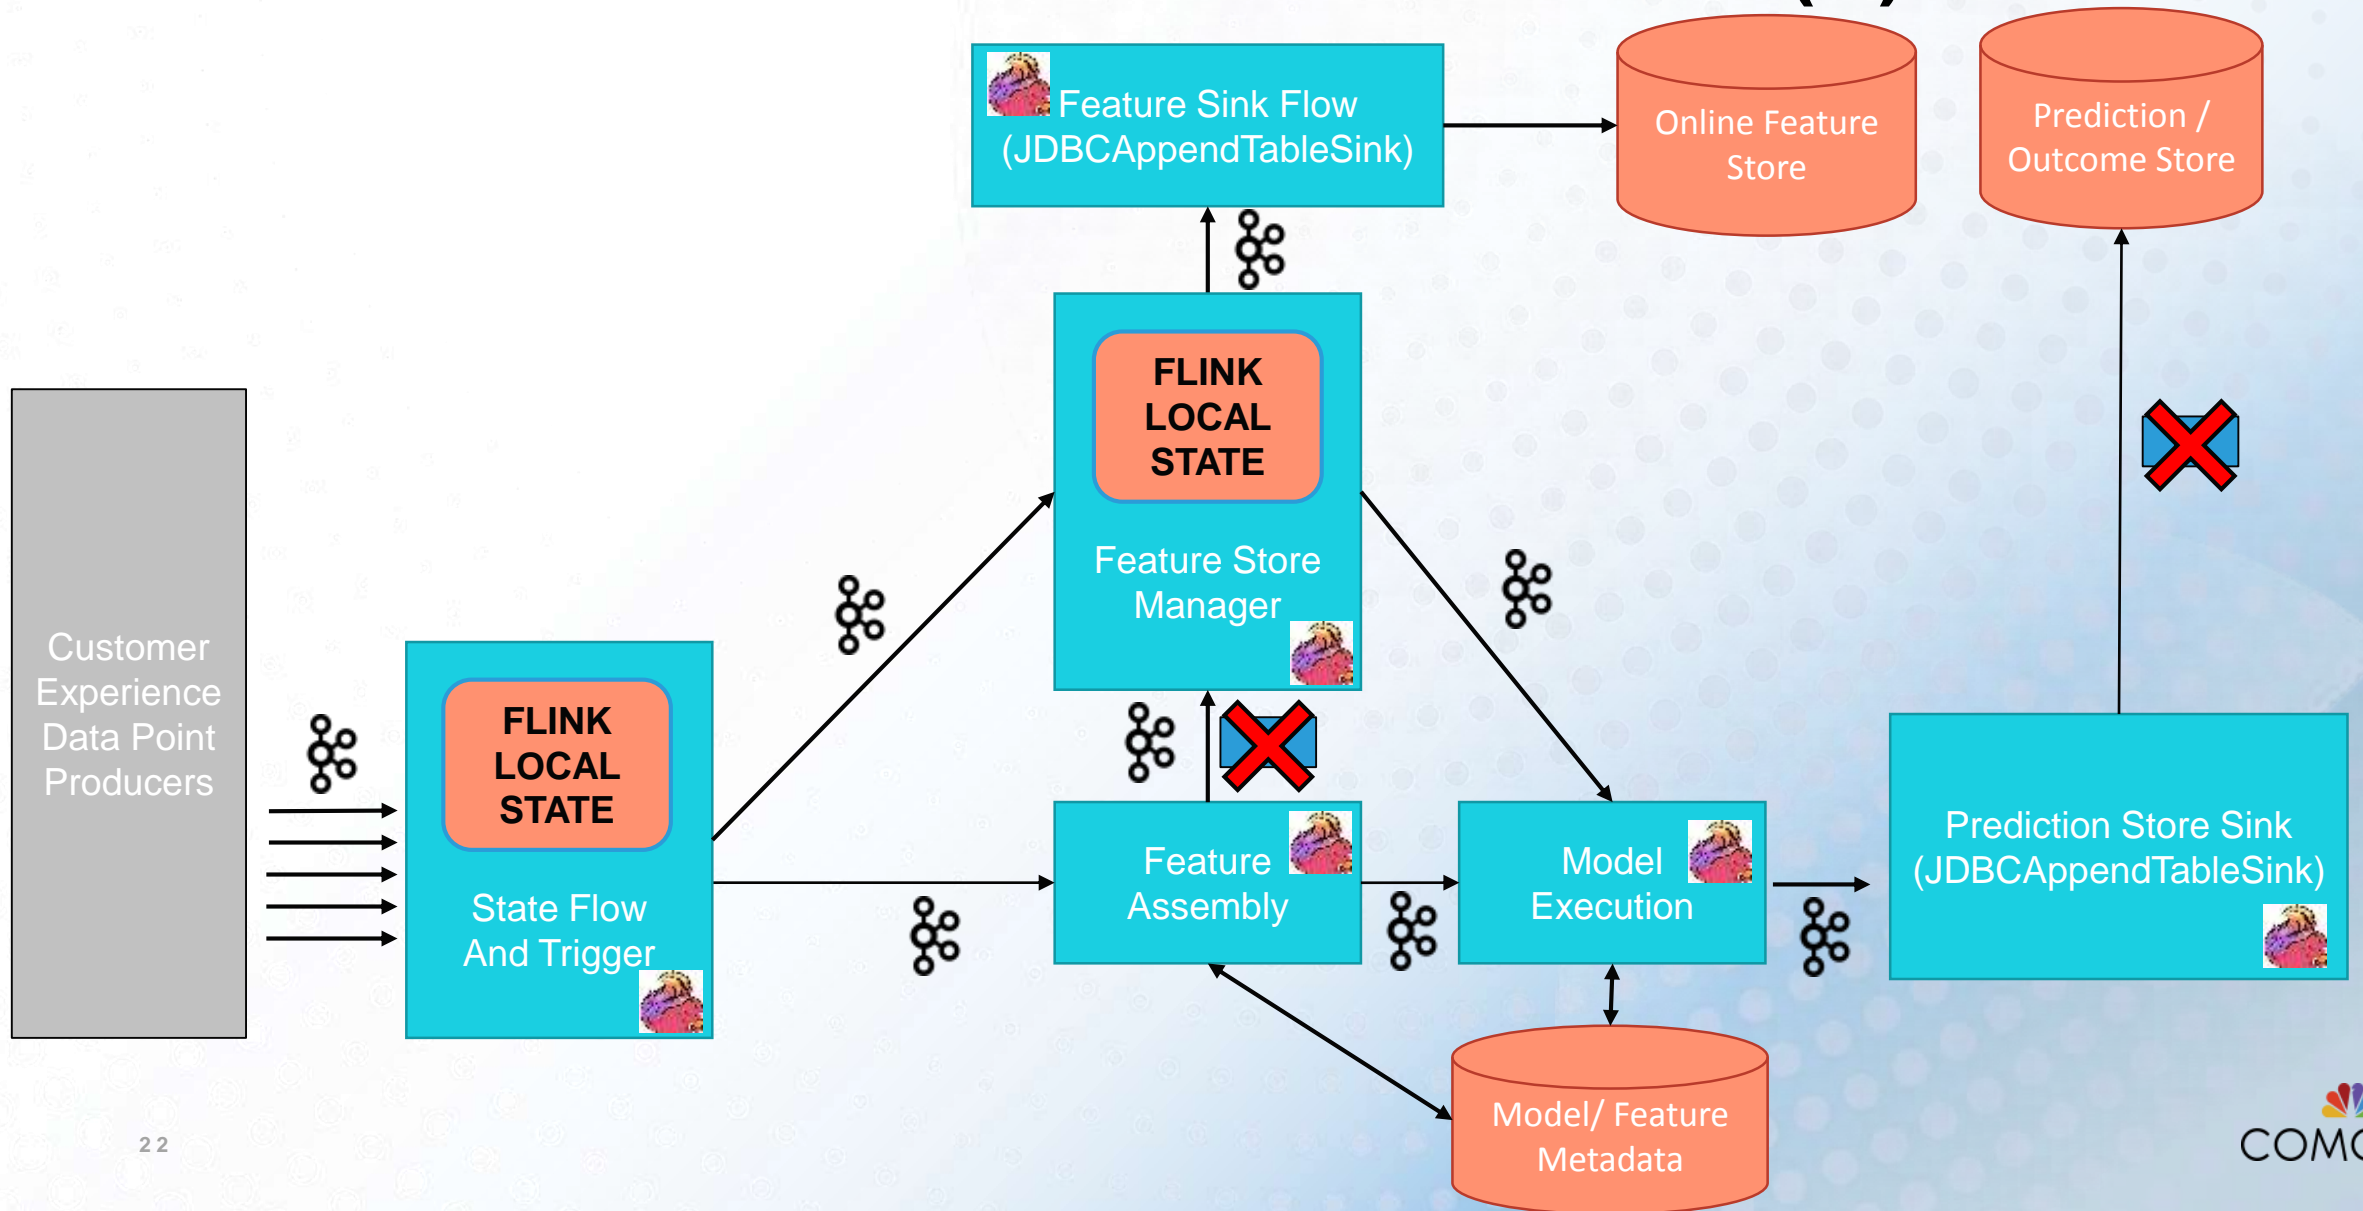
`numberOfRequestsPerPeriodPerProcess =
Math.max(1, numberOfRequestsPerPeriod /
numberOfProcesses);`

`}`

REPLACE REST CALLS WITH KAFKA (I)



REPLACE REST CALL WITH KAFKA (II)



THE INEFFICIENT OBJECT HANDLING PROBLEM

HAS THIS EVER HAPPENED TO YOU?

```
String transformedJson =  
new JoltJsonTransformer()  
    .transformJson(myJsonValue,  
        pTool.getRequired(SPEC_PATH));
```

```
JSONObject inputPayload = new  
JSONObject(transformedJson);  
try {  
    aType = inputPayload.getJSONObject("info")  
        .getJSONObject("aType").  
        getString("string");  
} catch (Exception e) {  
    // ignore, invalid object  
    aType = STATUS_INVALID;  
}
```

New “JOLT” JSON Transformer
(but don’t cache it)

→ New ObjectMapper()
(but don’t cache it)

→ Parse JSON string into Map

→ Load and parse “Transform” JSON
from the resource path

(but don’t cache it)

→ Serialize transformed Map *back*
into string result

HAS THIS EVER HAPPENED TO YOU?

```
String transformedJson =  
    new JoltJsonTransformer()  
        .transformJson(myJsonValue,  
            pTool.getRequired(SPEC_PATH));  
JSONObject inputPayload = new  
    JSONObject(transformedJson);  
try {  
    aType = inputPayload.getJSONObject("info")  
        .getJSONObject("aType").  
        getString("string");  
} catch (Exception e) {  
    // ignore, invalid object  
    aType = STATUS_INVALID;  
}
```

New “JOLT” JSON Transformer
(but don’t cache it)
→ New ObjectMapper()
(but don’t cache it)
→ Parse JSON string into Map
→ Load and parse “Transform” JSON
from the resource path
(but don’t cache it)
→ Serialize transformed Map *back*
into string result
Then let’s make a NEW JSONObject
from that string

HAS THIS EVER HAPPENED TO YOU?

```
String transformedJson =  
new JoltJsonTransformer()  
    .transformJson(myJsonValue,  
        pTool.getRequired(SPEC_PATH));  
JSONObject inputPayload = new  
JSONObject(transformedJson);  
try {  
    aType = inputPayload.getJSONObject("info")  
        .getJSONObject("aType").  
        getString("string");  
} catch (Exception e) {  
    // ignore, invalid object  
    aType = STATUS_INVALID;  
}
```

New “JOLT” JSON Transformer
(but don’t cache it)

→ New ObjectMapper()
(but don’t cache it)

→ Parse JSON string into Map

→ Load and parse “Transform” JSON
from the resource path

(but don’t cache it)

→ Serialize transformed Map *back*
into string result

Then let’s make a NEW JSONObject
from that string

Go deep into the path to get value

HAS THIS EVER HAPPENED TO YOU?

```
String transformedJson =  
new JoltJsonTransformer()  
    .transformJson(myJsonValue,  
        pTool.getRequired(SPEC_PATH));  
JSONObject inputPayload = new  
JSONObject(transformedJson);  
try {  
    aType = inputPayload.getJSONObject("info")  
        .getJSONObject("aType").  
        getString("string");  
} catch (Exception e) {  
    // ignore, invalid object  
    aType = STATUS_INVALID;  
}
```

New “JOLT” JSON Transformer
(but don’t cache it)
→ New ObjectMapper()
(but don’t cache it)
→ Parse JSON string into Map
→ Load and parse “Transform” JSON
from the resource path
(but don’t cache it)
→ Serialize transformed Map *back*
into string result
Then let’s make a NEW JSONObject
from that string
Go deep into the path to get value
Cast exception? STATUS_INVALID

SLIGHTLY MORE EFFICIENT SOLUTION

```
if (mapper == null)
    mapper = new ObjectMapper();
```

```
Map<?, ?> resultMap =
    JsonUtils.readValue(mapper, (String)
    result, Map.class);
```

```
MapTreeWalker mapwalker = new
    MapTreeWalker(resultMap);
```

```
final String aType =
    mapwalker.step("info")
    .step("aSchemaInfo")
    .step("aType")
    .step("string")
    .<String>get().
    orElse(STATUS_INVALID);
```

Cache our new ObjectMapper()

Parse JSON string into Map *ONCE*

Lightweight Map Traversal utility

Java Stream syntax

Optional.ofNullable

Optional.orElse

FLINK OBJECT EFFICIENCIES

```
StreamExecutionEnvironment env;  
env.getConfig().enableObjectReuse();
```

REDUCE OBJECT
CREATION AND
GARBAGE
COLLECTION

SEMANTIC ANNOTATIONS

@NonForwardedFields

@ForwardedFields

BE CAREFUL OF
SIDE EFFECTS
WITHIN
OPERATORS

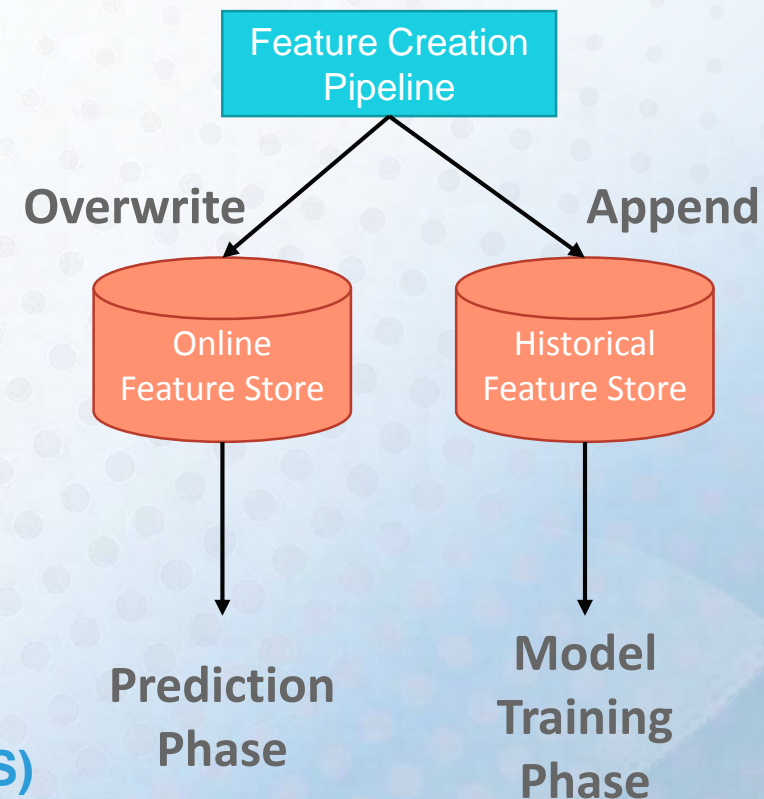
THE FEATURE STORE PROBLEM

WHAT IS A FEATURE STORE?

“FEATURE” IS A DATA POINT INPUT TO ML MODEL

WE NEED TO:

- **STORE ALL THE INPUT DATA POINTS**
- **SNAPSHOT AT ANY MOMENT IN TIME**
- **ASSOCIATE WITH A DIAGNOSIS (MODEL OUTPUT)**
- **HAVE ACCESS TO THE DATA POINTS (API FOR OTHER APPS)**
- **STORE ALL DATA POINTS FOR ML MODEL TRAINING**



FIRST TRY: AWS AURORA RDBMS

“HEY IT SHOULD GET US UP TO 10,000 / SECOND”

...THE UPSERT PROBLEM

DIDN'T WORK MUCH PAST 3,000 / SECOND

...SOON OUR INPUT RATE WAS 20,000 / SECOND



MITIGATION: STORE ONLY AT DIAGNOSIS TIME

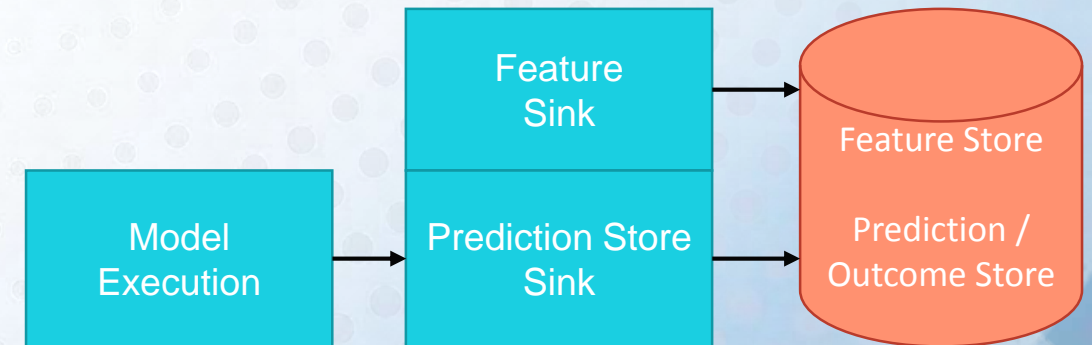
STOPPED STORING ALL RAW DATA POINTS

ONLY STORE DATA POINTS ALONGSIDE
DIAGNOSIS

CON: ONLY A SMALL % OF DATA POINTS

CON: DATA NOT CURRENT, ONLY AS OF
RECENT DIAGNOSIS

CON: LIMITED USEFULNESS



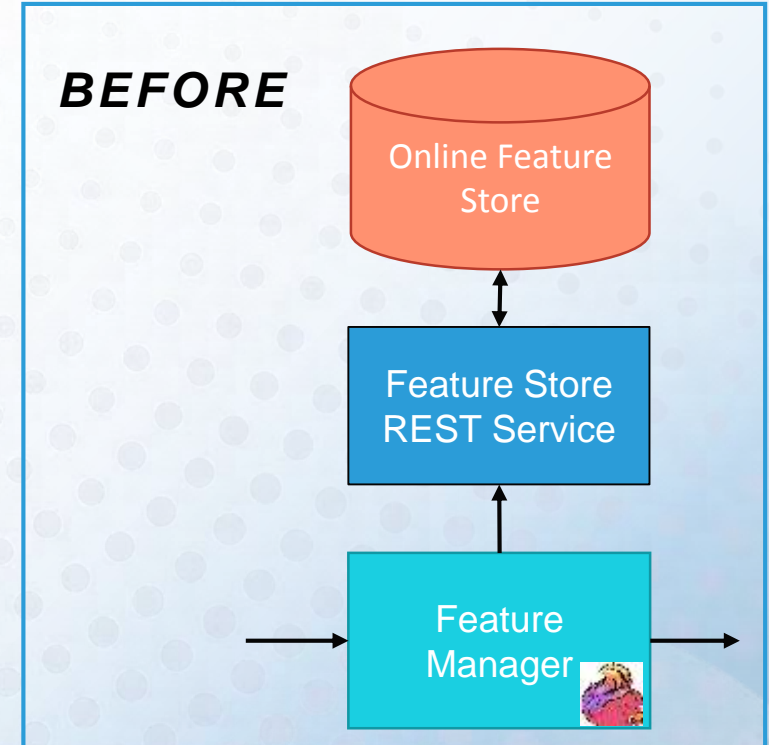
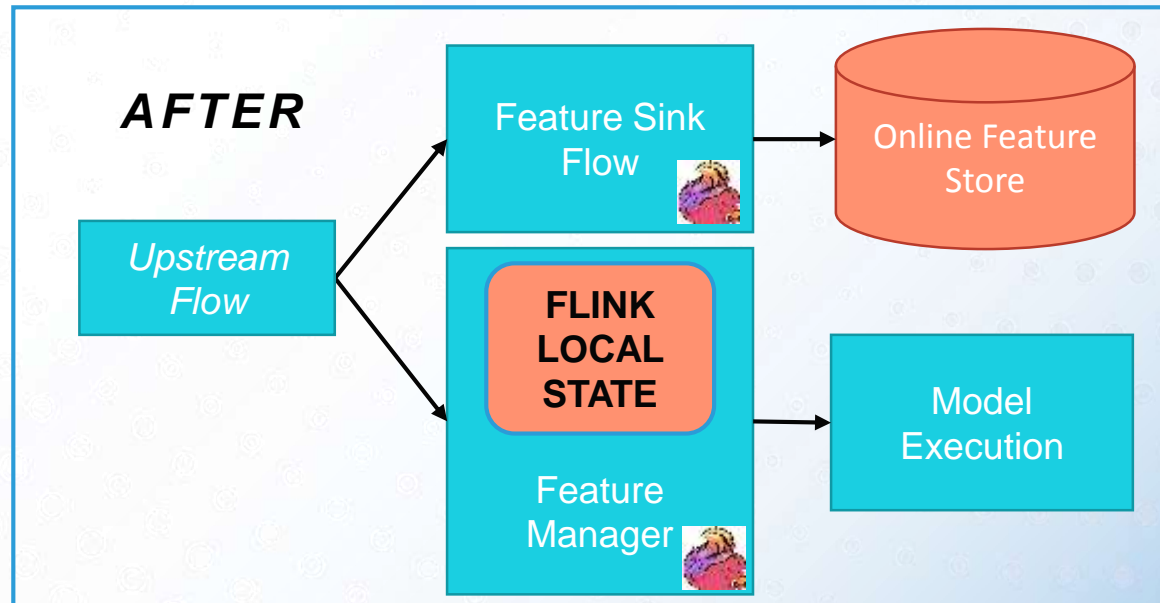
OPTIMIZING THE WRITE PATH

WHICH FLINK FLOW WRITES TO FEATURE STORE?

FEATURE STORE NOT NEEDED FOR TRIGGER

FEATURE STORE NOT NEEDED FOR MODEL EXECUTION

SEPARATE 'TRANSFORM AND SINK' FLOW



ALTERNATIVE DATA STORES

REDIS

STORE ONLY LATEST DATA POINTS?

CASSANDRA

COST AND COMPLEXITY?

FLINK QUERYABLE STATE

PRODUCTION HARDENING?

READ VOLUME?

TIME SERIES DB

DRUID

INFLUX

REDIS

WE ONLY NEED < 7 DAYS OF DATA, ONLY LATEST VALUE, SO USE REDIS AS THE KV STORE

ACCESSIBLE BY OTHER CONSUMERS OR VIA A SERVICE

COULDN'T PUT MORE THAN 8,000 / SECOND FOR A 'REASONABLE' CLUSTER SIZE

FLINK SINK CONNECTOR IS NOT CLUSTER-ENABLED

ONE OBJECT PUT PER CALL

JEDIS VS LETTUCE FRAMEWORK AND PIPELINING

SOME OPTIMIZATION IF ABLE TO BATCH REDIS SHARDS BY PRE-COMPUTING

FLINK QUERYABLE STATE

REUSE OUR FEATURE REQUEST/RESPONSE
FRAMEWORK

SCALABLE, FAST

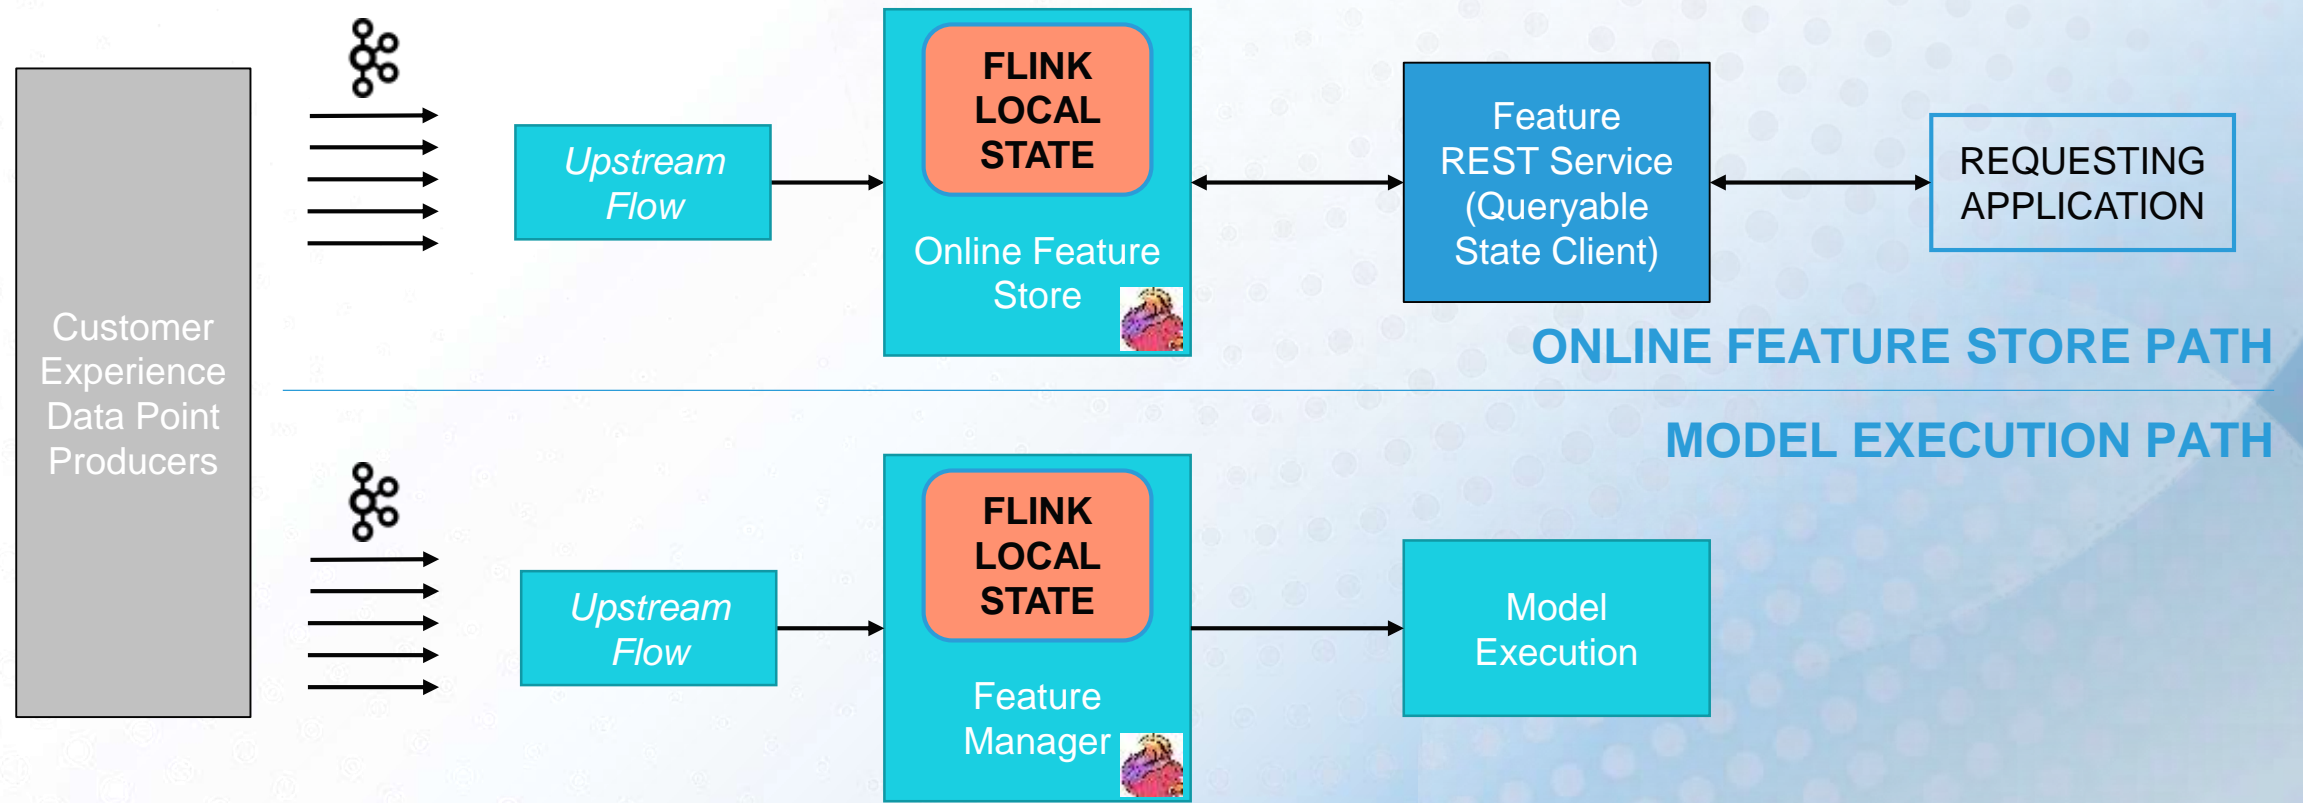
OUR READ LOAD IS LOW (10 / SECOND)

EXPENSIVE (RELATIVELY)

NOT A “PRODUCTION” CAPABILITY

NOT A “DATABASE” TECHNOLOGY


QUERYABLE STATE - ONLINE FEATURE STORE




The background features a complex, abstract design. On the left, there's a light beige area with a fine grid pattern. This transitions into a larger section on the right with a blue-to-white gradient. Overlaid on this gradient are several semi-transparent geometric shapes: a large circle, a square, and a triangle. These shapes contain different patterns, including a grid of small dots and a series of parallel lines. In the bottom right corner, a faint, stylized city skyline is visible, with several tall buildings. The overall aesthetic is modern and architectural.

THE VOLUME PROBLEM

RAMPED UP PRODUCTION – JULY 2018




ADDED NEW DATA POINT TYPES




INCREASED DATA POINT FREQUENCY

ONCE EVERY 24 HOURS




INCREASED DATA POINT FREQUENCY

ONCE EVERY 24 HOURS




ONCE EVERY 4 HOURS




INCREASED DATA POINT FREQUENCY


ONCE EVERY 24 HOURS



ONCE EVERY 4 HOURS




ONCE EVERY 5 MINUTES ???



SYMPTOM: KAFKA LATENCY INCREASING

FLINK CLUSTER COULDN'T KEEP UP WITH KAFKA VOLUME



KAFKA OPTIMIZATION

INCREASE PARTITIONS

1 KAFKA PARTITION
READ BY
1 FLINK TASK THREAD

MATCH PARALLELISM TO PARTITIONS (OR EVEN MULTIPLE)

100 KAFKA PARTITIONS
50 SLOT PARALLELISM

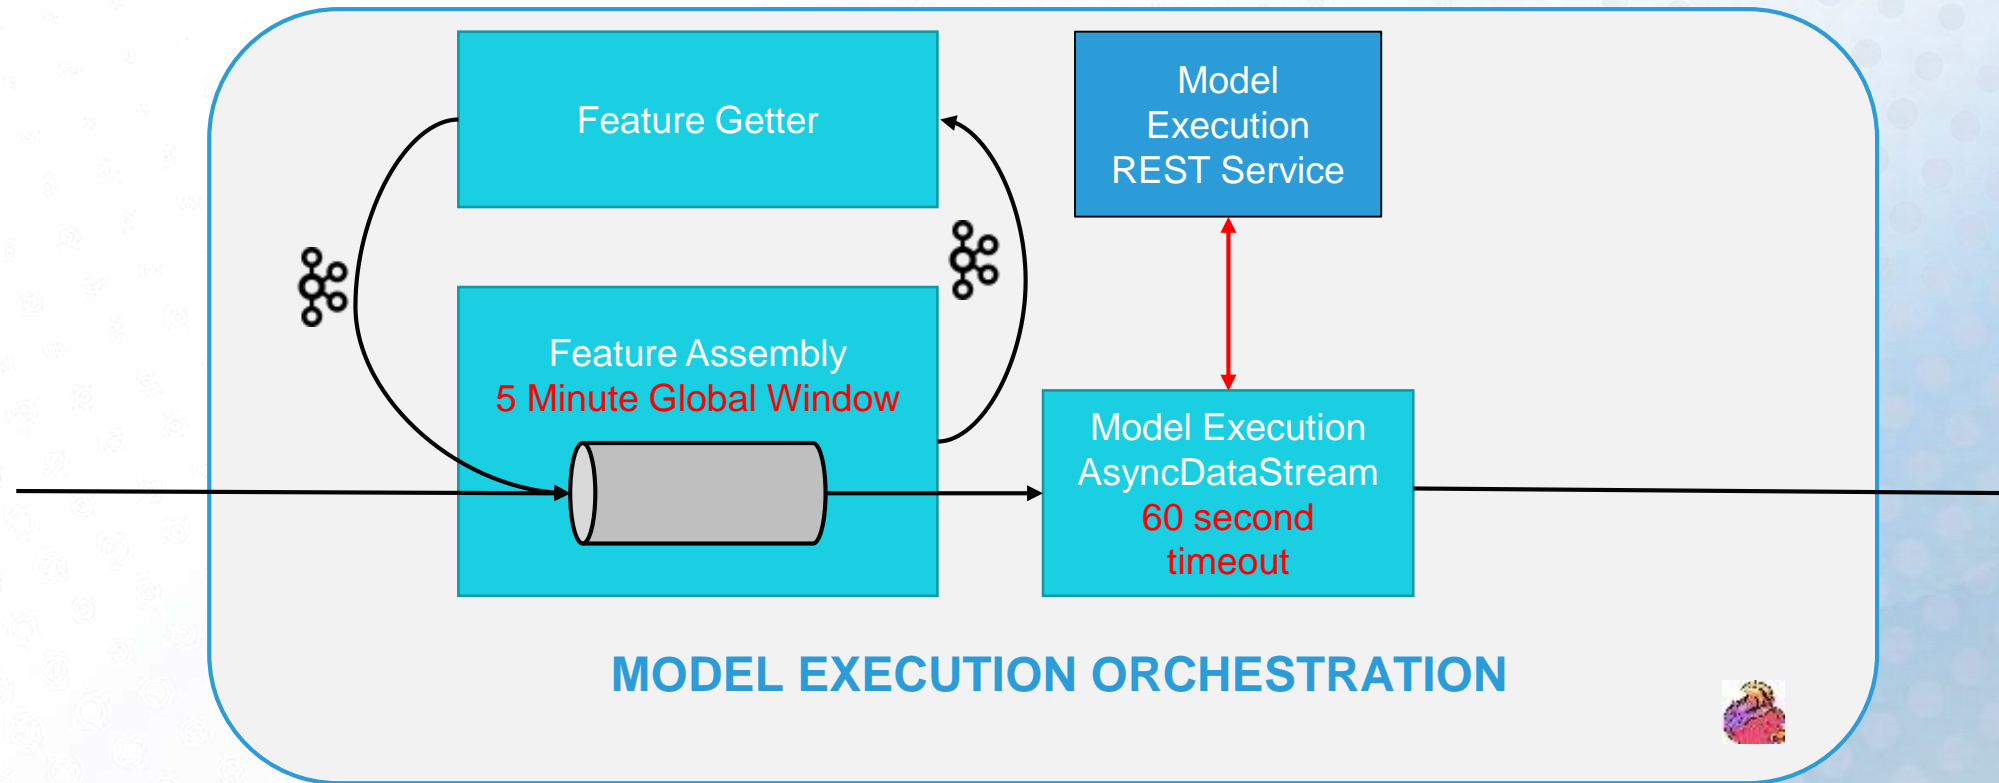
50 KAFKA PARTITIONS
50 SLOT PARALLELISM

ADJUST KAFKA CONFIGURATION

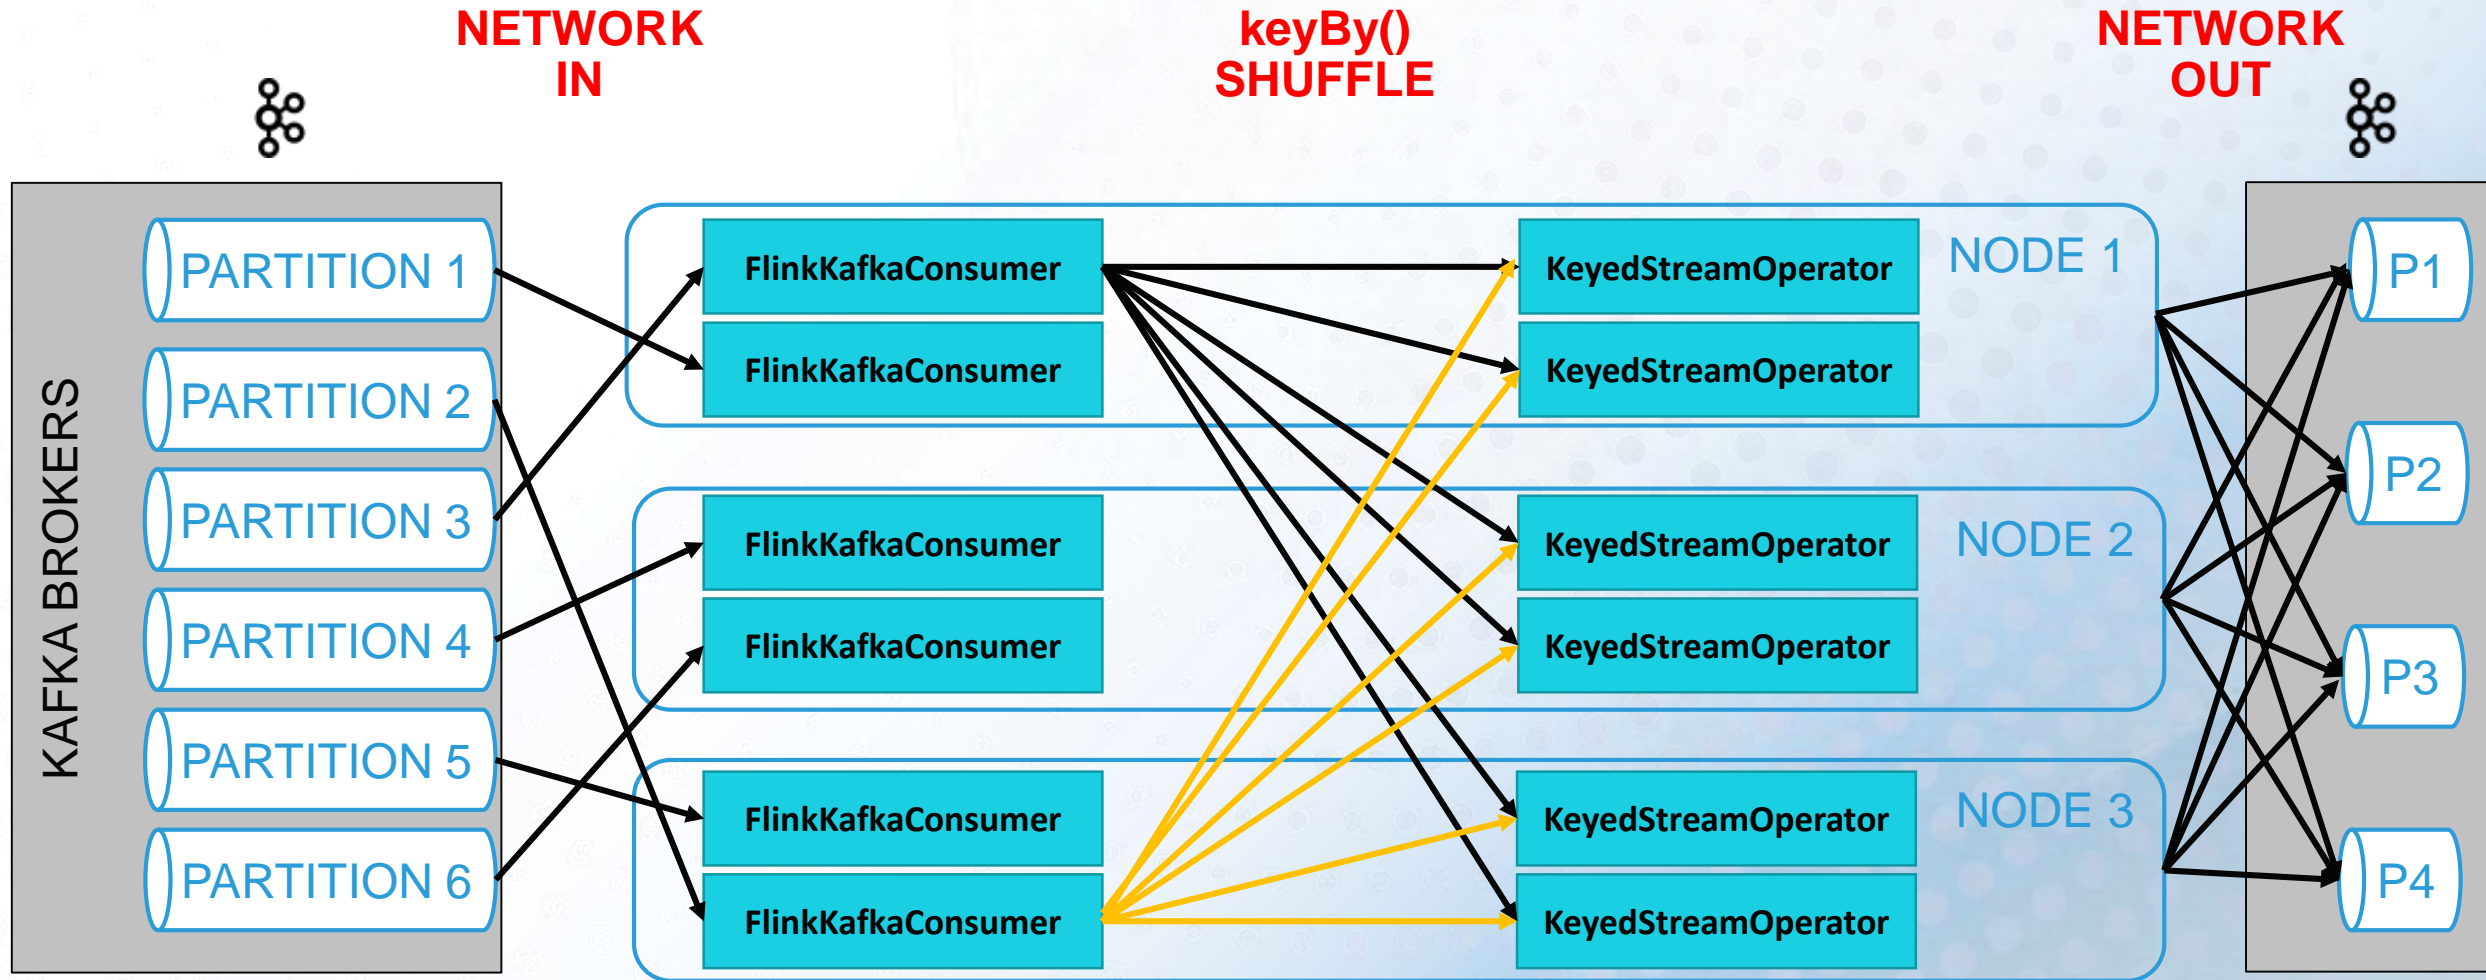
KAFKA CONSUMER SETTINGS
MAX.POLL.RECORDS = ~~500~~10000
FETCH.MIN.BYTES = ~~1~~102400
FETCH.MAX.WAIT.MS = ~~500~~1000
KAFKA PRODUCER SETTINGS
BUFFER.MEMORY = 268435456
BATCH.SIZE = ~~131072~~524288
LINGER.MS = ~~0~~50
REQUEST.TIMEOUT.MS = ~~90000~~
600000

MODEL EXECUTION BACKUP PROBLEM

GENERAL SLOWNESS BACKS UP MODEL EXECUTION PIPELINE



CLUSTER NETWORK I/O TRAFFIC




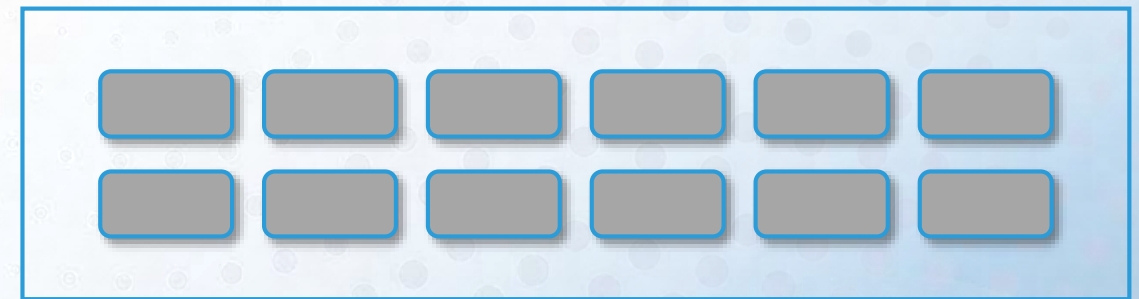
CLUSTER NETWORK I/O VOLUME

$(10,000 \text{ MESSAGES / SECOND}) * (1\text{KB / MESSAGE}) =$
10 MB / SECOND

100 MBITS/SEC NETWORK

SOLUTION – LARGER NODE TYPES WITH MORE
VCPU AND HIGHER PARALLELISM... FEWER
NODES IN CLUSTER

SWEET SPOT SEEMED TO BE ~20 NODES



GOOD READ: [HTTPS://WWW.VERVERICA.COM/BLOG/HOW-TO-SIZE-YOUR-APACHE-FLINK-CLUSTER-GENERAL-GUIDELINES](https://www.ververica.com/blog/how-to-size-your-apache-flink-cluster-general-guidelines)



THE CUSTOMER STATE PROBLEM

KEEPING TRACK OF DATA POINT STATE

25+ MILLION HIGH SPEED INTERNET CUSTOMERS

10+ DATA POINT TYPES

TWO FLINK STATES TO MANAGE:

RED / GREEN “TRIGGER”

MAPSTATE () / KEYBY (ENTITY ID)

[ENTITY ID] (DATA POINT TYPE, STATE)

25 MILLION CUSTOMERS

10+ DATA POINT TYPES

AVG 33 BYTES PER ENTITY ID

~1GB TOTAL

QUERYABLE STATE FEATURE STORE

UP TO 15KB (JSON) PER DATA POINT

10KB * 25 MILLION = 250 GB

1KB * 25 MILLION = 25 GB

APPROXIMATE TOTAL = 1 TB

THE CHECKPOINT PROBLEM

CHECKPOINT TUNING

FEATURE MANAGER – UP TO 1TB OF STATE

2.5 MINUTES WRITING A CHECKPOINT EVERY 10 MINUTES

TURNED OFF FOR AWHILE!

(RATIONALE: WOULD ONLY TAKE A FEW HOURS TO RECOVER ALL STATE)

INCREMENTAL CHECKPOINTING HELPED

LONG ALIGNMENT TIMES ON HIGH PARALLELISM

REDUCED THE SIZE OF STATE FOR 4 AND 24 HOUR WINDOWS USING FEATURE LOOKUP TABLE

The background features a faded, high-angle view of a city skyline with several skyscrapers. A large, semi-transparent blue circle is positioned on the right side of the image, partially overlapping the cityscape. The text 'THE HA PROBLEM' is centered on the left side of the image.

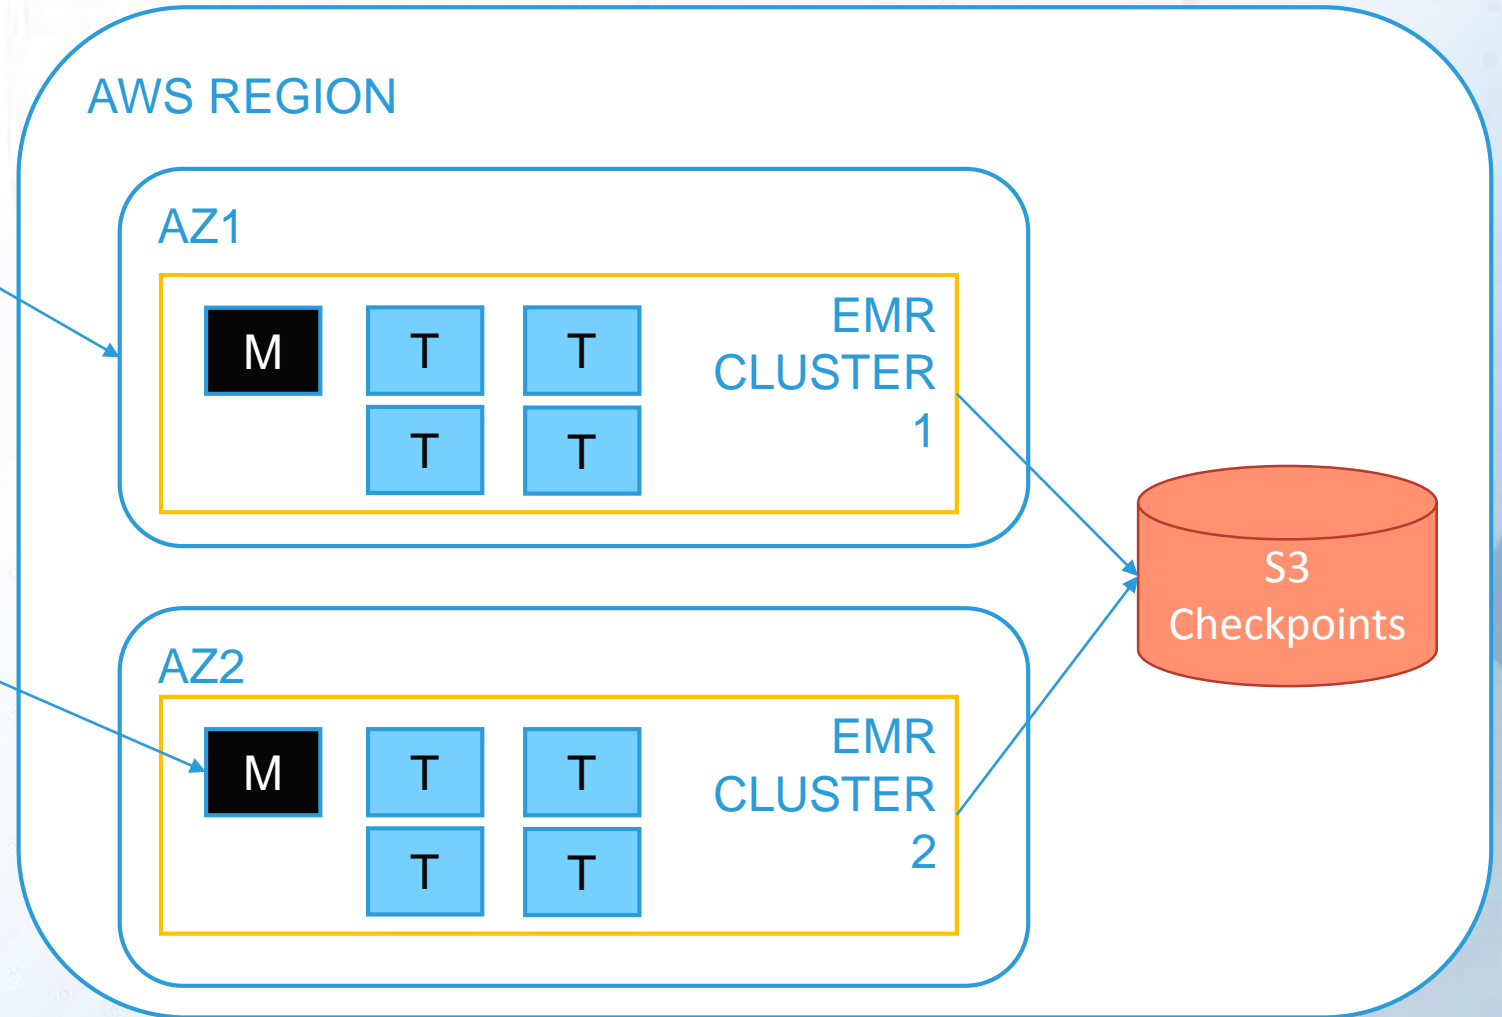
THE HA PROBLEM

AMAZON EMR AND HIGH AVAILABILITY

EMR FLINK CLUSTER

SINGLE AVAILABILITY ZONE (AZ)

SINGLE MASTER NODE
NOT HA



AMAZON AWS EMR

SESSIONS KEPT DYING

**TMP DIR IS PERIODICALLY CLEARED
BREAKS 'BLOB STORE' ON LONG RUNNING JOBS**

SOLUTION - CONFIGURE TO NOT BE IN /TMP

jobmanager.web.tmpdir

blob.storage.directory

EMR FLINK VERSION LAGS

APACHE RELEASE BY 1-2 MONTHS

TECHNOLOGY OPTIONS

FLINK ON KUBERNETES

AMAZON EKS (ELASTIC KUBERNETES SERVICE)

SELF-MANAGED KUBERNETES ON EC2

~~DATA ARTISANS~~ VERVERICA PLATFORM


Bonus:

Reduced Cost / Overhead by
having fewer Master Nodes

Less overhead for smaller (low
parallelism) flows

THE REALLY HIGH VOLUME AND FEATURE STORE PROBLEM #2

DATA POINT VOLUME MAY–NOVEMBER 2018



DATA POINT VOLUME MAY–NOVEMBER 2018




DATA POINT VOLUME MAY–NOVEMBER 2018



LET'S ADD ANOTHER 3 BILLION!

15 MILLION DEVICES

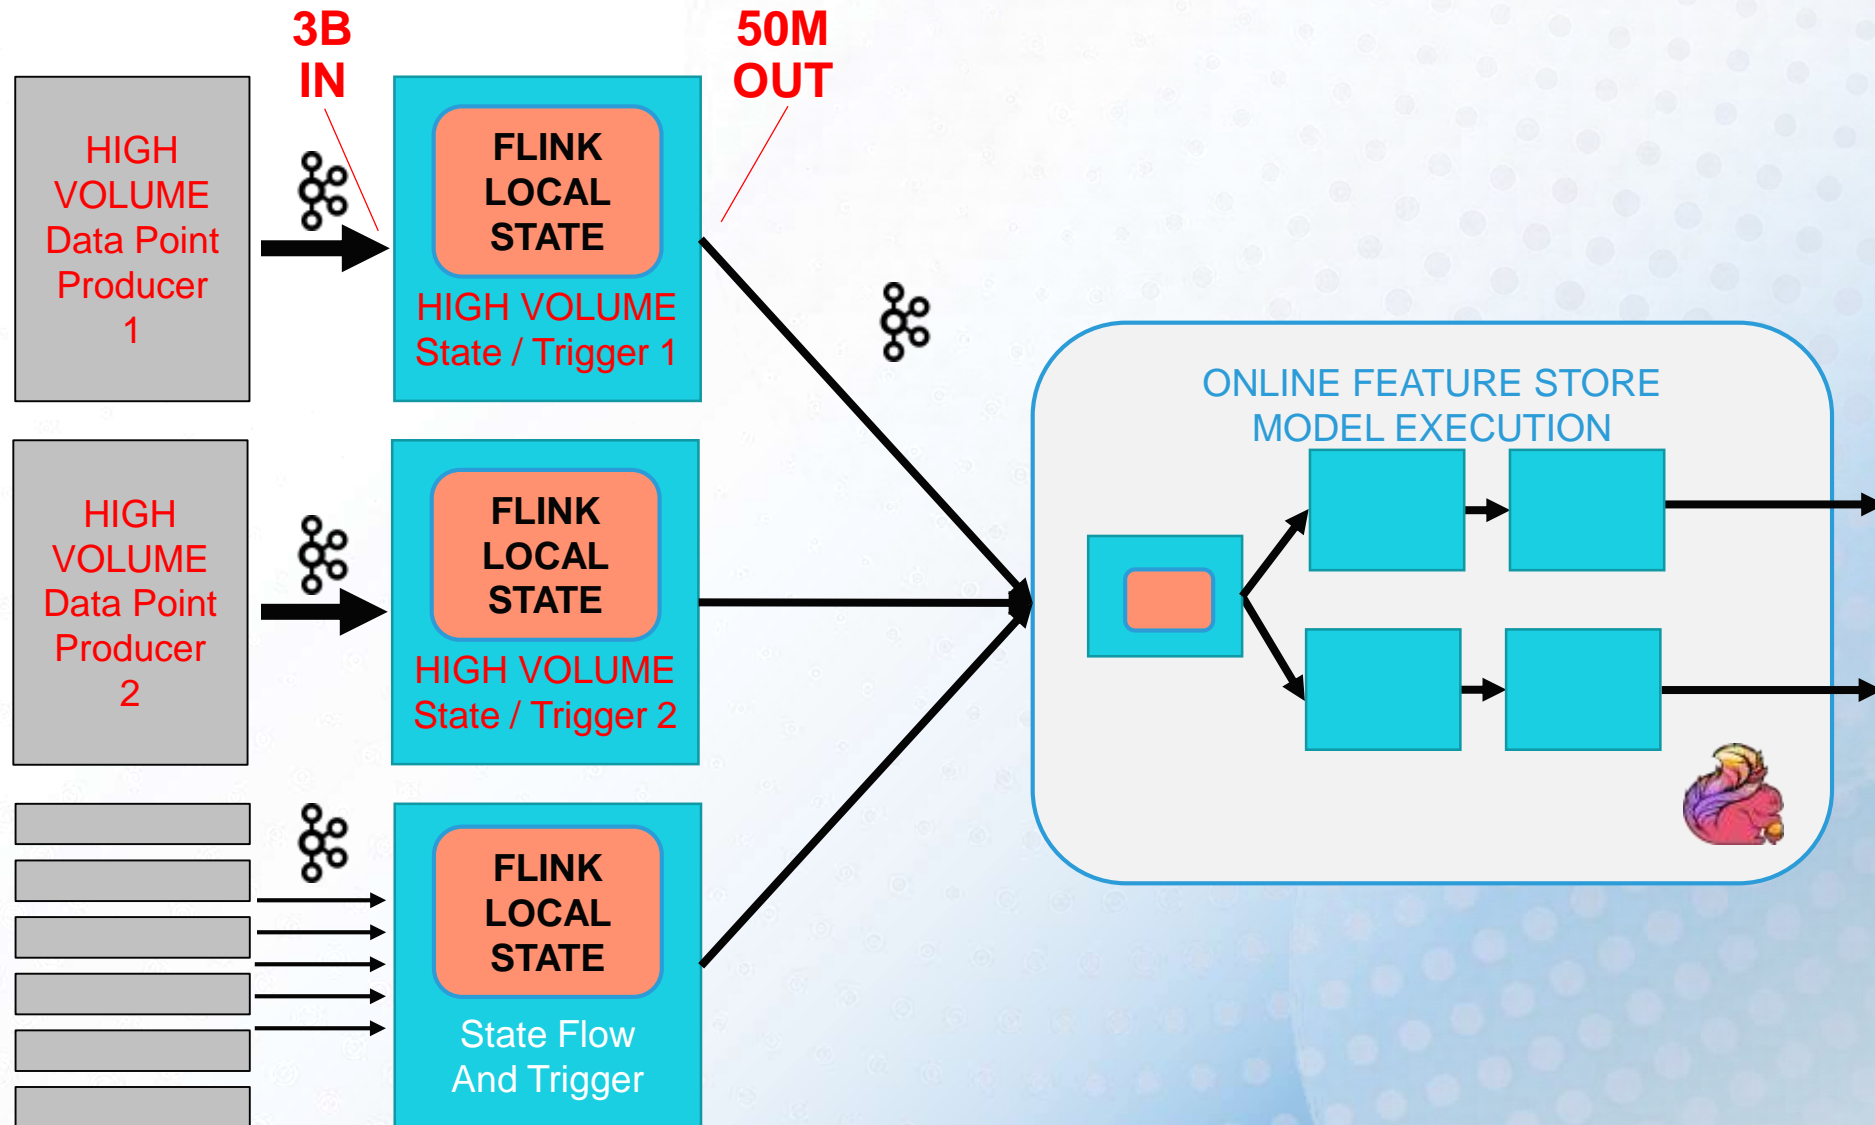
3 BILLION DATA POINTS



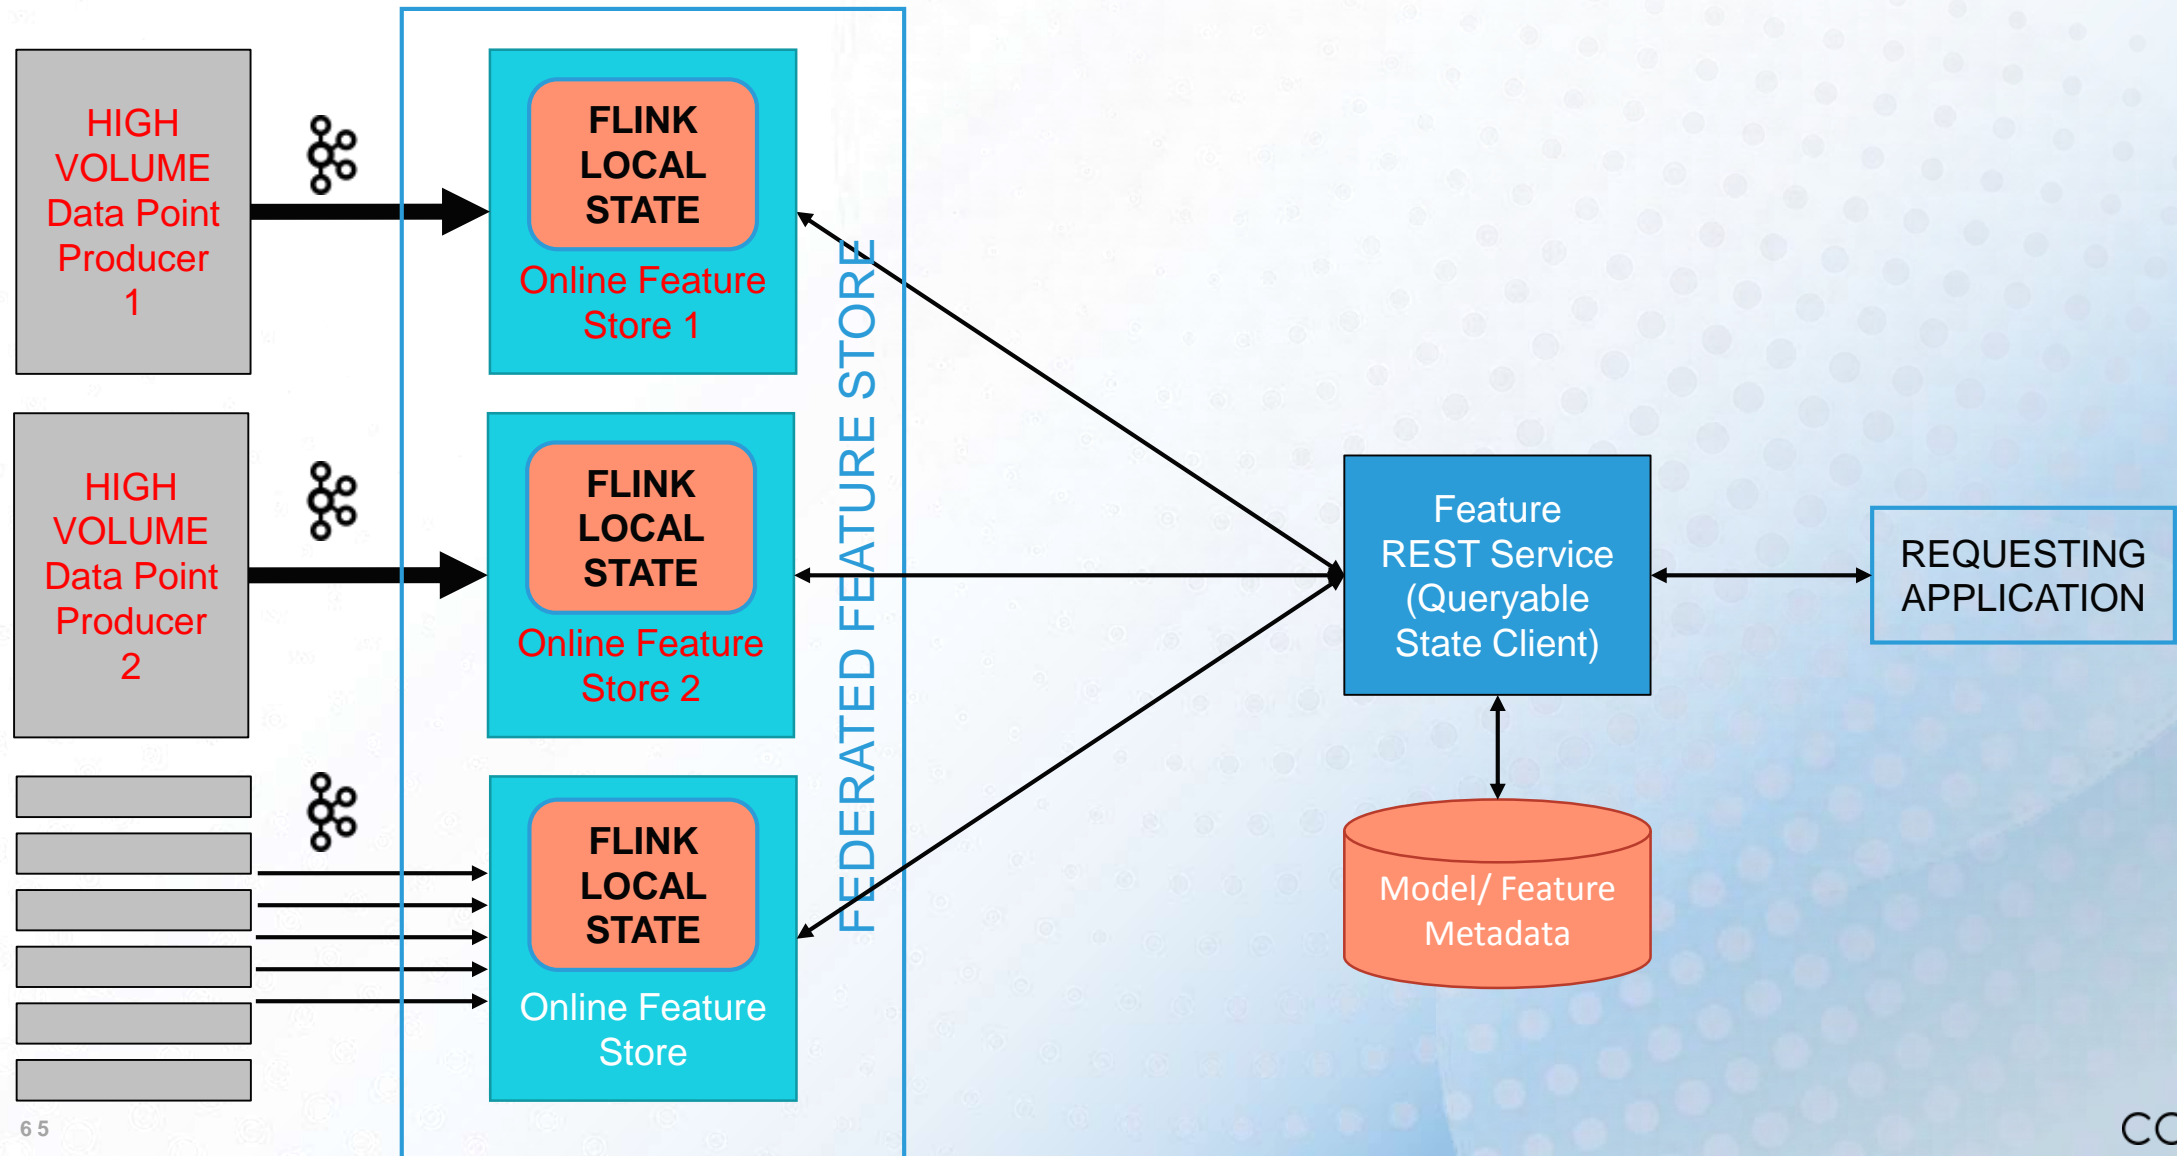
SOLUTION APPROACH

- 1 ISOLATE TRIGGER FOR HIGH-VOLUME DATA POINT TYPES**
- 2 ISOLATE ONLINE FEATURE STORE FOR HIGH- VOLUME DATA POINT TYPES**
- 3 ISOLATE DEDICATED FEATURE STORAGE FLOW TO HIGH-SPEED FEATURE STORE**
- 4 SEPARATE HISTORICAL FEATURE STORAGE (S3 WRITER) FROM REAL TIME MODEL EXECUTION FLOWS**

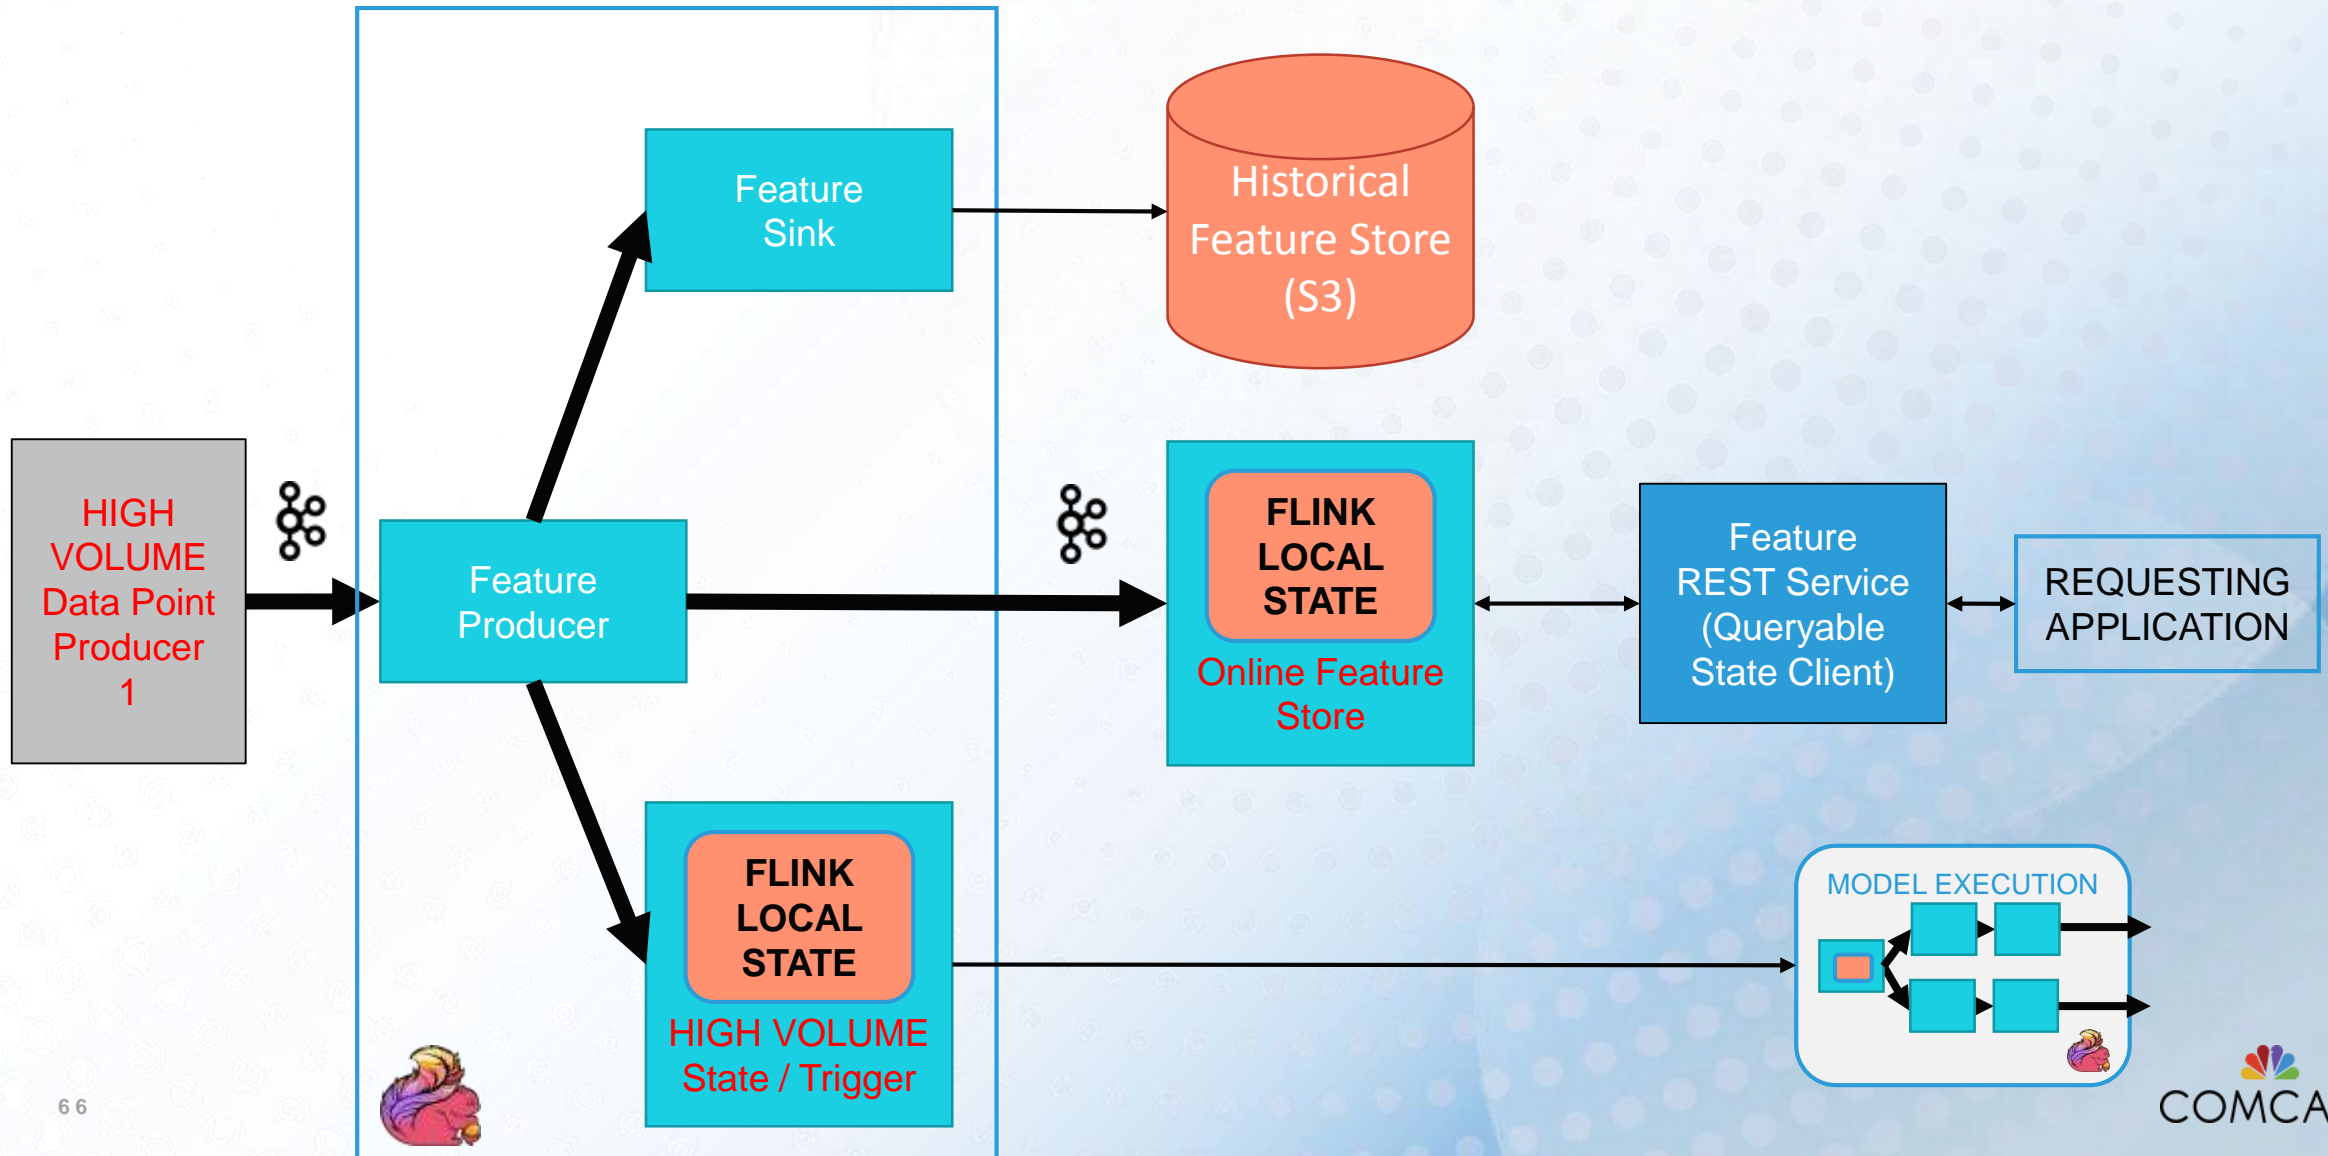
1: SEPARATE HIGH VOLUME TRIGGER FLOWS



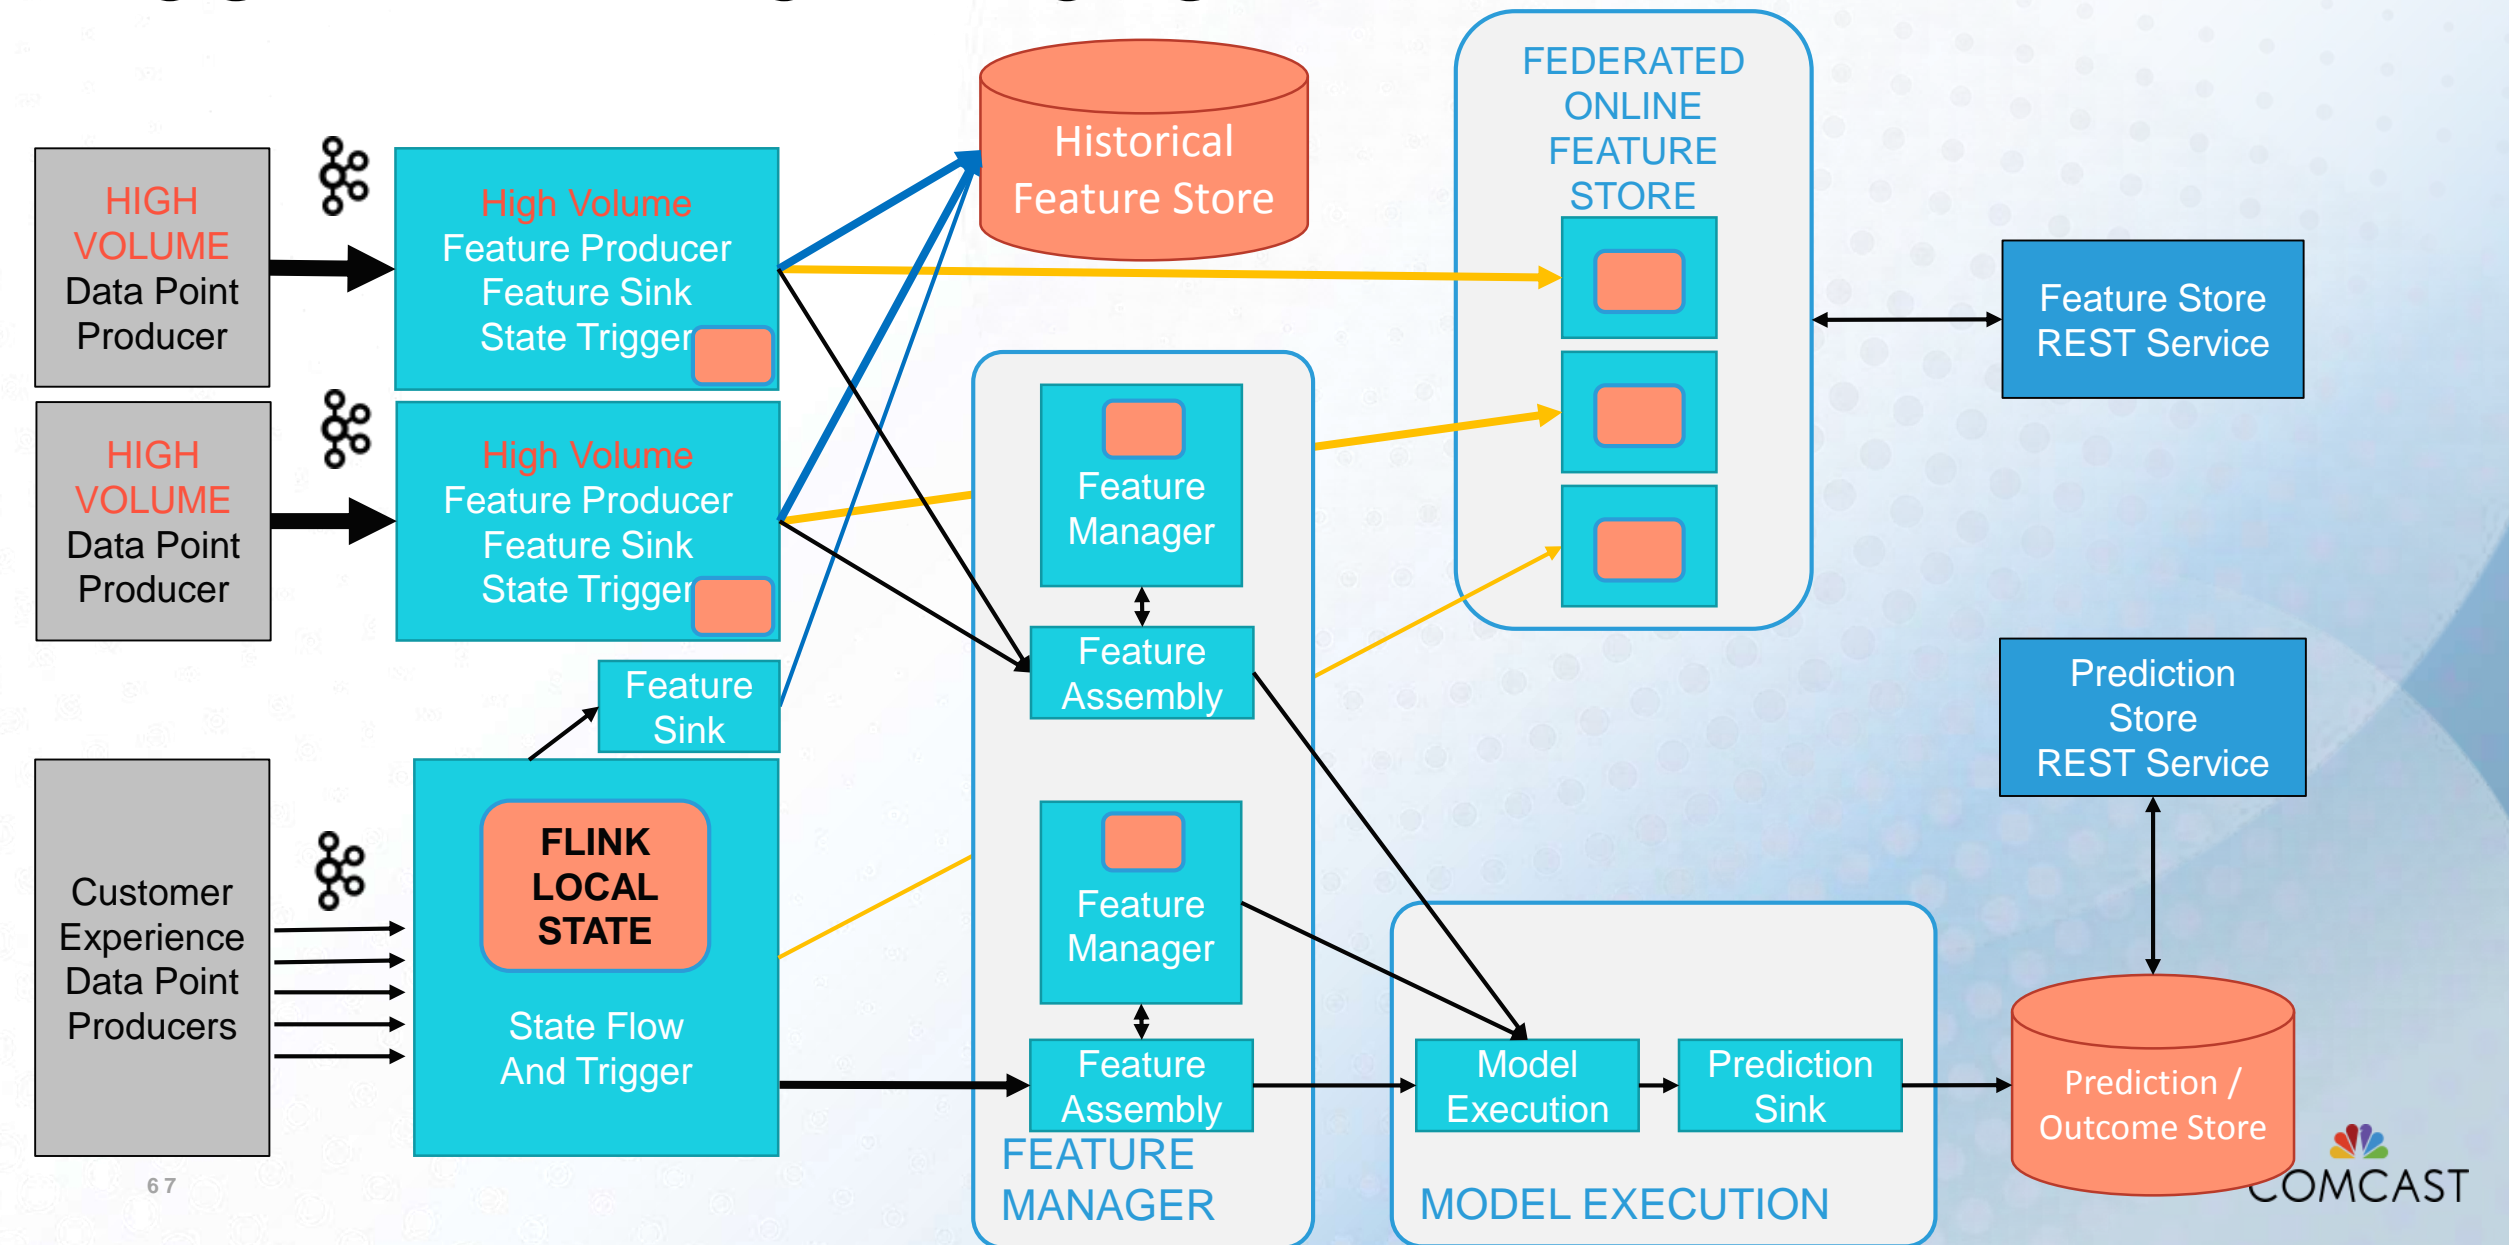
2: FEDERATED ONLINE FEATURE STORE



3+4: HISTORICAL FEATURE STORE



CURRENT ARCHITECTURE



WHERE ARE WE TODAY

INDICATORS

**725+
MILLION
DATA
POINTS
PER DAY**

HIGH-VOLUME INDICATORS

**6 BILLION
PER DAY**

TOTAL

**~7 BILLION
DATA POINTS
PER DAY**

WHERE ARE WE TODAY – FLINK CLUSTERS

CLUSTERS

14

VCPU

1100

INSTANCES

150

RAM

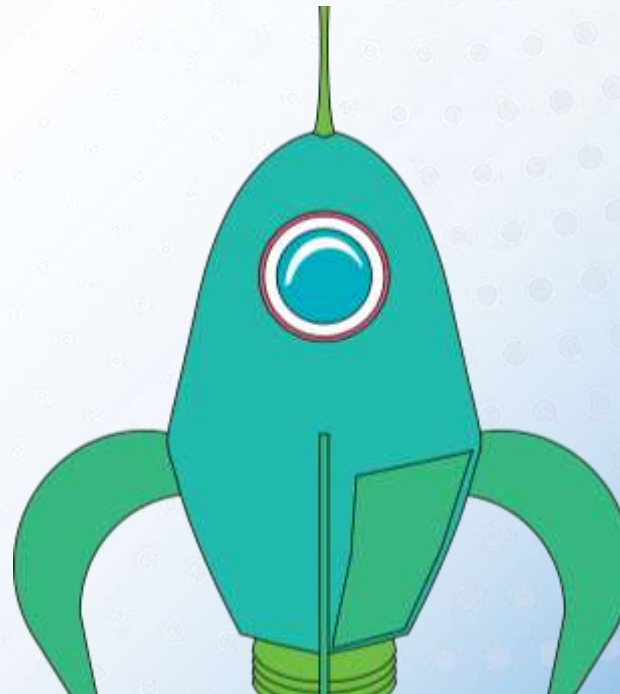
5.8 TB

FUTURE WORK

EXPAND
CUSTOMER EXPERIENCE
INDICATORS TO OTHER
COMCAST PRODUCTS

IMPROVED ML MODELS
BASED ON RESULTS AND
CONTINUOUS
RETRAINING

MODULAR
MODEL EXECUTION
VIA KUBEFLOW
AND GRPC CALLS



WE'RE HIRING!



**PHILADELPHIA
WASHINGTON, D.C.
MOUNTAIN VIEW
DENVER**



THANK YOU!



COMCAST