



Scaling Warehouse with Flink, Parquet & Kubernetes

Aditi Verma &
Ramesh Shanmugam



Aditi Verma

Sr Software Engineer

 @aditiverma89 <averma@branch.io>



Ramesh Shanmugam

Sr Data Engineer

 @rameshs01 <rshanmugam@branch.io>



branch

Agenda

- Background
- Moving data with Flink @ Branch
- Scale & Performance
- Flink on Kubernetes
- Auto Scaling & Failure Recovery

branch

12B requests per day (+70% y/y)

3B user sessions per day

50 TB of data per day

200K events per second

60+ Flink pipelines

5+ Kubernetes cluster



branch



kubernetes



Moving data with Flink @ Branch



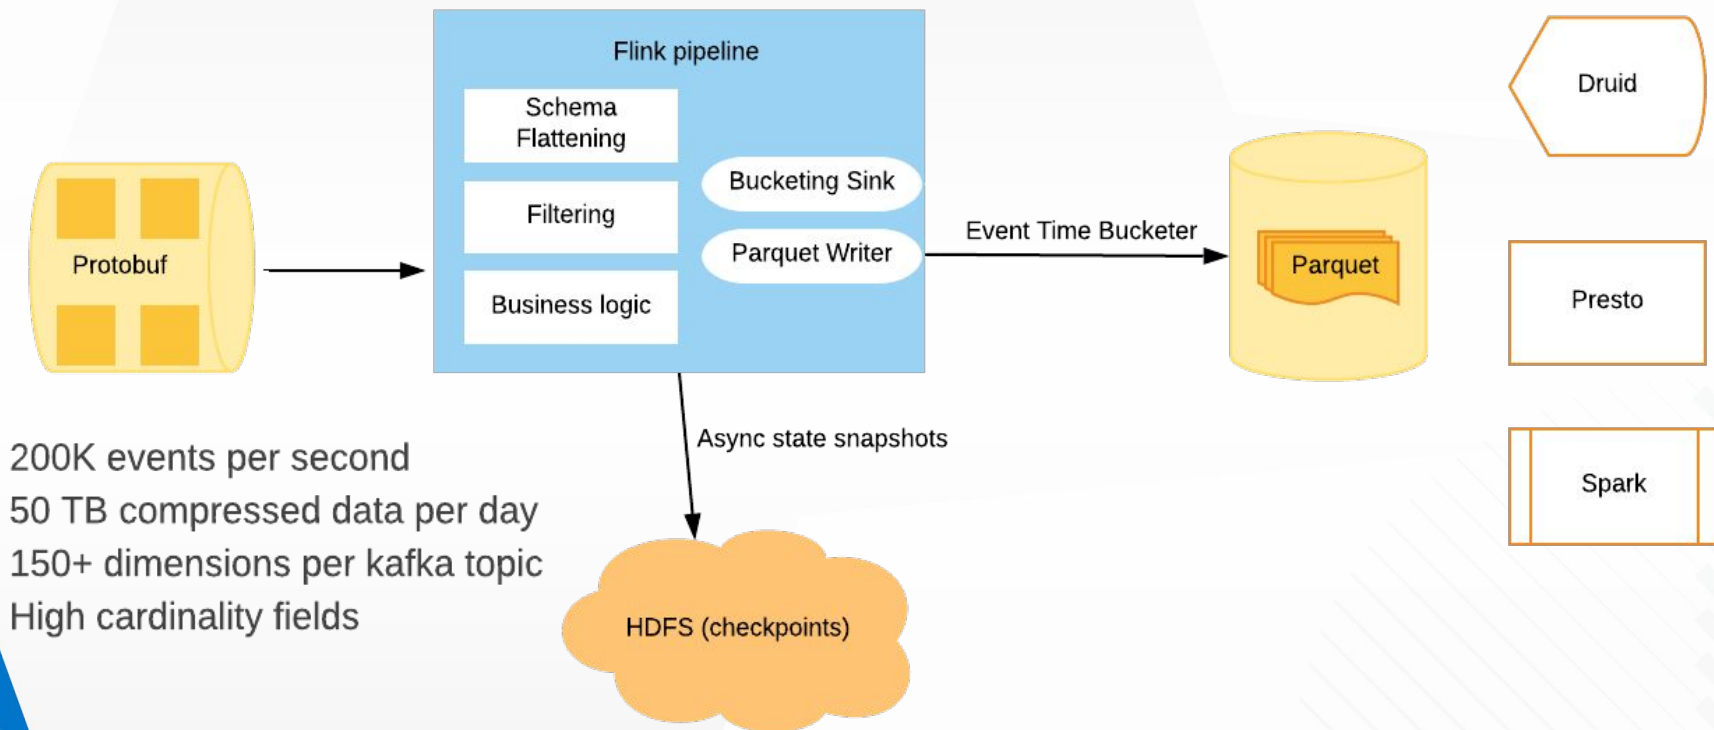
*"Life is 10% what happens to you and 90% how
you react to it."*

— **Charles R. Swindoll**


Receive information
Process it
React to it

FAST!!

Flink @ Branch



State Backend



- Relatively small state backend
- File system backed state



Parquet

- Higher compression
- Read heavy data set: ingested to Druid and Presto (3M+ queries/day)
- Avro data format
- Memory intensive writes



Writing parquet with Flink

Two approaches:

- 1) Close the file with checkpointing



Writing parquet with Flink

Two approaches:

- ~~a) Close the file with checkpointing~~
- b) Bucketing file sink
 - i) Configured with custom event-time bucketer, parquet writer and batch size
 - ii) Files are rolled out with a timeout of 10 min within a bucket



Performance and Scale

- 100% traffic increase each year
- Higher parallelism impacts application performance and state size
- Kafka partitions < Flink parallelism requires rebalance on the input stream
- Task manager timeouts



Message size

Network
buffers

Parquet page
size

Row group
size/Parquet
block size

State
backend
access
pattern

Checkpoint
frequency

Records per
second

MEMORY



Analyzing memory usage

- ❖ Network Buffers
- ❖ Memory Segments
- ❖ User code
- ❖ Memory and GC stats
- ❖ JVM parameters



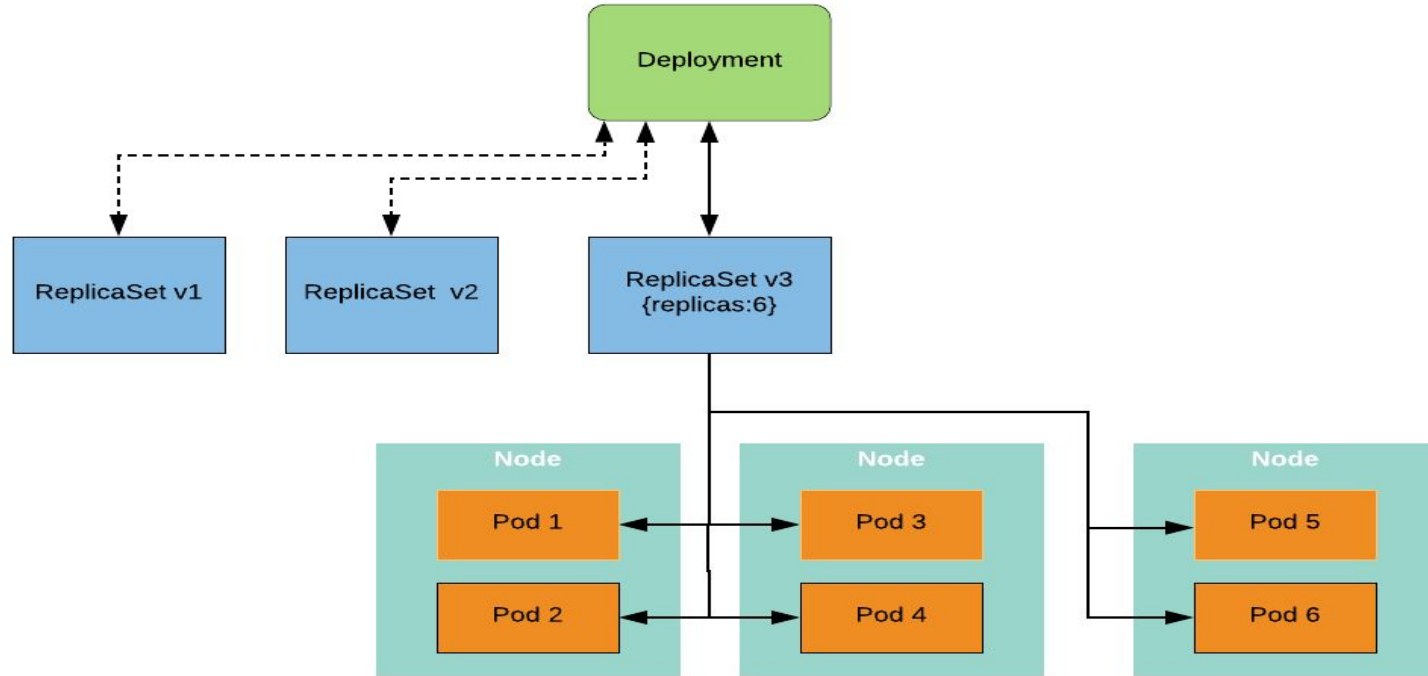
Containerizing Flink - Mesos

- Longer start-up time on Mesos
- Moved to containerizing Flink application on Kubernetes
- Kubernetes is resource oriented, declarative





Kubernetes Terms

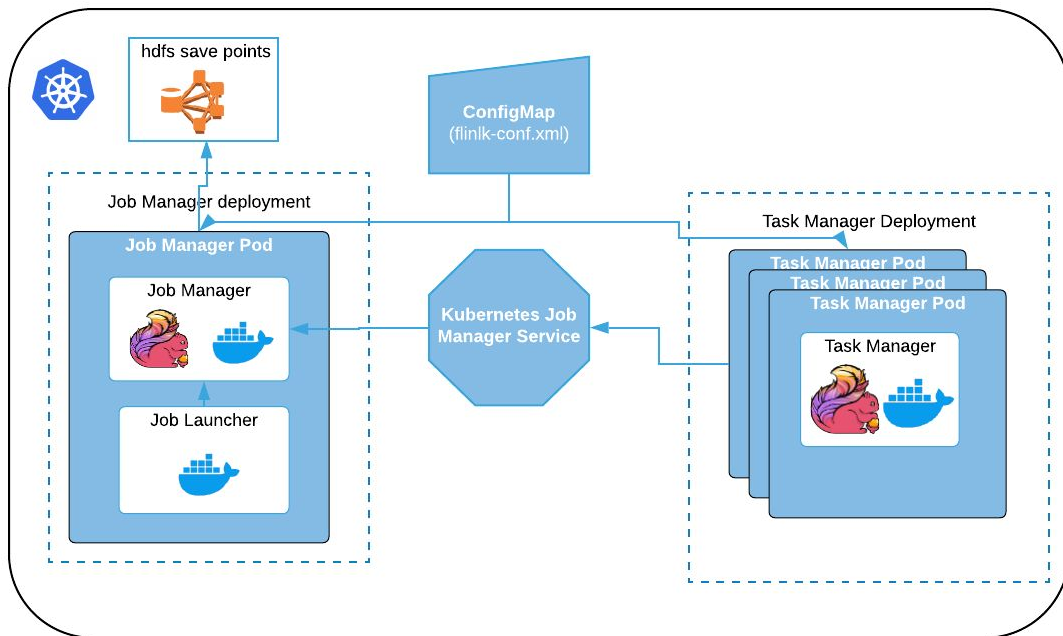




Flink



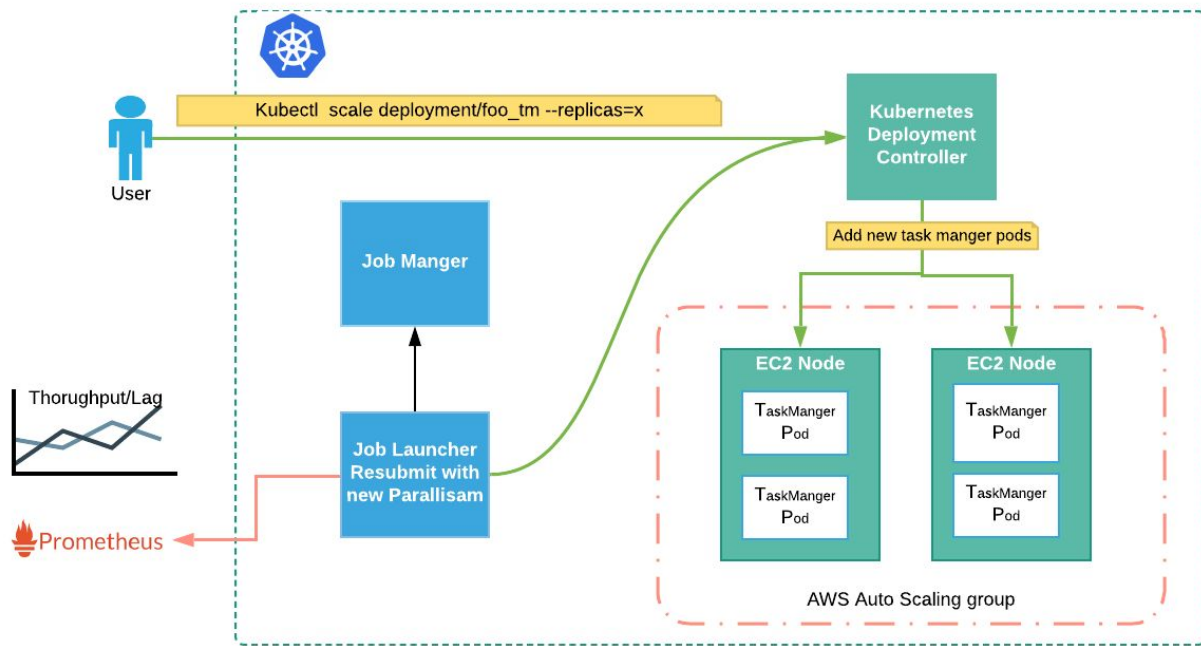
Flink on Kubernetes @ Branch



- Single job per cluster
- Docker image
 - flink image - Task manager + job manager
 - Job launcher - custom launcher + job jar
- Job launcher
 - Application jar
 - Uploads jar
- Config map - flink config.xml
 - jobmanager.rpc.address



Auto Scaling

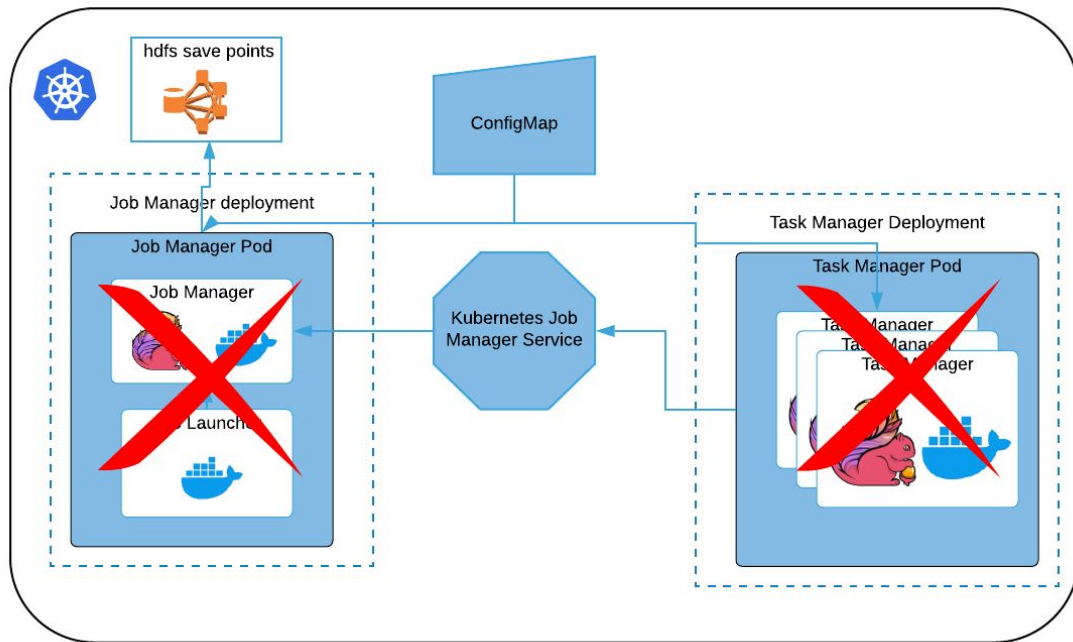


- When & How much scale
 - Auto - Joblauncher
- Scale
 - Replica Set
- Flink job with new parallelism

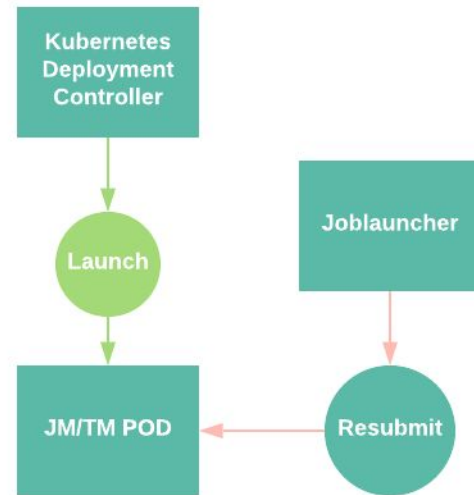
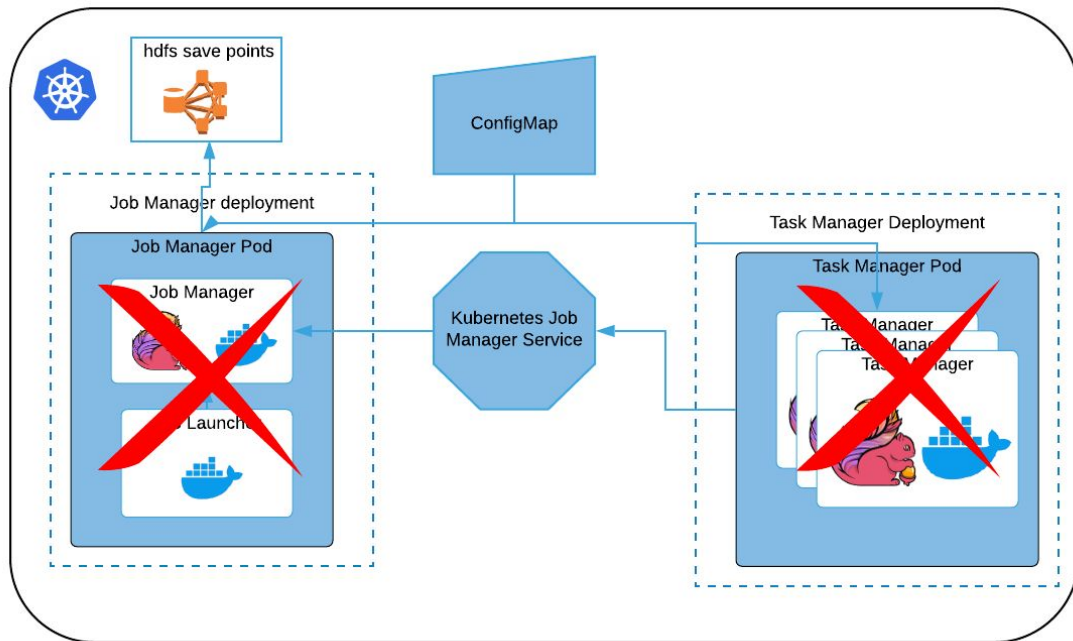


Failure Recovery

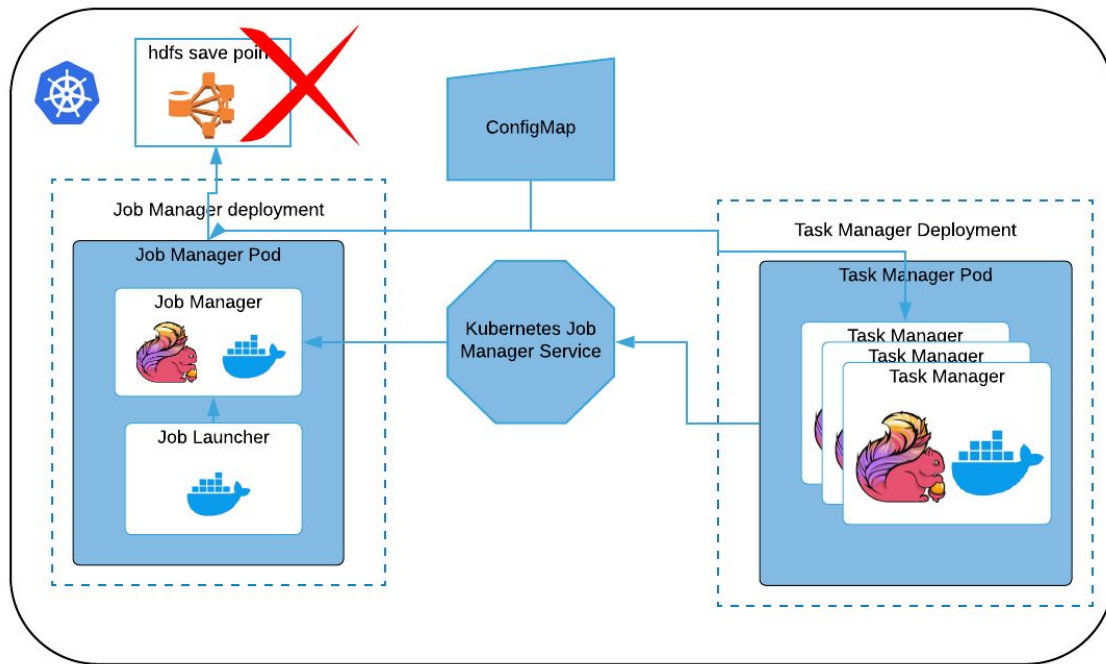
Job / Task Manager Goes Down?



Job / Task Manager Goes Down?



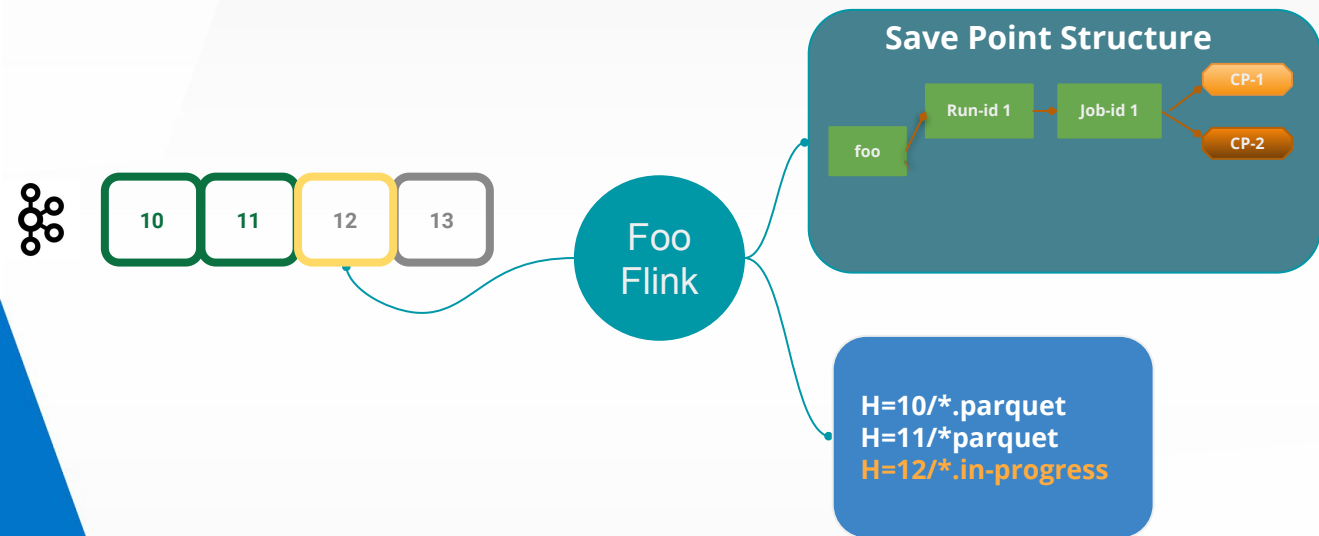
Savepoint Failure



- Reasons
 - Truncation
 - Schema mismatch
 - Hdfs outage

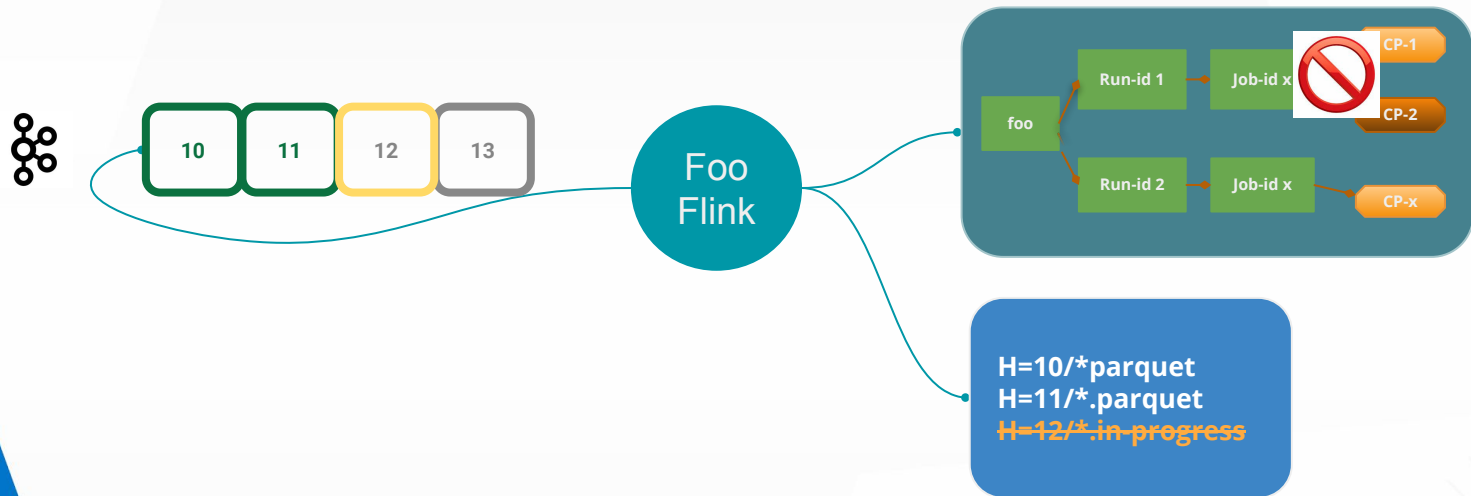


Savepoint Structure

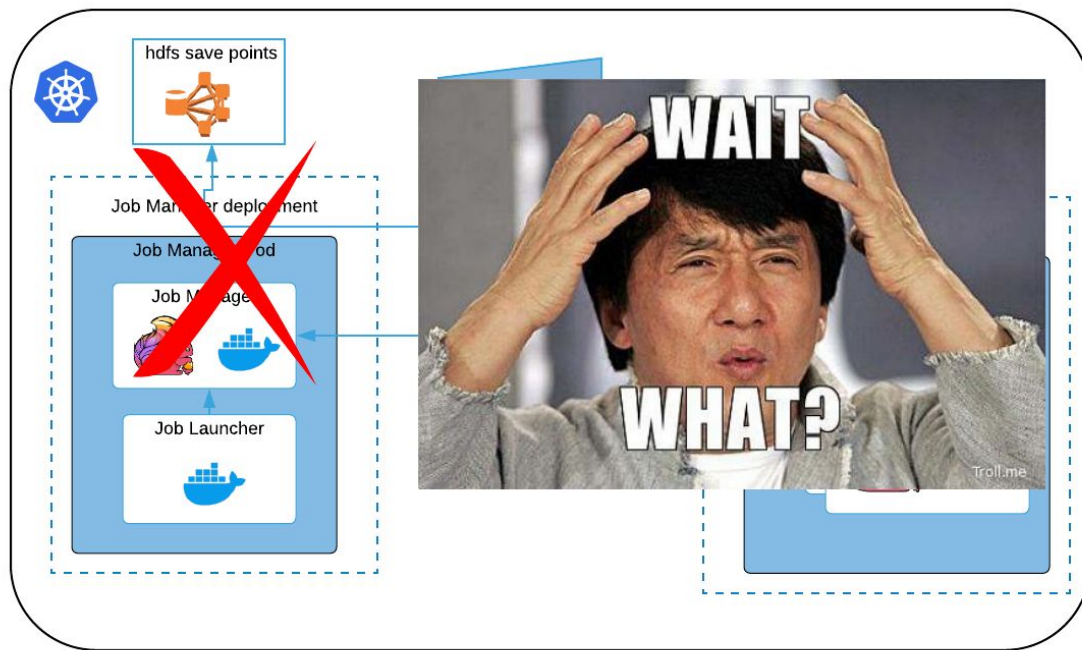


- job/run-id/flink-job-id/cp-x
- Run id - incremental number
- Job id - flink job name

Savepoint failure recovery



Auto Recovery does not work?



- Continuous monitoring and proper alerts
- start job from latest offset
- Have different backfill route



Next Steps....

- Parquet memory consumption (when too many buckets open)
 - Window + Rocks db => Parquet
 - Two stage process
 - row oriented streaming
 - batch to convert columnar



Q & A