

# Talk Python to Me

Stream Processing in Your Favourite Language with Beam on Flink



Apache Beam



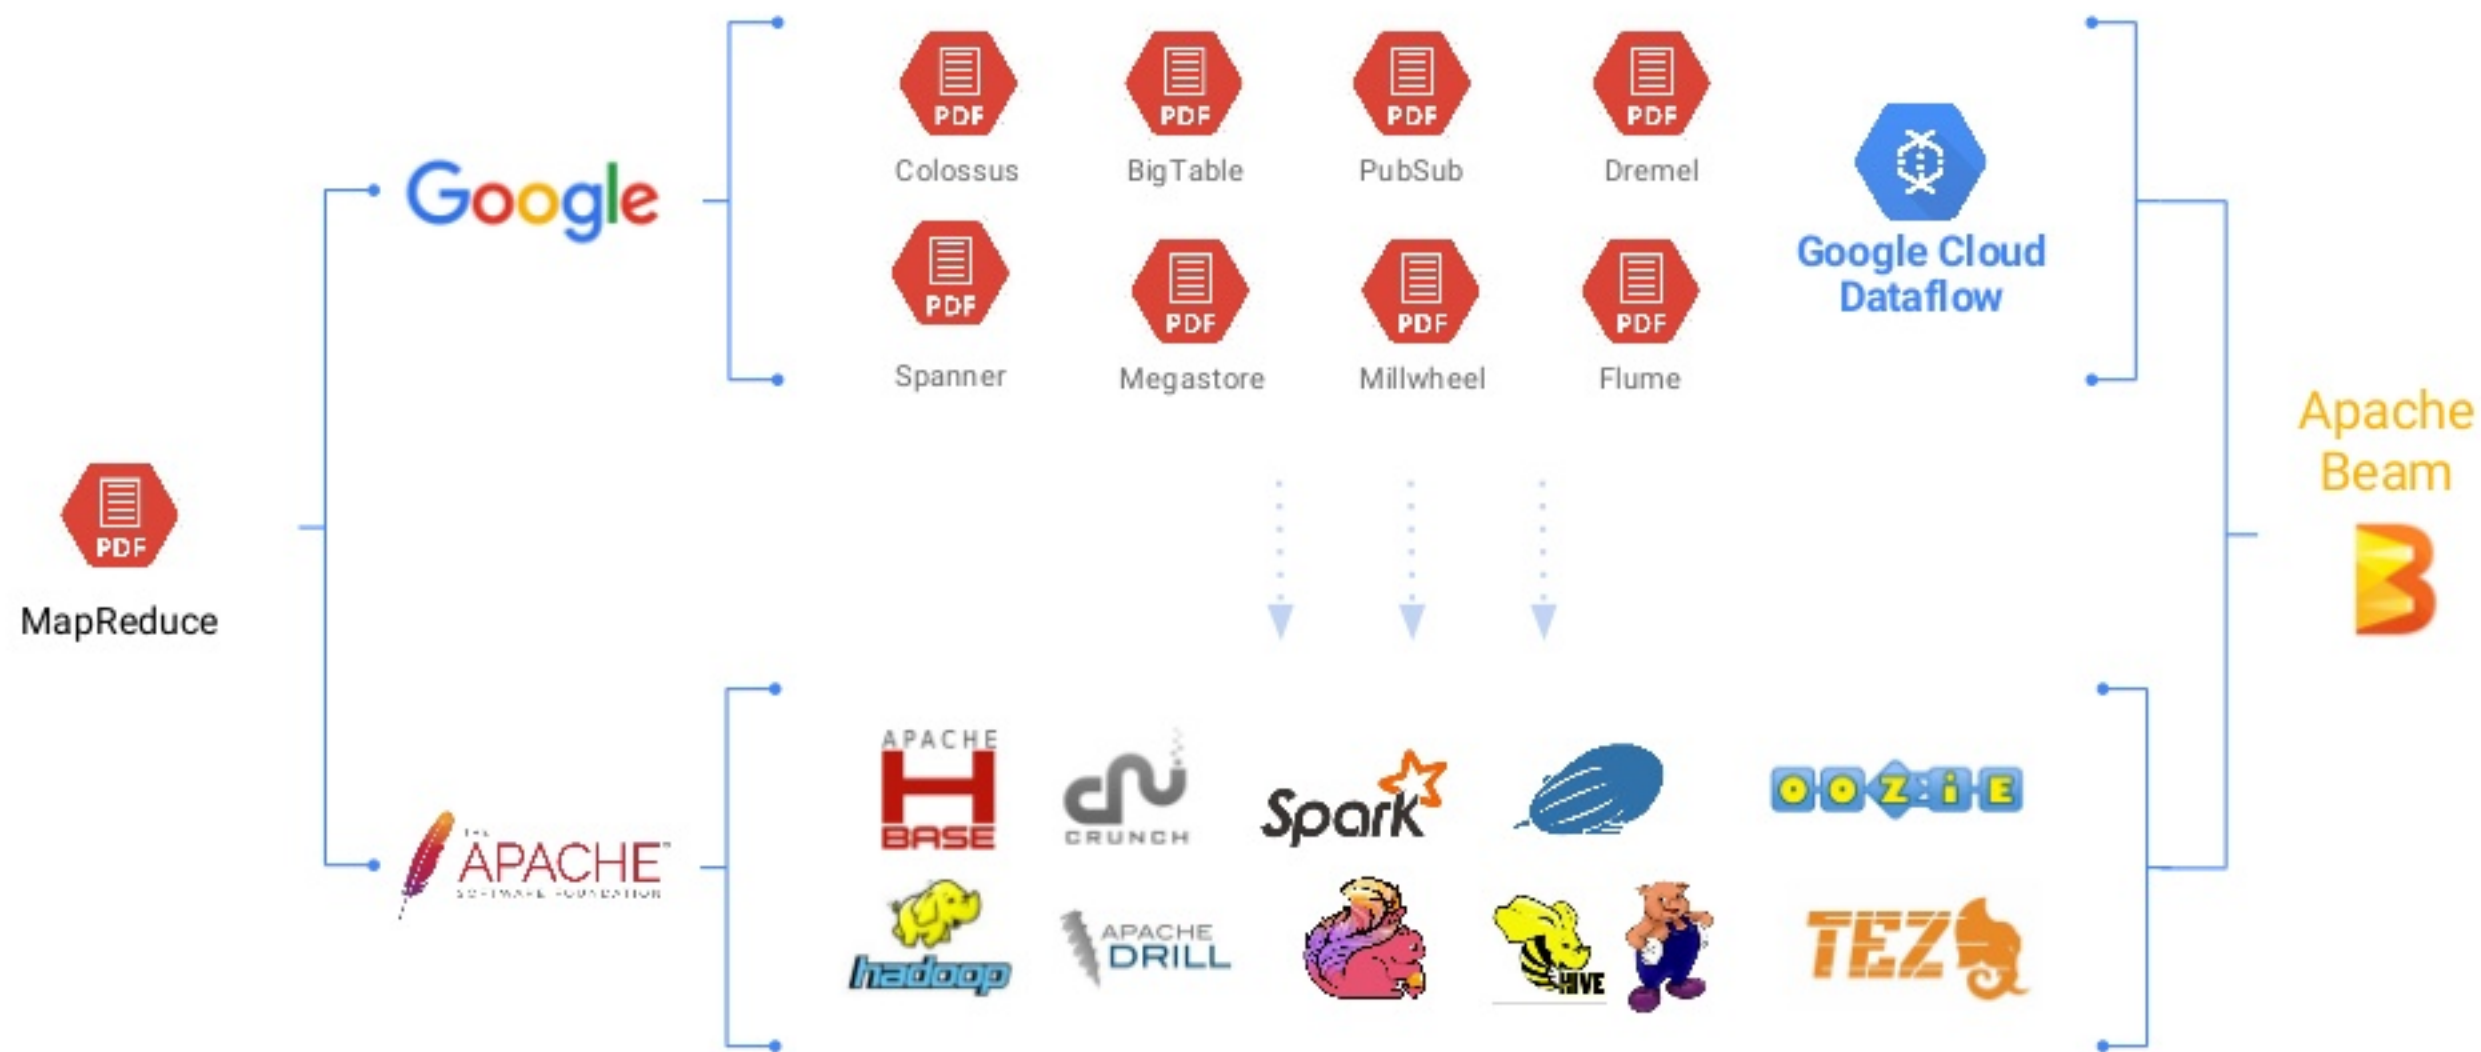
Apache Flink

# Agenda

1. What is Beam?
2. The Beam Portability APIs (Fn / Pipeline)
3. Executing Pythonic Beam Jobs on Flink
4. The Future

What is Beam?

# The Evolution of Apache Beam




# Beam Model: Generations Beyond MapReduce

Improved abstractions let you focus on your application logic

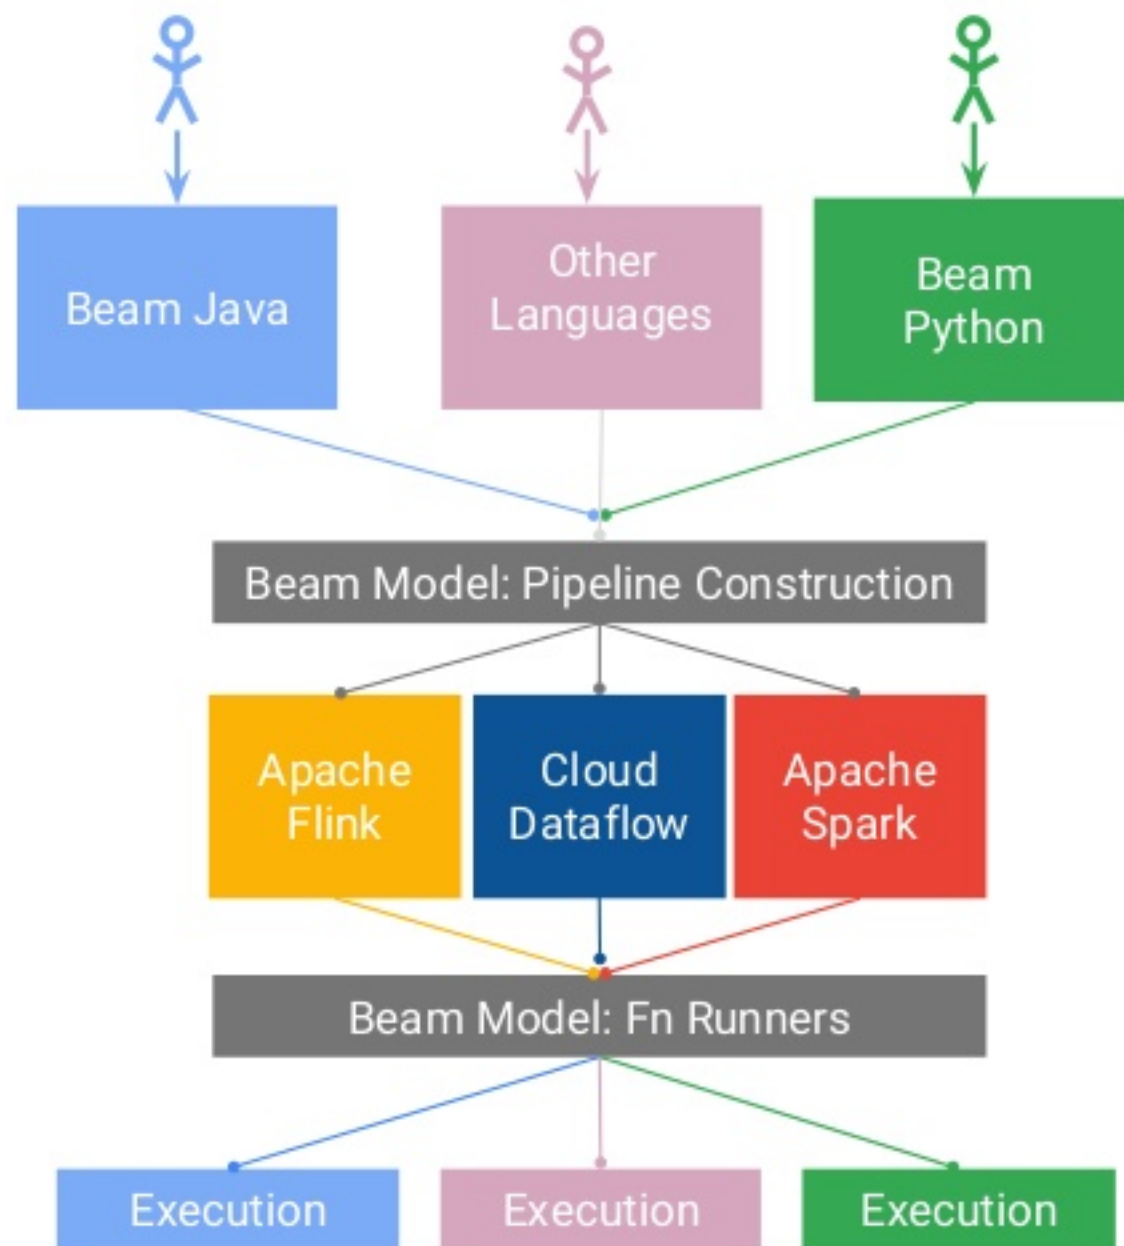
Batch and stream processing are *both* first-class citizens -- no need to choose.

Clearly separates event time from processing time.





# The Apache Beam Vision

1. **End users:** who want to write pipelines in a language that's familiar.
2. **SDK writers:** who want to make Beam concepts available in new languages.
3. **Runner writers:** who have a distributed processing environment and want to support Beam pipelines




# The Beam Model



(Flink draws it more like this)



# The Beam Model





# Beam Model: Asking the Right Questions

**What** results are calculated?

**Where** in event time are results calculated?

**When** in processing time are results materialized?


**How** do refinements of results relate?

# The Beam Model: **What** is Being Computed?

```
PCollection<KV<String, Integer>> scores = input  
    .apply(Sum.integersPerKey());
```

```
scores = (input  
    | Sum.integersPerKey())
```

# The Beam Model: **What** is Being Computed?

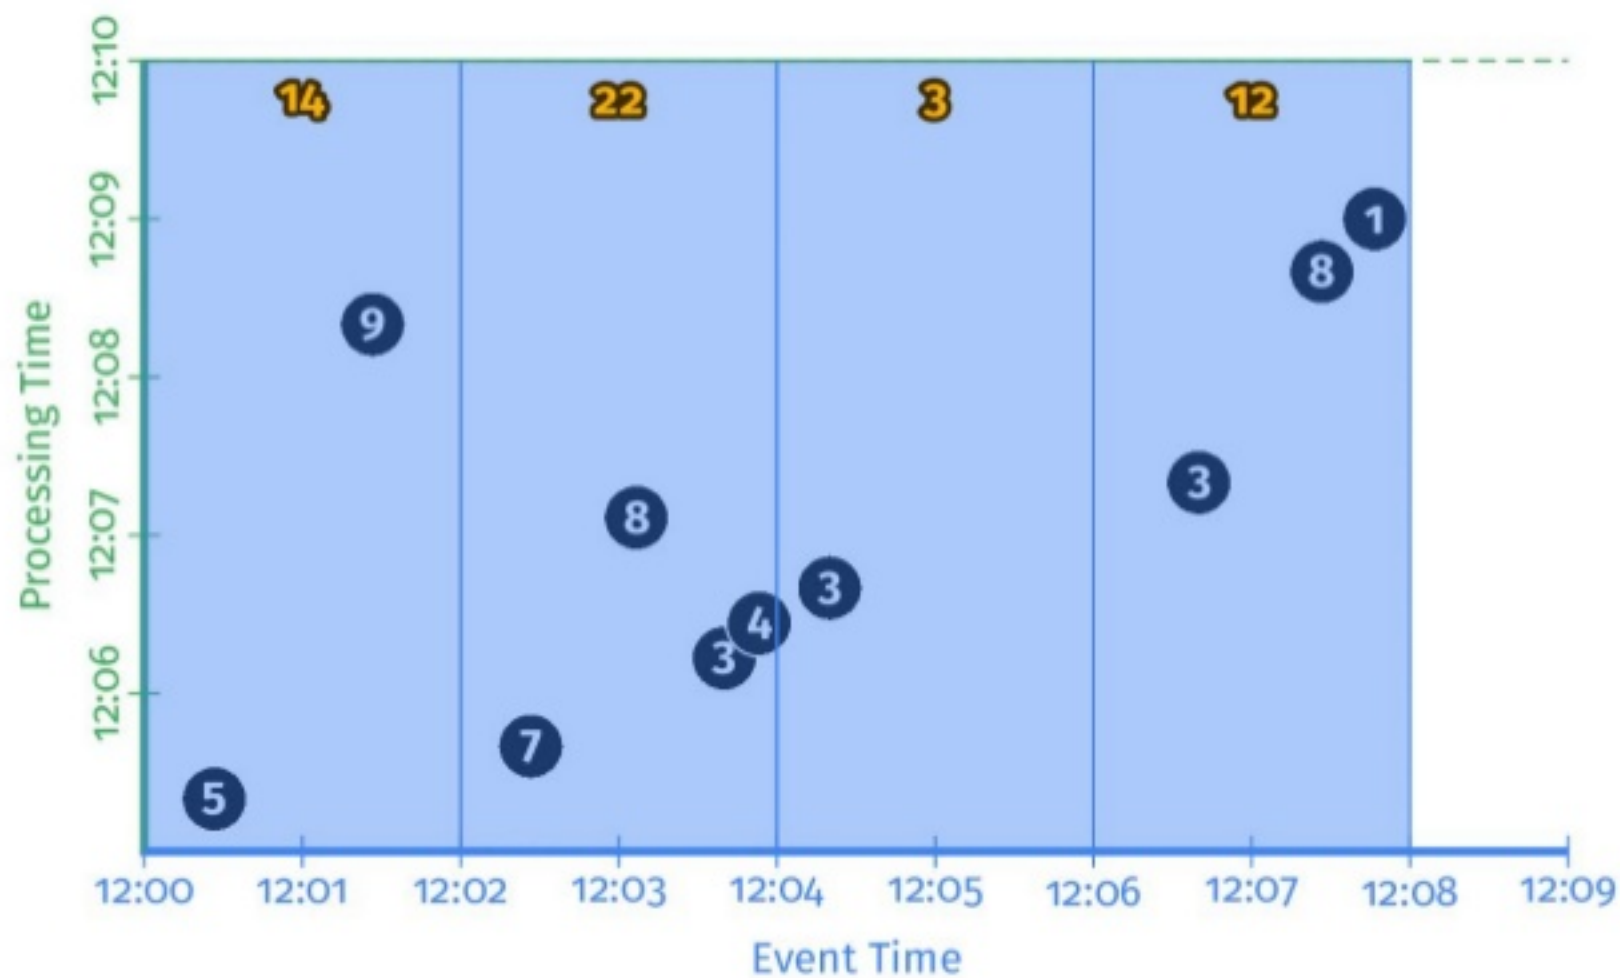


# The Beam Model: **Where** in Event Time?

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))))
    .apply(Sum.integersPerKey());
```

```
scores= (input
| beam.WindowInto(FixedWindows(2 * 60))
| Sum.integersPerKey())
```

# The Beam Model: **Where** in Event Time?




# The Beam Model: **When** in Processing Time?

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))
        .triggering(AtWatermark())))
    .apply(Sum.integersPerKey());
```

```
scores = (input
    | beam.WindowInto(FixedWindows(2 * 60)
        .triggering(AtWatermark()))
    | Sum.integersPerKey())
```

# The Beam Model: **When** in Processing Time?



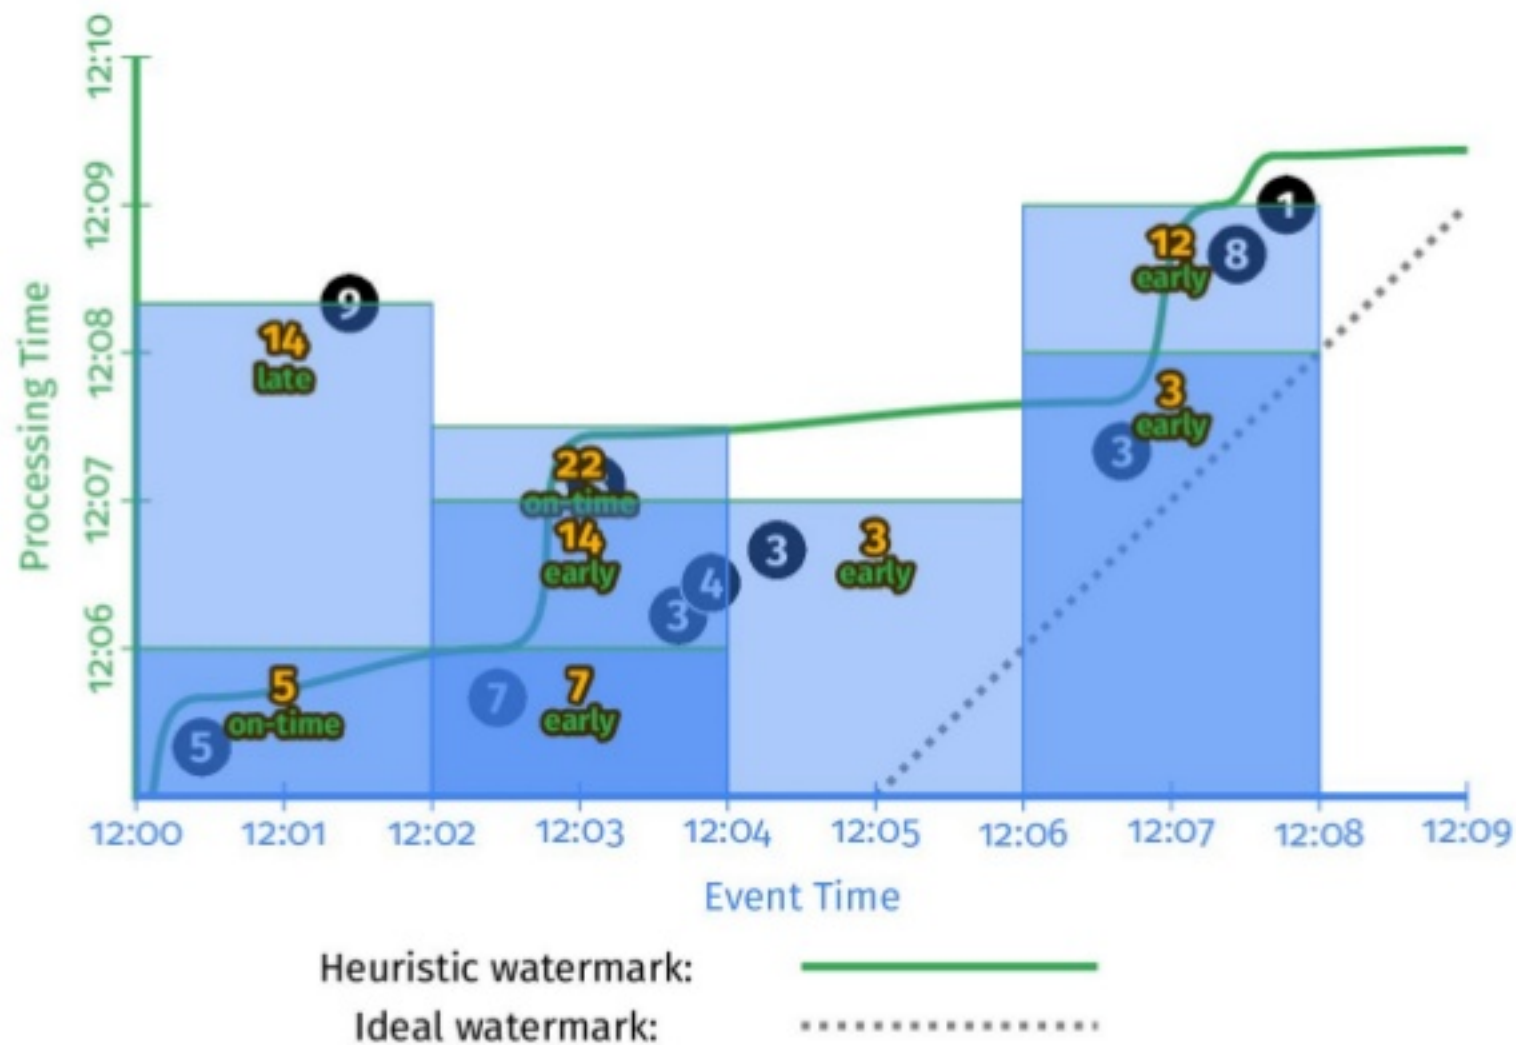


# The Beam Model: **How** Do Refinements Relate?

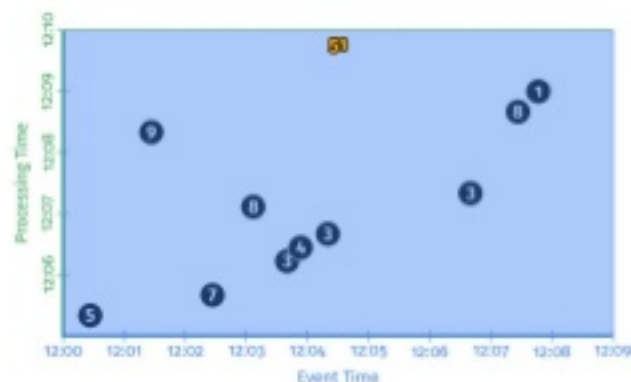
```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Duration.standardMinutes(1)))
            .withLateFirings(AtCount(1)))
        .accumulatingFiredPanels()))
    .apply(Sum.integersPerKey());
```

```
scores = (input
    | beam.WindowInto(FixedWindows(2 * 60)
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(1 * 60))
            .withLateFirings(AtCount(1)))
        .accumulatingFiredPanels())
    | Sum.integersPerKey())
```


# The Beam Model: **How** Do Refinements Relate?




# Customizing **What** **Where** **When** **How**



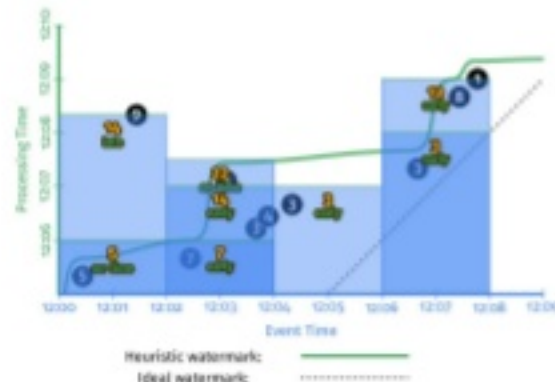
**1**  
**Classic  
Batch**



**2**  
**Windowed  
Batch**



**3**  
**Streaming**



**4**  
**Streaming  
+ Accumulation**

For more information see <https://cloud.google.com/dataflow/examples/gaming-example>

# A Complete Example of Pythonic Beam Code

```
import apache_beam as beam
with beam.Pipeline() as p:
    (p
     | beam.io.ReadStringsFromPubSub("twitter_topic")
     | beam.WindowInto(SlidingWindows(5*60, 1*60))
     | beam.ParDo(ParseHashTagDoFn())
     | beam.combiners.Count.PerElement()
     | beam.ParDo(BigQueryOutputFormatDoFn())
     | beam.io.WriteToBigQuery("trends_table"))
```

# What is Apache Beam?

1. The Beam Model: **What** / **Where** / **When** / **How**
2. SDKs for writing Beam pipelines
3. Runners for Existing Distributed Processing Backends
  - Apache Apex
  - Apache Flink
  - Apache Spark
  - Google Cloud Dataflow
  - Local (in-process) runner for testing



# Beam Portability APIs (Pipeline / Job / Fn)



# What are we trying to solve?

- Executing user code written in an arbitrary language (Python) on a Runner written in a different language (Java)
- Mixing user functions written in different languages (Connectors, Sources, Sinks, ...)



# Terminology

## **Beam Model**

Describes the API concepts and the possible operations on PCollections.

## **Pipeline**

User-defined graph of transformations on PCollections. This is constructed using a Beam SDK. The transformations can contain UDFs.

## **Runner**

Executes a Pipeline. For example: *FlinkRunner*.

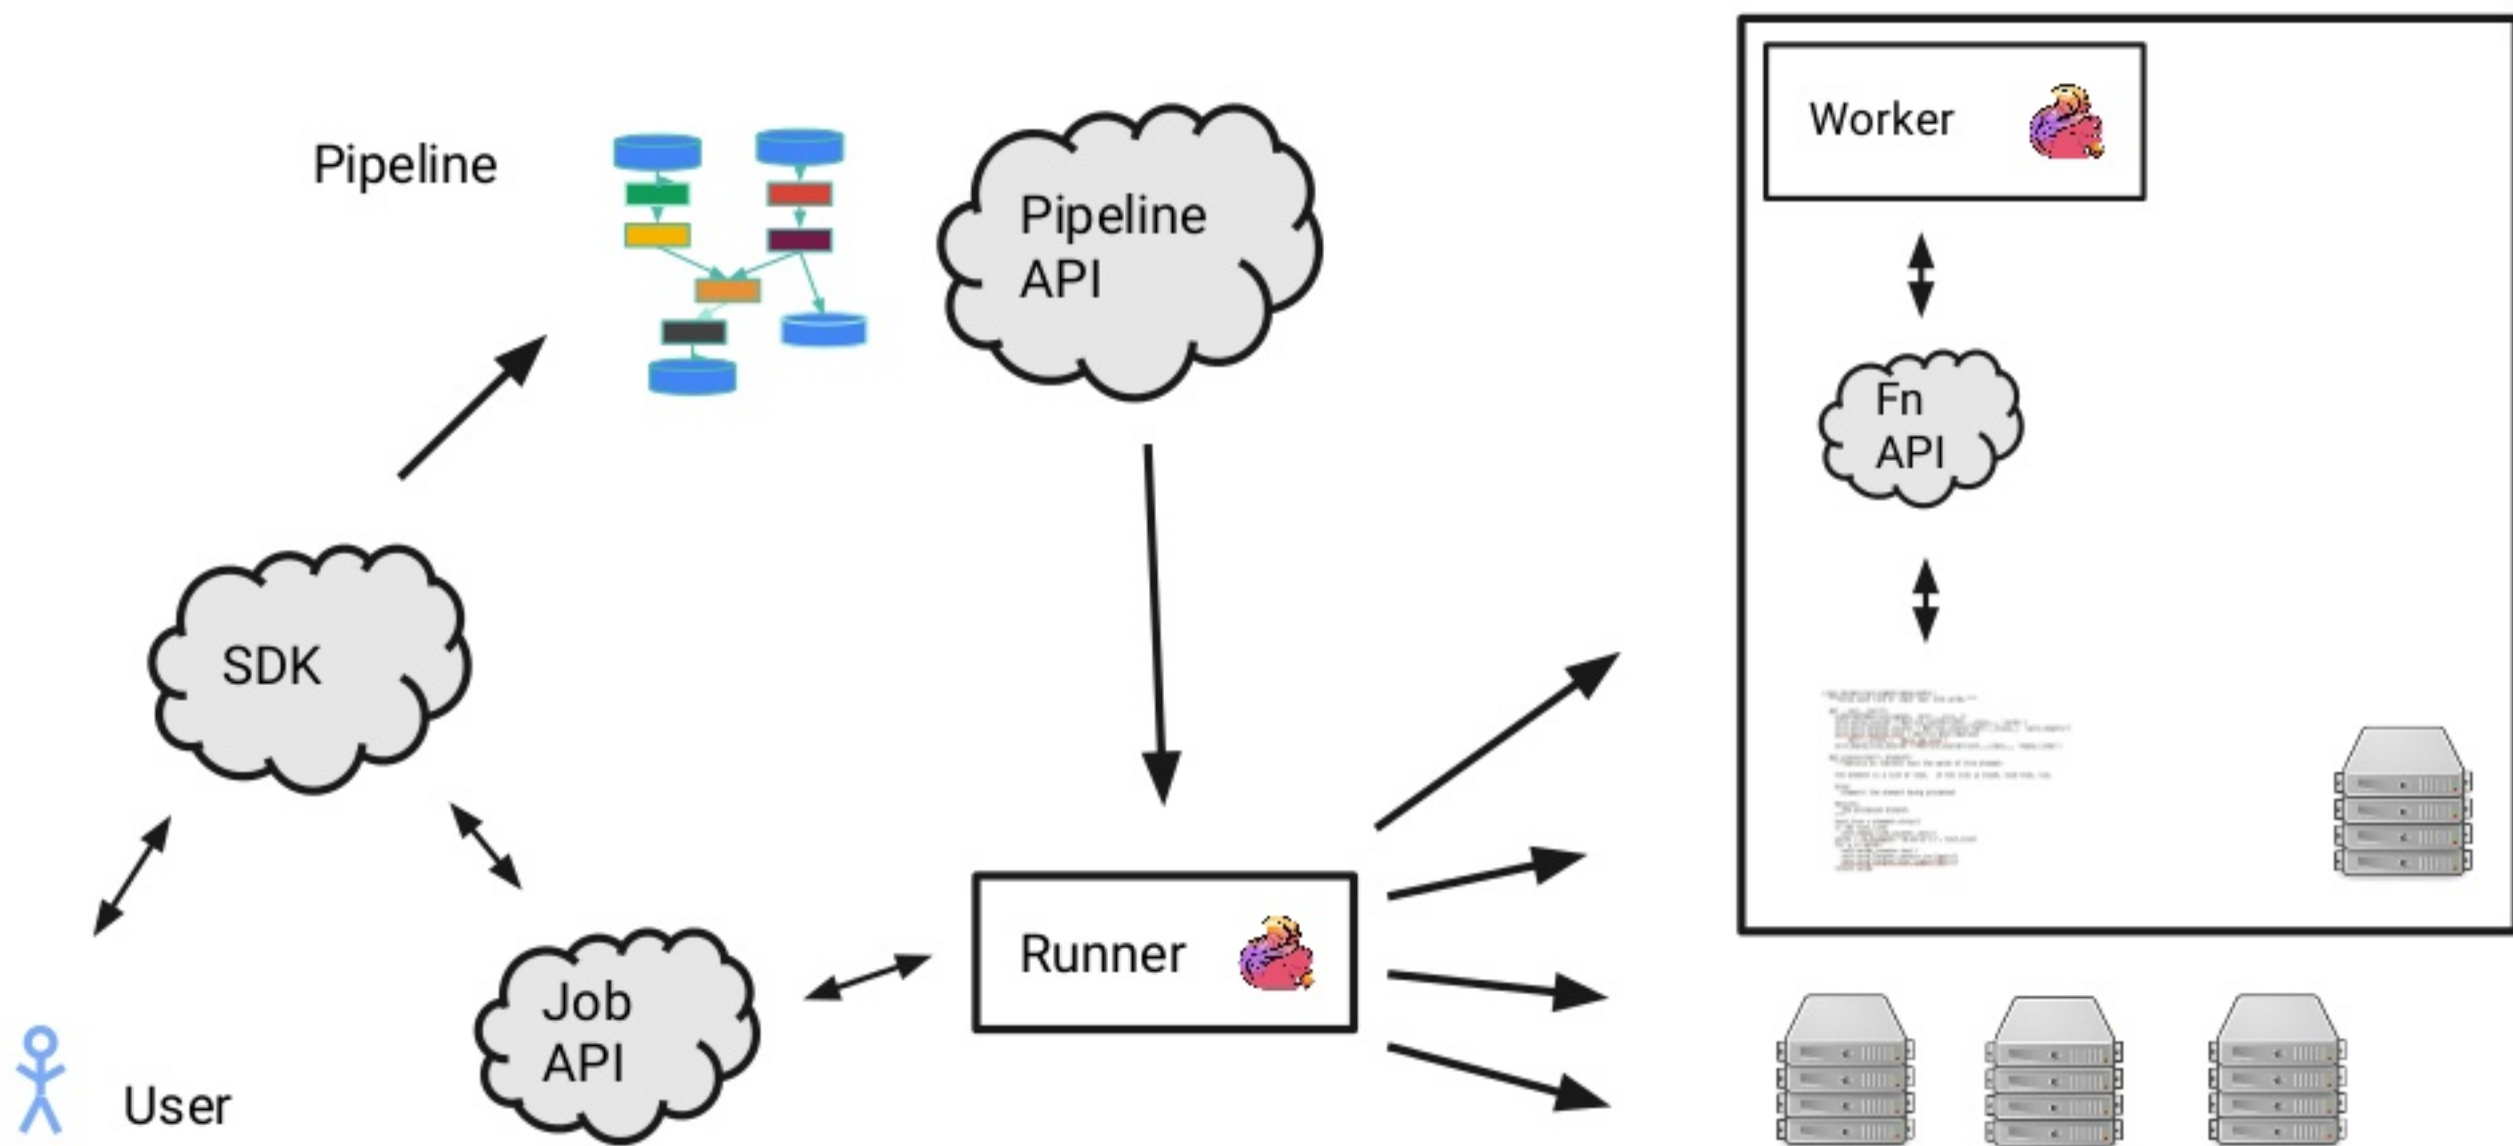
## **Beam SDK**

Language specific library/framework for creating programs that use the Beam Model. Allows defining Pipelines and UDFs and provides APIs for executing them.

## **User-defined function (UDF)**

Code in Java, Python, ... that specifies how data is transformed. For example *DoFn* or *CombineFn*.

# Executing a Beam Pipeline - The Big Picture



# APIs for Different Pipeline Lifecycle Stages

## Pipeline API

- Used by the SDK to construct SDK-agnostic Pipeline representation
- Used by the Runner to translate a Pipeline to runner-specific operations

## Fn API

- Used by an SDK harness for communication with a Runner
- User by the Runner to push work into an SDK harness

## Job API

- (API for interacting with a running Pipeline)

# Pipeline API (simplified)

- Definition of common primitive transformations (Read, ParDo, Flatten, Window.into, GroupByKey)
- Definition of serialized Pipeline (protobuf)

```
Pipeline = {PCollection*, PTransform*, WindowingStrategy*,  
Coder*}
```

```
PTransform = {Inputs*, Outputs*, FunctionSpec}
```

```
FunctionSpec = {URN, payload}
```


# Job API

```
public interface JobApi {  
    State getState(); // RUNNING, DONE, CANCELED, FAILED ...  
    State cancel() throws IOException;  
    State waitUntilFinish(Duration duration);  
    State waitUntilFinish();  
    MetricResults metrics();  
}
```

# Fn API

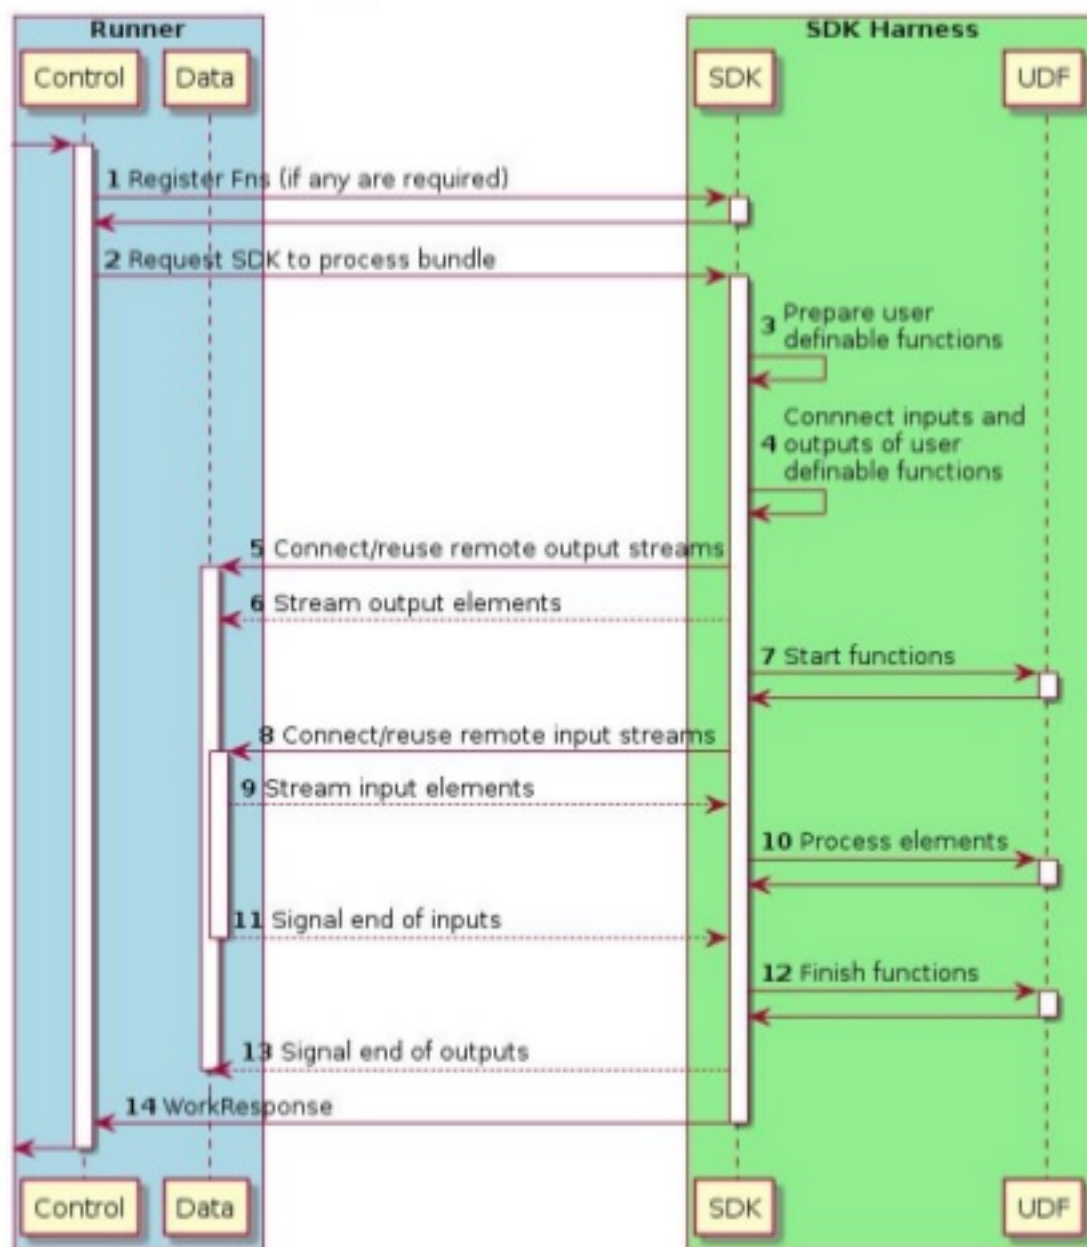
- gRPC interface definitions for communication between an SDK harness and a Runner
- **Control:** Used to tell the SDK which UDFs to execute and when to execute them.
- **Data:** Used to move data between the language specific SDK harness and the runner.
- **State:** Used to support user state, side inputs, and group by key reiteration.
- **Logging:** Used to aggregate logging information from the language specific SDK harness.

## Fn API (continued)







# Fn API - Bundle Processing



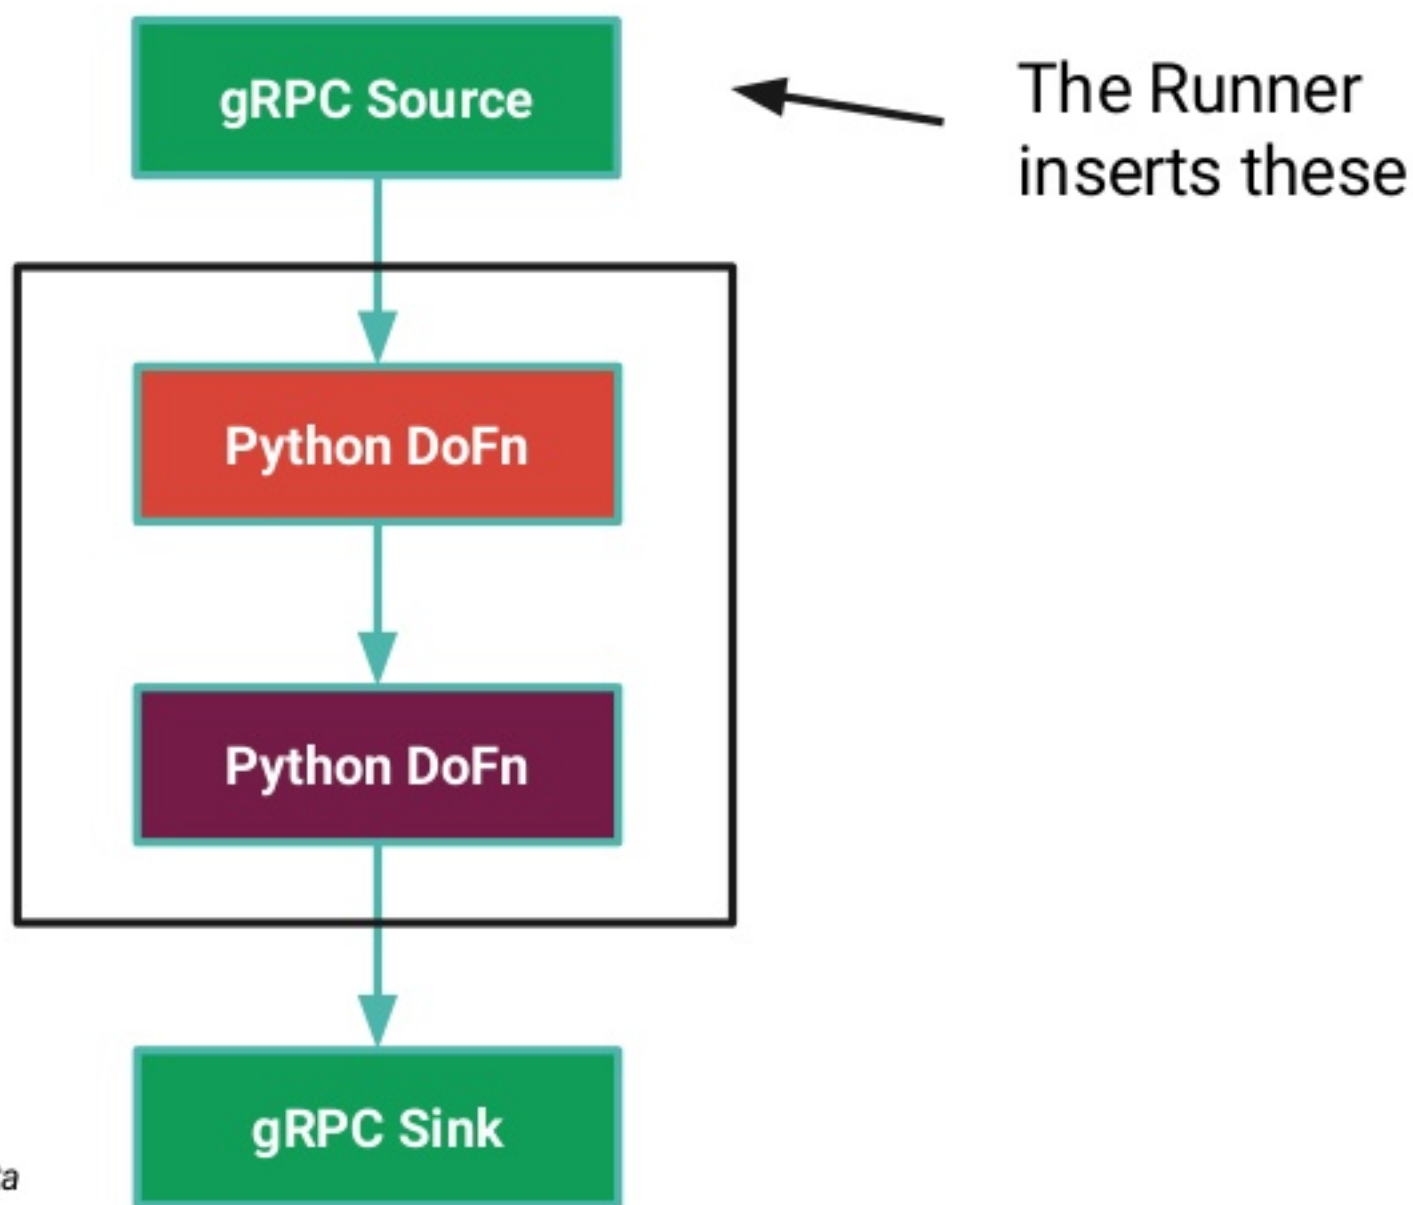
# Fn API - Processing DoFns



# Fn API - Processing DoFns

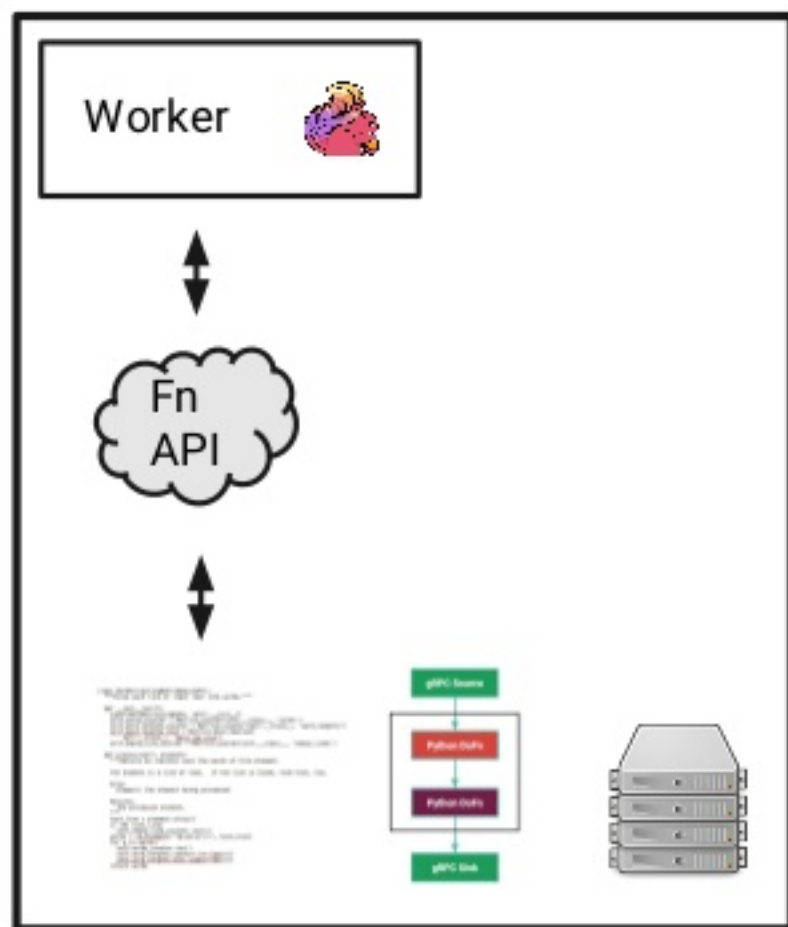


# Fn API - Processing DoFns (Pipeline manipulation)



# Fn API - Executing the user Fn using a SDK Harness

- We can execute as a separate process
- We can execute in a Docker container
- Repository of containers for different SDKs
- We inject the user code into the container when starting
- Container is user-configurable



# Executing Pythonic\* Beam Jobs on Fink

\*or other languages

# What is the (Flink) Runner/Flink doing in all this?

- Analyze/transform the Pipeline (Pipeline API)
- Create a Flink Job (DataSet/DataStream API)
- Ship the user code/docker container description
- In an operator: Open gRPC services for control/data/logging/state plane
- Execute arbitrary user code using the Fn API

**Easy, because Flink state/timers map well to Beam concepts!**



# Advantages/Disadvantages

- Complete isolation of user code
  - Complete configurability of execution environment (with Docker)
  - We can support code written in arbitrary languages
  - We can mix user code written in different languages
- Slower (RPC overhead)
  - Using Docker requires docker 😊

# The Future

# Future work

- Finish what I just talked about
- Finalize the different APIs (not Flink-specific)
- Mixing and matching connectors written in different languages
- Wait for new SDKs in other languages, they will just work 😊

# Learn More!



## Apache Beam/Apache Flink

<https://flink.apache.org> / <https://beam.apache.org>

## Beam Fn API design documents

<https://s.apache.org/beam-runner-api>

<https://s.apache.org/beam-fn-api>

<https://s.apache.org/beam-fn-api-processing-a-bundle>

<https://s.apache.org/beam-fn-state-api-and-bundle-processing>

<https://s.apache.org/beam-fn-api-send-and-receive-data>

<https://s.apache.org/beam-fn-api-container-contract>

## Join the mailing lists!

[user-subscribe@flink.apache.org](mailto:user-subscribe@flink.apache.org) / [dev-subscribe@flink.apache.org](mailto:dev-subscribe@flink.apache.org)

[user-subscribe@beam.apache.org](mailto:user-subscribe@beam.apache.org) / [dev-subscribe@beam.apache.org](mailto:dev-subscribe@beam.apache.org)

## Follow @ApacheFlink / @ApacheBeam on Twitter

Thank you!

# Backup Slides

# Processing Time vs. Event Time

