



Efficient Distributed R Dataframes on Apache Flink

Andreas Kuntz, Jens Meiners, Tilmann Rabl, Volker Markl



- R got huge traction
 - Open source
 - Rich support for analytics & statistics
- But, standalone not well suited for out of core data loads
- Multiple extensions for distributed execution
 - Hadoop + R
 - Spark + R
 - SystemML

Our Goals

- 1 Provide API with natural feeling
 - `df <- select(df, f = df$flights, df$distance)`
 - `df$km <- df$miles * 1.6`
 - `df <- apply(df, key = id, aggFunc)`
- 2 Achieve comparable performance as native dataflow system

General Approach

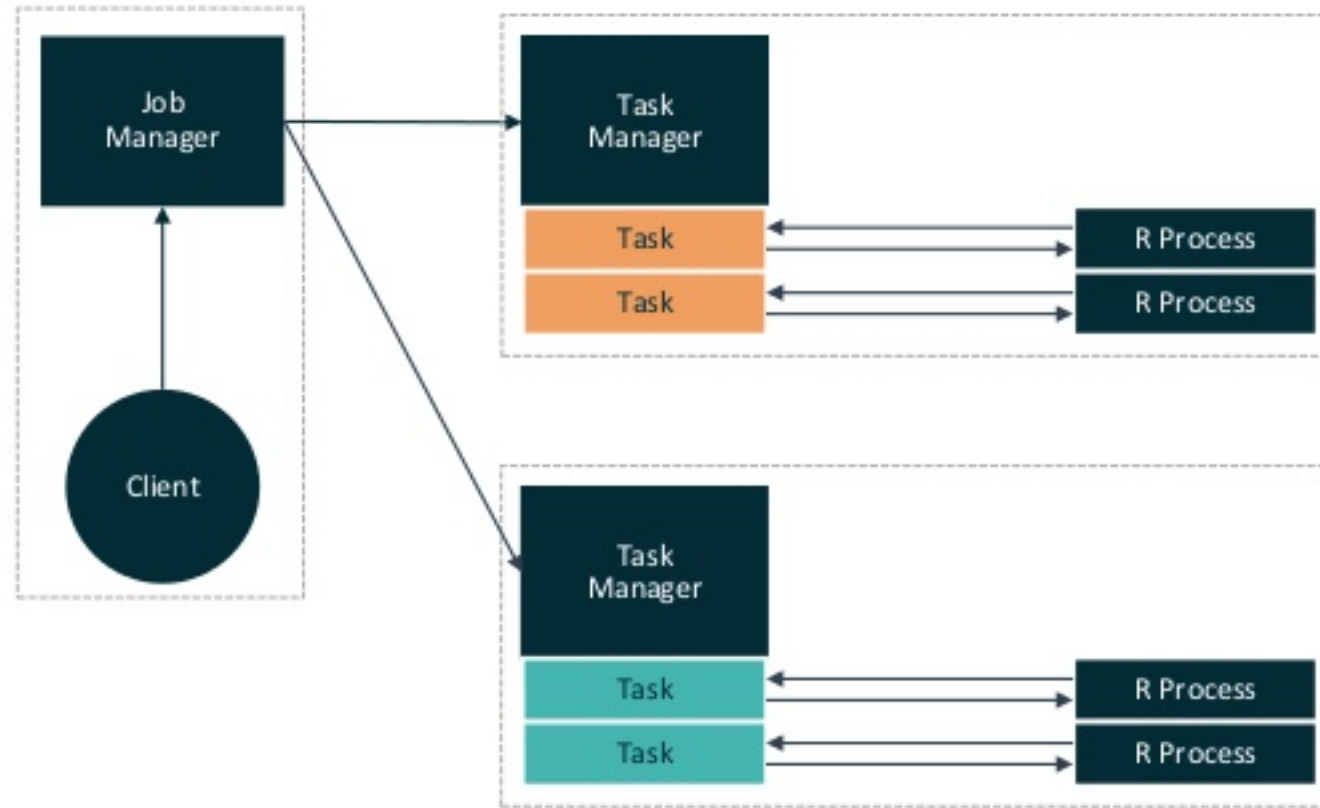
- R dataframe(T_1, T_2, \dots, T_N) as $\text{DataSet}\langle \text{TupleN}\langle T_1, T_2, \dots, T_N \rangle \rangle$
- Create execution plan
 - Map R dataframe functions to the native API whenever possible
e.g., select to projections
 - Call **user defined R functions** within the worker nodes

General Approach


- R dataframe(T_1, T_2, \dots, T_N) as $\text{DataSet}\langle \text{TupleN}\langle T_1, T_2, \dots, T_N \rangle \rangle$
- Create execution plan
 - Map R dataframe functions to the native API whenever possible
e.g., select to projections
- Call **user defined R functions** within the worker nodes

Handling user defined R functions

Inter Process Communication



Inter Process Communication



- 1 Communication + Serialization
- 2 Java and R compete for memory

Source-to-Source Translation

- Translate restrict set of operations to native dataflow API
- Operations are executed natively

```
df <- filter(  
  df,  
  df$language == 'R'  
)
```



```
val df = df.filter($"language" === "R")
```


```
df$km <- df$miles * 1.6
```




```
val df = df.withColumn("km", $"miles" * 1.6)
```

Flink + fastR

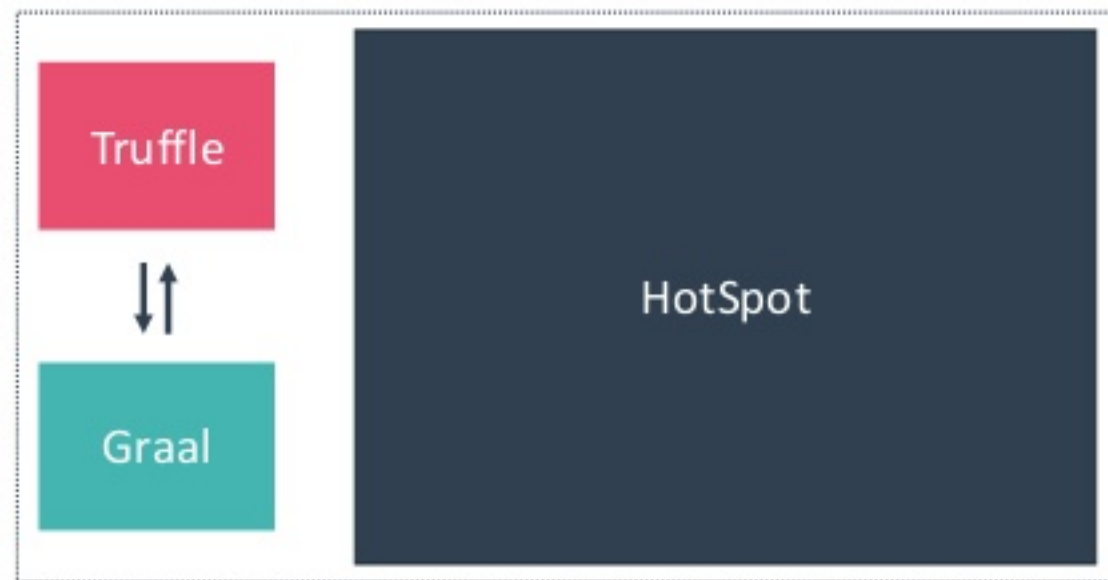
Truffle/Graal



Truffle/Graal

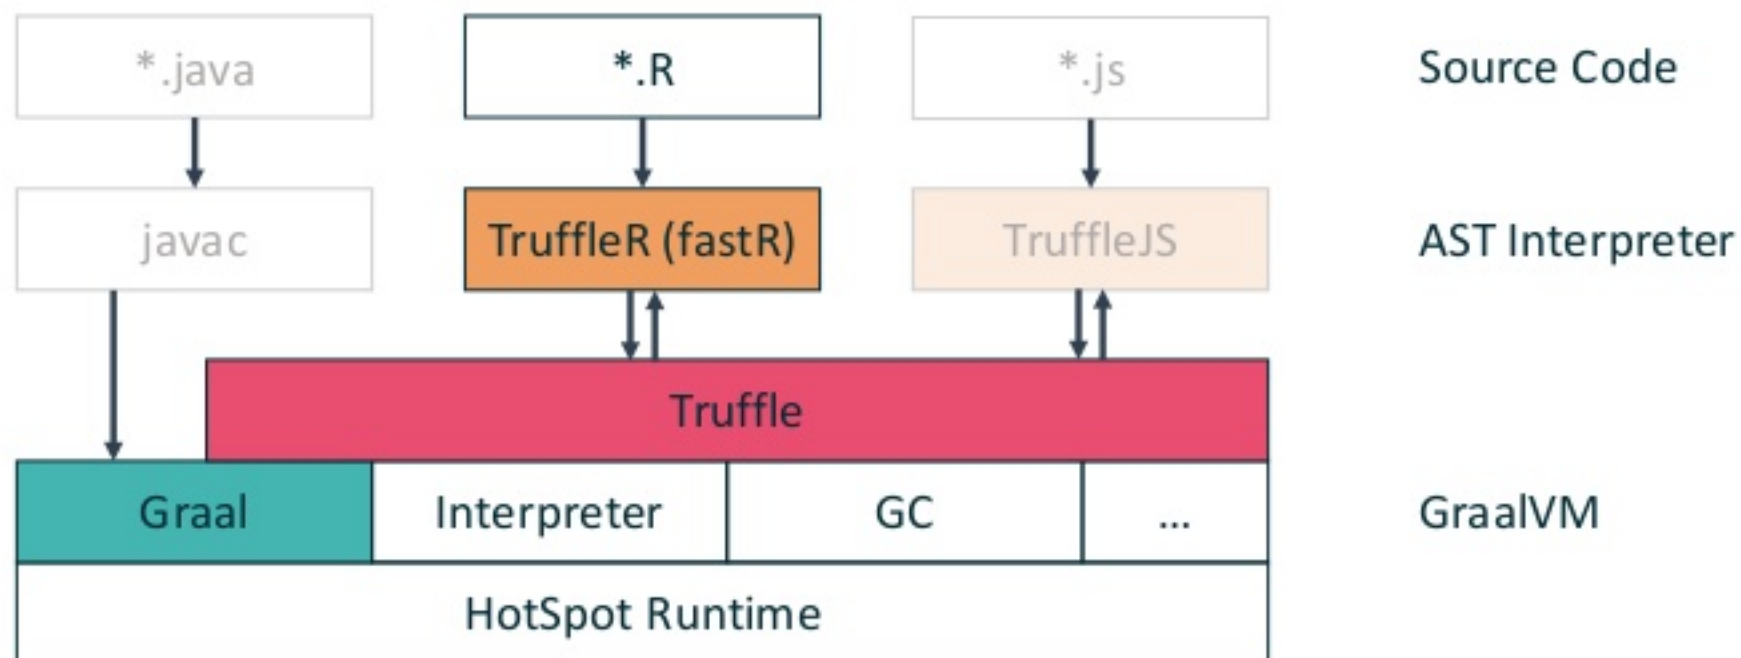


Truffle/Graal



GraalVM

Truffle/Graal



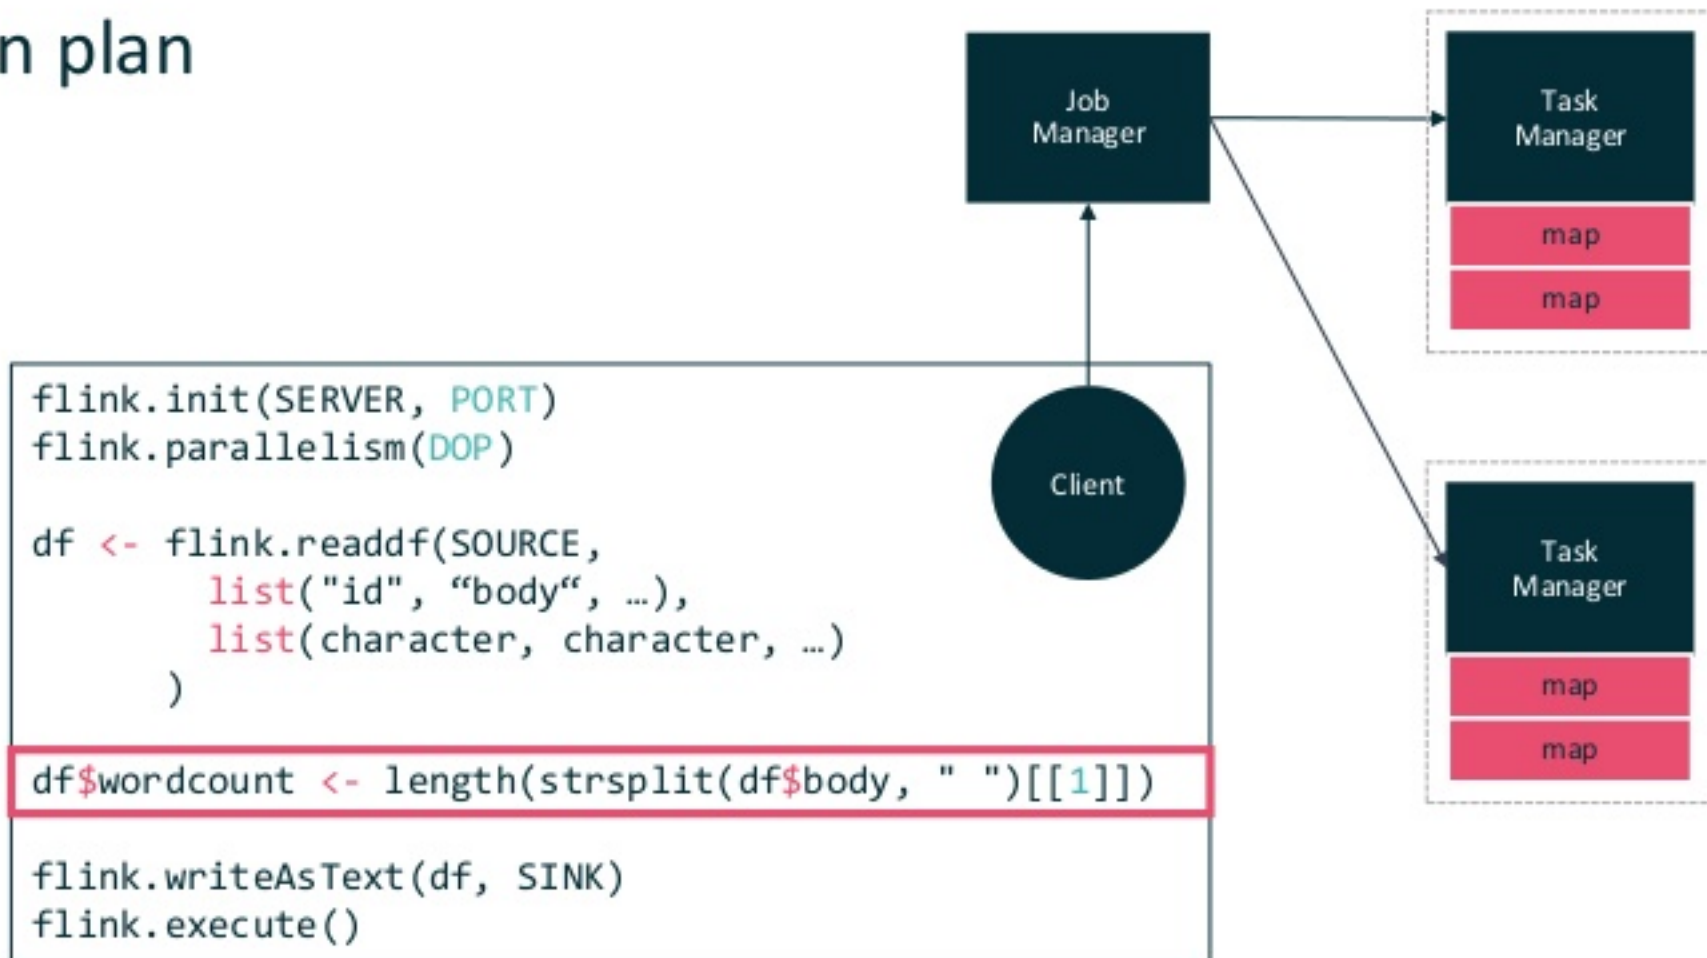
Flink + fastR

fastR: R implementation on top of Truffle/Graal

- Allows us to execute R code in the same VM as Flink
- Infer result types of R functions
- Access Java (Flink) data types in R

Client:

1. Dataframe rows to Flink tuples
2. Determine return types of UDFs
3. Create execution plan



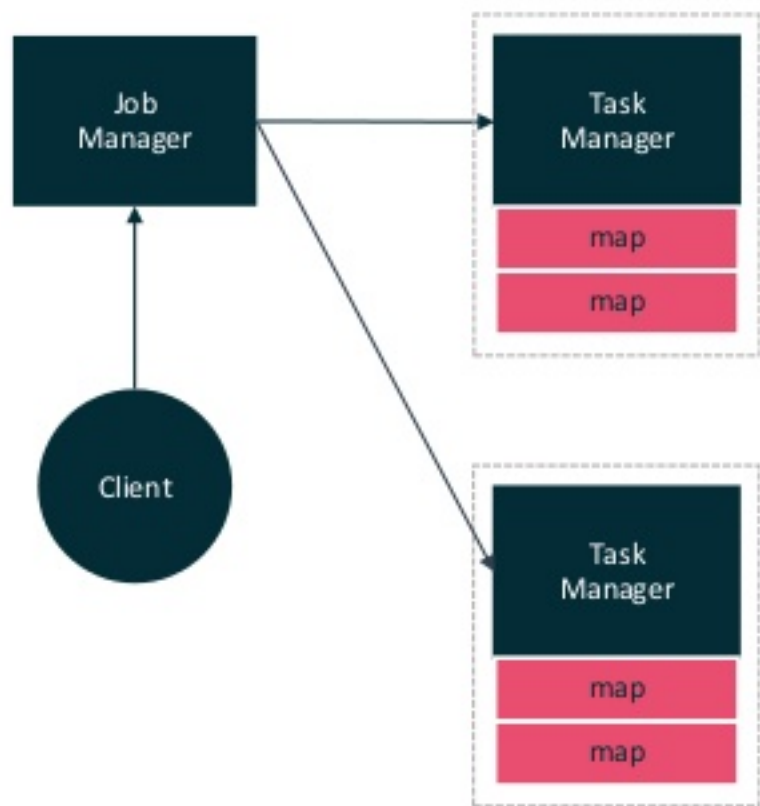

```
df$wordcount <- length(strsplit(df$body, " ")[[1]])
```



```
function(tuple) {  
  .fun <- function(tuple) { length(strsplit(tuple[[2]], " ")[[1]]) }  
  flink.tuple(tuple[[1]], tuple[[2]], .fun(tuple))  
}  
  3 2
```

- 1 Dataframe proxy keeps track of columns, provides efficient access
- 2 Can be extended with new columns
- 3 Rewrite to directly use Flink tuples


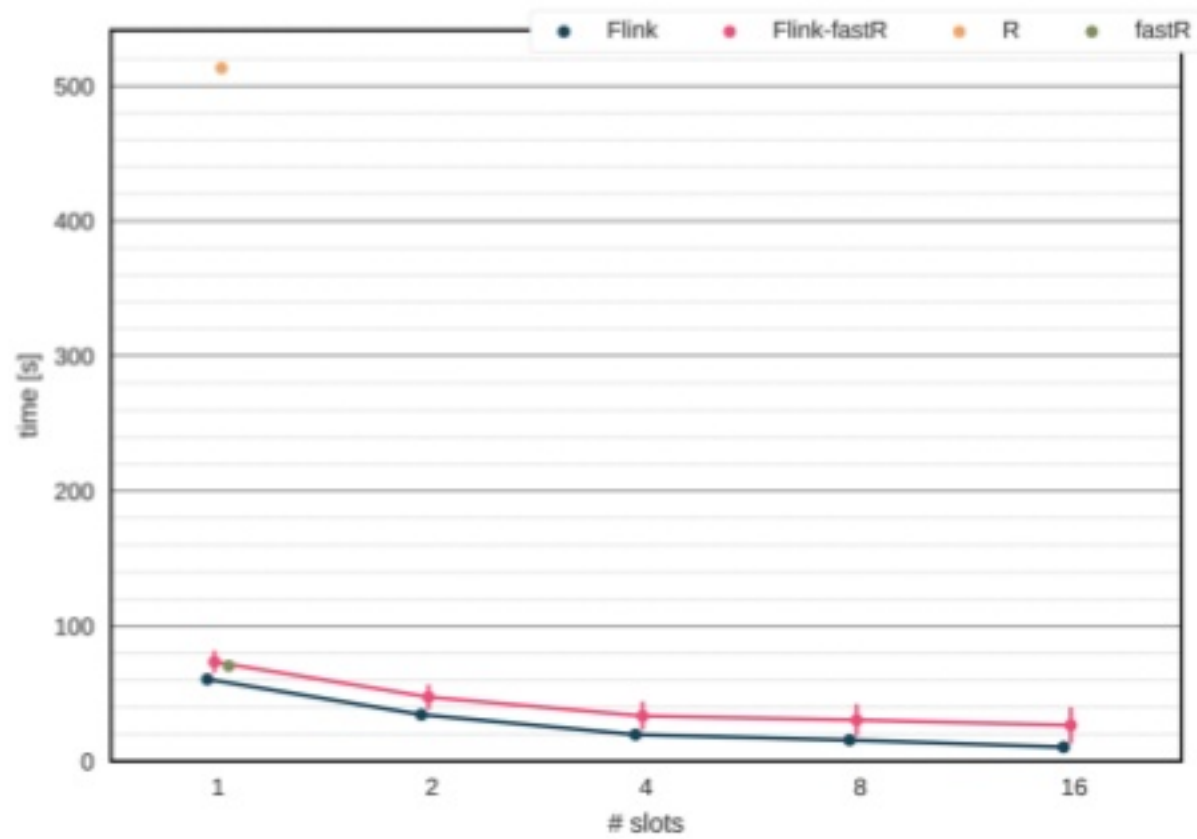
- Task Manager:
Evaluate R UDF & Execute





```
map { tuple =>  
    executeRFunction(func, tuple)  
}
```

```
map { tuple =>  
    executeRFunction(func, tuple)  
}
```


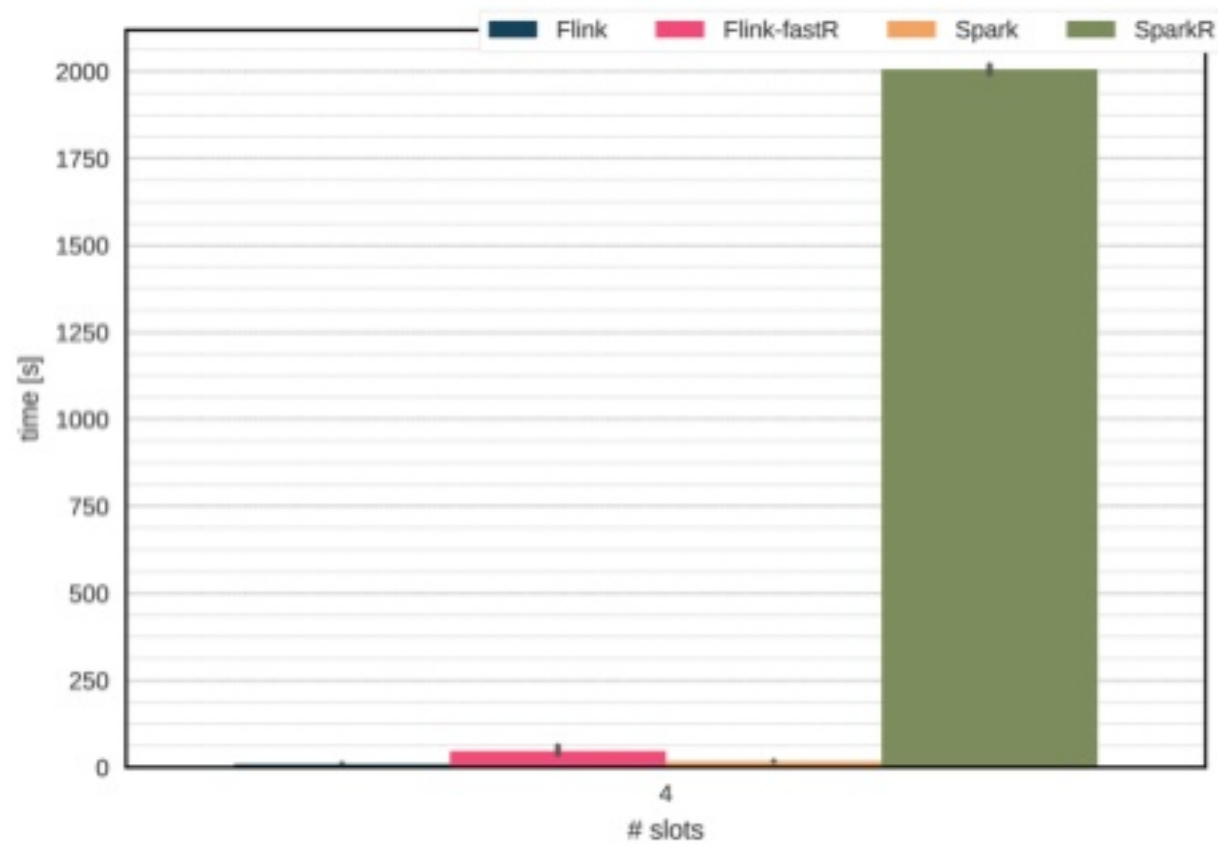
Local - 1.4GB




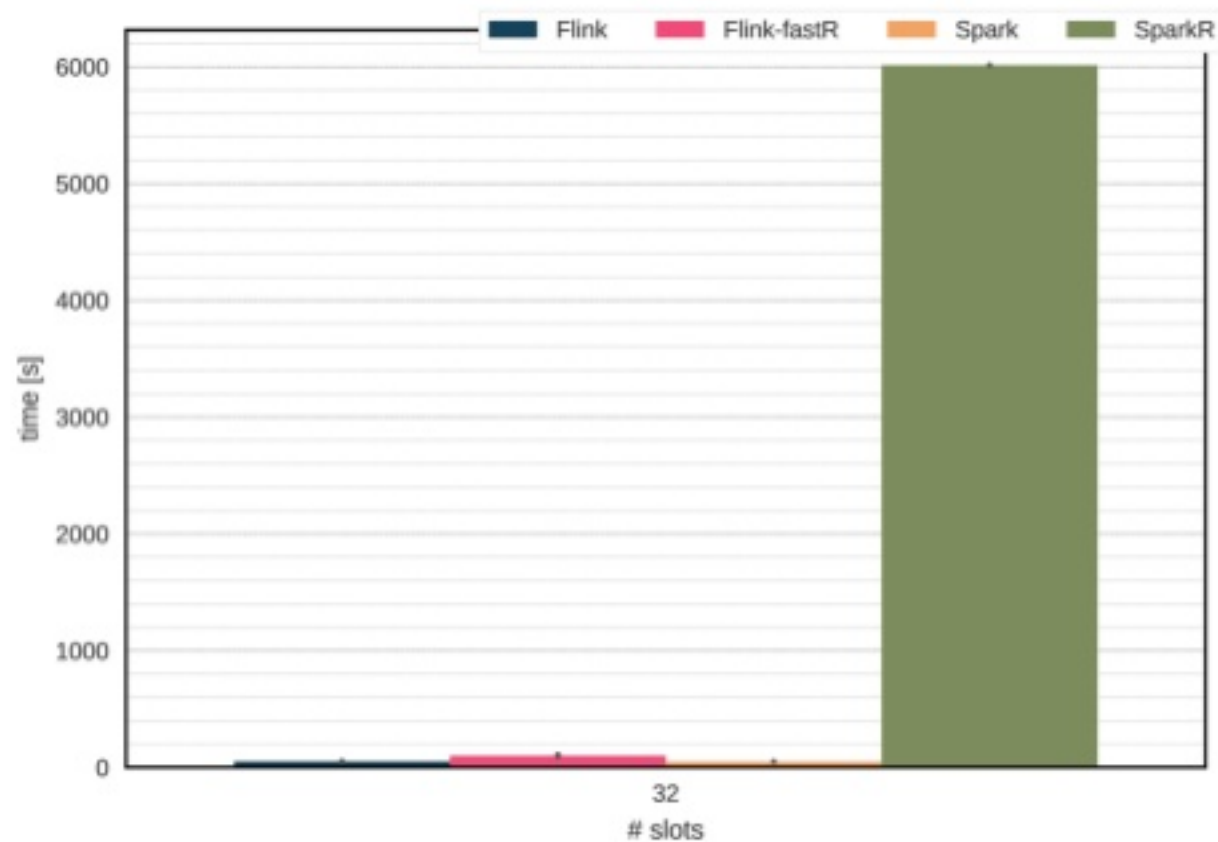
Local - 14GB



Local – 1GB



Cluster – 10GB



fastR + Flink

- R dataframe abstraction for distributed computation
- Performance gains even on single node (local mode)
- Approaches native performance even for R UDFs
- Interesting opportunities for:
 - Streaming
 - Other dynamic languages
 - Dynamic Re-optimization

Thank you for your attention!