

King

Building & Operating Large-scale Streaming Applications

Gyula Fóra
gyula.fora@king.com

Background

About King

We make awesome mobile games

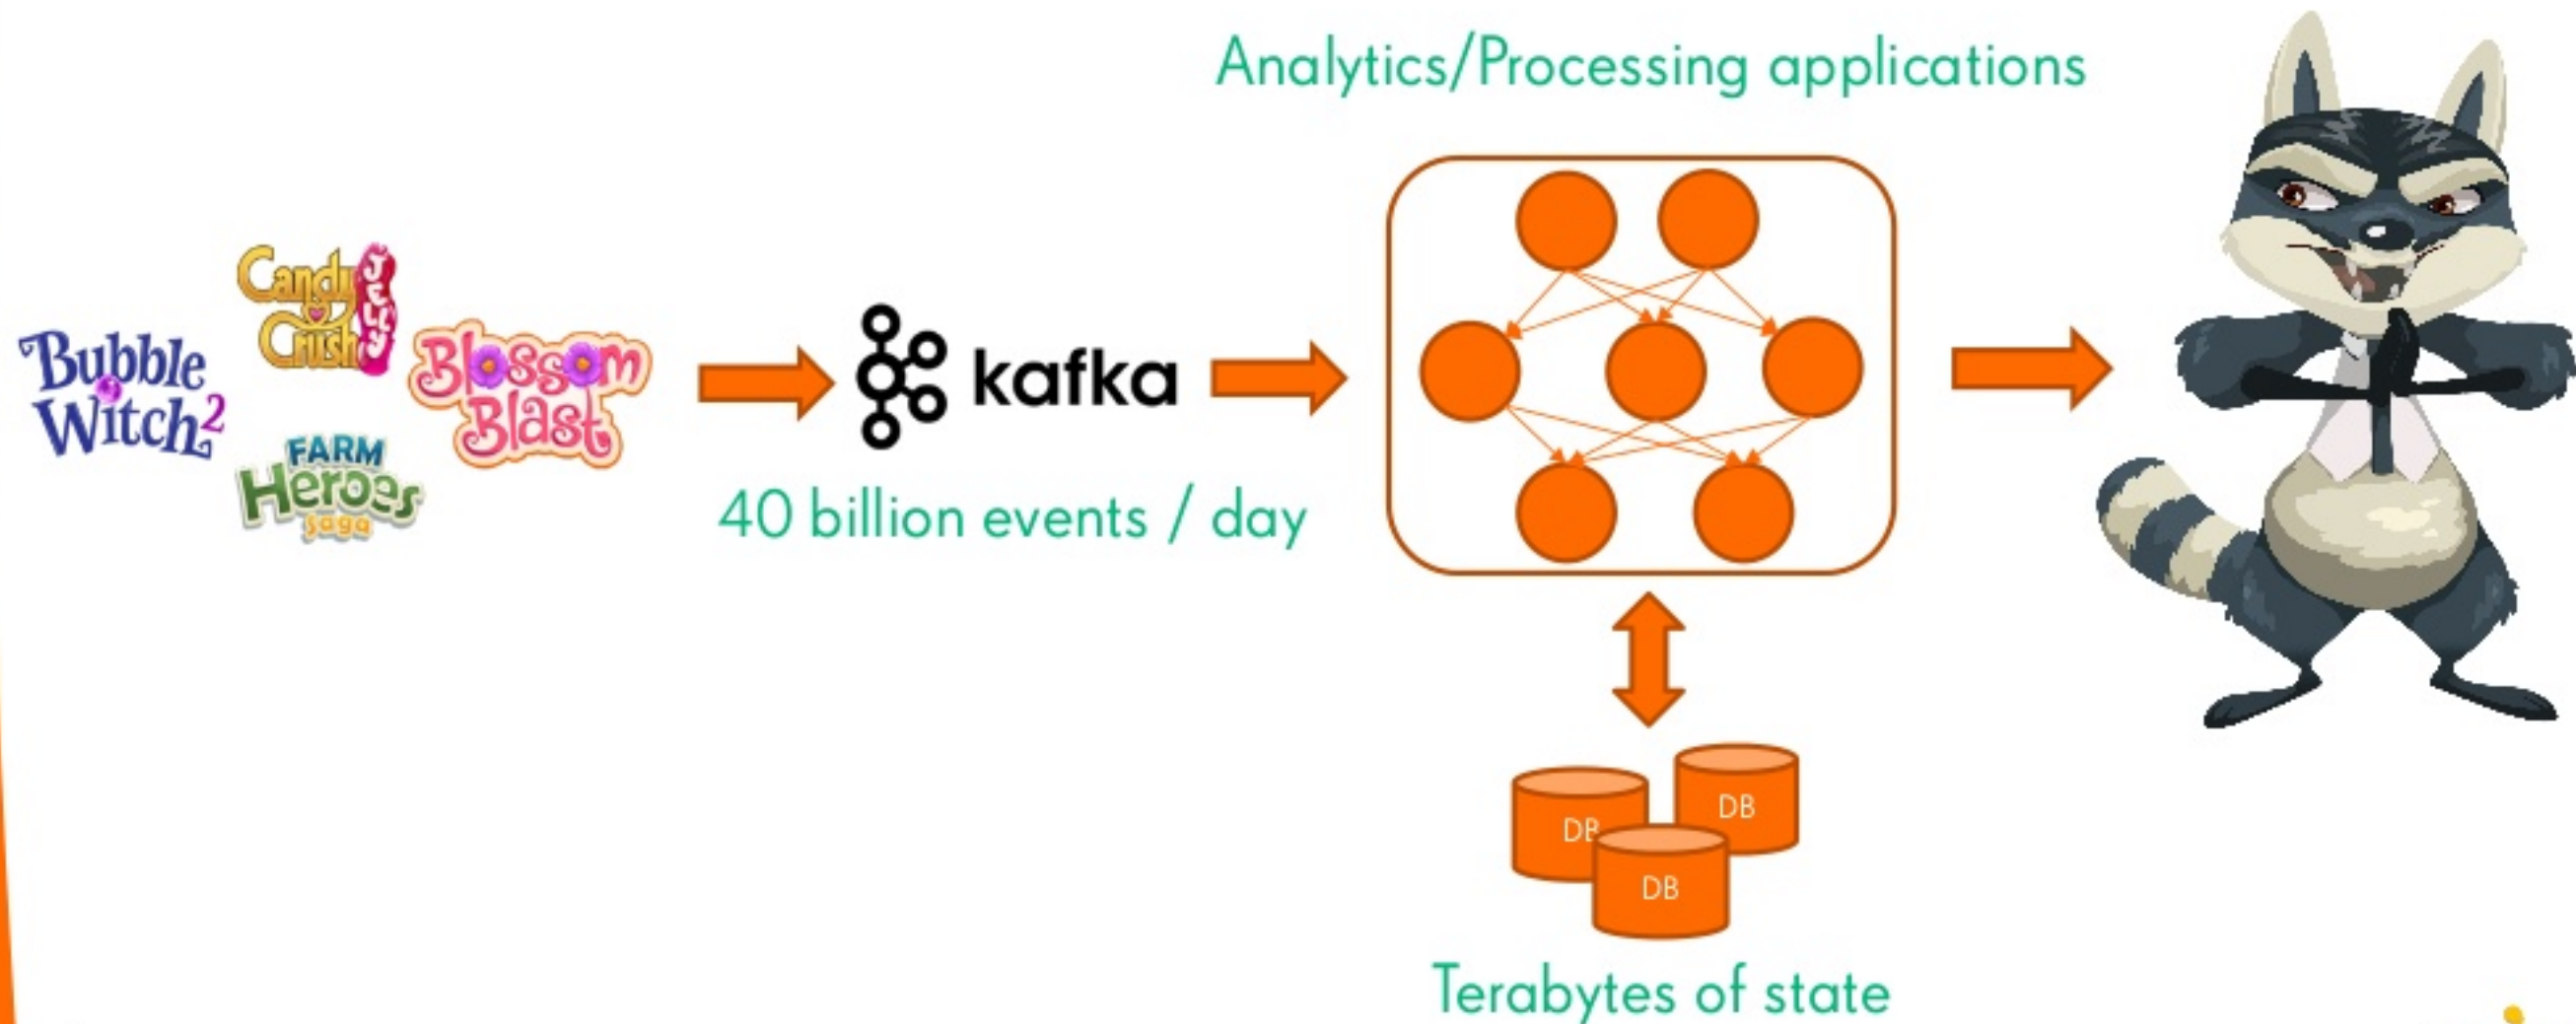
Hundreds of millions of players

40 billion events per day

And a lot of data...



From a streaming perspective...



Stream processing at King



Real-time dashboards



Kafka Consumers



King Streaming SDK & Libraries



Real-time analytics platform

The RBea platform

Powered by Apache Flink

Scripting on the live streams

Window aggregates, Timers, Sessions

Complex stateful computations

Scalable + fault tolerant



RBea in production

"Stable" since summer 2016

Running for 20+ games

Live/QA environments

250+ live scripts

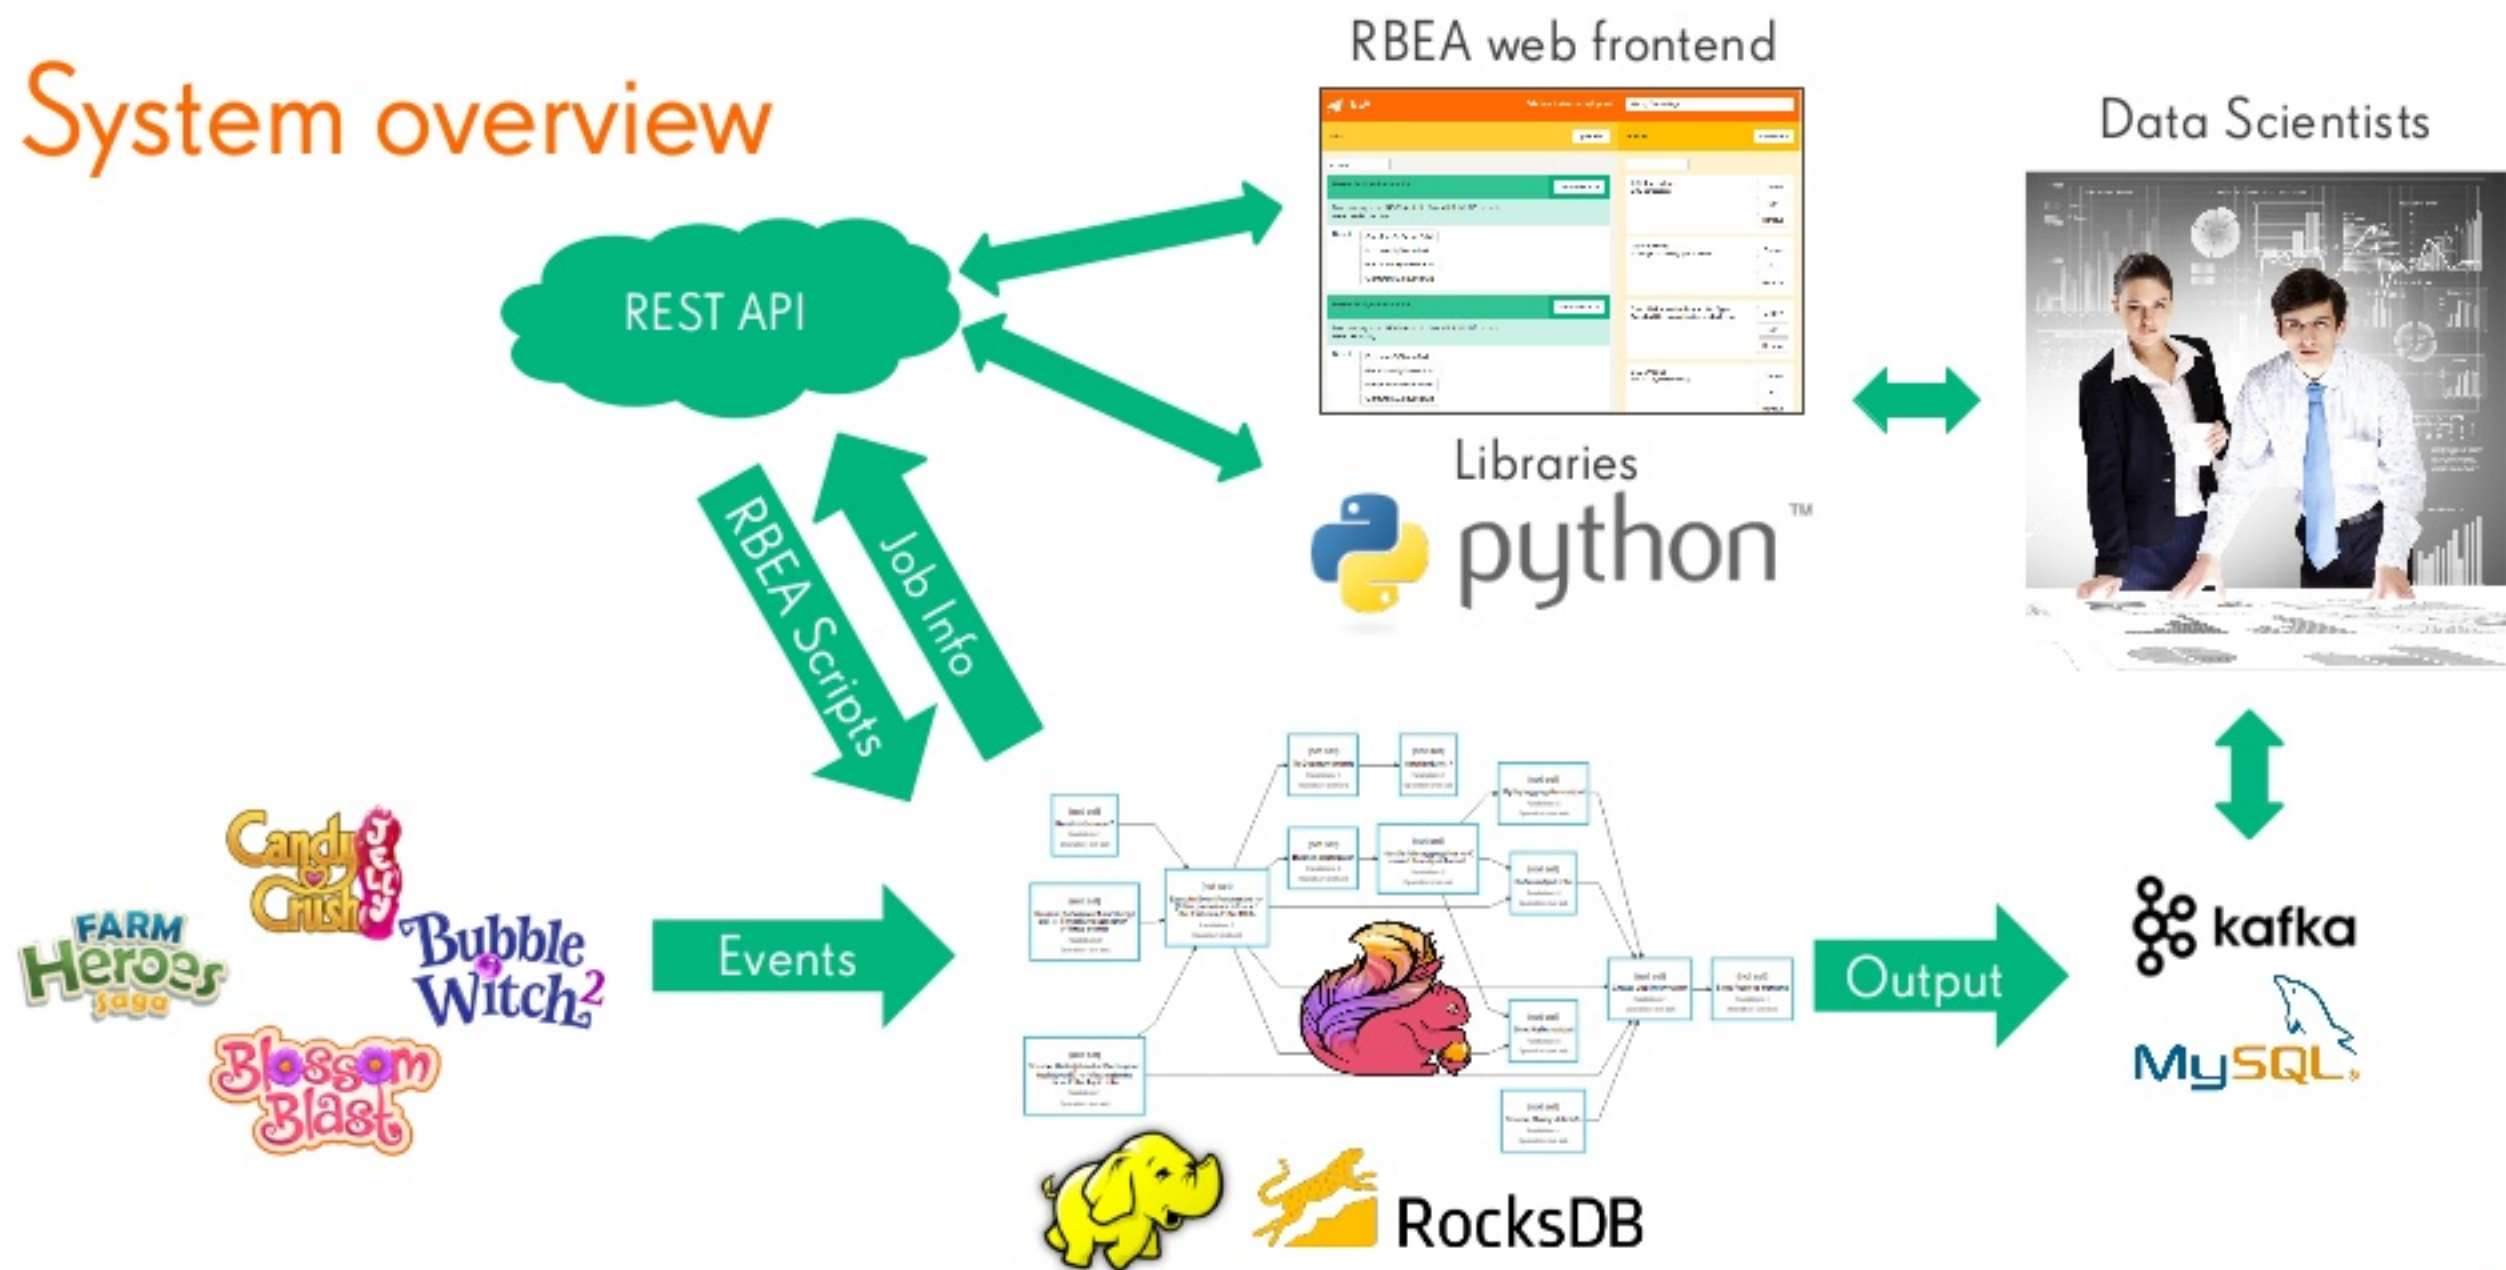
5+ TBs of user state (and growing)



Building a streaming platform

A.K.A the coolest Flink program I have ever written...

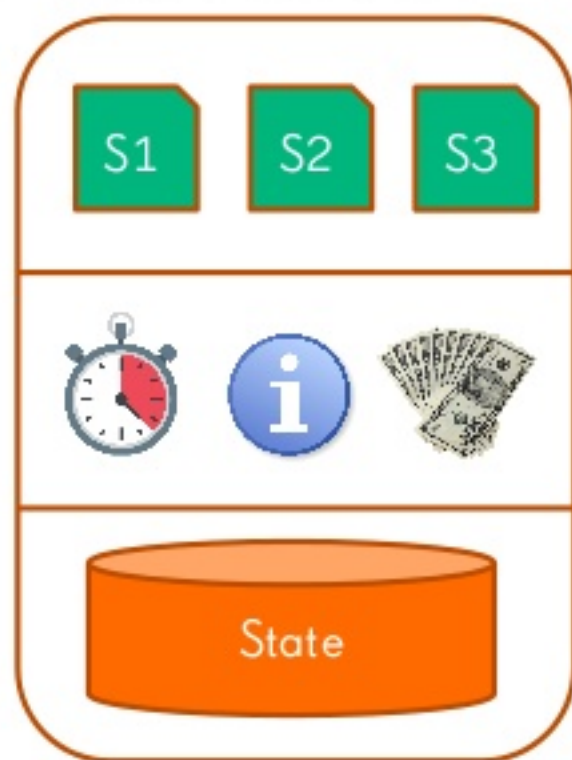

System overview




RBea backend architecture

Partitioned Event Stream

```
for(script: deployedScripts)  
  script.process(nextEvent)
```



Output based
on API calls



Broadcasted control stream

Dynamic Time Windows

User states and state backends


Roughly 1 billion keys

< 10 GB

< 100 GB

< 1-2 TB

< ?? TB




- Too much memory

- Poor read performance
- Operationally complex

- Poor recovery speed
- Large checkpoints hurt the cluster

- Poor recovery speed

"Incremental" improvements




Disk IO



Network traffic



Window aggregates and dashboards



Problems

Connection issues

Throughput problems

Duplicated logic

Sneaky bugs



Window aggregates and dashboards



Pipeline testing

```
testJob = RBEATestPipeline
    .startWithDeployment(1000, new TestProcessor1())
    .thenDeploy(6000, new TestProcessor2())
    .thenEvent(3, "1500")
    .thenRemoveProcessor(1000)
    .thenWatermark(800)
    .thenFailAndRestoreJob()
    .thenEvent(2, "1000")
    .thenWatermark(1600)
    .thenDeploy(2000, new ProcWithFailingTimer());

outputs = runTestPipeline(testJob)

// Do some output verification...
```



Running a streaming platform

Now that we are done with the easy part...

Apache Flink infrastructure



Standalone



16 cores
24G RAM
HDD

x6



hadoop
YARN



x12



hadoop
YARN



x36

Checkpoints



ceph



ceph



hadoop
HDFS

Flink application management

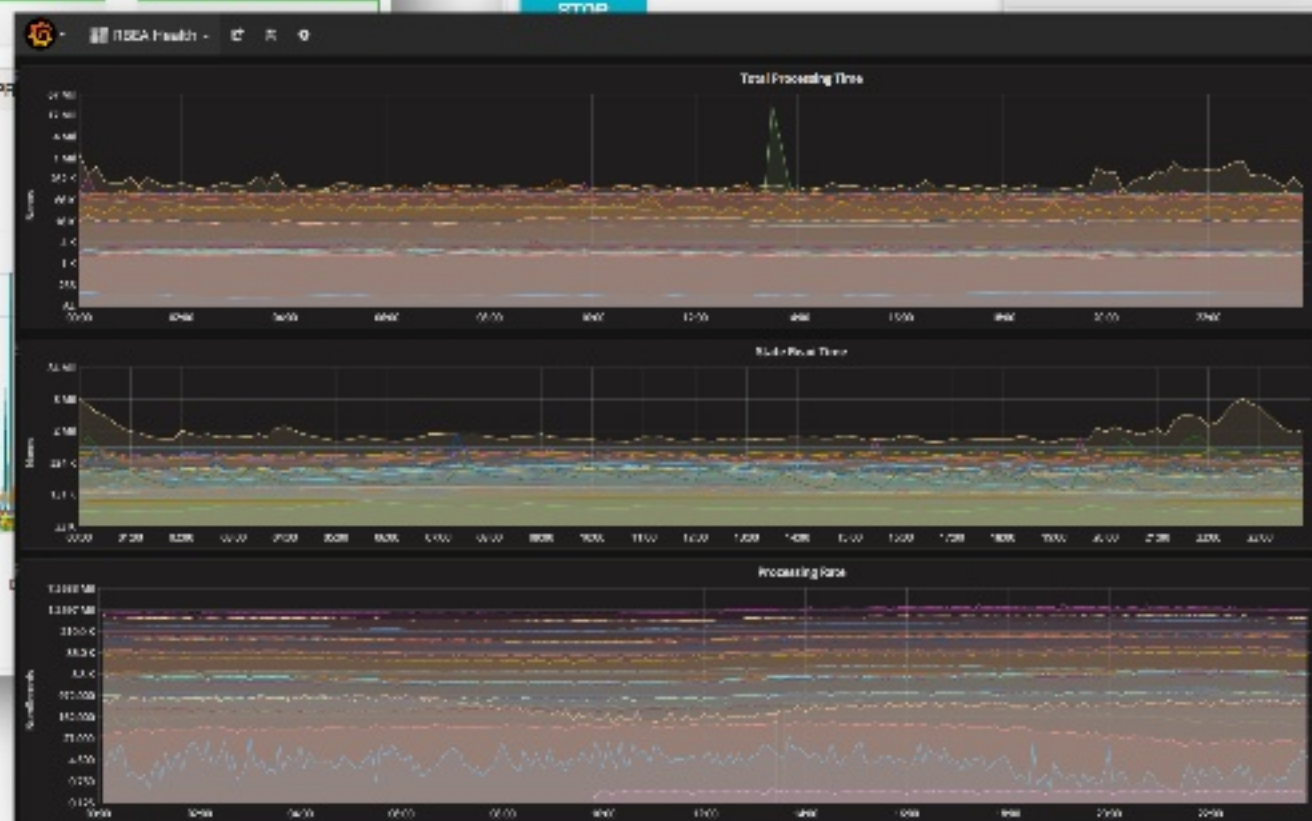
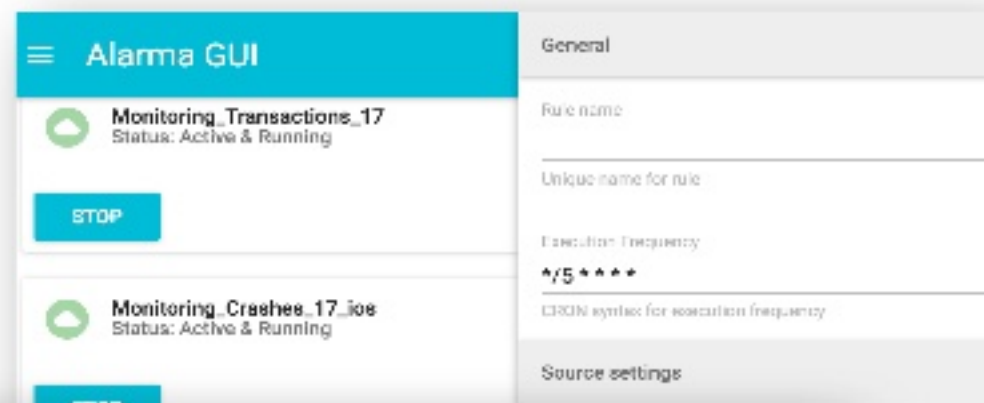
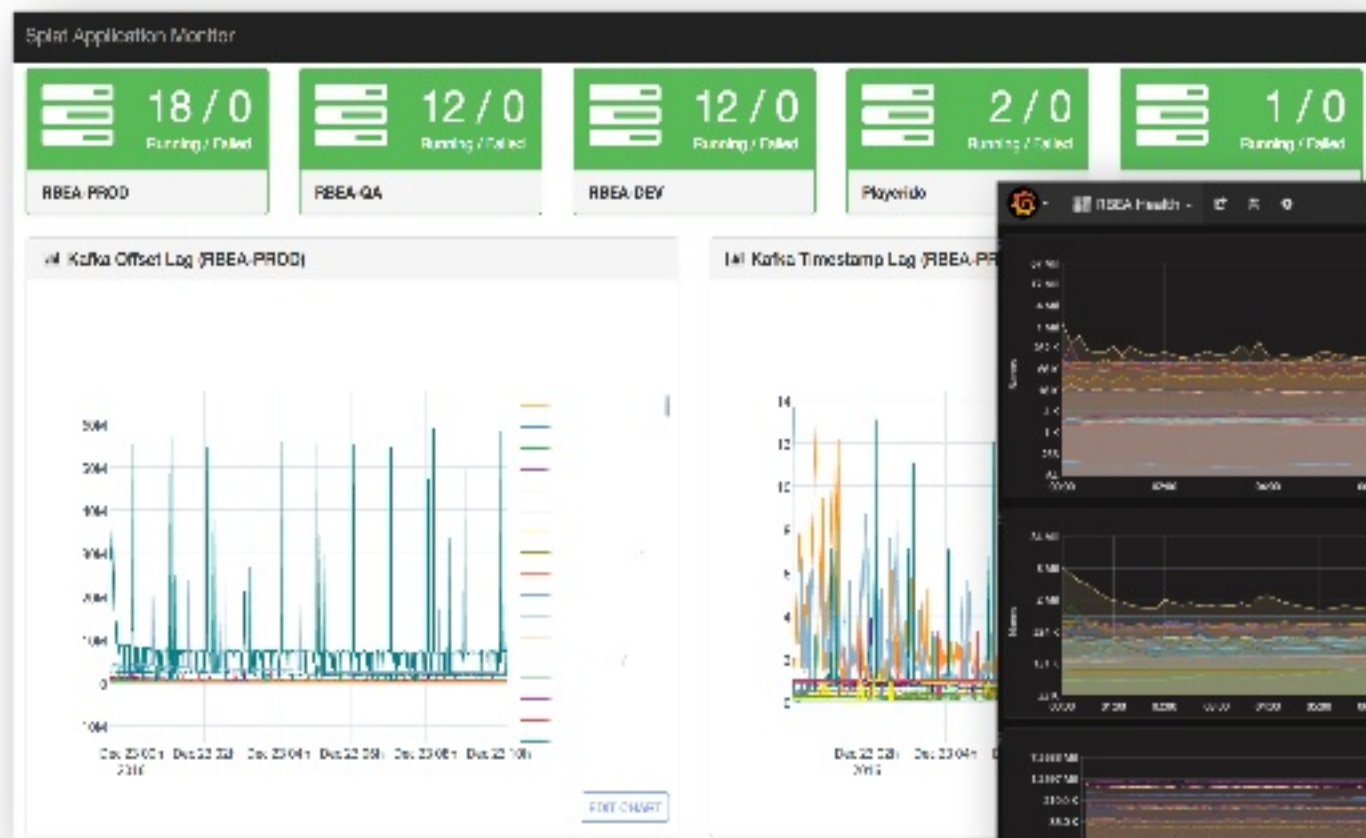
The heavy lifting is done by Flink: Checkpoints, Savepoints, Rescaling

But we have a lot of Python scripts to make it nice

- Lookup and monitor applications (Application IDs)
- Savepoint management (restore from-latest/at, fork)
- Application versioning, automatic fallback to stable
- Continuous deployment (through Jenkins)
- Across several clusters/environments



Metrics, metrics everywhere...



from the mean

Metrics, metrics everywhere...

Use custom Flink metrics to measure all API interactions

- State read/write times, sizes
- Cache statistics
- Script execution stats

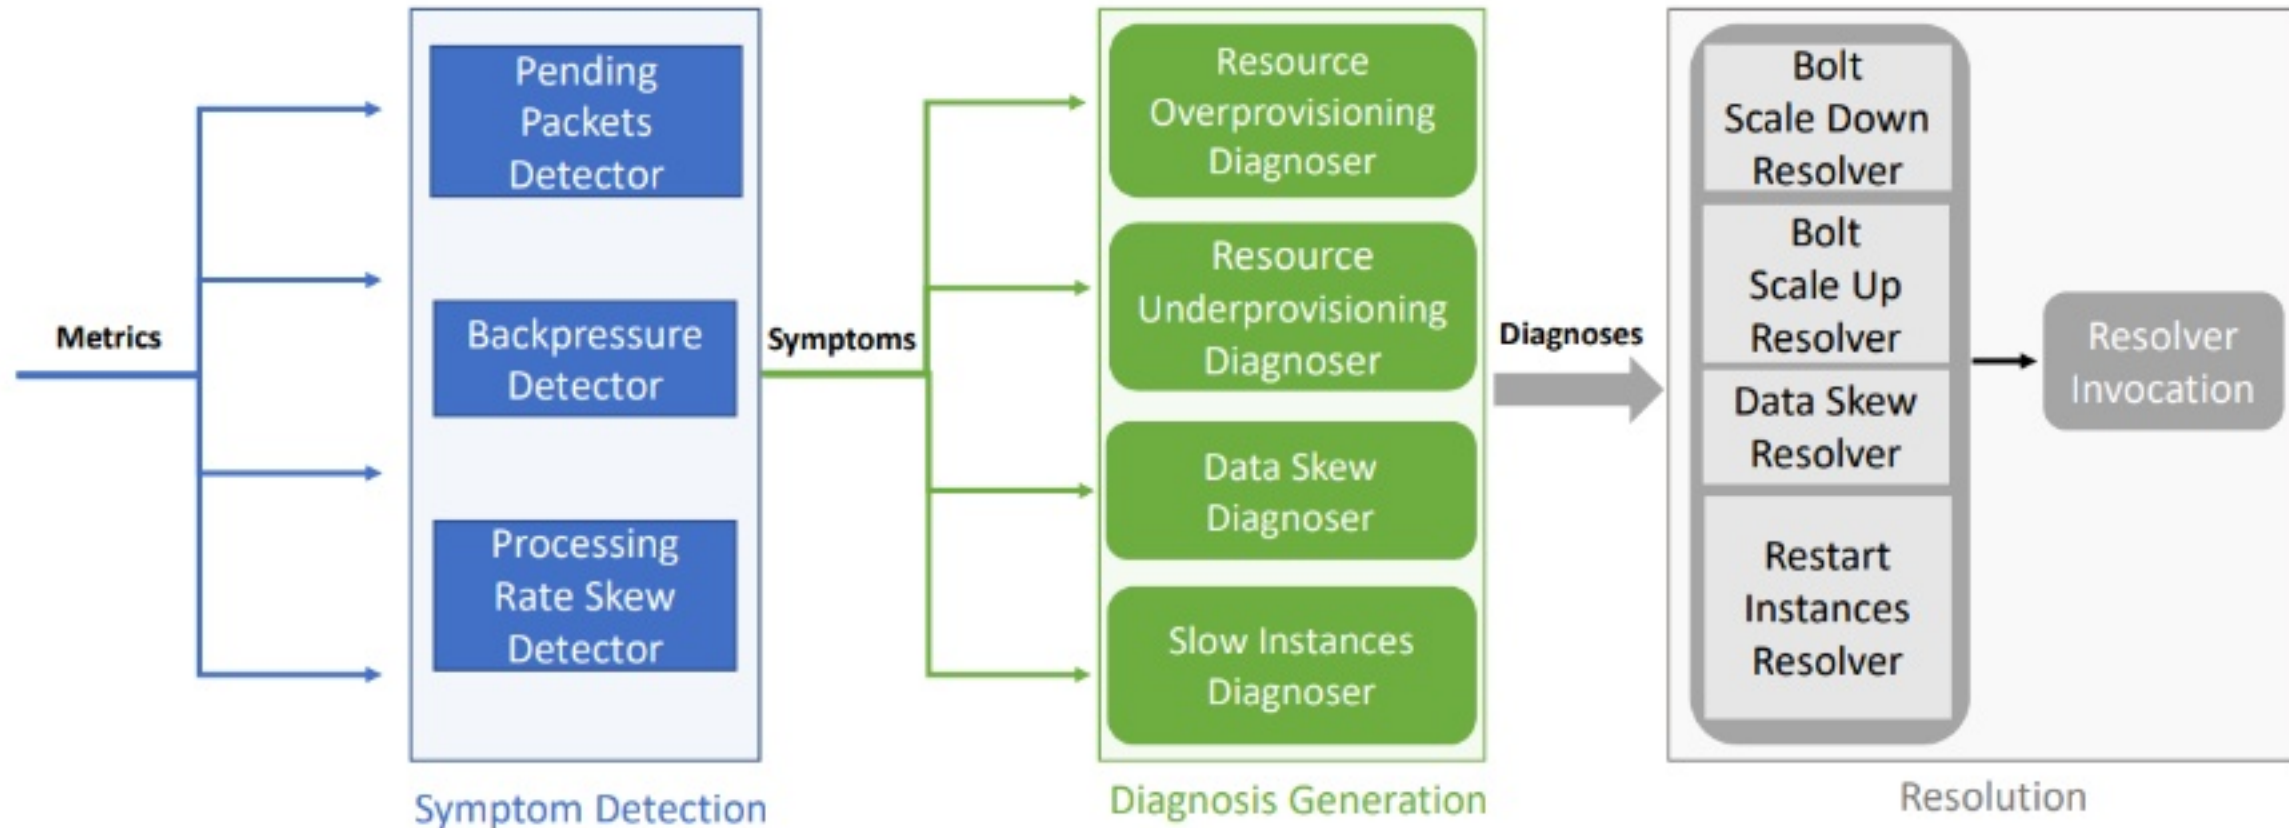
Exponential Moving Average (Gauge) for most metrics

Leverage windowing mechanism for dynamic stats

We also do a lot of CPU profiling 😊



Microsoft Dhalion



<https://www.microsoft.com/en-us/research/project/dhalion/>
<https://github.com/Microsoft/Dhalion>

Join the Streaming Platform Team!

We are looking for passionate developers to help us transform the gaming industry!

Talk to us at the conference or search all available roles at <https://jobs.king.com/>

Thank you!