

A MATERIALIZATION ENGINE FOR DATA INTEGRATION WITH FLINK



MIHAIL VIERU

13-09-2017



AGENDA

- Microservices Architecture
- Data Integration Challenge
- Materialization Engine
- Flink Backend
- Stream Compaction
- Advantages over Legacy Approach

ABOUT ME

Mihail Vieru

Big Data Engineer
Team "Flux" Stream Processing
Data Engineering Dept.



Europe's leading online fashion platform

15 countries

~21 million active customers

~3.6 billion € revenue 2016

250,000+ products

2,000 brands

13,000+ employees in Europe



WE ARE CONSTANTLY INNOVATING TECHNOLOGY

**HOME-BREWED,
CUTTING-EDGE
& SCALABLE**

technology solutions



help our brand to
WIN ONLINE



~ 1,800

employees from



77

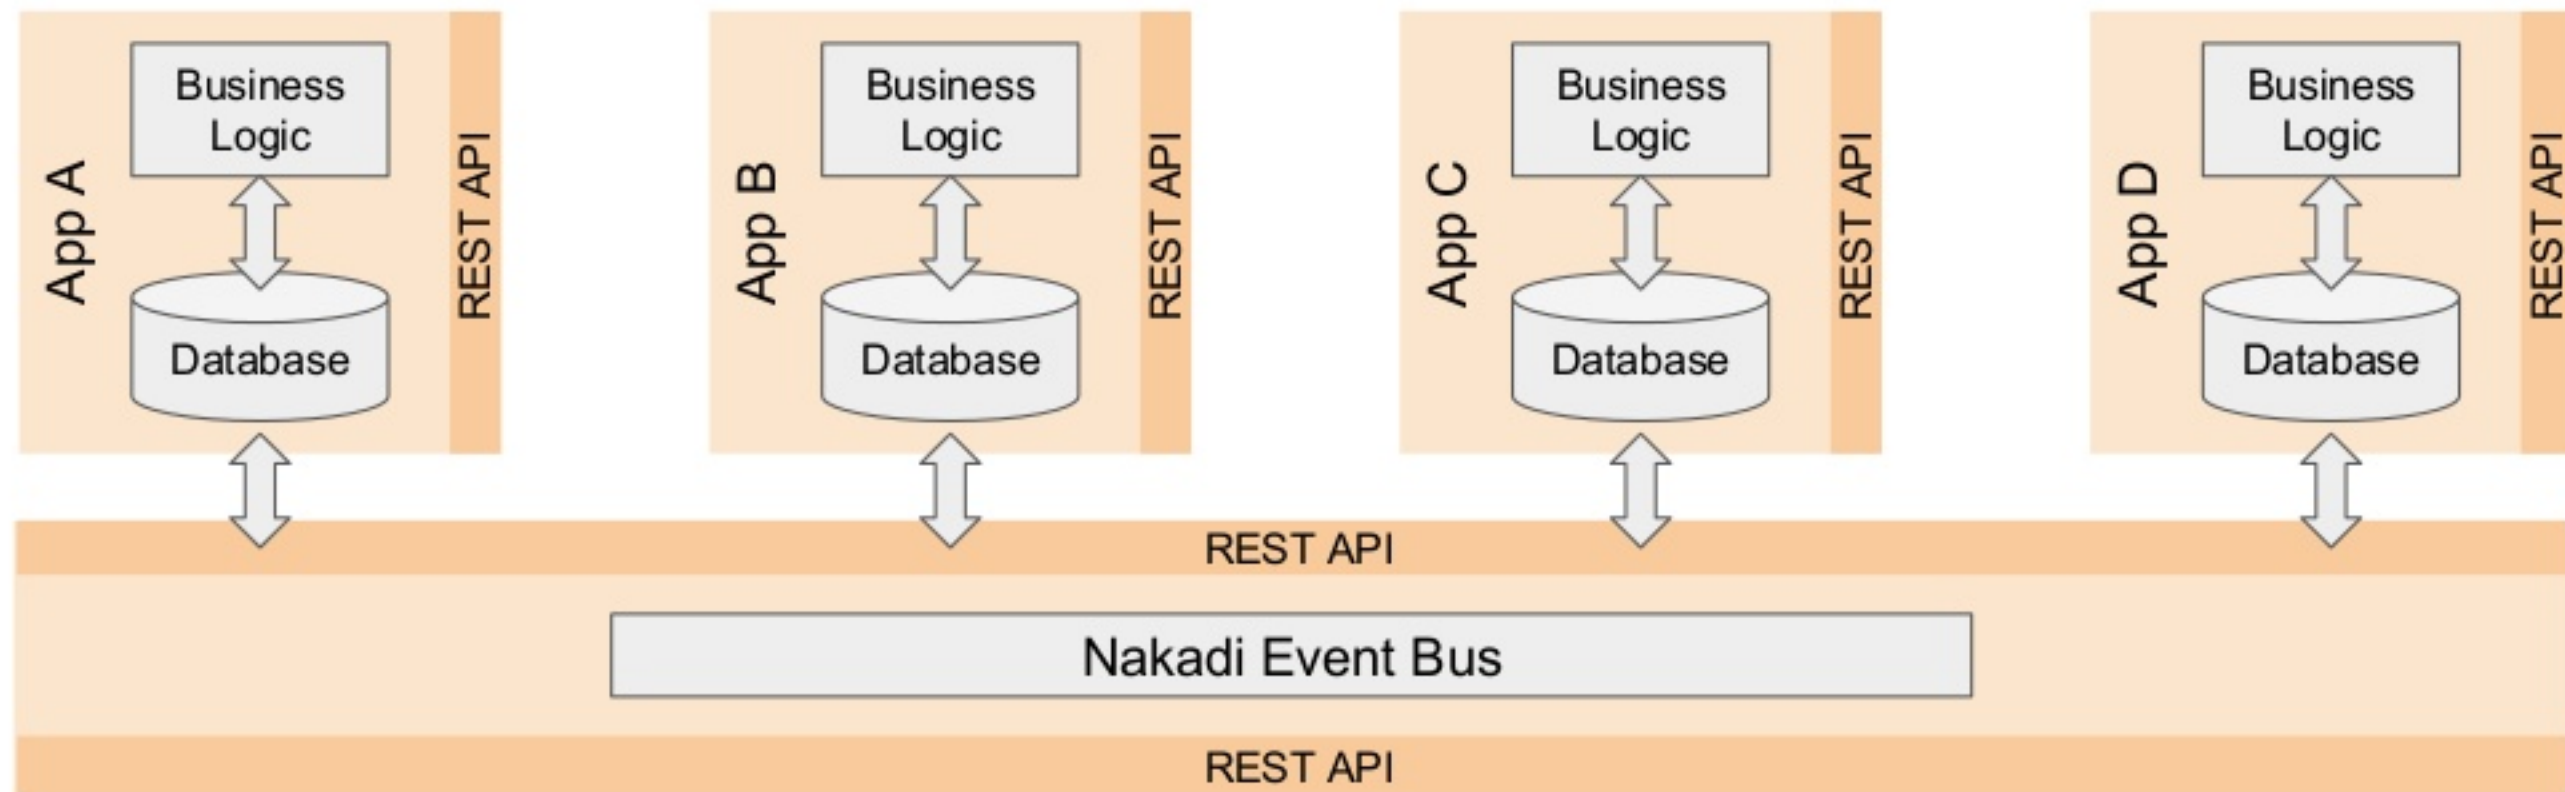
nations



6 tech locations
+ HQs in Berlin

MICROSERVICES ARCHITECTURE

MICROSERVICES ARCHITECTURE



Everything runs on Amazon Web Services

NAKADI - CENTRAL EVENT BUS

A distributed event bus that implements a RESTful API abstraction over Kafka-like queues.

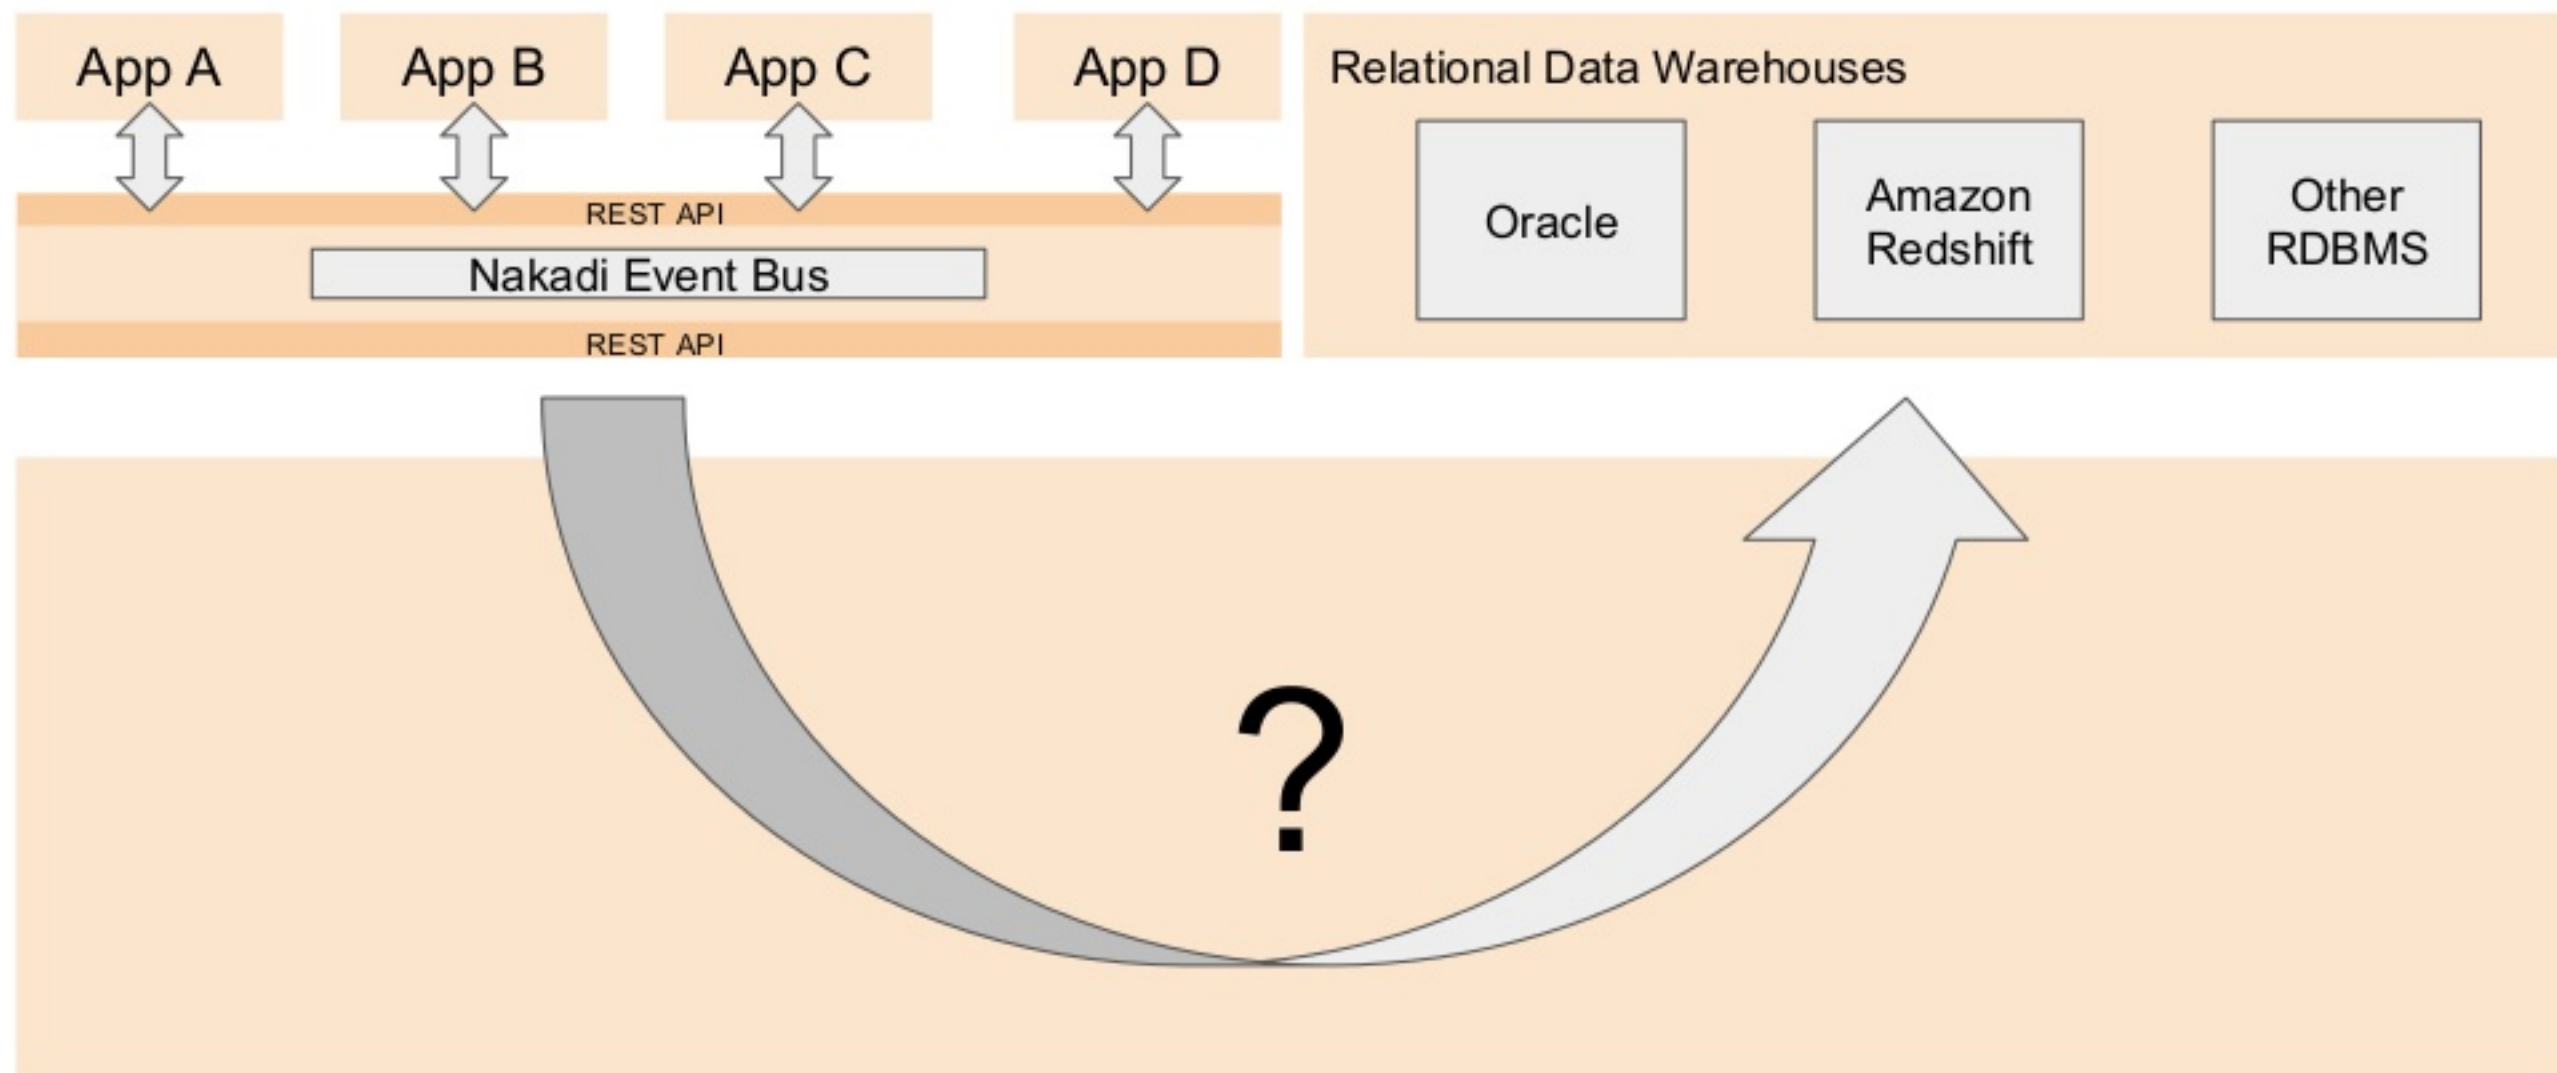
- 1000s of Kafka topics
- high variance in event structure, size & throughput
- consumer can specify partition offsets and batch size in the GET request

<https://github.com/zalando/nakadi>



DATA INTEGRATION CHALLENGE

DATA INTEGRATION CHALLENGE



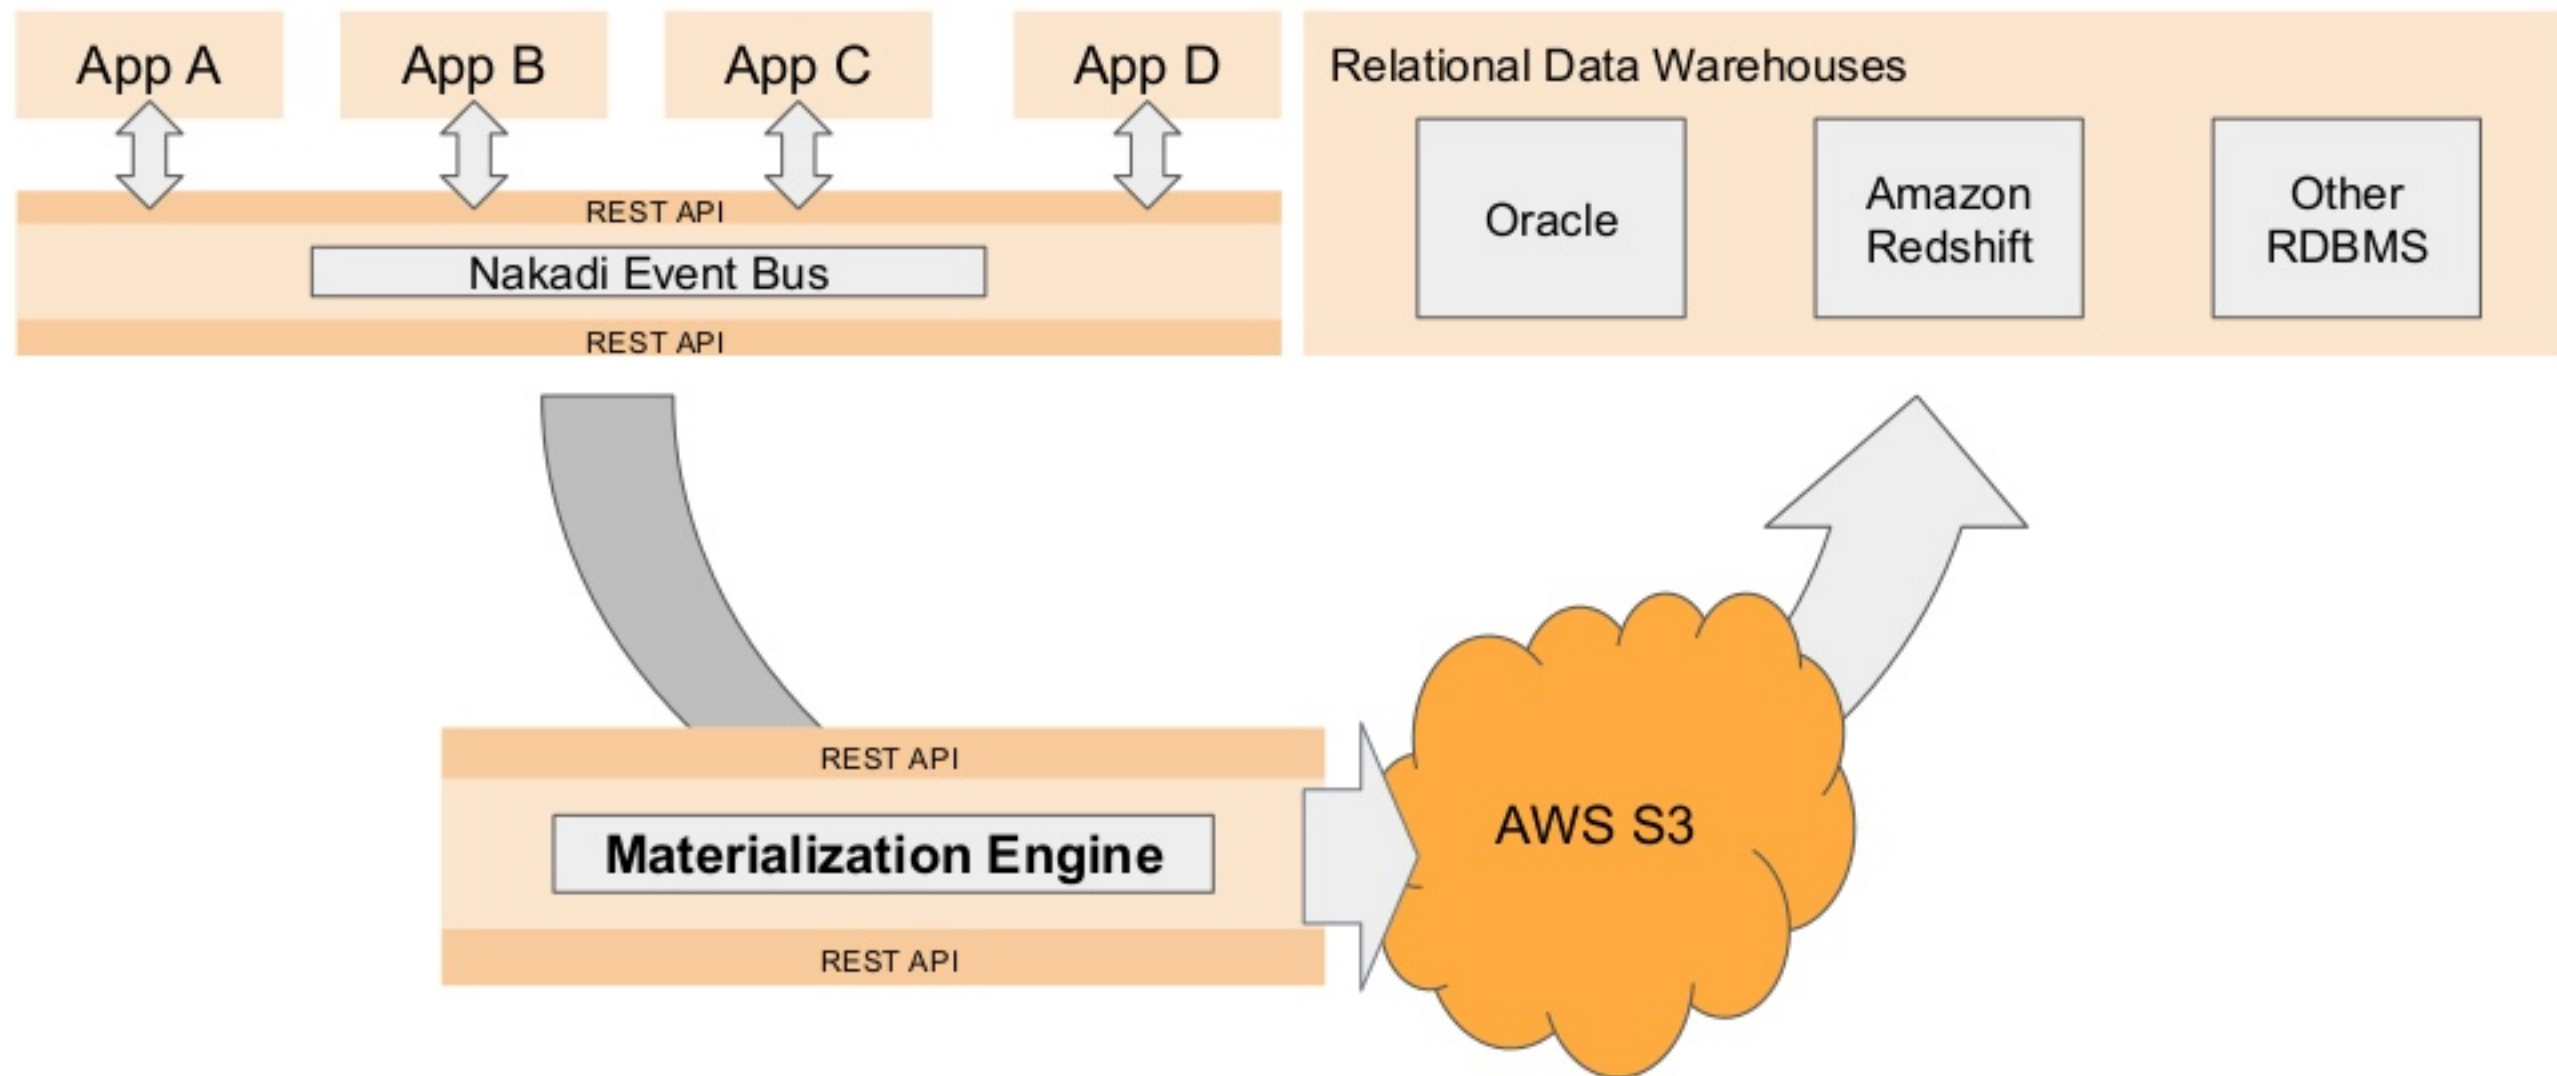
REQUIREMENTS

Goal: Consume data streams in a relational database friendly way

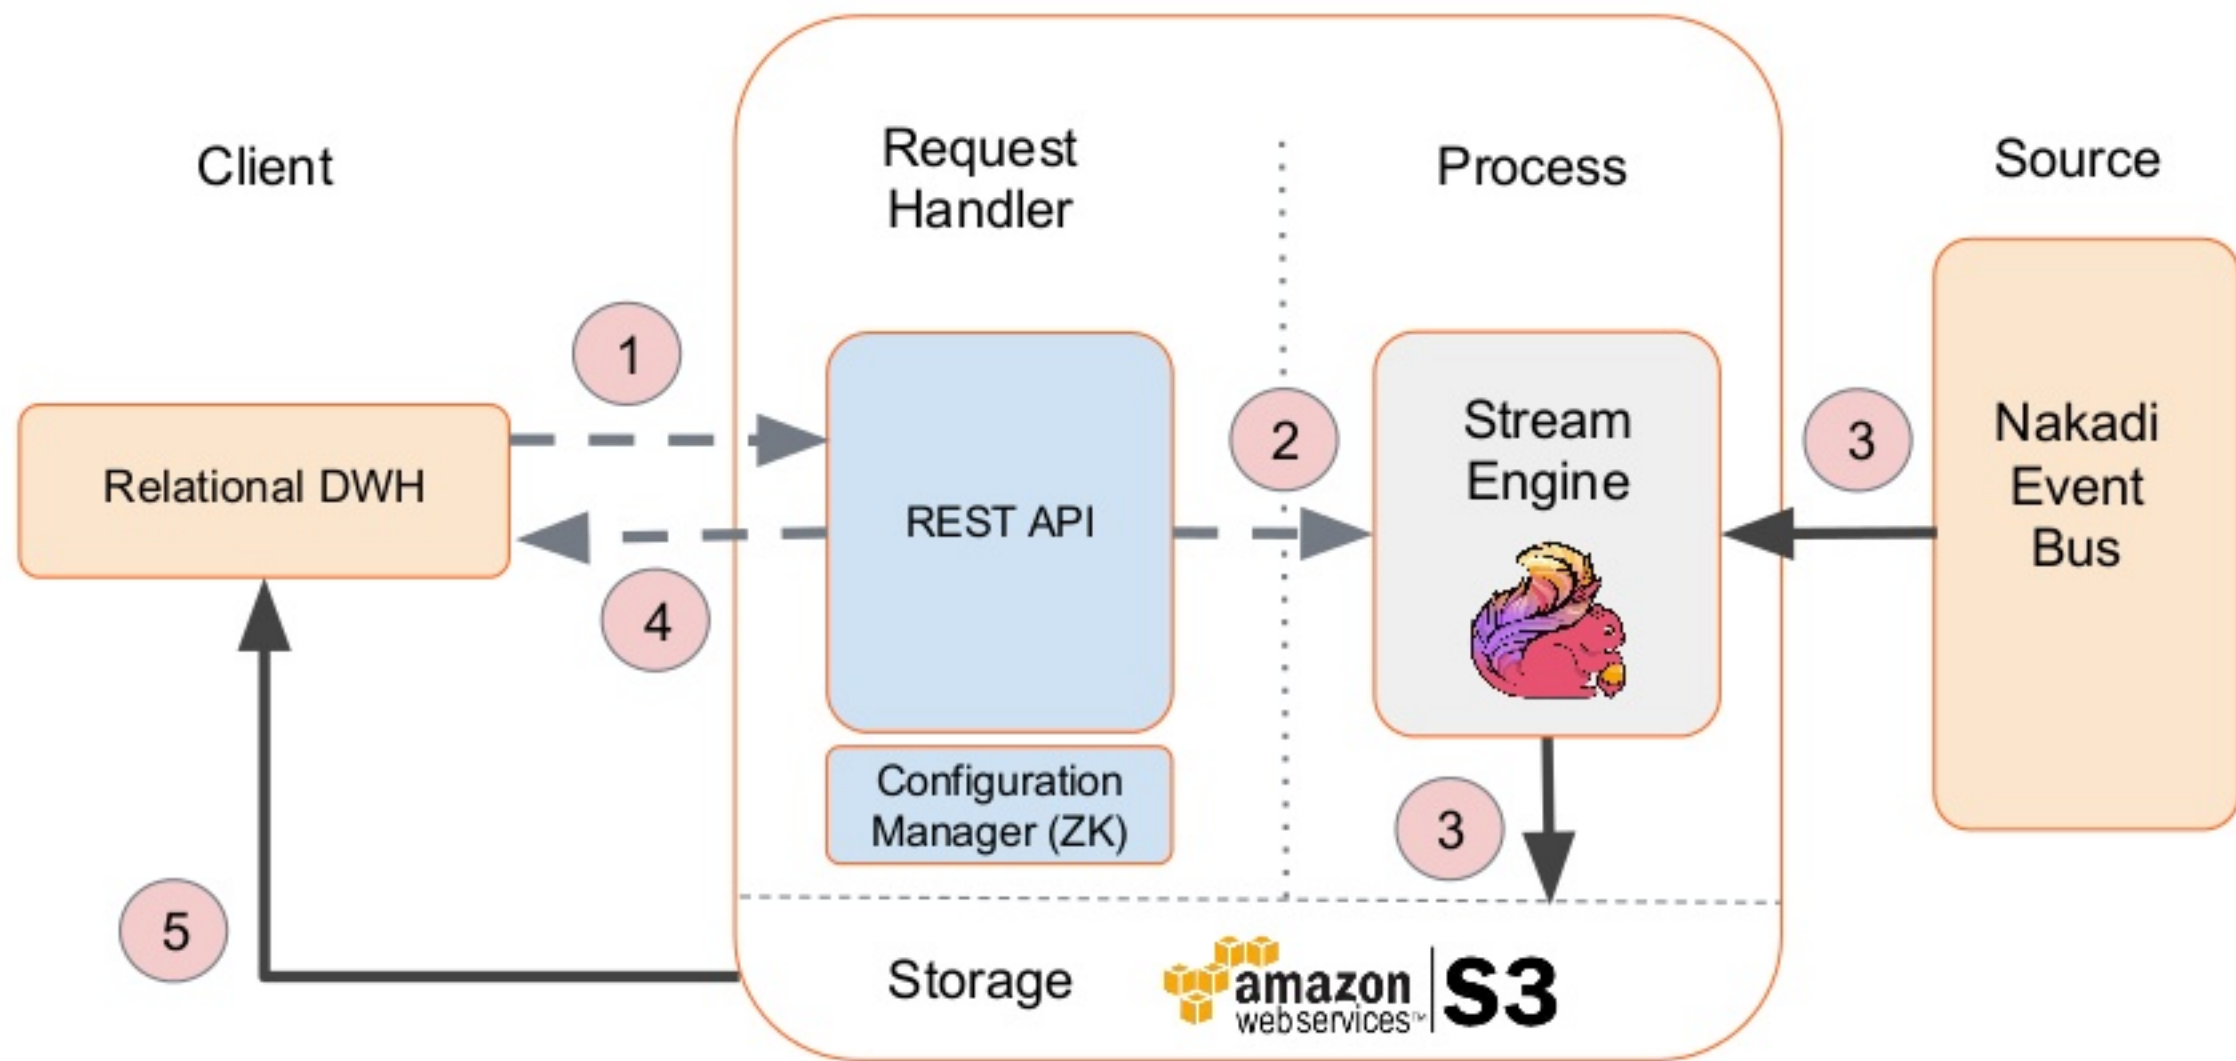
- Materialize data from Nakadi event bus into cloud storage
- Transform complex JSON events into easily ingestible CSV flat files, incl. the flattening of arrays
- Relieve load on the (monolithic) data warehouse by compacting event streams according to event properties

MATERIALIZATION ENGINE

DATA INTEGRATION CHALLENGE



MATERIALIZATION ENGINE



MATERIALIZATION ENGINE

- Materialization API as another abstraction layer over Flink's REST API: deploy, monitor, control jobs
- Configuration Manager (ZooKeeper) stores per Nakadi topic:
 - JSON-to-CSV mapping
 - partitioning key
 - ordering key (for compaction)

FLINK BACKEND

FLINK BACKEND: STREAM TO BATCH



- Short-lived stream-to-batch jobs: 1 Job / Call
- Stoppable Nakadi Stream Source
- Modified BucketingSink for S3 Frankfurt region
 - Writes files to TaskManagers' attached EBS storage
 - Moves them to persistent S3 storage
 - Circumvents [HADOOP-13324](#) of Flink's S3 Connector
- Batch processing using a Streaming API

FLINK BACKEND: ACCUMULATORS



- Flink jobs expose progress through accumulators
 - # Nakadi consumers finished
 - # events read
 - # files delivered to S3
- Materialization API periodically queries Flink's REST API to measure progress.
Issues job stop request upon completion

FLINK BACKEND: CLUSTER



- Flink in standalone mode inside Docker containers on AWS EC2 t2.large instances
- Cluster specifics:
 - Overprovisioning of TaskManagers/ TaskSlots
 - UpScaling on TaskSlots via CloudWatch

STREAM COMPACTION

STREAM COMPACTION: MOTIVATION

- Relinquish DWH resources by reducing the size of the data to ingest
- Can be applied to events which represent changes to the same resource, i.e. having the same partitioning key, but different ordering keys:

{		:	{	
"article_id":	123,	:	"article_id":	123,
"brand":	"Nike",	:	"brand":	"Nike",
"model":	"Air Max",	:	"model":	"Air Max",
"version":	1,	:	"version":	2,
"available_qty":	50,	:	"available_qty":	30,
}		:	}	

STREAM COMPACTION: IMPLEMENTATION



Deploy Flink jobs with config as parameter from Request Handler.

Compact the stream according to the partitioning and ordering keys.



Compaction rates up to 70%

STREAM COMPACTION: OUT-OF-ORDER EVENTS

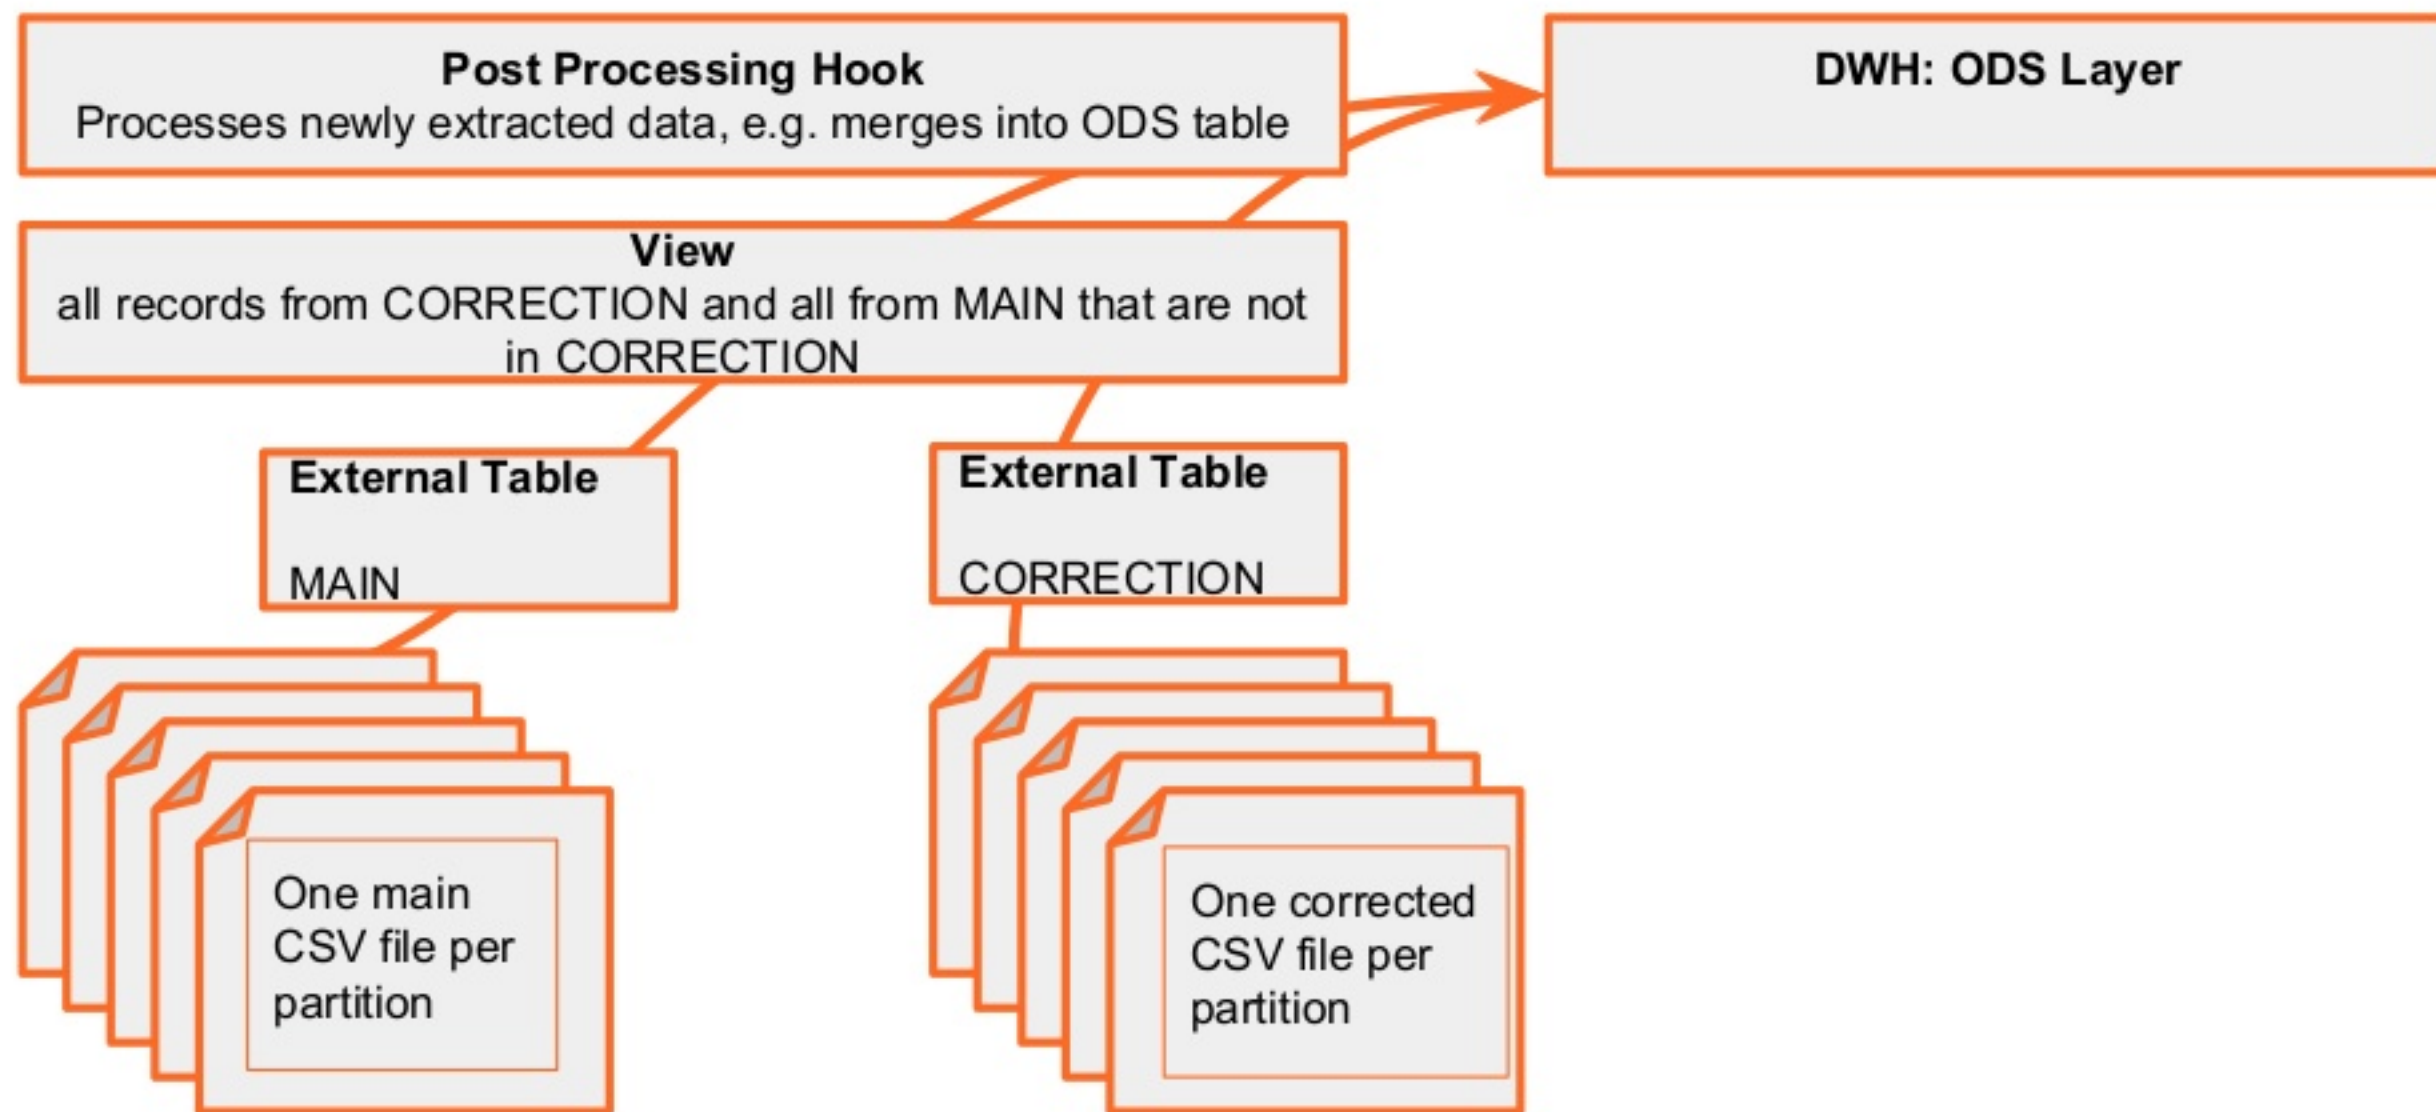


Out-of-order events = Events which were received not in the order we expected, e.g. higher version value first

- Materialize out-of-order events in a different CSV file (correction file)
 - Generated as a side output from main stream
- Rel. DWH creates a view from main and corrected files before merging into Operational Data Store (ODS) table

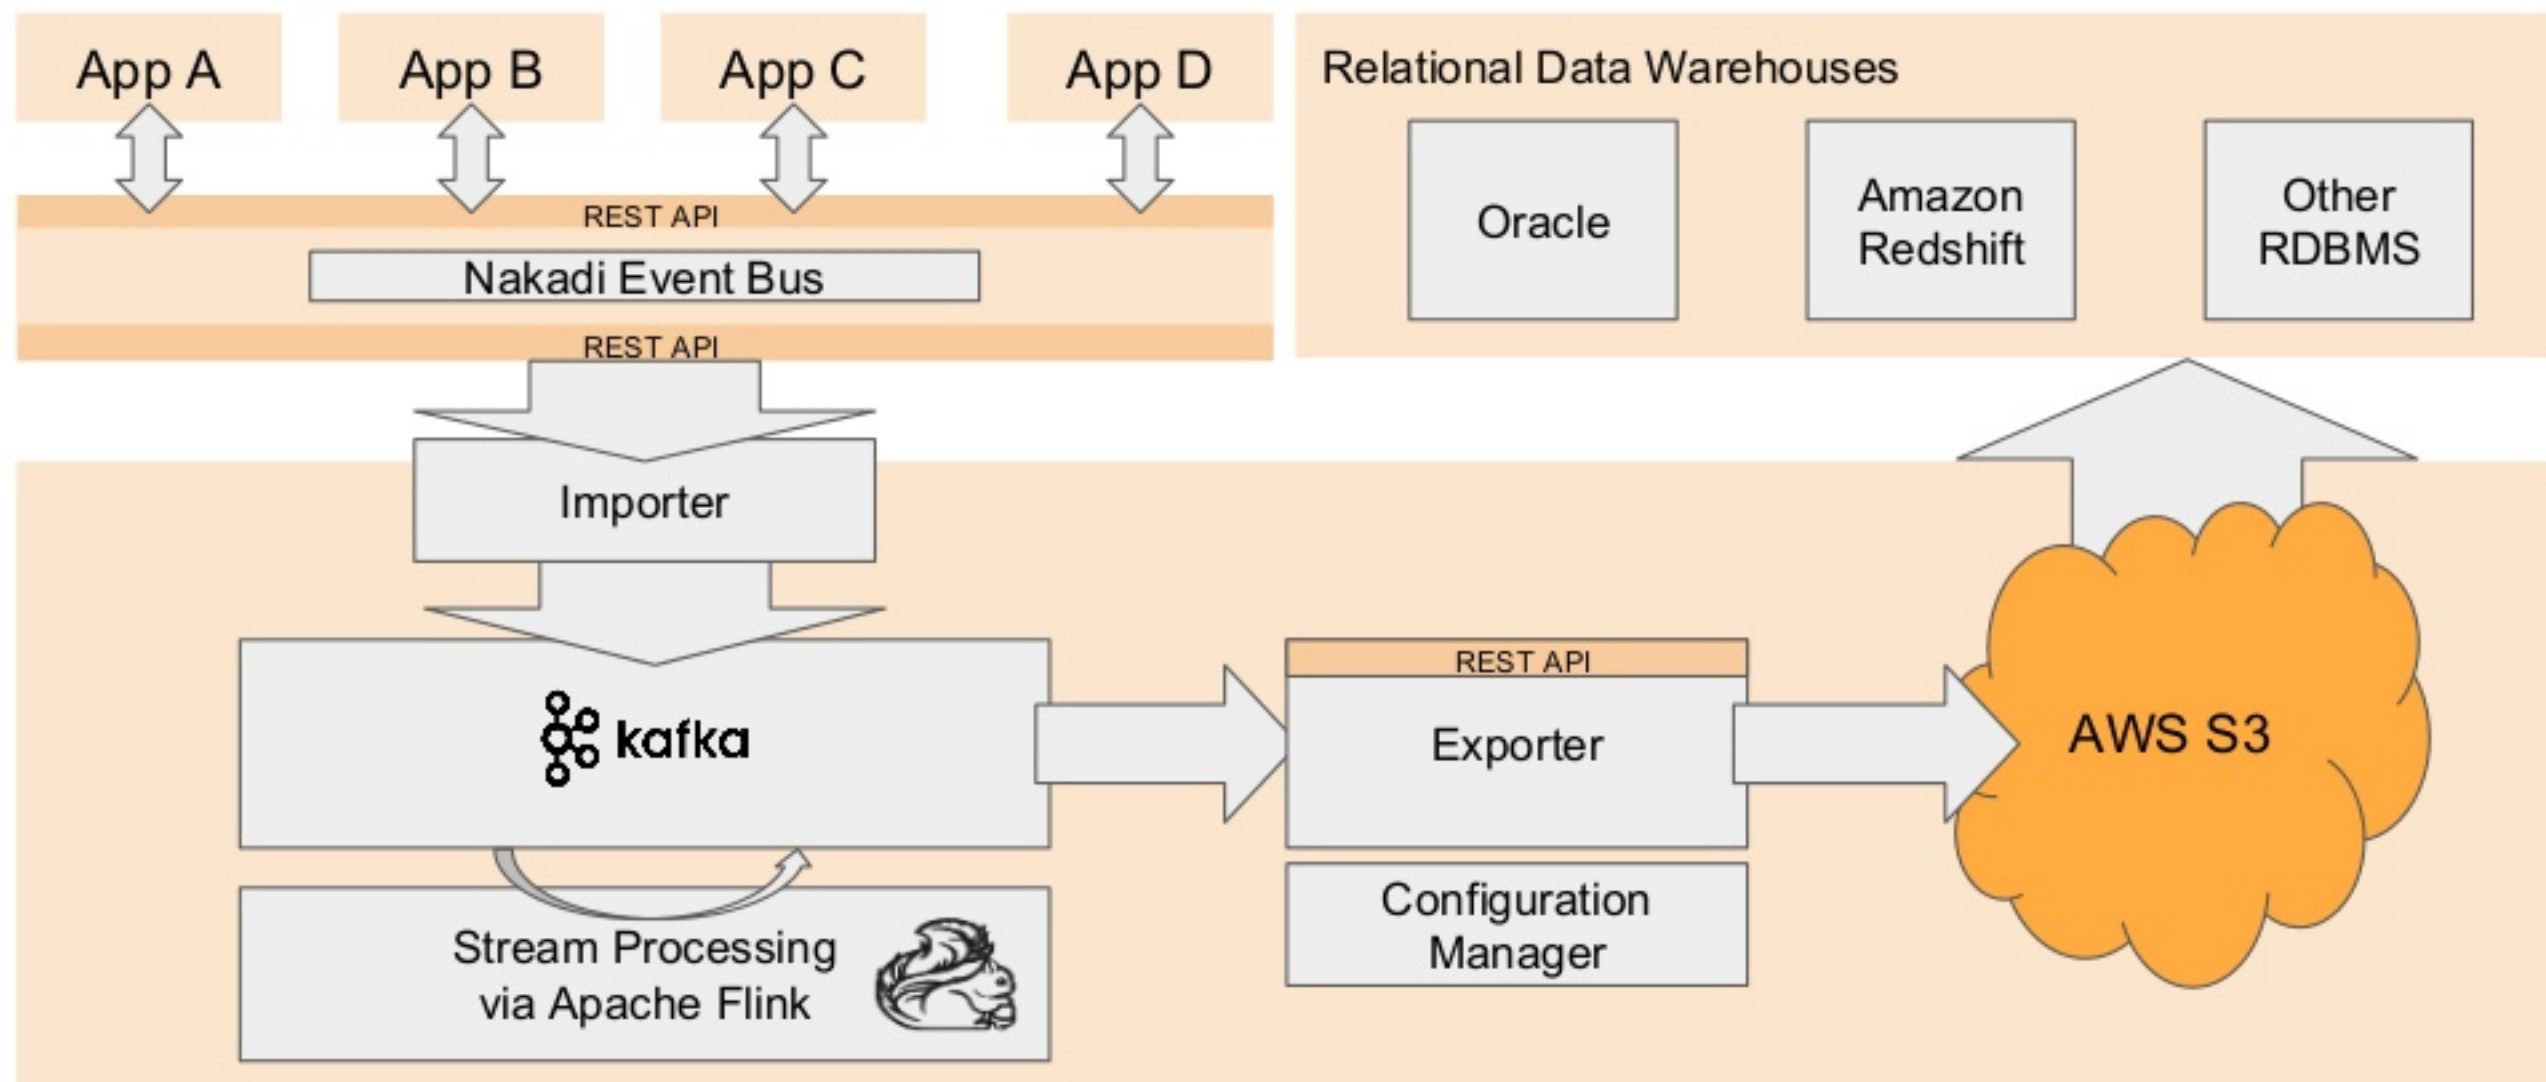
DWH CLIENT'S IMPORT PROCESS

ORACLE®



ADVANTAGES OVER LEGACY APPROACH

LEGACY DATA INTEGRATION ARCHITECTURE



ADVANTAGES OVER LEGACY APPROACH

Fewer stacks: Flink + Materialization API + ZooKeeper
instead of Importer + Kafka + Flink + Exporter API + ZooKeeper

- Reduced AWS costs
- Decreased operational overhead
 - No data redundancy (Nakadi + Kafka), no Importer setup
 - Far less maintenance, e.g. no streaming array flattening jobs
- Easier reasoning and implementation through Flink's API
 - Compaction
 - Extensible feature set (more ETL in the future)

THANK YOU

mihail.vieru@zalando.de

BACKUP SLIDES

ORACLE CLIENT'S IMPORT PROCESS

