



# RADICALBIT

Deep, Different, Disruptive.

## **NSDb (Natural Series Db)**

A time series streaming oriented database  
optimized for the serving layer

Roberto Bentivoglio - @robbenti

Saverio Veltri - @save\_veltri

Berlin – 13 September 2017

## WHO WE ARE

**Radicalbit** is a highly **specialized software firm**, founded in Milan, Italy, in 2015.

We are focussed on the design and development of **Fast Data products and solutions**, using the best-of-breed **streaming technologies** (and of course **Flink**).

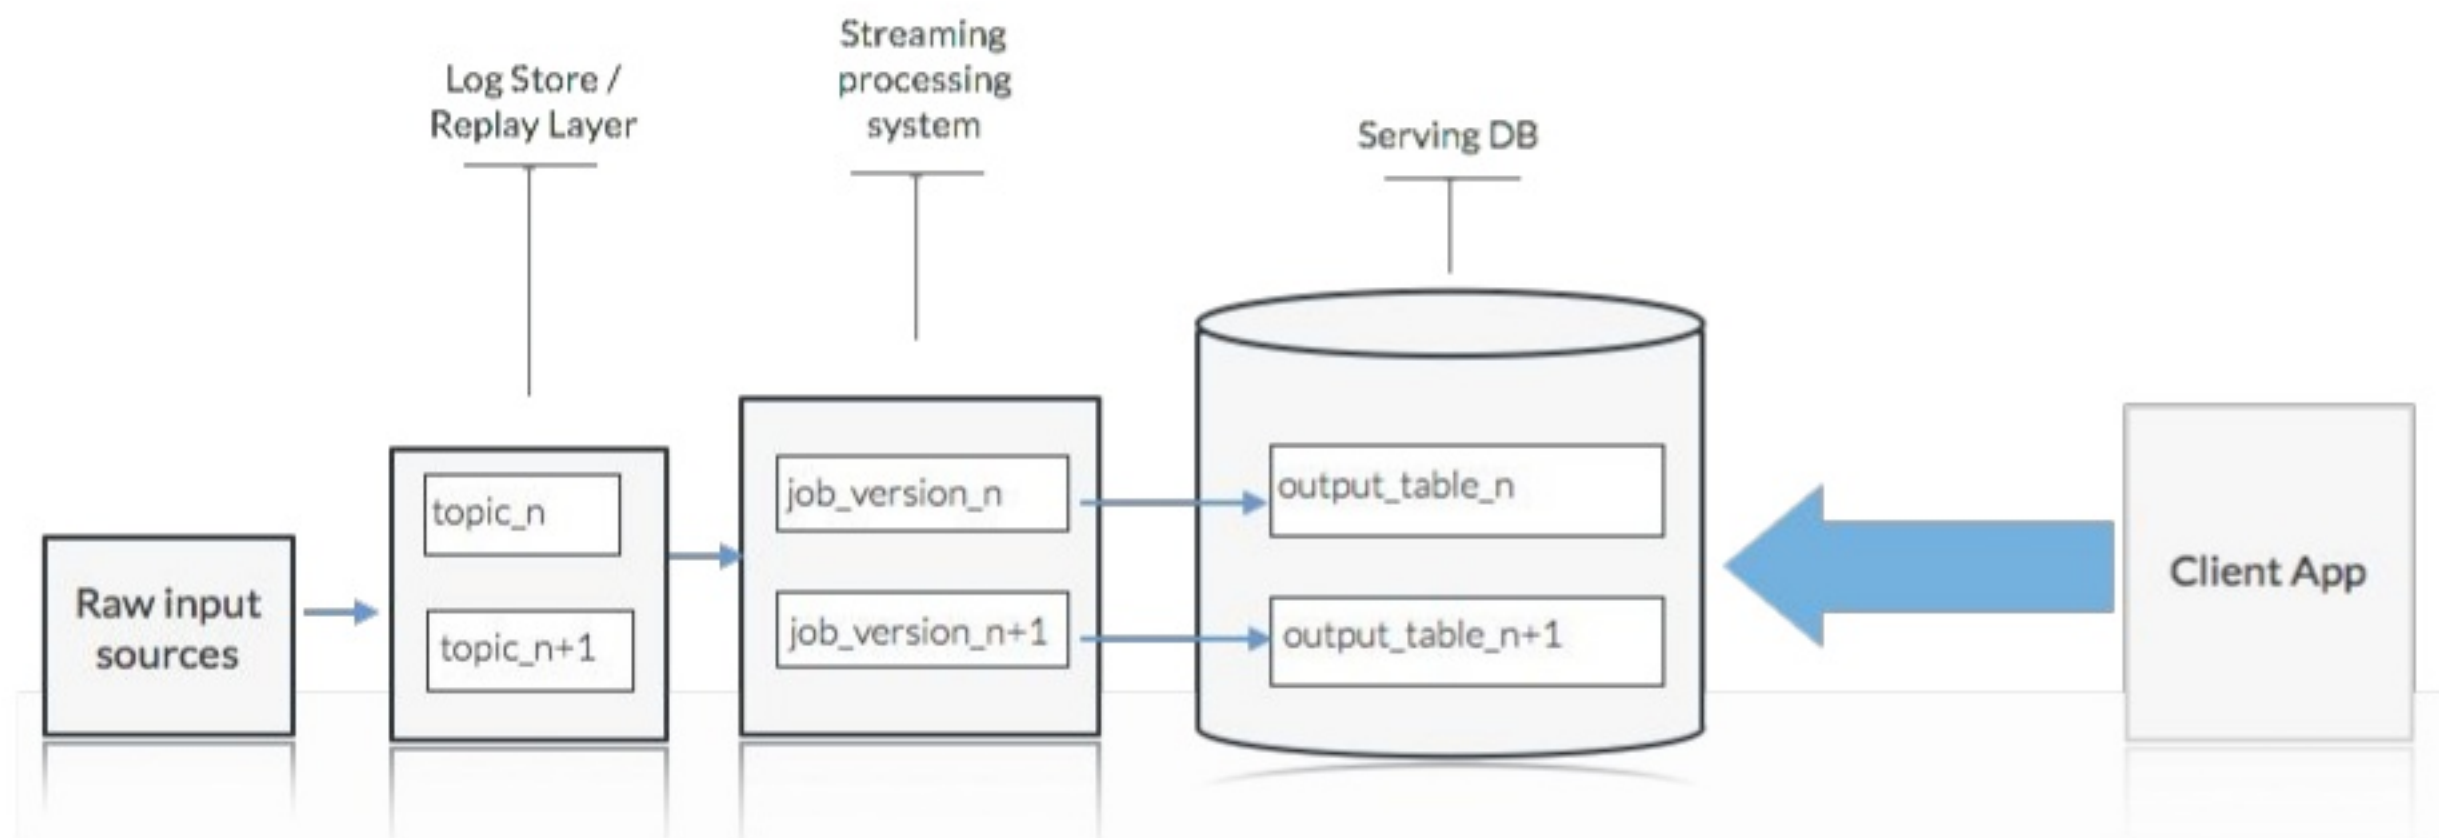
We're rapidly growing, offering our enterprise solutions to data intensive companies.

# NSDb - MOTIVATIONS

---

- Our Fast Data distribution lacked of Time Series functionalities
- Too many licensing and pricing issues exploring third-party OEM solutions
- Third-party solutions don't completely fit
  - our non-functional requirements
  - our streaming real-time analytics needs
- Have a deep technical ownership of the solution
- Publish / Subscribe model
- Push capabilities
- Create a community around Radicalbit products
- We'd like to grow up our Scala codebase (we're Scala guys)

# KAPPA ARCHITECTURE



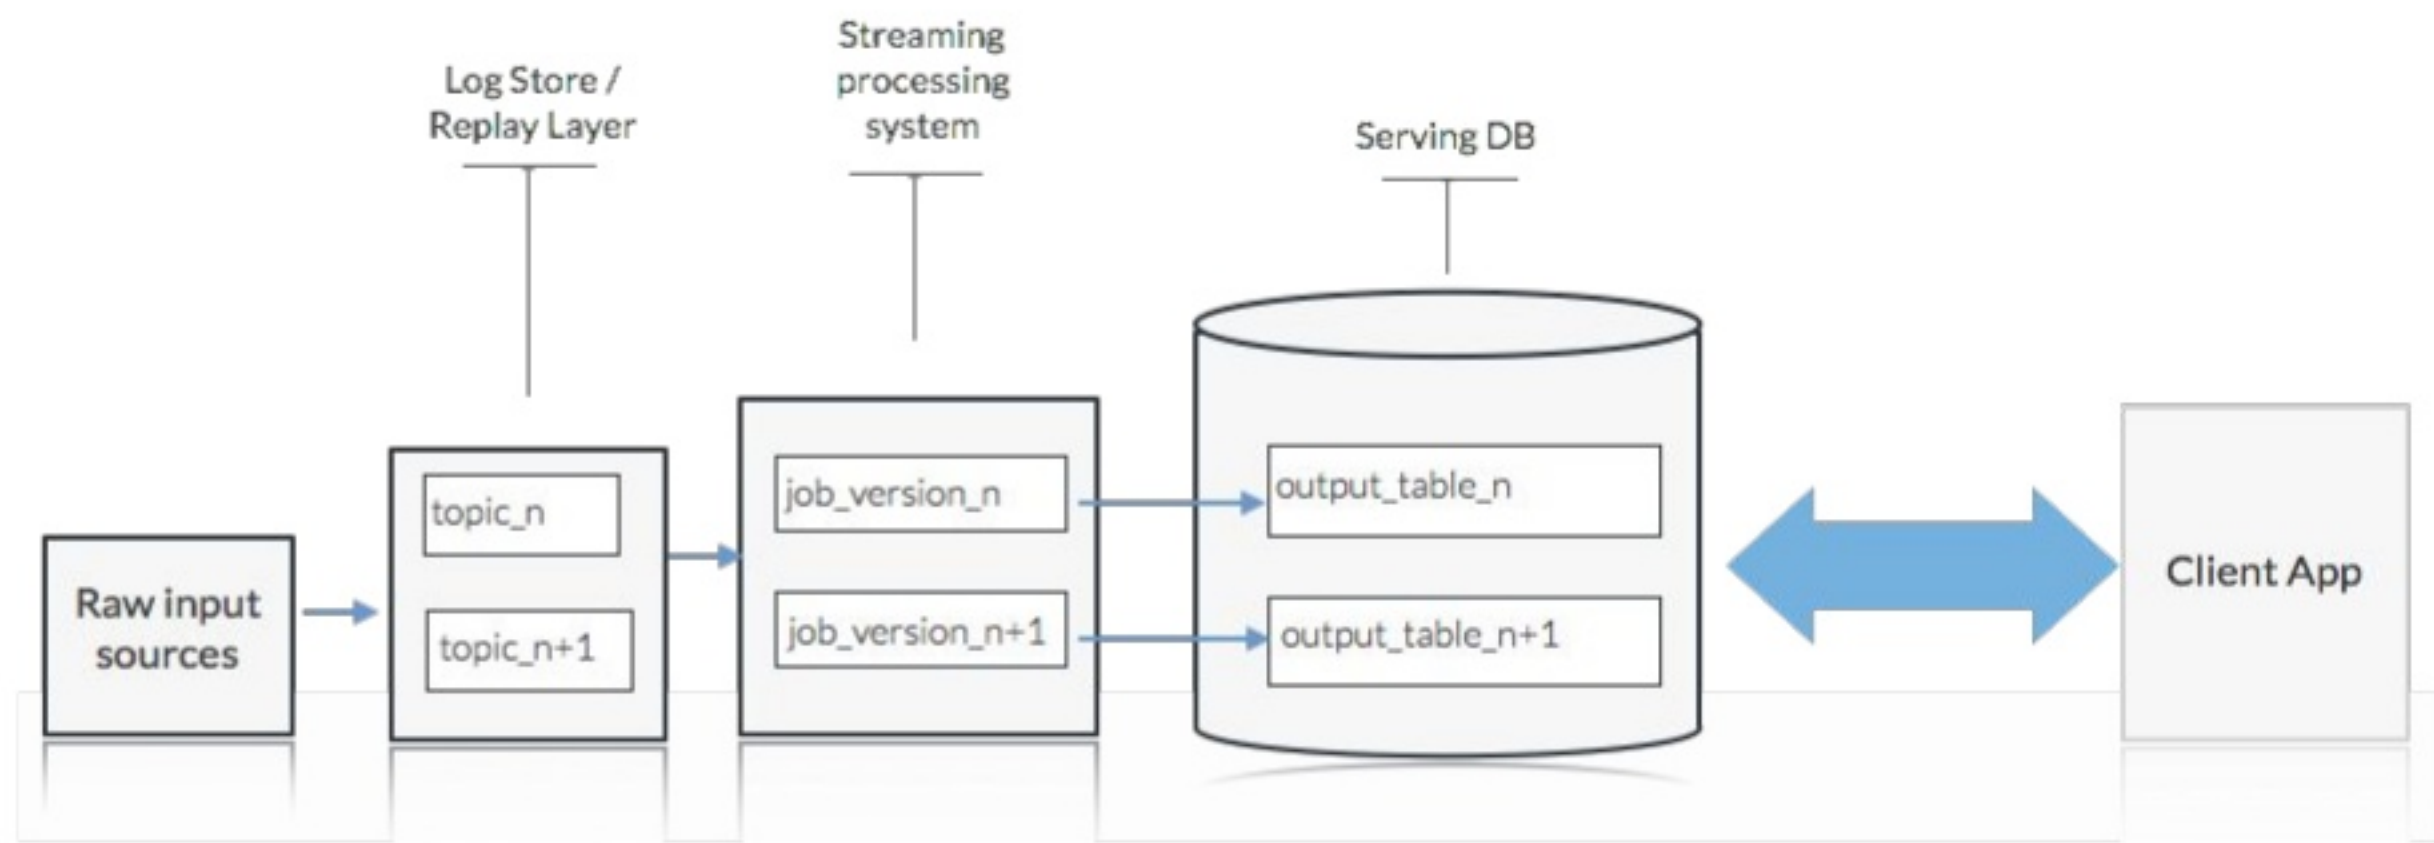
# CHALLENGES IMPLEMENTING KAPPA ARCHITECTURE

---

- Everything is handled as a **streaming event**. This is great!
- There are plenty of solutions for handling real time writes / journaling  
(Some of them are being presented concurrently to our talk :))
- If a traditional database is being used, any possible **real time capability** will be lost on the **serving layer**
- Any application must query periodically the DB for gathering data in a **pull fashion**



# KAPPA ARCHITECTURE REVISITED



# NSDb - MAIN FEATURES

- Optimized time series management
- Focus on read performance
- High availability and clustering
- Ad-hoc publish-subscribe streaming feature (using WebSockets)
- SQL like query language
- Fluent Java / Scala Api

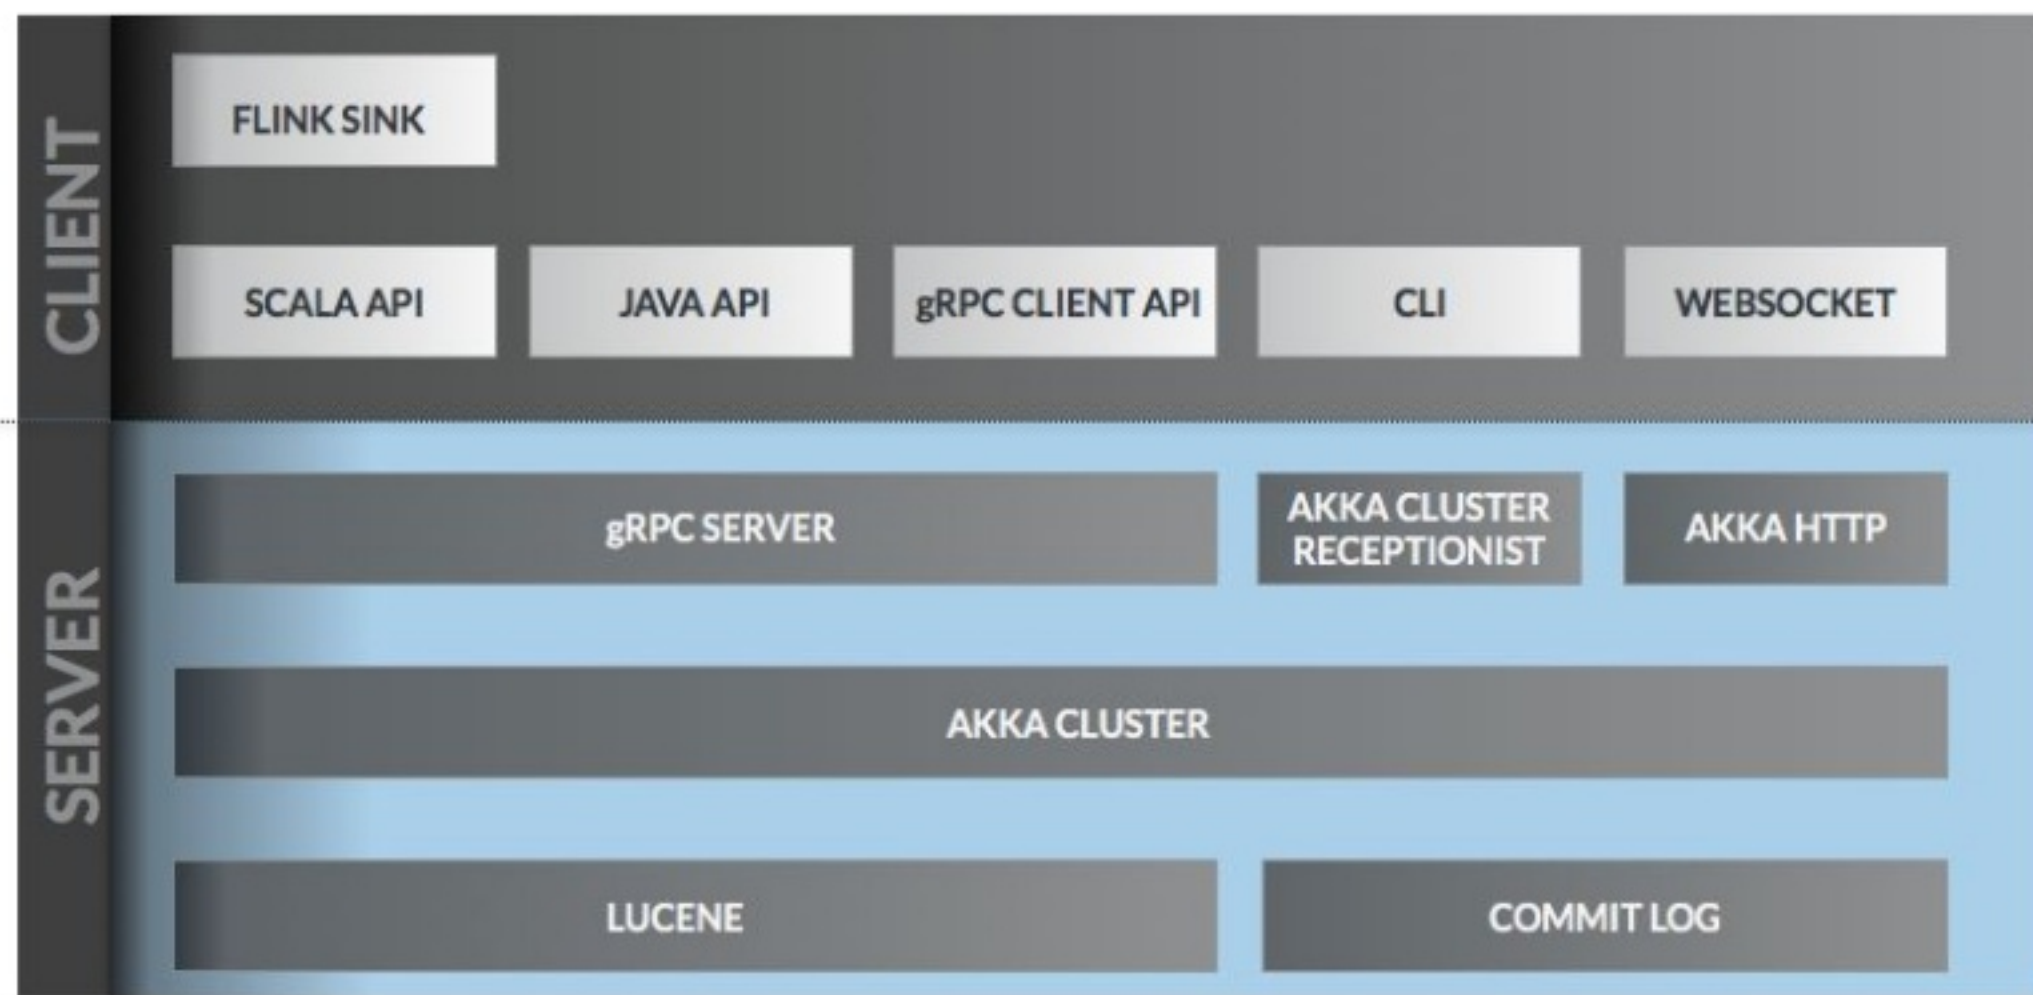
# NSDb - MAIN FEATURES

---

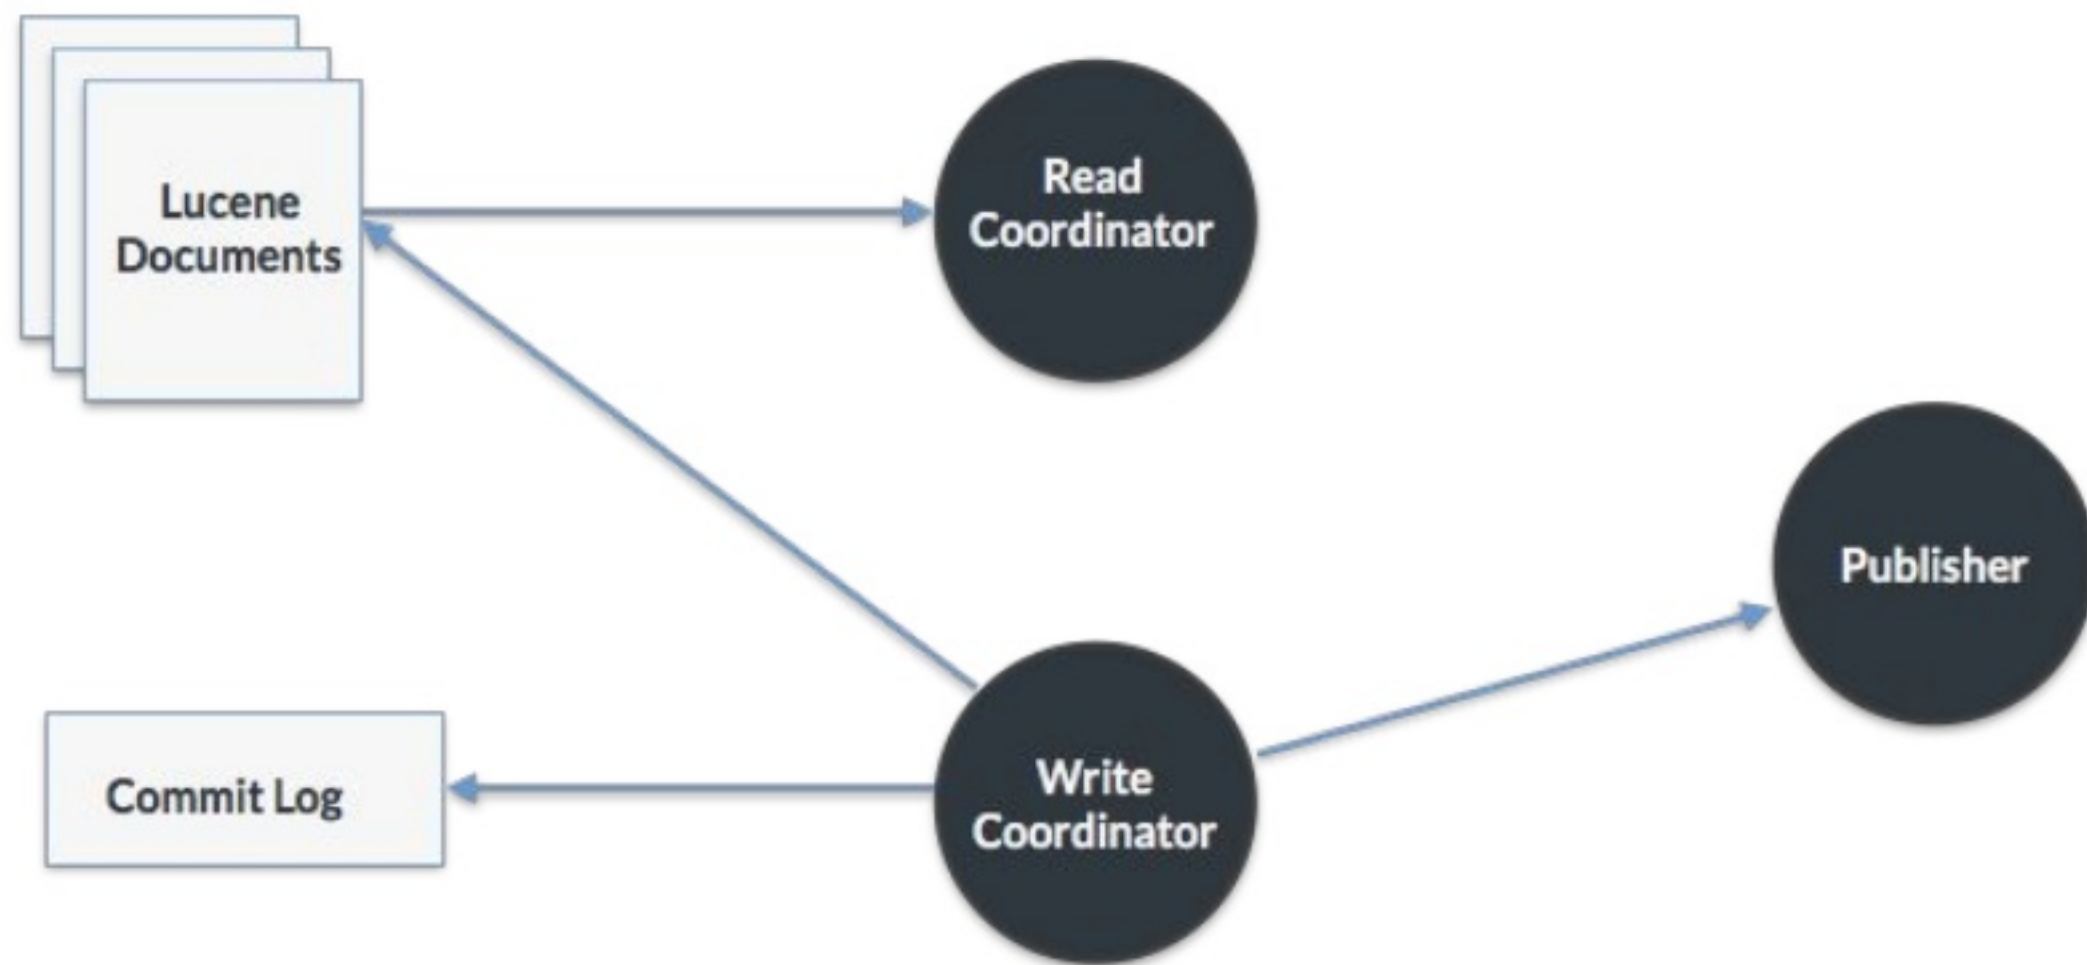
- **Apache Flink** integration (sink)
- Integration with many languages thanks to gRPC (Go, Python, C++, ...)
- Metrics versioning
- Native support for Apache Avro (ingestion, schema publishing and versioning)
- JDBC support
- Import / Export features



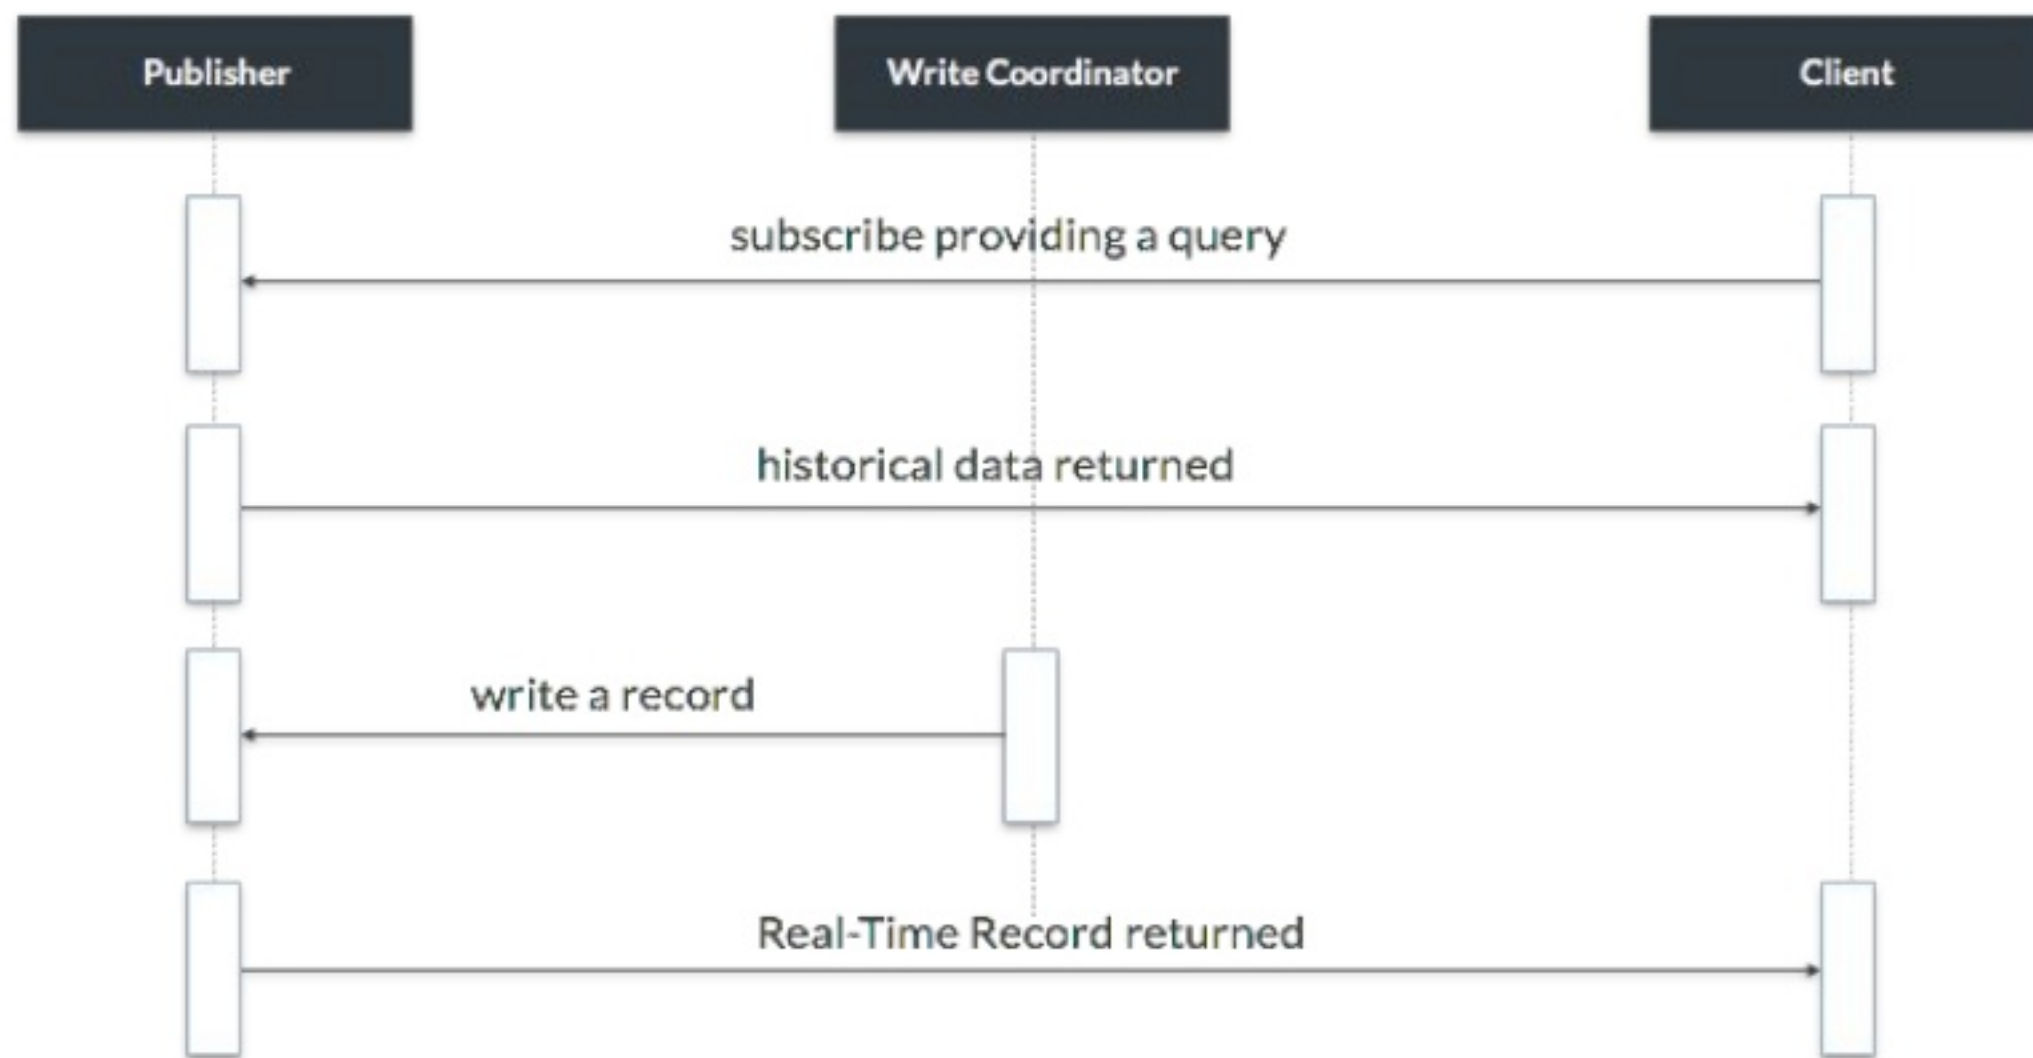
# NSDb - OVERALL ARCHITECTURE



# NSDb - CORE ARCHITECTURE

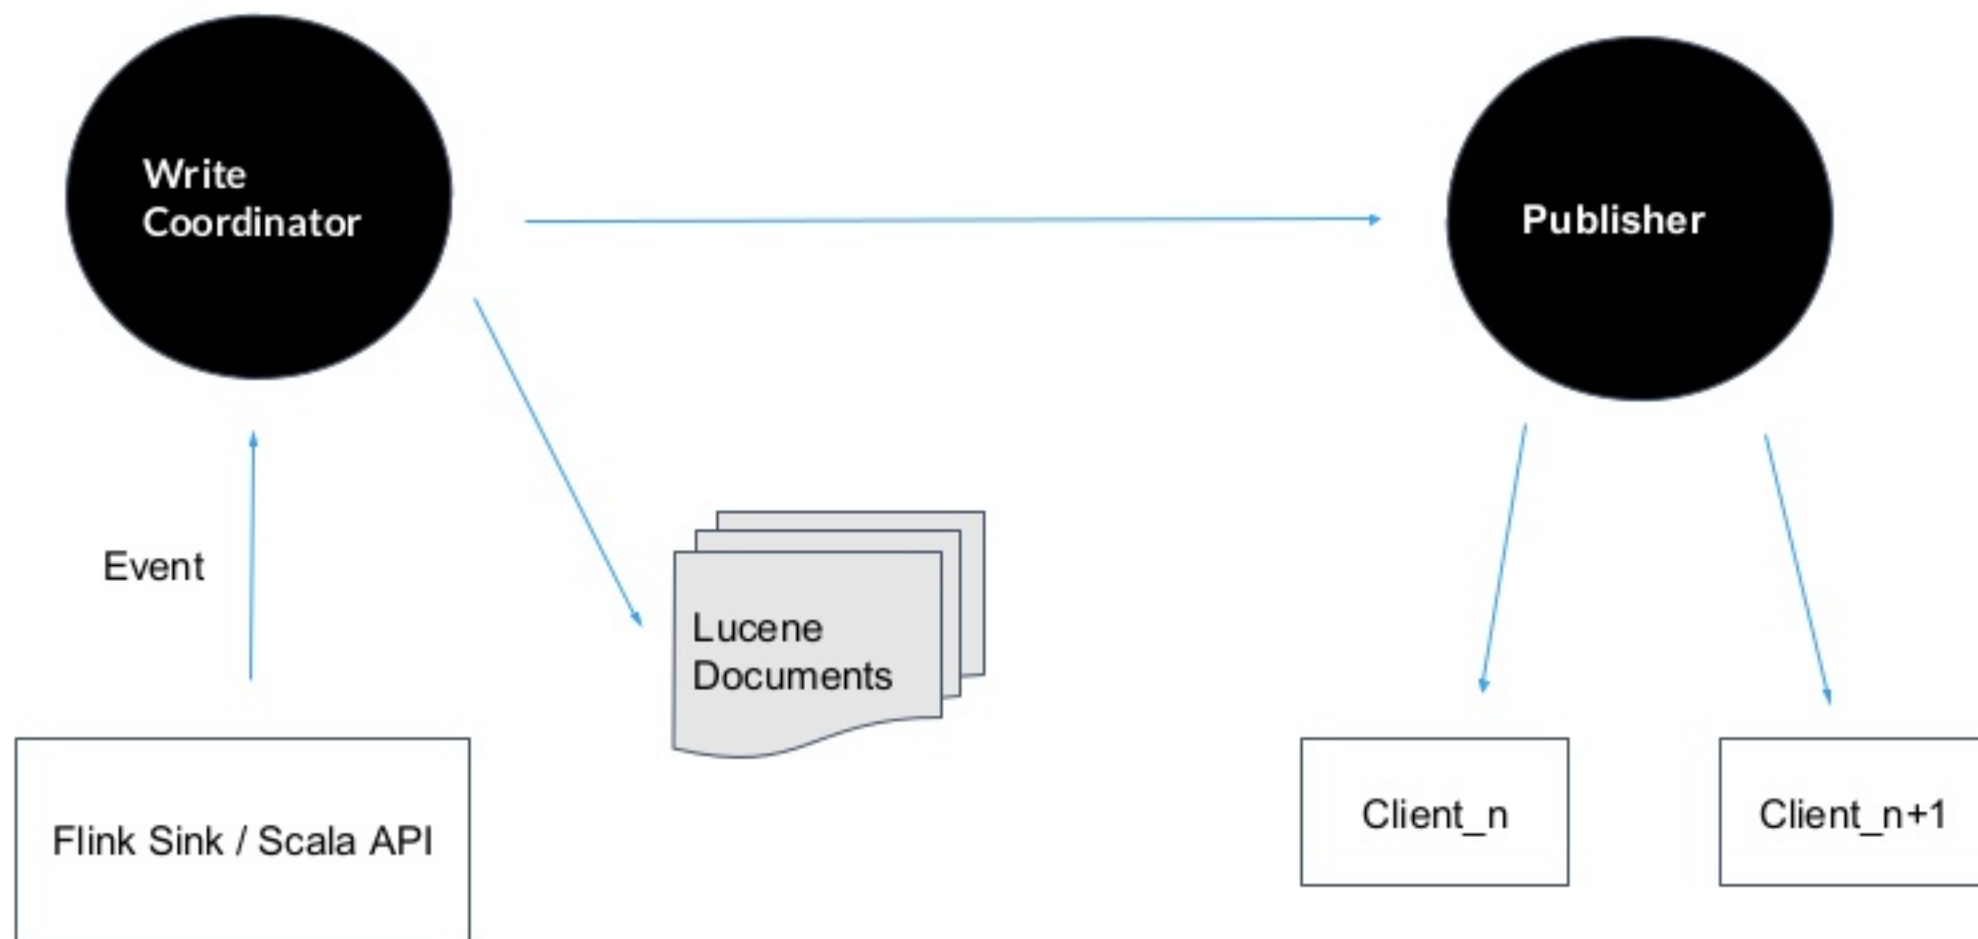


# NSDb | PUBLISH-SUBSCRIBE



# NSDb - CONSISTENCY

- Eventual consistency
- Real time delivery for subscribed client



# SCALING


- Scaling out is a challenging topic
- How would it be possible to scale a generic Sql or NoSql workload ?

¯\\_(ツ)\_/¯

- In most of the well known NoSql Databases data partitioning is something demanded to the user
- Unless boundaries are being defined

# NSDb - BIT

MultiDimensional time series value






## NSDb - SCALING

---

- NSDB's Bits are immutable. New data continuously arrives, and will be always inserted and never updated.
- Time series data are naturally ordered by time



Time is Totally a Gentleman!

# NATURAL TIME SHARDING

- Workloads can be partitioned across time and space
- space partitioning will be demanded to Lucene indexes
- Time partitioning leverages intrinsic time linearity



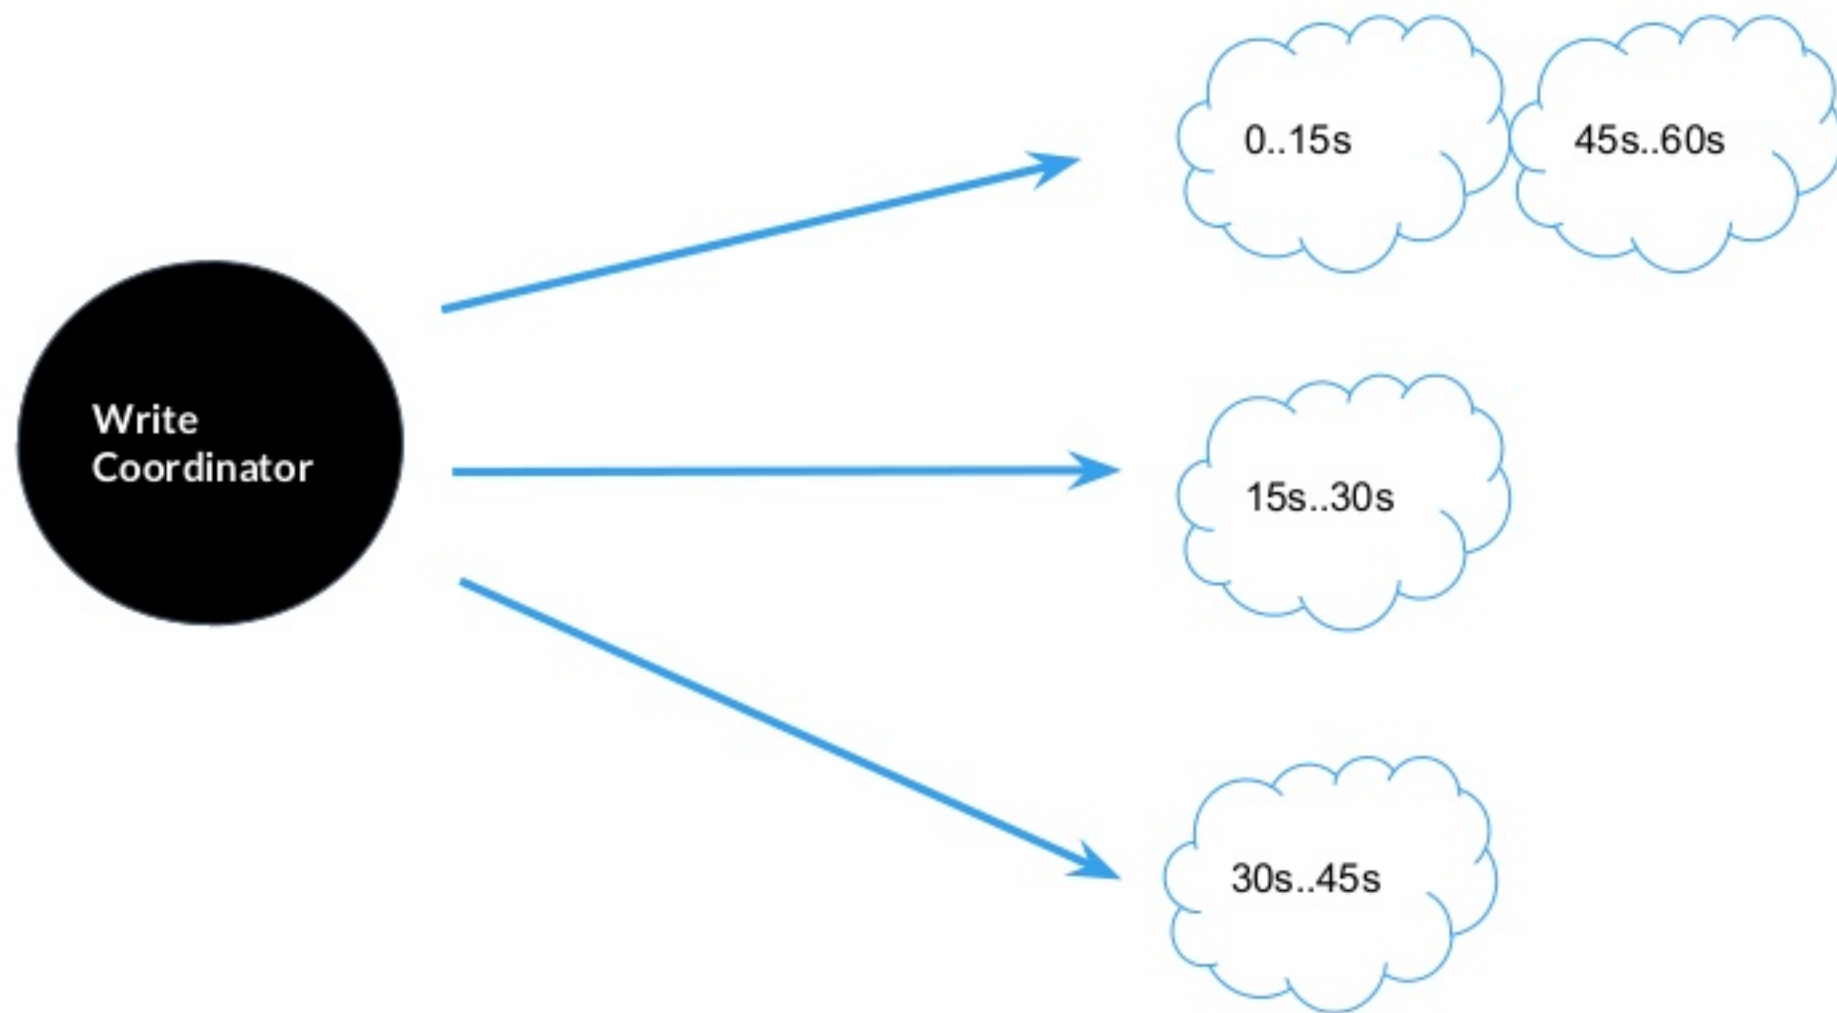
# NSDb - NATURAL TIME SHARDING

```
sharding {  
  ...  
  interval = 15s  
}
```

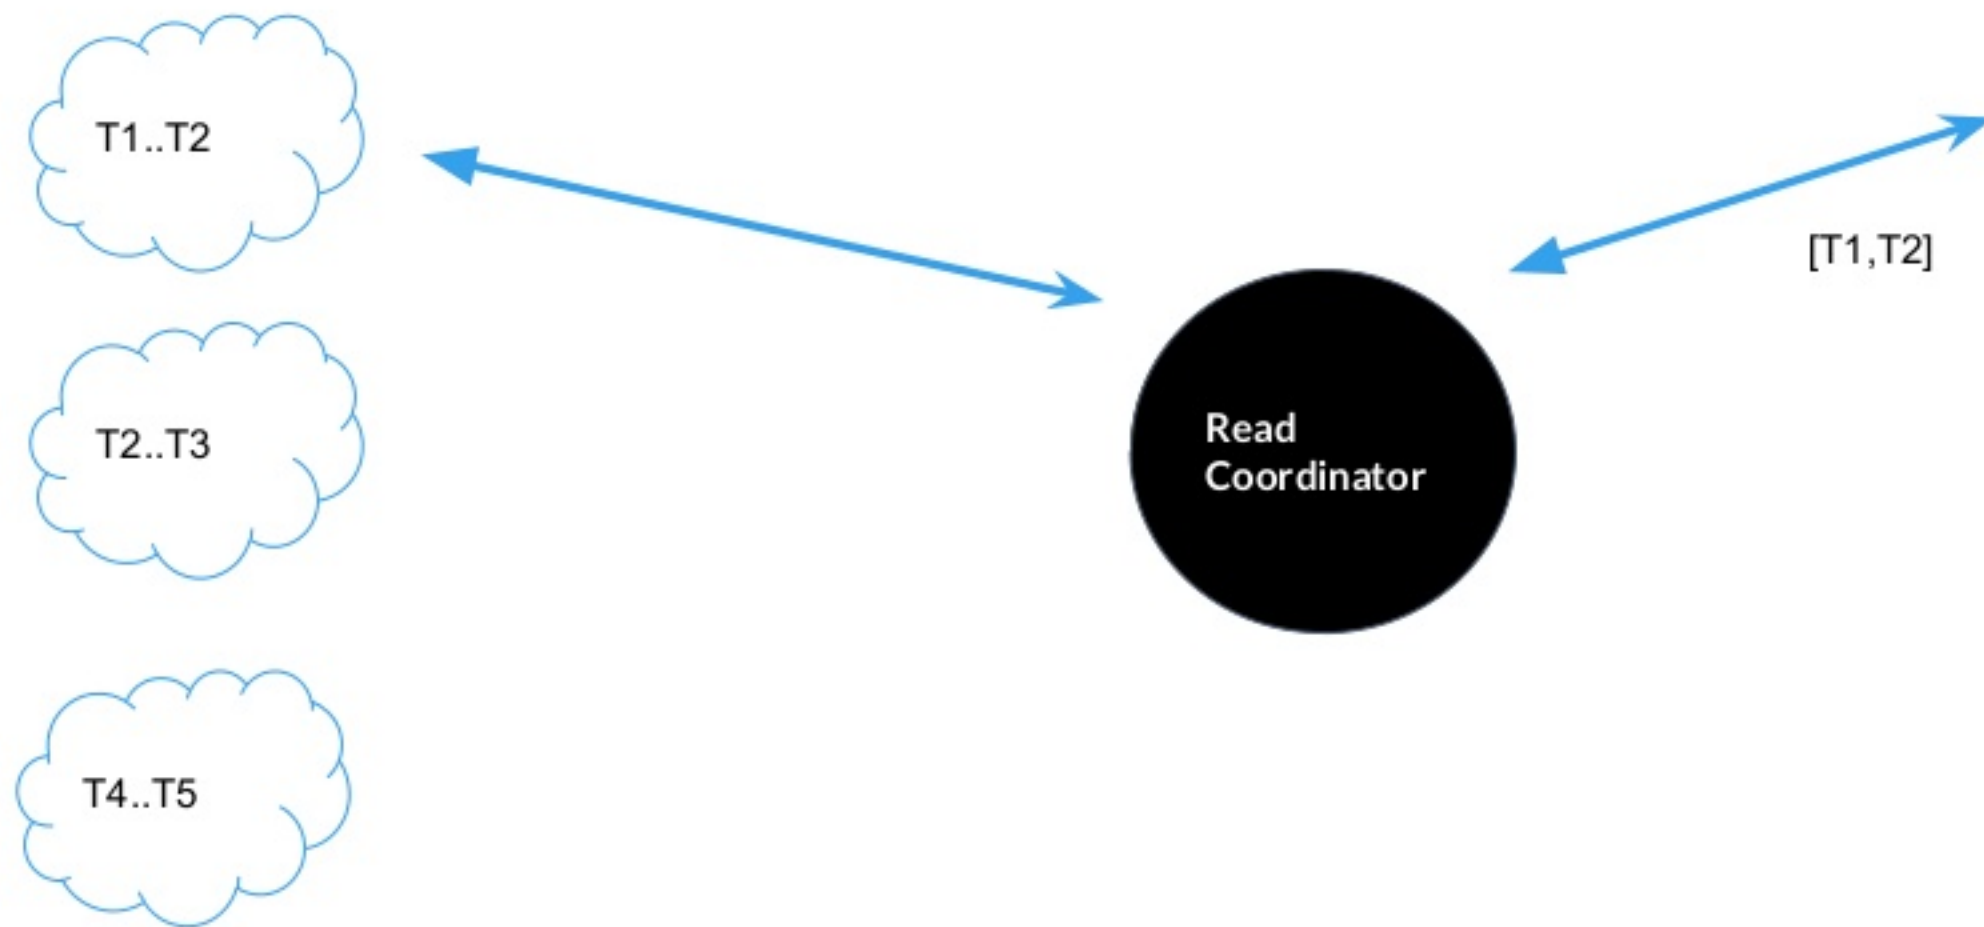
- A node is picked up for writing until time length reaches configured interval
- Autoscaling is supported

# Natural Time Sharding - Data Partitioning

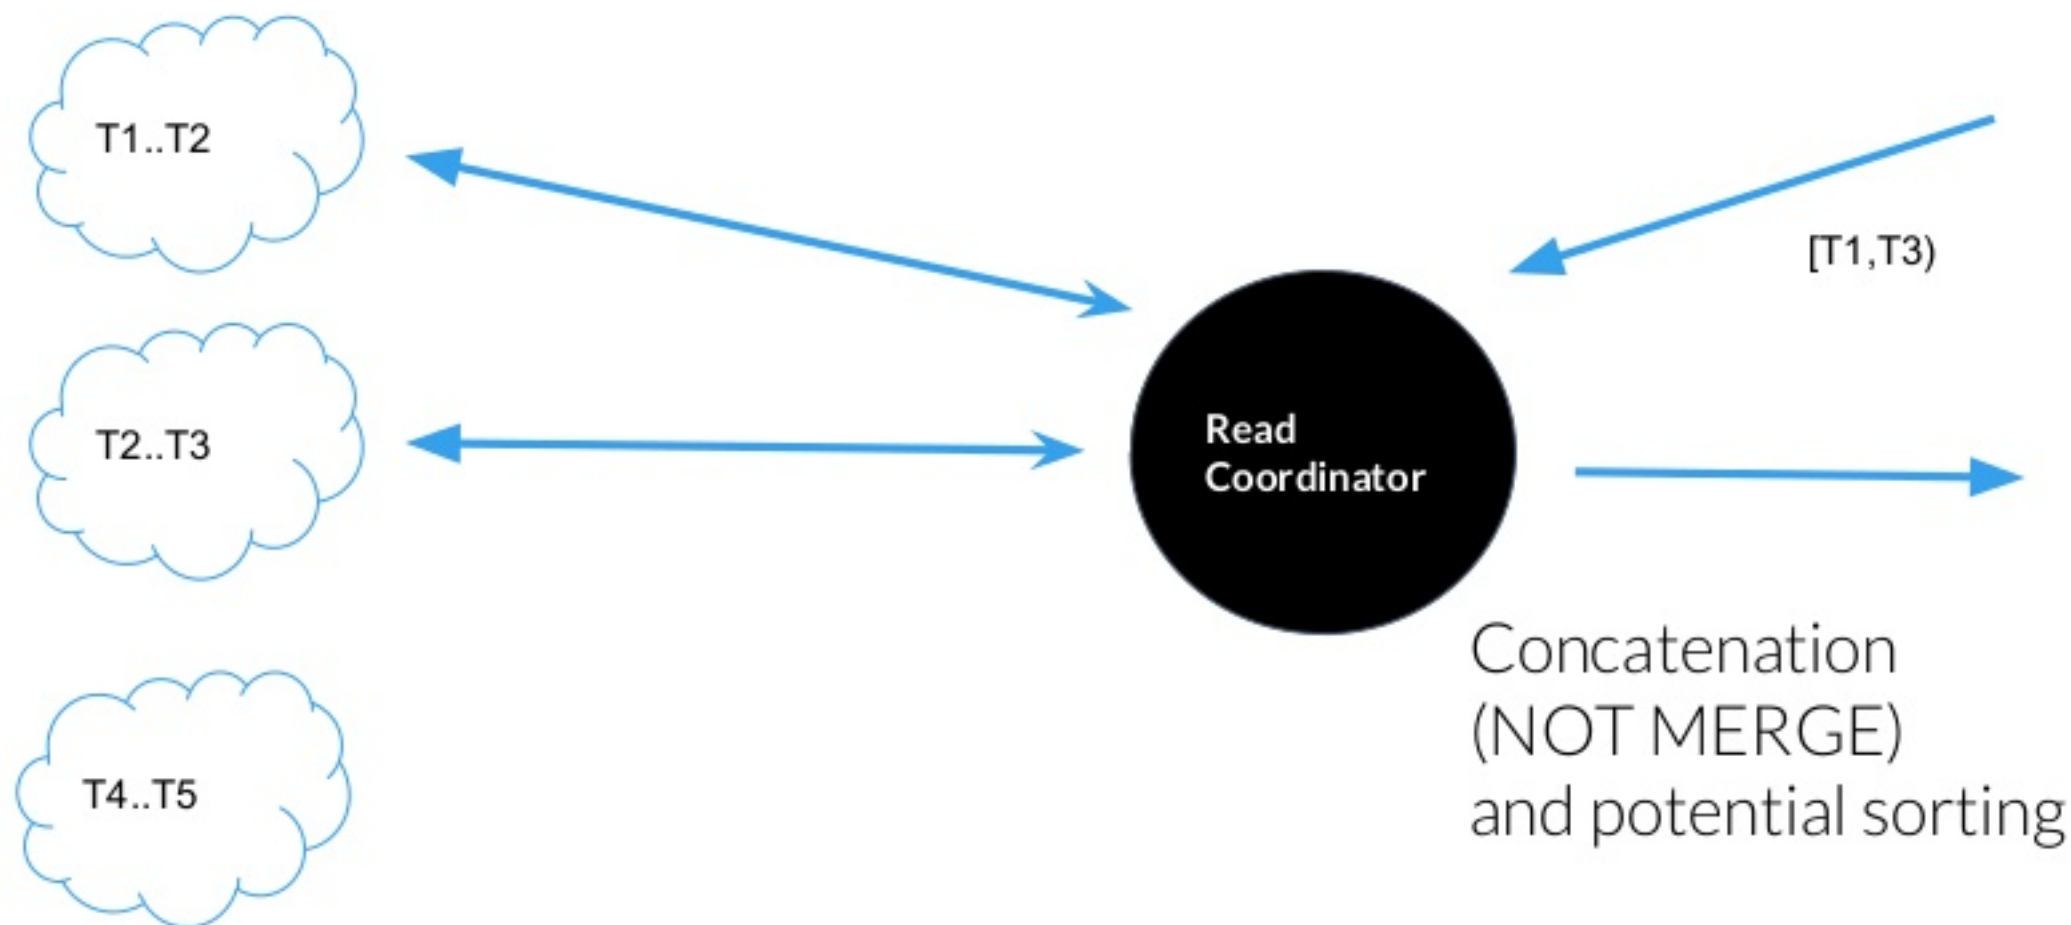
`interval = 15s #nodes = 3`



# Natural Time Sharding - Best Case

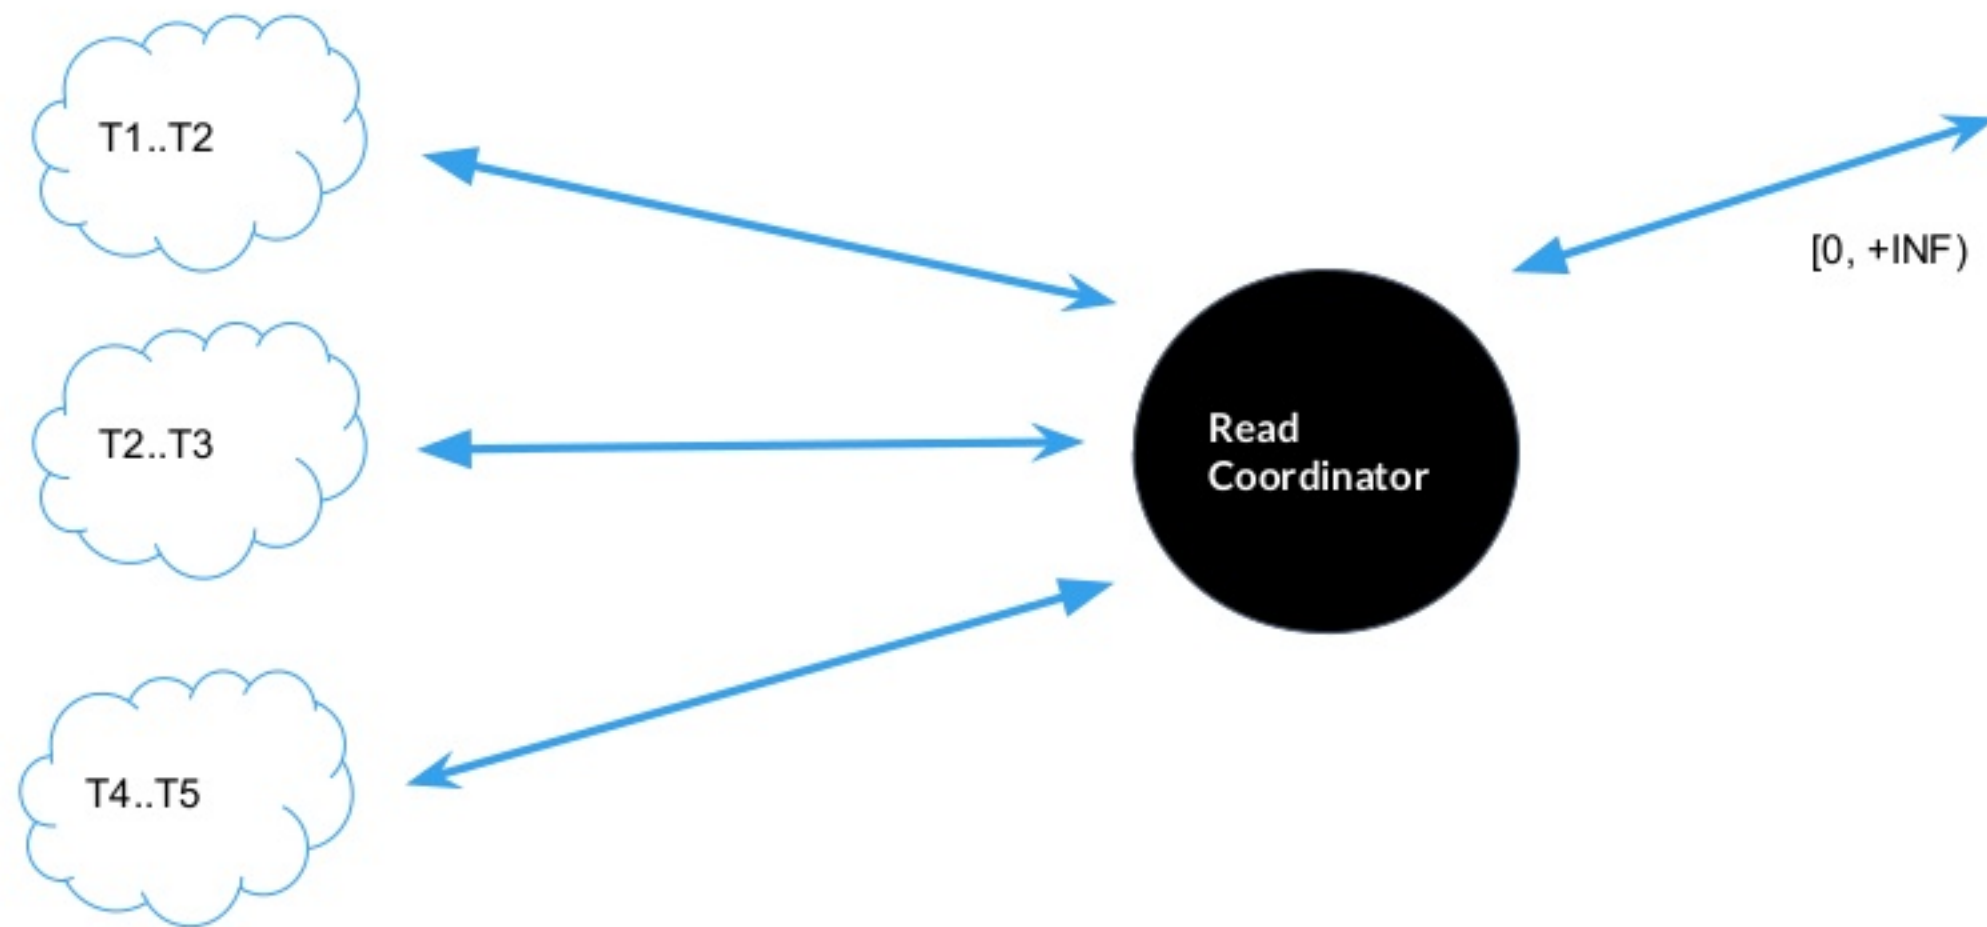


## Natural Time Sharding - Average Case





# Natural Time Sharding - Worst Case



# NSDb - TIME SERIES MANAGEMENT

---

- **Namespace**: high level structure grouping homogeneous metrics
- **Metric**: a series of Bit (Point)
  - **Bit**: logical construct used into the API. It groups together:
    - **Timestamp**: the event time
    - **Value**: the numerical value being measured
    - **Dimensions**: a dynamic list of queryable String -> Value pairs

## NSDb - SCALA API

```
val nsdb = NSDB.connect(host = "127.0.0.1", port =  
7817) (ExecutionContext.global)
```

```
val series = nsdb.namespace("registry")  
  .bit("order amount")  
  .value(34.20)  
  .dimension("city", "Berlin")  
  .dimension("gender", "M")
```

```
val res: Future[RPCInsertResult] = nsdb.write(series)  
res onComplete {  
  case Success(r) => // do something here...  
  case Failure(t) => // manage the errors here...  
}
```

## NSDb - FLINK SINK

```
val sentiment: DataStream[TweetSentiment] =  
sentimentalAnalysisByHashtag.flatMap(new  
ScoreText[TweetByHashtag])
```

```
val sentimentSink = new NSDBSink[TweetSentiment]  
("localhost", 7817)  
{  
  (s: TweetSentiment) =>  
    Bit(namespace = "twitter", metric = "TweetSentiment")  
      .dimension("key", s.score.toString) // Positive,  
Negative, Neutral  
      .value(1)  
}  
sentiment.addSink(sentimentSink)
```

Live demo showing a stream of tweets



# NSDb - ROADMAP

---

The first Release Candidate will be published around November / December 2017 (Apache License v2.0) and it will offer:

- Clustering
- Java and Scala + gRPC API
- **Flink connector** (sink)
- WebSocket
- CLI

The following points will be released in the following versions:

- Integration to stream seamlessly from Pravega and Apache Kafka
- Native Apache Avro management (schema publishing)
- JDBC API thanks to Apache Calcite
- SSL and Kerberos support
- many others



# Q&A

---

# DANKE!

---

**<radicalbit.team/>**

info@radicalbit.io