



# From Apache Flink® 1.3 to 1.4

Till Rohrmann  
till@data-artisans.com  
@stsffap

**dataArtisans**

# dataArtisans



Original creators of Apache  
Flink®



PLATFORM

Providers of  
dA Platform 2, including  
open source Apache Flink +  
dA Application Manager

# Overview

---



- ✦ Apache Flink 1.3 – Previously on Apache Flink
- ✦ Apache Flink 1.4 – What's happening now?
- ✦ Apache Flink 1.5+ – Next on Apache Flink



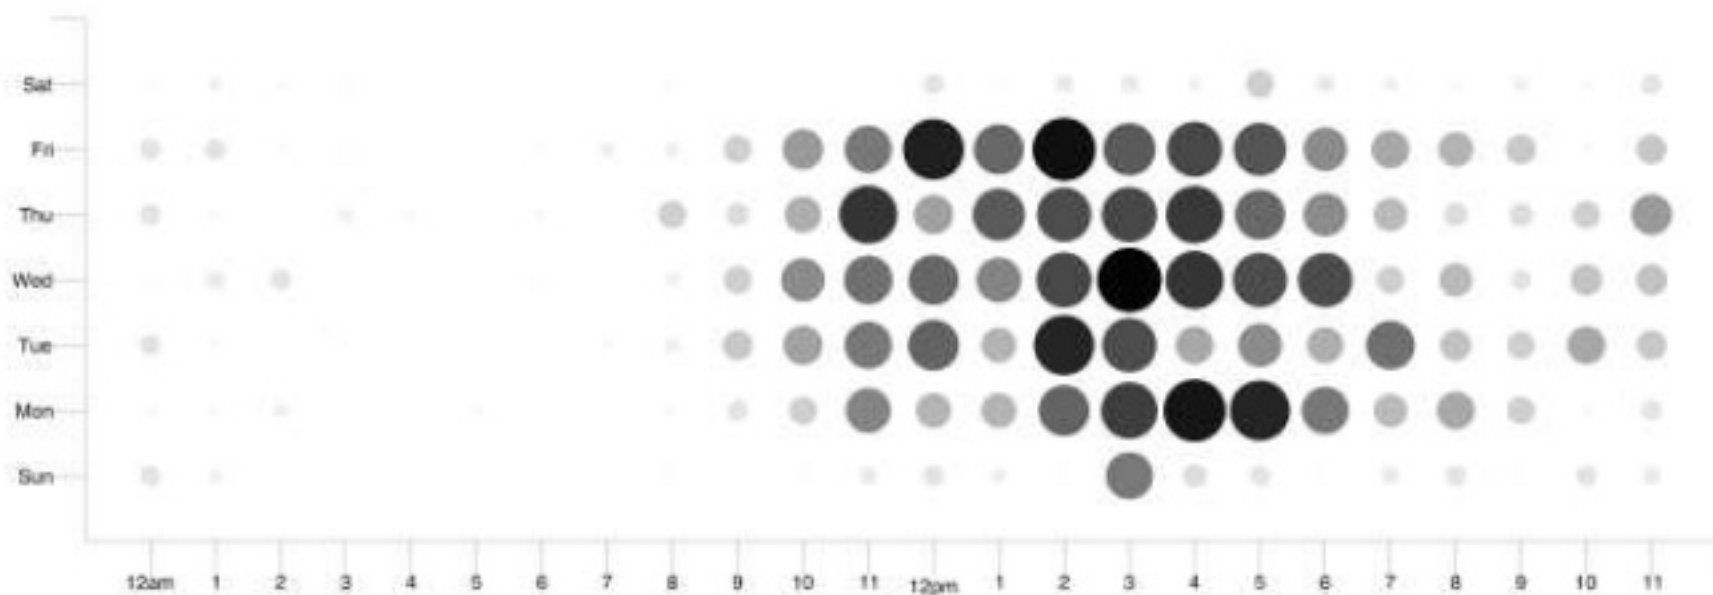
Apache Flink 1.3

# Previously on Apache Flink

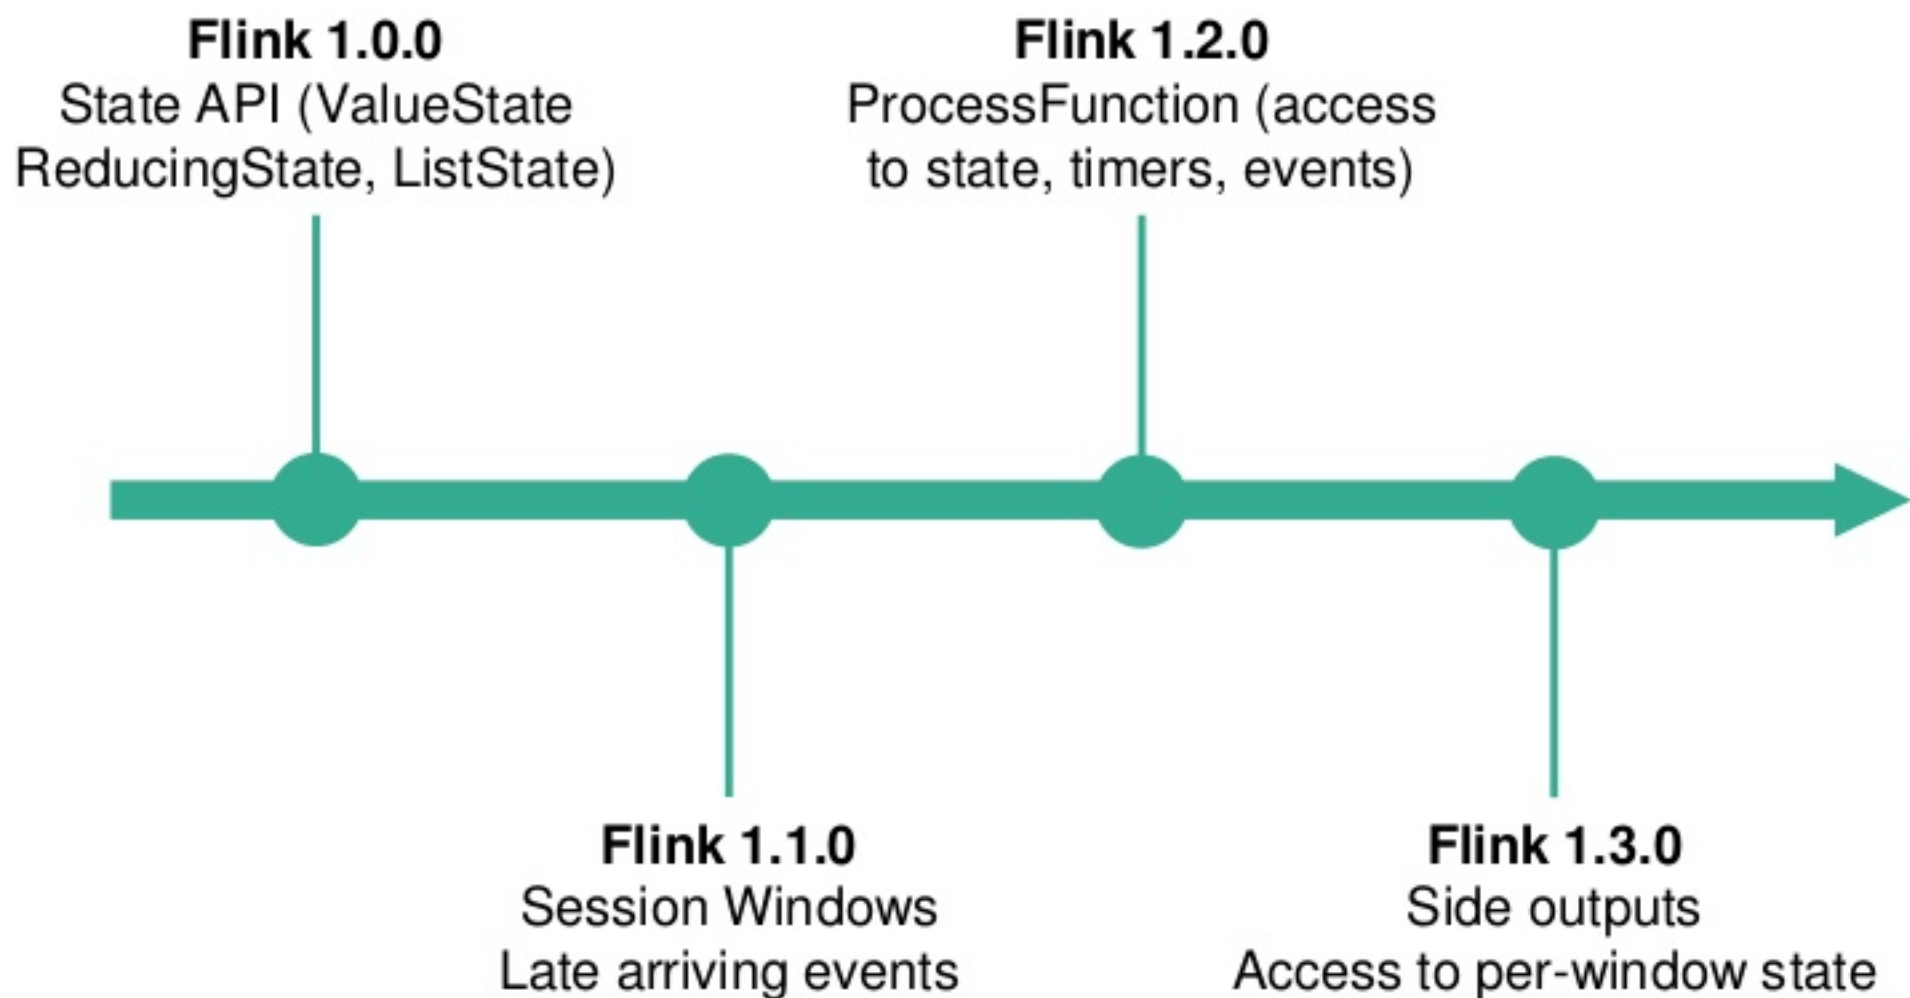
# Apache Flink 1.3 in Numbers



- ★ 141 contributors (no deduplication)
- ★ 1400 commits
- ★  $\geq 680$  resolved JIRA issues
- ★ +261813 / -65646 LOC



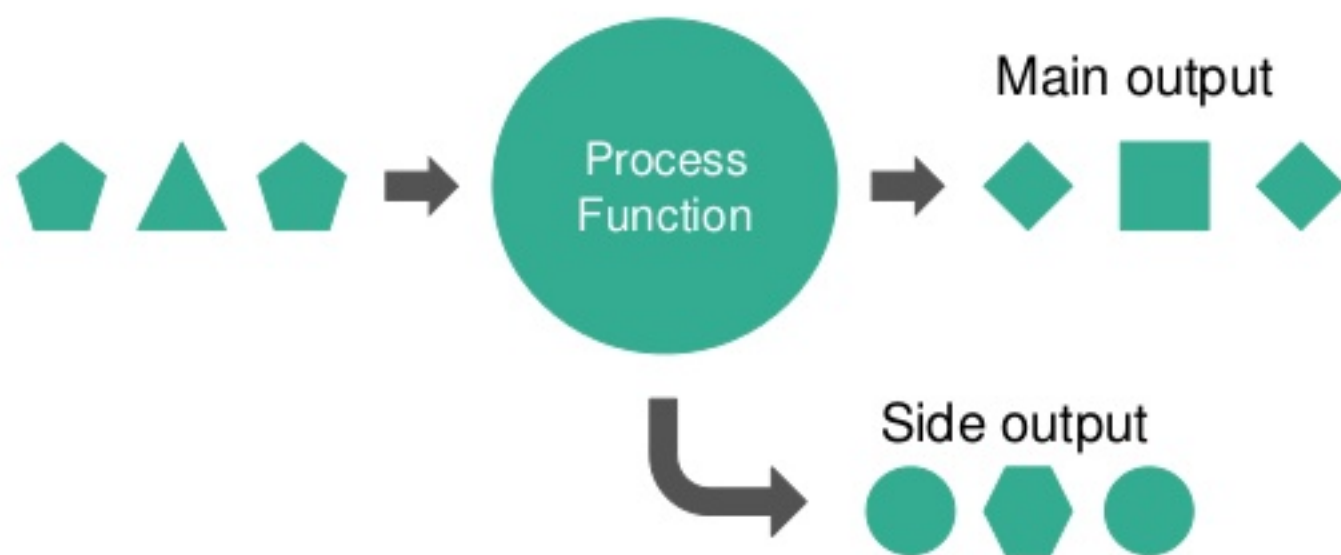
# Evolution of Flink's API



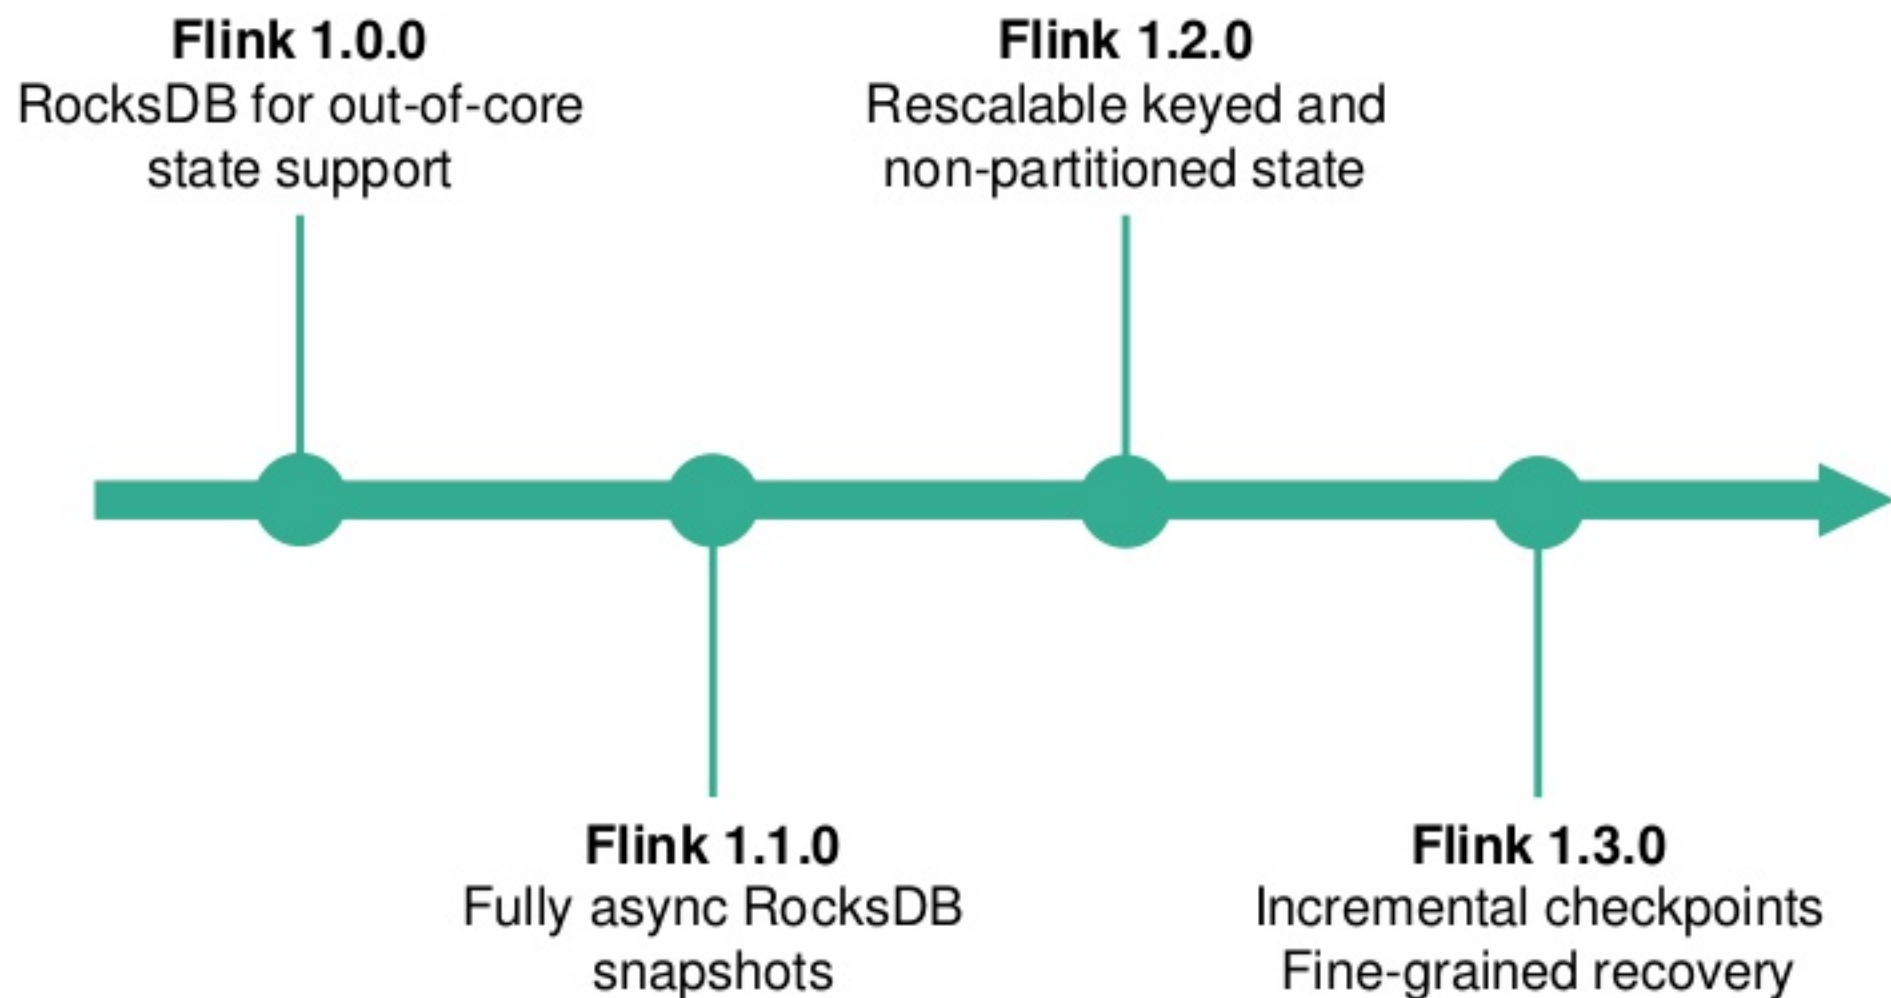
# Side Outputs



- ✦ Additional outputs for a stream
  - ⑩ Late events
  - ⑩ Corrupted input data
- ✦ More expressive APIs
- ✦ [FLINK-4460](#)

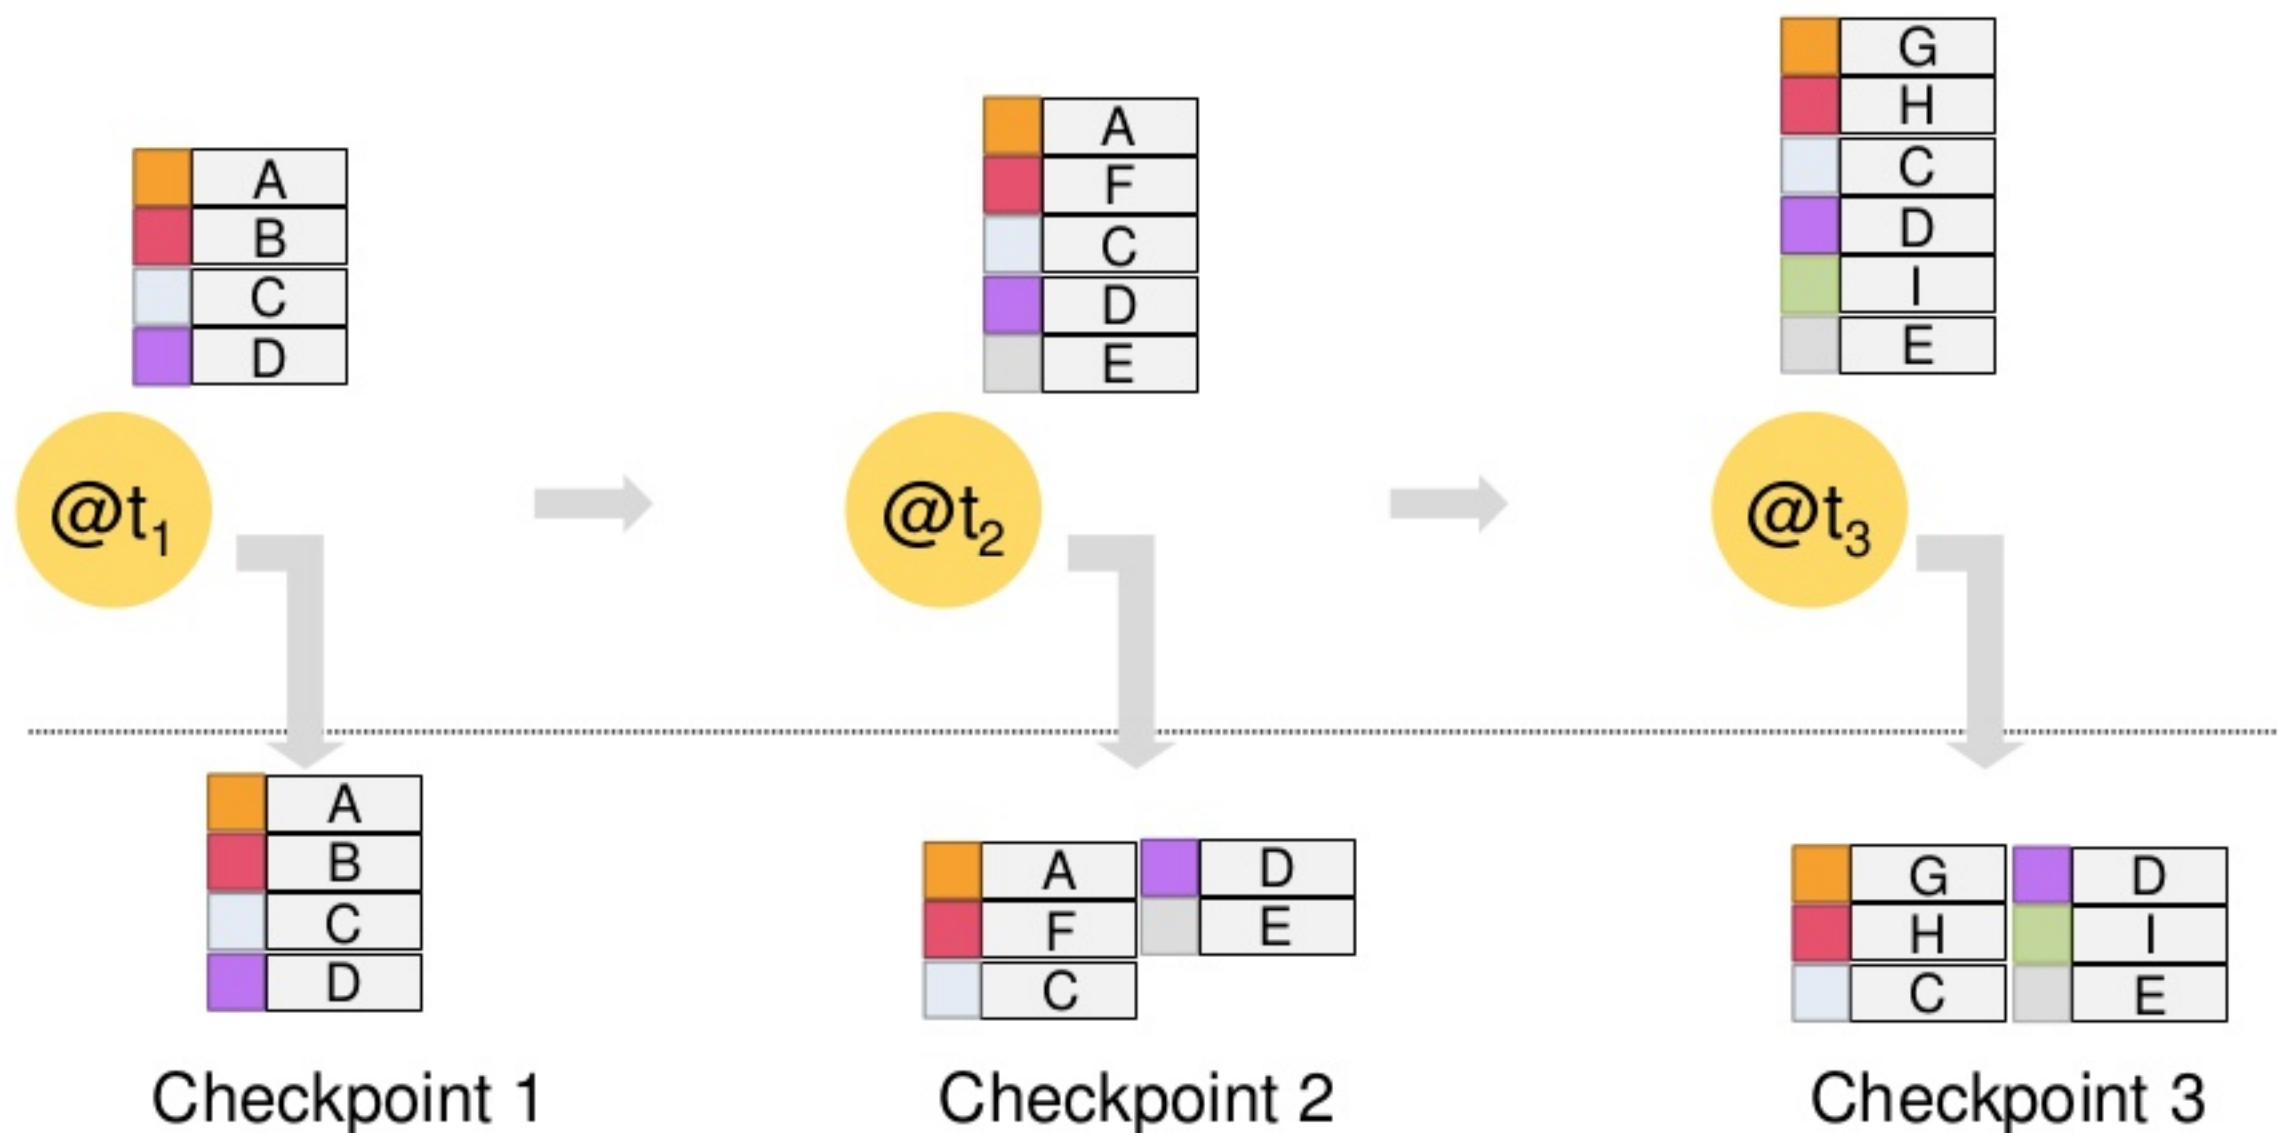


# Evolution of Large State Handling

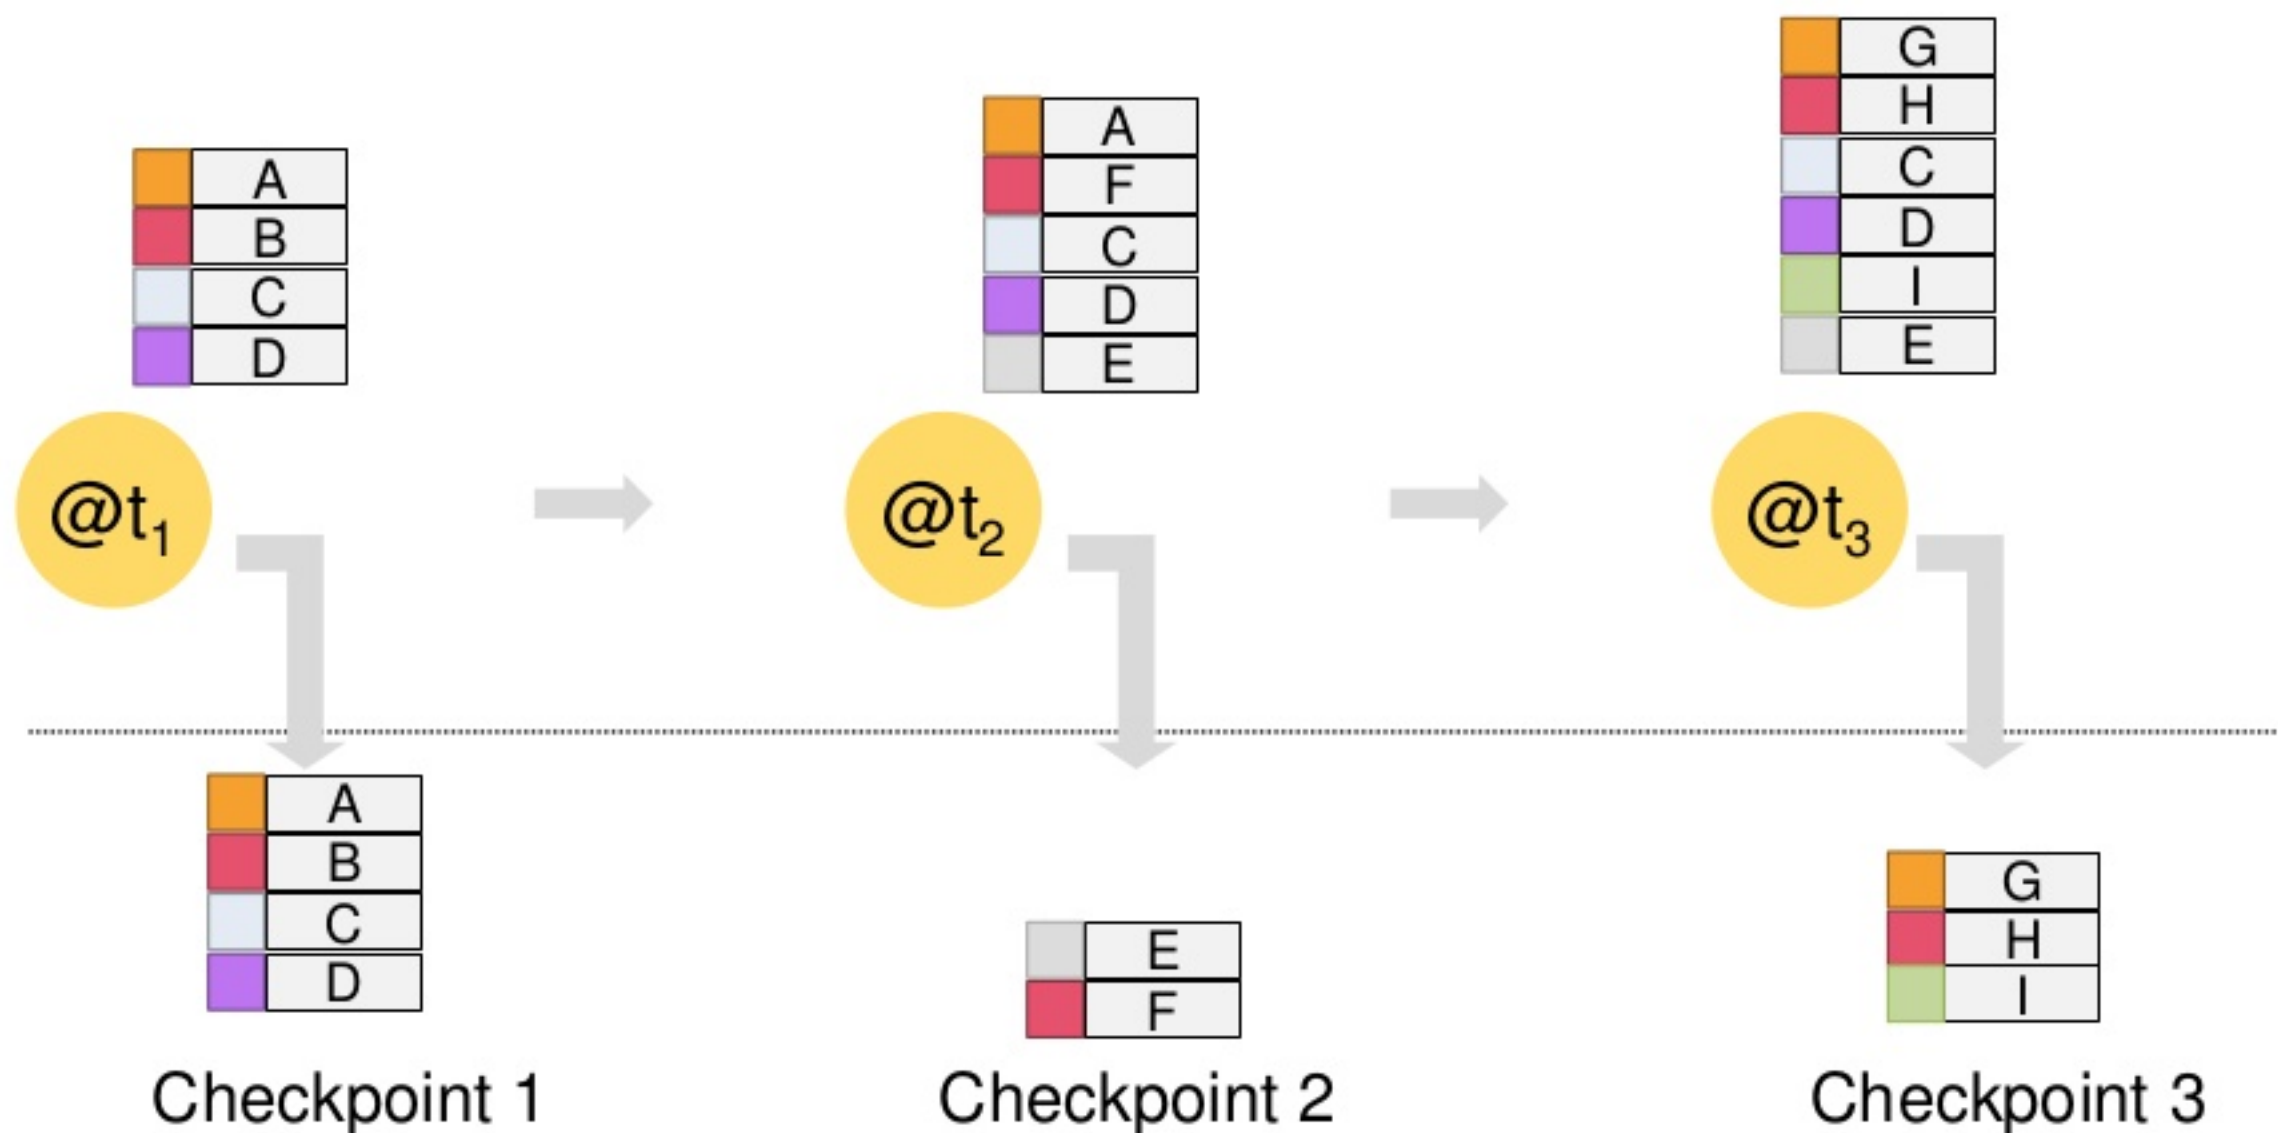




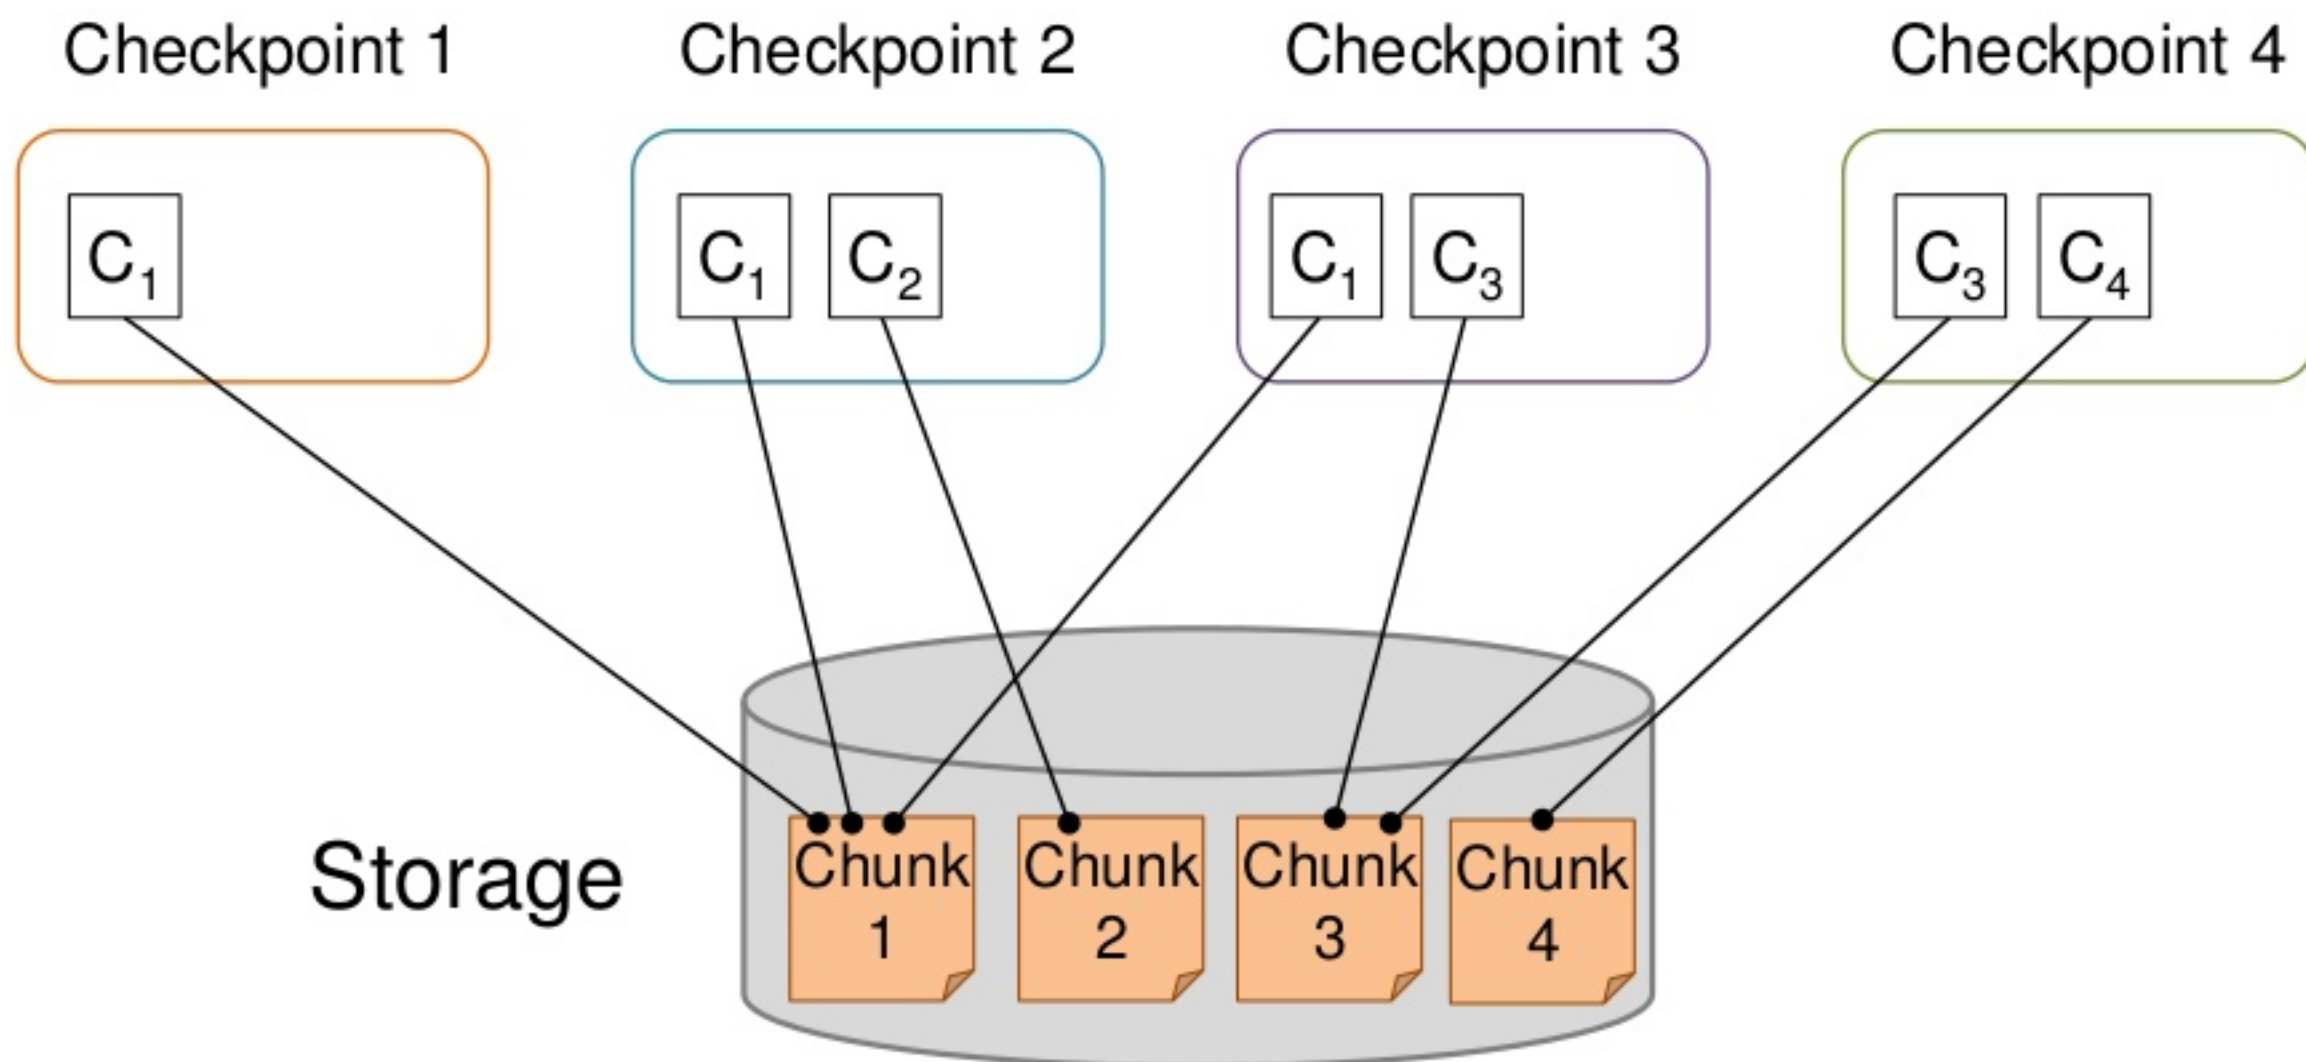
# Full Checkpoints



# Incremental Checkpoints



# Incremental Checkpoints



# Incremental Checkpointing Contd.

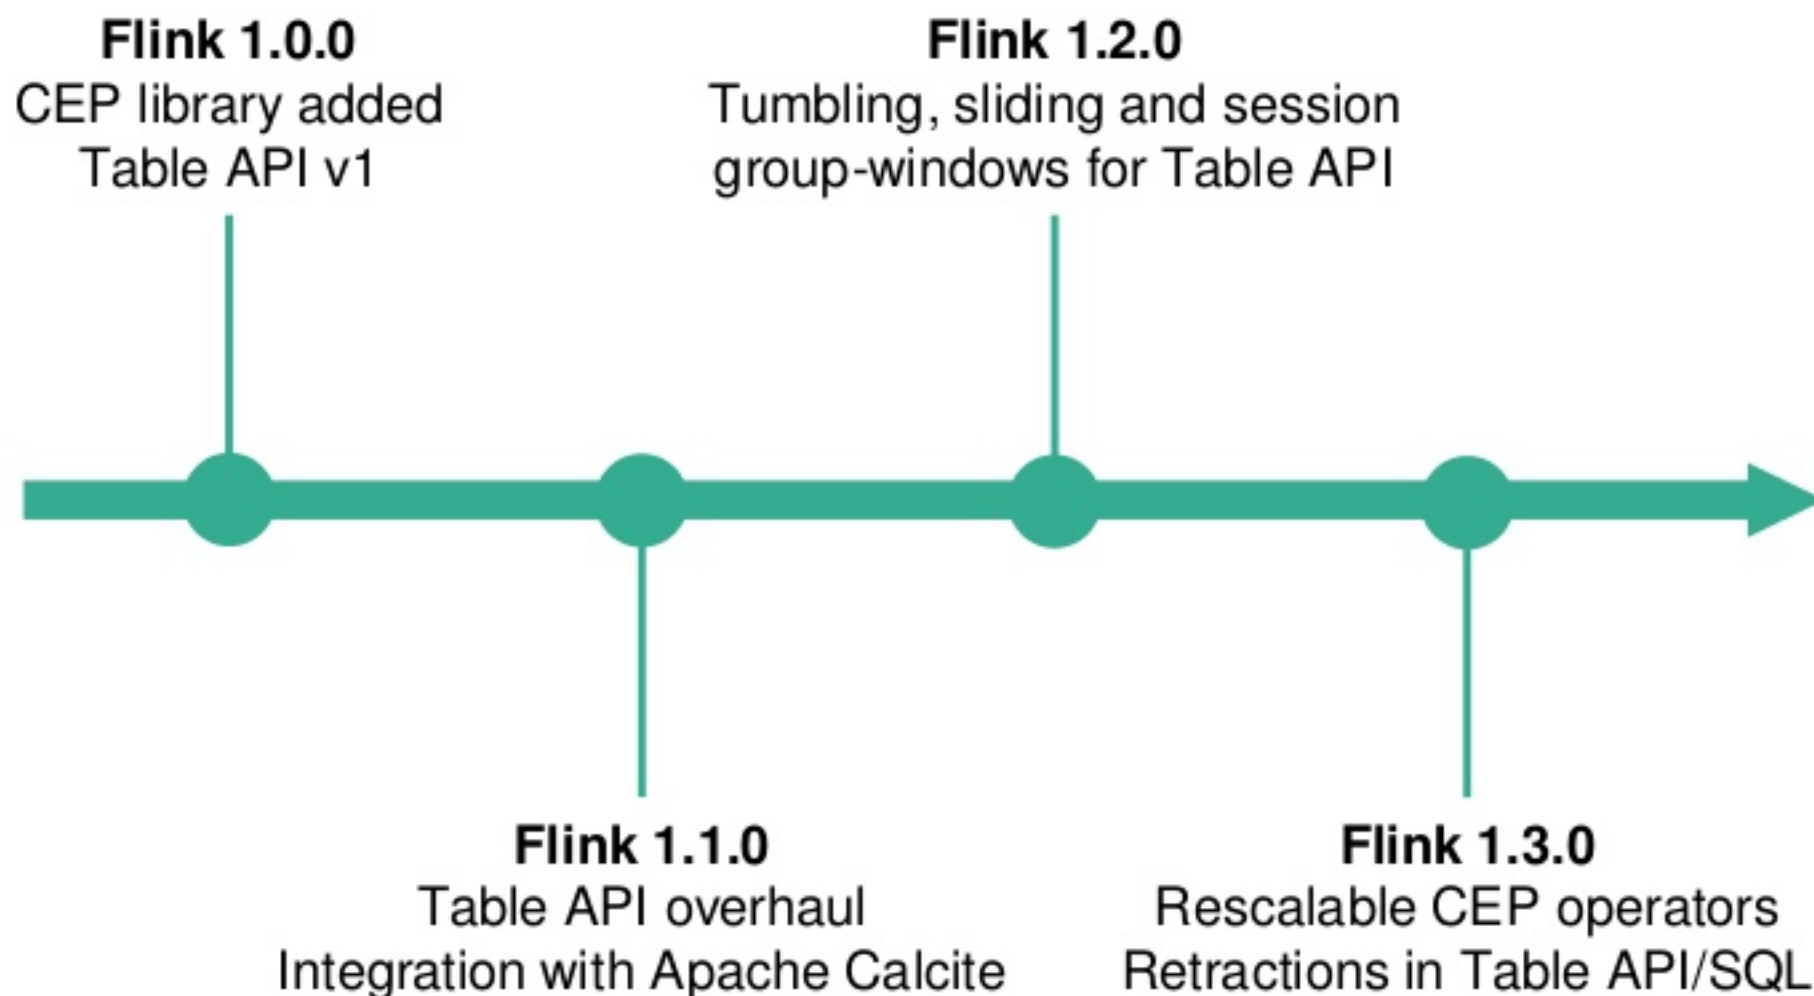


- ★ Currently supported for RocksDB state backend
- ★ [FLINK-5053](#)
- ★ Faster and smaller checkpoints

*“A Look at Flink’s Internal Data Structures and Algorithms for Efficient Checkpointing” by Stefan Richter, Tomorrow @ 12:20 pm Maschinenhaus*

	Full checkpoint	Incremental checkpoint
Size	60 GB	1 – 30 GB
Time	180 s	3 – 30 s

# Evolution of High Level APIs



# Enriched CEP Language



✦ Support for quantifiers (+, \*, ?)

⑩ [FLINK-3318](#)

✦ Iterative conditions

⑩ [FLINK-6197](#)

✦ Not operator

⑩ [FLINK-3320](#)

*“Complex Event Processing With Flink: The State of FlinkCEP” by Kostas Kloudas, Today @ 2:30 pm Maschinenhaus*





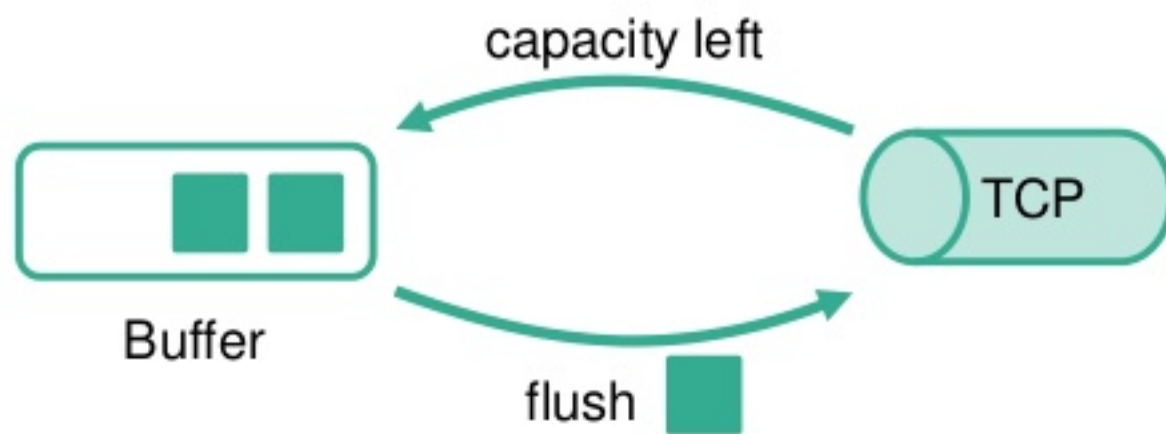
Apache Flink 1.4

# What's Happening Now?

# Event Driven I/O



- ✦ Rework of Flink's network stack
- ✦ Event driven network I/O
  - ⑩ Use full available capacity
  - ⑩ Near perfect latency behaviour





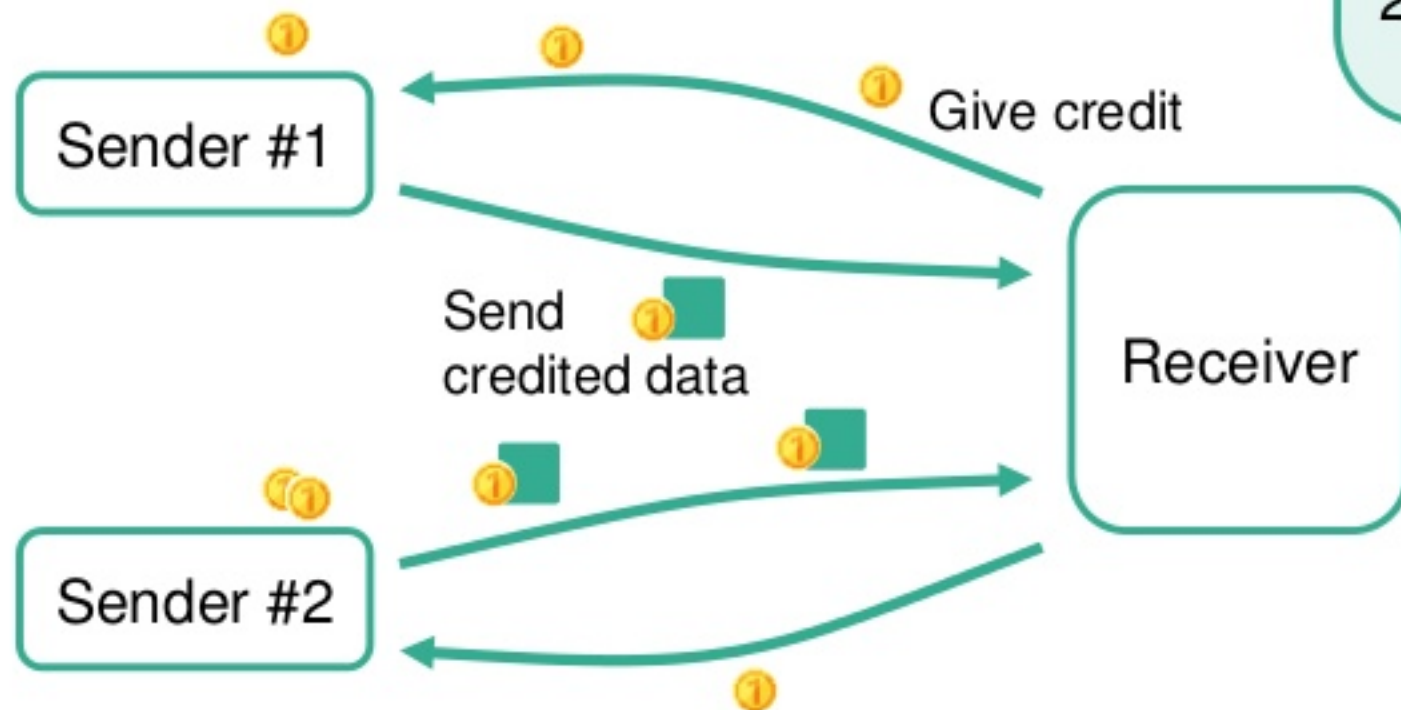
# Flow Control



## ★ Flow control for TaskManager communication

- ⑩ Single channel no longer stalls other multiplexed channels
- ⑩ Fine-grained backpressure control
- ⑩ Improves checkpoint alignments

*"Building a Network Stack for Optimal Throughput / Low-Latency Trade-Offs"*  
by Nico Kruber, Today @ 2:00 pm Palais Atelier

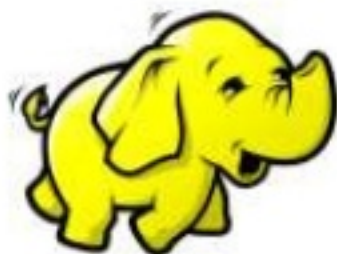


# New Deployment Model



- ★ Rework of Flink's distributed architecture
- ★ Ready for multitude of deployment scenarios
- ★ Support for dynamic scaling

*"Flink in Containerland"* by  
Patrick Lucas, Tomorrow  
@ 3:20 pm Maschinenhaus



MESOS

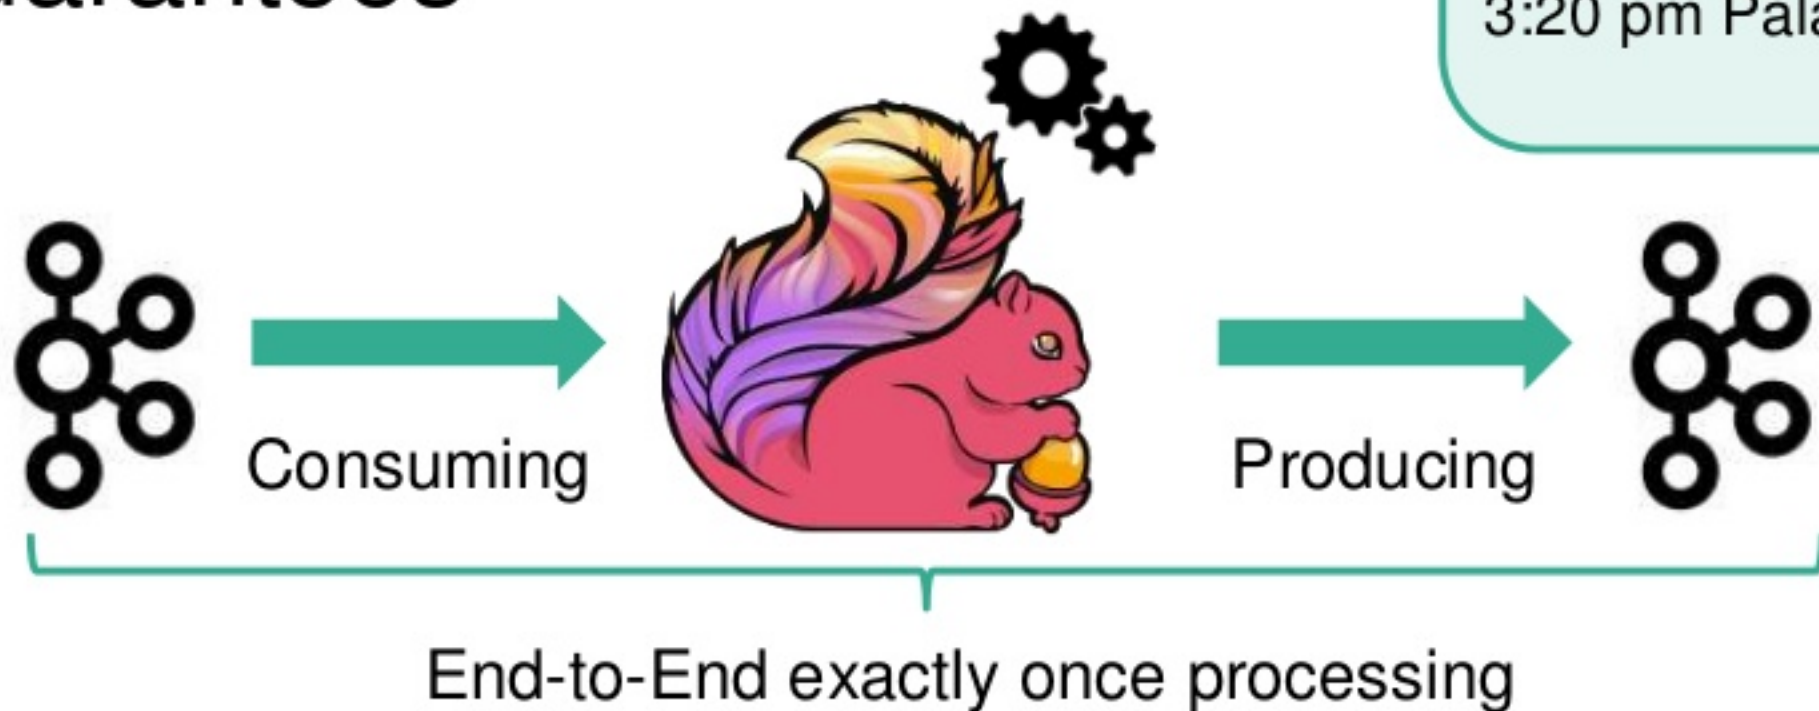


# Producing Exactly Once with Kafka 0.11



- ✦ Support for Kafka 0.11
- ✦ First Kafka producer with *exactly once processing* guarantees

*"Hit Me, Baby, Just One Time  
– Building End-to-End Exactly  
Once Applications With Flink"*  
by Piotr Nowowski, Today @  
3:20 pm Palais Atelier





# Operational Robustness



- ✦ Drop Java 7
- ✦ Support Scala 2.12
- ✦ Avoid dependency hell
  - ⑩ Child first class loading
  - ⑩ Relocation of dependencies
- ✦ De-Hadoopification





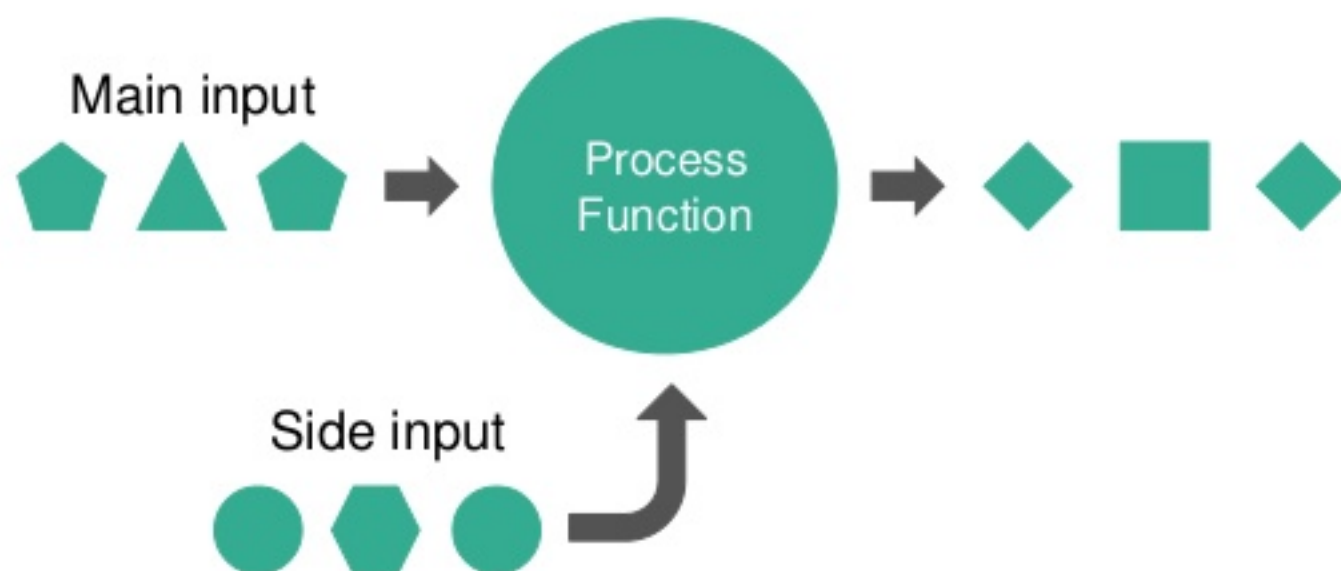
Apache Flink 1.5+

# Next on Apache Flink

# Side Inputs



- ✦ Additional input for operator
  - ⑩ Join with static data set
  - ⑩ Feeding of externally trained ML model
  - ⑩ Window joins
- ✦ Flip-17 design document: <https://goo.gl/W4yMEu>



# State Management & Evolution



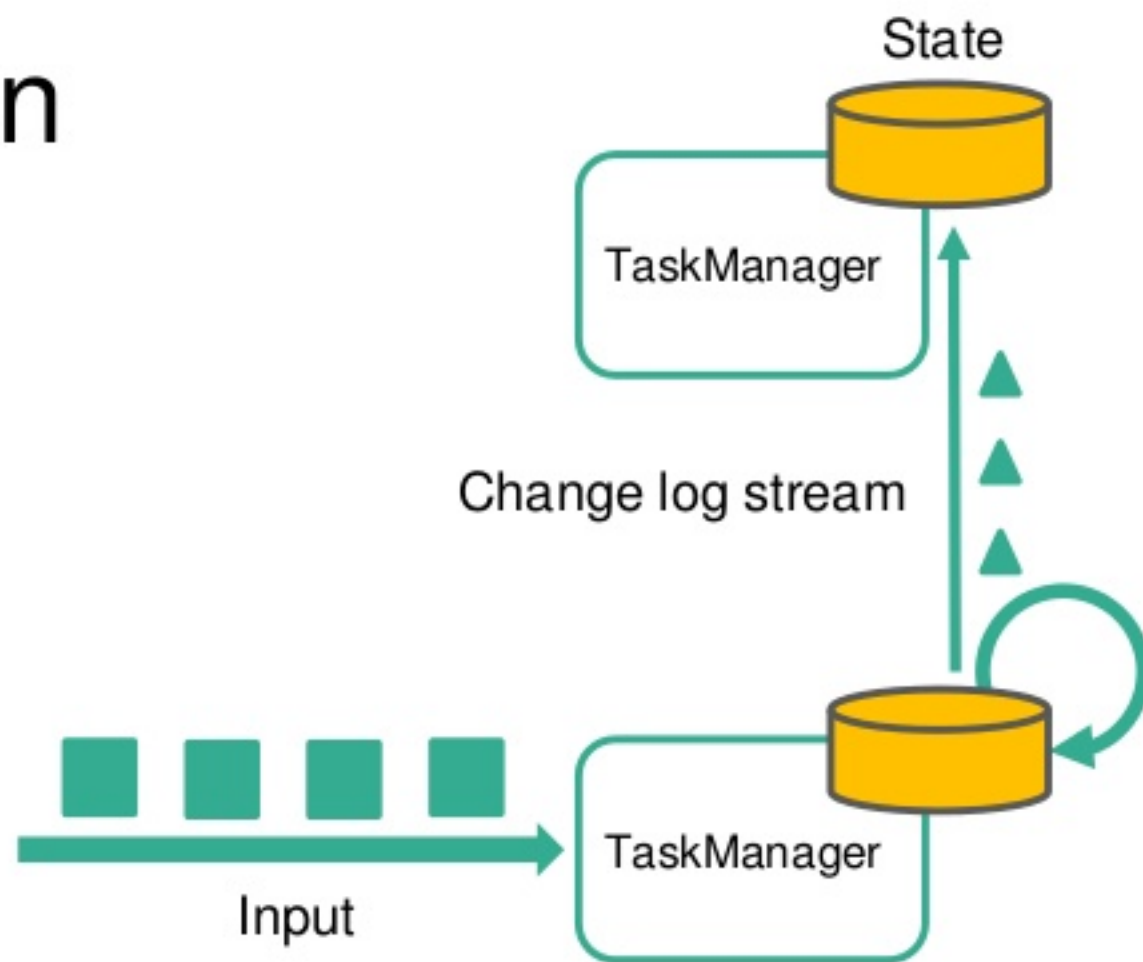
- ★ Eager state declaration
  - ⑩ State type, serializer and name known at pre-flight time
  - ⑩ Flip-22 design document: <https://goo.gl/trFiSi>
- ★ Evolving existing state
  - ⑩ Schema updates
  - ⑩ Serializer upgrades

*“Managing State in Apache Flink”* by  
Tzu-Li Tai, Today @  
4:30 pm Kesselhaus

# State Replication



- ✦ Replicate state between TaskManagers
  - 10 Faster recovery in case of failures
  - 10 High throughput queryable state





# Programmatic Job Control

---



- ✦ Improve client to give better job control
- ✦ Run concurrent jobs from the same program
- ✦ Trigger savepoints programmatically
- ✦ Better testing facilities

# JobClient & ClusterClient



```
StreamExecutionEnvironment env = ...;
// define program

JobClient jobClient = env.execute();

CompletableFuture<Acknowledge> savepointFuture = jobClient.takeSavepoint(savepointPath);

// wait for the savepoint completion
savepointFuture.get();

CompletableFuture<JobExecutionResult> resultFuture = jobClient.getResultFuture();

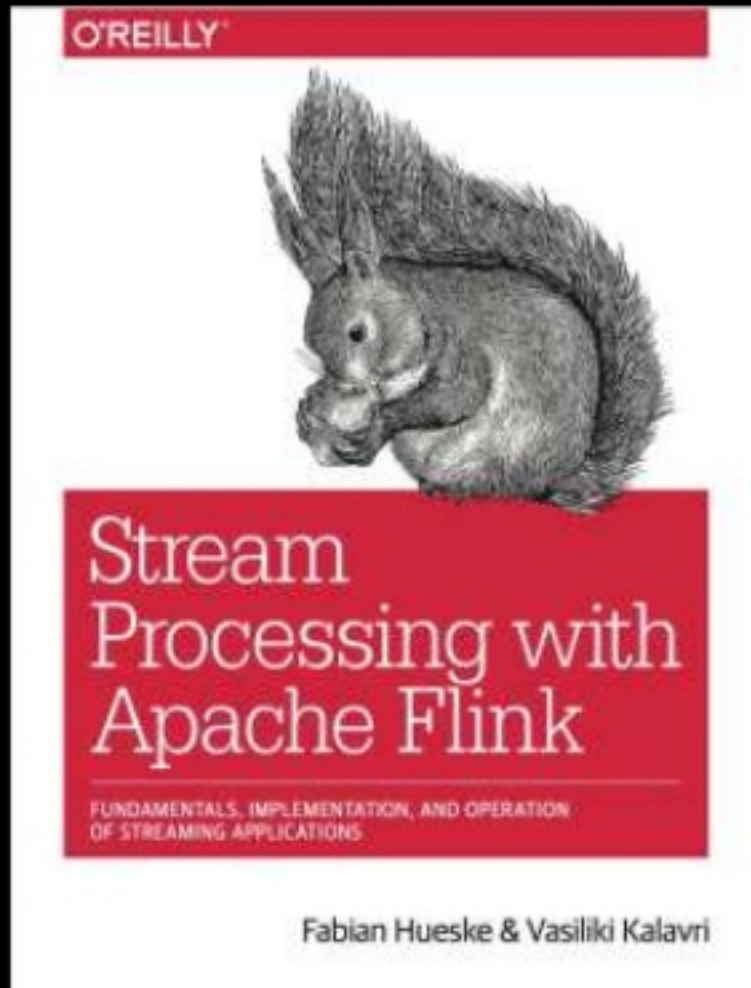
// cancel the job
jobClient.cancelJob();

// get the execution result --> should be canceled
JobExecutionResult result = resultFuture.get();

// get list of all still running jobs on the cluster
ClusterClient clusterClient = jobClient.getClusterClient();
CompletableFuture<List<JobInfo>> jobInfosFuture = clusterClient.getJobInfos();
List<JobInfo> jobInfos = jobInfosFuture.get();
```



- ✦ Apache Flink one of the most innovative open source stream processing platforms
- ✦ Stay tuned what's happening next 😊
- ✦ Visit the in depths talks to learn more about Flink's internals



# Thank you!

@stsffap

@ApacheFlink

@dataArtisans

# dataArtisans

We are hiring!

[data-artisans.com/careers](https://data-artisans.com/careers)